

MODULE 13 · DOMAIN 1 - AI & ML FUNDAMENTALS

Telling Similar Services Apart

The confusion pairs, the perpetual distractors, and how to eliminate fast.

DOMAIN all five EXAM WEIGHT n/a MODULE 13 of 14

This exam is not a recall test. It is a **discrimination** test — telling two plausible things apart.

Four shapes, and they repeat

- ◆ **Right data, wrong verb.** Comprehend offered for a scanned form. Both involve documents.
- ◆ **Right verb, wrong data.** Rekognition offered for invoice fields. Both extract text.
- ◆ **The heavier option.** Build on SageMaker when a managed service already does it.
- ◆ **The adjacent service.** Kendra offered when the task is classification, not search.
- ◆ A fifth: **true but irrelevant.** A correct statement that does not answer the question asked.

ON THE EXAM

Learn to name the shape. Once you can say 'that is the heavier option', you stop being tempted by it.

Ground Truth or A2I?

SAGEMAKER GROUND TRUTH

Humans label data

- Happens BEFORE training
- Creates the labelled dataset
- Input is raw, unlabelled data
- Purpose: make training possible

AMAZON A2I

Humans review predictions

- Happens AFTER a prediction
- Checks low-confidence outputs
- Input is a model's uncertain result
- Purpose: catch errors before they reach users

THE TIEBREAKER

Both use human workforces, which is what makes them confusable. Ask when in the lifecycle the humans appear.

Everything that reads text

SERVICE	THE ONE THING IT DOES
Textract	Pulls text and form structure out of scanned documents.
Comprehend	Finds meaning in text: sentiment, entities, language, PII.
Comprehend Medical	The same, but for clinical text and terminology.
Translate	Converts text between languages.
Rekognition	Reads text in photos and scenes, not structured documents.

ON THE EXAM

Detecting which language a document is written in is Comprehend, not Translate. Translate converts; Comprehend identifies.

Everything that searches or retrieves

SERVICE	CHOOSE IT WHEN
Kendra	Managed enterprise search returning direct answers from documents.
Bedrock Knowledge Bases	Managed RAG feeding a foundation model. AWS's successor to Kendra.
Amazon Q Business	A finished assistant with connectors and a UI already built.
OpenSearch	You are operating the vector or keyword index yourself.

ON THE EXAM

These overlap genuinely. Ask what you are building: a search box, a RAG pipeline, a finished product, or an index.

Everything that watches something

SERVICE	WATCHES
SageMaker Model Monitor	A deployed model, for drift and quality decay.
SageMaker Clarify	A model or dataset, for bias and explainability.
Amazon CloudWatch	Metrics, logs and alarms. Performance right now.
AWS CloudTrail	API calls. Who did what, when, from where.
AWS Config	Resource configuration against policy.
AWS Trusted Advisor	Your account, for cost and best-practice advice.

ON THE EXAM

Six services that all 'monitor'. The question is always what they monitor, never whether they monitor.

Everything that finds sensitive data

SERVICE	WHERE IT LOOKS
Amazon Macie	Inside objects stored in S3.
Amazon Comprehend	Inside text you pass to it, and can redact it.
Bedrock Guardrails	Inside prompts and model responses, and can mask it.
Amazon Inspector	Nowhere — it scans workloads for vulnerabilities.
Amazon GuardDuty	Nowhere — it detects threats and unusual activity.

ON THE EXAM

The last two appear in the same option lists and never find sensitive data. They are the security-sounding distractors.

Guardrails or Clarify?

BEDROCK GUARDRAILS

Controls what is said

- Blocks harmful or off-topic content
- Operates at runtime, on every request
- Denied topics, PII masking, grounding checks
- For generative applications

SAGEMAKER CLARIFY

Analyses what the model is

- Measures bias and explains predictions
- Operates on datasets and trained models
- Feature importance, SHAP, bias reports
- For models you built

THE TIEBREAKER

Guardrails prevents bad output in the moment. Clarify diagnoses the model's properties. Prevention versus analysis.

Everything that documents or proves

ARTEFACT	WHO WRITES IT, AND ABOUT WHAT
SageMaker Model Card	You write it, about a model you built.
AWS AI Service Card	AWS writes it, about an AWS service.
AWS Artifact	AWS's own compliance certifications, for you to download.
AWS Audit Manager	Evidence about your workloads, collected continuously.
Model Registry	Versioned models with approval status before deployment.

ON THE EXAM

Two questions separate all five: is it about you or about AWS, and is it documentation or evidence?

The customisation ladder, in order of cost

APPROACH	USE WHEN
Prompt engineering	The model can do it but needs clearer instruction. Free.
RAG	The model lacks facts. Ground it in your documents.
Fine-tuning	The model lacks a behaviour. Needs labelled pairs.
Continued pre-training	The model lacks vocabulary. Needs unlabelled domain text.
Train from scratch	Almost never. Months and millions.

ON THE EXAM

An answer that skips down this ladder without justification is usually wrong. The exam rewards trying the cheap option first.

Services that are nearly always distractors

- ◆ **Amazon Inspector** and **GuardDuty** — offered for anything security-adjacent.
- ◆ **AWS Trusted Advisor** — offered whenever a question mentions monitoring or cost.
- ◆ **AWS Config** — offered whenever a question mentions auditing.
- ◆ **Amazon Forecast** — offered for anything involving future values.
- ◆ Each is real and useful. Each answers a question the exam is not asking.

ON THE EXAM

If one of these appears alongside a service purpose-built for the described task, the purpose-built one wins.

How to eliminate quickly

- ◆ **Name the verb** the scenario needs. Extract, classify, generate, search, monitor, document.
- ◆ **Name the input.** Scanned image, digital text, audio, tabular rows, a question.
- ◆ Those two alone eliminate two options on most questions.
- ◆ **Then check effort.** If two remain, prefer the managed one unless control was requested.
- ◆ If an option is true but does not answer the question asked, it is wrong.

ON THE EXAM

Verb, then input, then effort. In that order. Most questions resolve before you reach the third step.

Carry these into the exam

- ◆ Ground Truth labels before training. A2I reviews after a prediction.
- ◆ Detecting a language is Comprehend. Converting it is Translate.
- ◆ Guardrails prevents bad output. Clarify analyses model properties.
- ◆ Inspector, GuardDuty, Trusted Advisor and Config are usually distractors.
- ◆ Model Card is yours, AI Service Card is Amazon's, Artifact is AWS's certifications.
- ◆ Verb, then input, then effort.

A logistics company has two hundred thousand unlabelled photographs of damaged packages. It wants to train a model to categorise damage type.

It needs people to categorise a sample first, so there is something to train on.

Which service does this?

- | | |
|--|---|
| ☐ A Amazon A2I | ☐ B SageMaker Ground Truth |
| ☐ C SageMaker Clarify | ☐ D Amazon Rekognition |

◆ PAUSE AND DECIDE

Which service does this?

- A

 Amazon A2I
Reviews predictions. There is no model yet.
- B

 Ground Truth
Humans label raw data so training can begin.
- C

 Clarify
Analyses bias and explains predictions. No labelling.
- D

 Rekognition
Pre-trained vision. It cannot label your custom categories.

ON THE EXAM

'So there is something to train on' places this before training. A2I cannot apply because no model exists yet.

A global publisher receives manuscripts in many languages. Before routing each one to the correct editorial team, it needs to identify which language the document is written in.

No translation is required at this stage.

Which service should it use?

- | | |
|---|--|
| ☐ A Amazon Translate | ☐ B Amazon Comprehend |
| ☐ C Amazon Textract | ☐ D Amazon Polly |

◆ PAUSE AND DECIDE

Which service should it use?

- A

 Amazon Translate
Converts between languages. Nothing is being converted.
- B

 Amazon Comprehend
Dominant language detection is a Comprehend feature.
- C

 Amazon Textract
Extracts text from scans. The text is already digital.
- D

 Amazon Polly
Text to speech. Wrong direction entirely.

ON THE EXAM

The sentence 'no translation is required' is placed there deliberately. Read it as the answer key.

A generative assistant occasionally produces responses that mention competitor brands and, on rare occasions, mildly offensive language.

The company wants these blocked before any user sees them.

Which service applies?

- ☐ A SageMaker Clarify
- ☐ B Bedrock Guardrails
- ☐ C SageMaker Model Monitor
- ☐ D Amazon Comprehend

◆ PAUSE AND DECIDE

Which service applies?

A Clarify

Analyses bias and explainability. Does not block output.

C Model Monitor

Watches SageMaker models for drift. Not content.

B Guardrails

Word filters and content filters block before delivery.

D Comprehend

Analyses text. It reports; it does not block.

ON THE EXAM

'Before any user sees them' means prevention at runtime. Every other option analyses or reports after the fact.

A model that scores loan applications has been in production for a year. The risk team wants to know whether its decisions have started to disadvantage any age group since launch. They already have bias reports from before deployment.

What should they use now?

- ☐ A SageMaker Clarify again
- ☐ B SageMaker Model Monitor
- ☐ C Amazon CloudWatch
- ☐ D AWS Config

◆ PAUSE AND DECIDE

What should they use now?

- A

 Clarify again
Point-in-time analysis. They already have that.
- B

 Model Monitor
Continuous monitoring of a live model, including bias drift.
- C

 CloudWatch
Latency and error rates. Not fairness.
- D

 AWS Config
Resource configuration. Not model behaviour.

ON THE EXAM

Clarify and Model Monitor overlap on bias. Clarify is a snapshot; Model Monitor watches continuously in production.

A support team wants a foundation model to answer in a consistent house style: second person, no jargon, always ending with a next step.

They have no labelled examples but can write a few model answers by hand.

What should they try first?

- ☐ A Fine-tune the model
- ☐ B Few-shot prompting with those examples
- ☐ C Continued pre-training
- ☐ D Build a RAG pipeline

◆ PAUSE AND DECIDE

What should they try first?

- A

 Fine-tune
Right for style, but skips the free option and needs a dataset.
- B

 Few-shot prompting
A few hand-written examples demonstrate the style immediately.
- C

 Continued pre-training
Teaches vocabulary. Style is not vocabulary.
- D

 RAG
Supplies facts. The problem is style, not knowledge.

ON THE EXAM

Fine-tuning is genuinely the right tool for style — at scale. 'A few examples' and 'try first' both point down the ladder.

A company's board asks for a single document describing what a model it built is for, how it was trained, how well it performs, and where it should not be used.

They are not asking about AWS's own certifications.

Which artefact is this?

- ☐ A AWS Artifact report
- ☐ B AI Service Card
- ☐ C SageMaker Model Card
- ☐ D Audit Manager assessment

◆ PAUSE AND DECIDE

Which artefact is this?

- A

 AWS Artifact
AWS's certifications. Explicitly excluded by the scenario.
- B

 AI Service Card
AWS documenting an AWS service, not your model.
- C

 SageMaker Model Card
Purpose, training, performance, limitations. Yours.
- D

 Audit Manager
Compliance evidence, not model documentation.

ON THE EXAM

Four artefacts, one question: is it about your model, or about AWS? Everything else follows.

An operations team needs to detect when personally identifiable information has been uploaded by mistake into the S3 bucket that feeds a knowledge base.

Four security services are proposed.

Which one does this?

- | | |
|---|--|
| ☐ A Amazon Inspector | ☐ B Amazon Macie |
| ☐ C Amazon GuardDuty | ☐ D AWS Trusted Advisor |

◆ PAUSE AND DECIDE

Which one does this?

- A

 Inspector
Software vulnerabilities and network exposure.
- B

 Macie
Discovers and classifies sensitive data inside S3 objects.
- C

 GuardDuty
Threat detection from account and network activity.
- D

 Trusted Advisor
Account-level cost and best-practice recommendations.

ON THE EXAM

This exact option set recurs across the exam. Only Macie opens the objects and reads what is inside them.

A small retailer with no technical staff wants to predict weekly demand for around four hundred products, using three years of its own sales spreadsheets.

Four approaches are proposed and all four would technically work.

Which should they choose?

- ☐ A Train a custom model with SageMaker training jobs
- ☐ B SageMaker Canvas
- ☐ C A Bedrock foundation model with prompt engineering
- ☐ D Build a RAG pipeline over the spreadsheets

◆ PAUSE AND DECIDE

Which should they choose?

- A

 Custom training jobs
The heavier option. No technical staff exist.
- B

 SageMaker Canvas
No-code forecasting on their own tabular data.
- C

 Bedrock with prompting
Generative models are poor at numeric forecasting.
- D

 RAG over spreadsheets
Retrieval answers questions. It does not forecast.

ON THE EXAM

When several options work, effort and fit decide. Tabular forecasting plus no technical staff is Canvas, every time.