

PRACTICE EXAM 1 OF 3 · EXAM STRATEGY - ALL DOMAINS

Practice Exam 1

Twenty questions, mixed across all five domains — as the real exam presents them.

QUESTIONS 20 MIX Weighted PART 1 of 3

How to use this

- ◆ Pause on each question. Commit to an answer before the walkthrough.
- ◆ Give yourself about **80 seconds** per question, matching real exam pace.
- ◆ Write down your answers. Score at the end and note the domain of each miss.
- ◆ Watching the walkthrough without answering first teaches you almost nothing.

ON THE EXAM

The domain marker on the left edge shows which domain each question belongs to. Track where your misses cluster.

A car manufacturer's assistant answers dealer questions about warranty terms. Terms are revised quarterly and differ by market.

The assistant gives outdated answers. Current warranty documents exist in S3; there is no labelled training data.

What should the company implement?

- ☐ A Fine-tune a foundation model on the warranty documents
- ☐ B Implement RAG over the warranty documents
- ☐ C Train a model from scratch
- ☐ D Increase the model's context window

◆ PAUSE AND DECIDE

What should the company implement?

- A

 Fine-tune
Needs labelled pairs, and freezes quarterly-changing terms.
- B

 Implement RAG
Grounds answers in current documents; updates when they do.
- C

 Train from scratch
Months and millions for a knowledge problem.
- D

 Bigger context window
More room for text the model still does not have.

ON THE EXAM

Knowledge problem, changing content, documents but no labelled data. Three signals, all pointing at RAG.

A property firm has ten years of sales records with location, size, condition and final sale price.

It wants to estimate the likely selling price of new listings.

Which type of ML problem is this?

- ☐ A Binary classification
- ☐ C Regression

- ☐ B Clustering
- ☐ D Anomaly detection

◆ PAUSE AND DECIDE

Which type of ML problem is this?

- A

 Binary classification
Output would be a category. Price is a quantity.
- B

 Clustering
Unsupervised. The data already has known prices.
- C

 Regression
Labelled history, and the output is a number.
- D

 Anomaly detection
Finds unusual sales. Does not estimate price.

ON THE EXAM

Price is a number. Number out means regression, every time.

A compliance officer must determine which IAM user invoked a specific Bedrock model last month, and from which IP address.

The information is needed for an internal investigation.

Which service provides this?

- | | |
|--|--|
| ☐ A Amazon CloudWatch | ☐ B AWS Config |
| ☐ C AWS CloudTrail | ☐ D AWS Trusted Advisor |

◆ PAUSE AND DECIDE

Which service provides this?

- A

 CloudWatch
Performance metrics and application logs, not caller identity.
- B

 AWS Config
Resource configuration state.
- C

 CloudTrail
Records who called which API, when, and from where.
- D

 Trusted Advisor
Cost and best-practice recommendations.

ON THE EXAM

'Which user' and 'which IP' are identity questions. Only CloudTrail records those.

Users of a customer service chatbot complain that it forgets details they mentioned earlier in the same conversation.

The development team confirms each request is sent to the model independently.

What explains this?

- ☐ A The model needs fine-tuning on conversation data
- ☐ B Foundation models have no memory between calls; history must be resent
- ☐ C The temperature is set too high
- ☐ D The model is hallucinating

◆ PAUSE AND DECIDE

What explains this?

- A

 Fine-tune on conversations
Teaches style, not memory of this conversation.
- C

 Temperature too high
Affects variation, not recall.
- B

 No memory between calls
The application must resend history in the context window.
- D

 Hallucinating
Inventing facts is different from forgetting them.

ON THE EXAM

'Each request is sent independently' is the answer stated in the scenario. The model is stateless by design.

A regulatory reporting assistant must produce identical wording when asked the same question twice.

Currently, repeated questions return differently phrased answers.

Which inference parameter should be adjusted?

- | | |
|--|--|
| ☐ A Increase max tokens | ☐ B Lower the temperature |
| ☐ C Increase top-k | ☐ D Raise the temperature |

◆ PAUSE AND DECIDE

Which inference parameter should be adjusted?

- A

 Increase max tokens
Allows longer answers, not more consistent ones.
- B

 Lower the temperature
Concentrates sampling on likely tokens, increasing consistency.
- C

 Increase top-k
Widens the candidate pool. Makes variation worse.
- D

 Raise the temperature
Directly the wrong direction.

ON THE EXAM

Options C and D both move away from consistency. Two of the four are wrong in direction, not just in choice.

A lender's model approves applicants from some postcodes at markedly lower rates. The team removes postcode from the features and retrains.

The disparity persists.

What is the most likely explanation?

- ☐ A The model needs more training epochs
- ☐ B Other features act as proxies for postcode
- ☐ C The test set is too small
- ☐ D The model is underfitting

◆ PAUSE AND DECIDE

What is the most likely explanation?

- A

 More epochs
Trains harder on the same biased data.
- B

 Proxy features
Employer, school and shopping patterns still encode location.
- C

 Test set too small
Would cause noisy metrics, not a consistent disparity.
- D

 Underfitting
Would show as poor accuracy overall.

ON THE EXAM

Deleting the sensitive attribute is the intuitive fix and does not work. Proxies survive.

A council receives thirty thousand handwritten planning objection forms a year. Staff retype the applicant name, address and objection category into a case system.

The council has no ML engineers.

Which service should it use?

- | | |
|--|---|
| ☐ A Amazon Comprehend | ☐ B Amazon Rekognition |
| ☐ C Amazon Textract | ☐ D Amazon Translate |

◆ PAUSE AND DECIDE

Which service should it use?

- A

 Comprehend
Cannot read a scanned image. Needs digital text.
- B

 Rekognition
Reads text in photos, but not structured form fields.
- C

 Textract
Reads handwriting and returns form fields with values.
- D

 Translate
Converts languages. Nothing is being translated.

ON THE EXAM

Follow the input. Handwritten paper forms means Textract, and named form fields confirm it.

A shipping company finds a foundation model misreads its documentation, treating terms like demurrage, laytime and bill of lading as ordinary words.

It holds forty years of shipping contracts and manuals, but no question-and-answer datasets.

Which customisation approach fits?

- ☐ A Instruction fine-tuning
- ☐ B RLHF
- ☐ C Continued pre-training
- ☐ D Prompt engineering alone

◆ PAUSE AND DECIDE

Which customisation approach fits?

- A

 Instruction fine-tuning
Needs labelled pairs. The company has documents.
- B

 RLHF
Aligns tone and preference. The gap is vocabulary.
- C

 Continued pre-training
Unlabelled domain text teaching specialist vocabulary.
- D

 Prompt engineering
Cannot teach words the model has never learned.

ON THE EXAM

Follow the data. Unlabelled documents plus a vocabulary gap equals continued pre-training.

An airline wants an assistant that answers baggage policy questions and can also check a live booking and change a seat assignment.

Policies are in S3. Booking and seat changes are available through internal APIs.

What architecture is required?

- ☐ A RAG over the policy documents
- ☐ A fine-tuned model
- ☐ An agent with action groups plus a knowledge base
- ☐ A larger context window

◆ PAUSE AND DECIDE

What architecture is required?

- A

 RAG alone
Answers the policy half. Cannot change a seat.
- B

 Fine-tuned model
Neither retrieves current policy nor performs actions.
- C

 Agent plus knowledge base
Retrieval for policy, action groups for booking changes.
- D

 Larger context window
Still read-only.

ON THE EXAM

Two requirements. Changing a seat modifies state, and only an agent can do that.

A government department must call Amazon Bedrock from applications in a private subnet. Policy prohibits any traffic reaching the public internet. There is no internet gateway in the VPC.

What should be configured?

- ☐ A A NAT gateway
- ☐ B An IAM policy granting bedrock:InvokeModel
- ☐ C A VPC endpoint using AWS PrivateLink
- ☐ D AWS Shield

◆ PAUSE AND DECIDE

What should be configured?

- A

 NAT gateway
Routes to the internet. Exactly what policy forbids.
- B

 IAM policy
Necessary for access, but creates no network path.
- C

 VPC endpoint with PrivateLink
Reaches Bedrock without leaving the AWS network.
- D

 AWS Shield
DDoS protection. Wrong direction.

ON THE EXAM

Option B is necessary but not sufficient. Permission and connectivity are separate problems.

A demand forecasting model achieves 96% accuracy on its training data and 58% on held-out test data.

The team is deciding what to do next.

What is happening?

- ☐ A Underfitting; train for longer
- ☐ B Overfitting; simplify the model or add data
- ☐ C The metric is wrong; switch to RMSE
- ☐ D Data leakage; remove the target column

◆ PAUSE AND DECIDE

What is happening?

- A

 Underfitting
Training accuracy would also be poor. It is 96%.
- B

 Overfitting
Excellent on seen data, weak on unseen. Textbook.
- C

 Wrong metric
Changing the metric does not close the gap.
- D

 Data leakage
Leakage inflates both scores, not just training.

ON THE EXAM

Read the two numbers. Training high, test low, means overfitting. Both low means underfitting.

An insurance company wants an application where an assessor uploads a photograph of vehicle damage and asks questions about it in text.

The response should be a written assessment. No images need to be generated.

Which model type is required?

- ☐ A A text-only model such as Nova Micro
- ☐ B An image generation model such as Nova Canvas
- ☐ C A multimodal model such as Nova Pro
- ☐ D An embeddings model

◆ PAUSE AND DECIDE

Which model type is required?

- A

 Nova Micro
Text only. Cannot accept an image.
- B

 Nova Canvas
Generates images. The scenario says none are needed.
- C

 Nova Pro
Accepts image plus text, returns text.
- D

 Embeddings model
Produces vectors, not written assessments.

ON THE EXAM

Check input and output separately. Image in, text out is multimodal input — not image generation.

A publisher summarises sixty thousand archived articles each night. Editors read the summaries the next morning.

The team has been asked to reduce cost. Which change is MOST cost-effective?

What should they do?

- ☐ A Buy provisioned throughput
- ☐ B Switch to batch inference
- ☐ C Lower the temperature
- ☐ D Use a larger model for better summaries

◆ PAUSE AND DECIDE

What should they do?

- A

 Provisioned throughput
Hourly cost for capacity idle most of the day.
- C

 Lower temperature
Changes style. Costs exactly the same.
- B

 Batch inference
Roughly half price, and nobody is waiting.
- D

 Larger model
Increases cost. The qualifier asked to reduce it.

ON THE EXAM

'Read the next morning' plus a cost qualifier equals batch. Two signals, one answer.

A retail chatbot occasionally discusses competitor pricing and has once produced mildly offensive phrasing.

The company wants both blocked before any customer sees a response.

What should be configured?

- ☐ A SageMaker Clarify
- ☐ B A longer system prompt
- ☐ C Bedrock Guardrails
- ☐ D SageMaker Model Monitor

◆ PAUSE AND DECIDE

What should be configured?

- A

 Clarify
Analyses bias and explainability. Does not block output.
- B

 Longer system prompt
Guidance, not enforcement. Defeated by injection.
- C

 Bedrock Guardrails
Denied topics and content filters block before delivery.
- D

 Model Monitor
Watches SageMaker models for drift.

ON THE EXAM

'Before any customer sees a response' means prevention at runtime. Only Guardrails intervenes in the response path.

A regional retailer wants to predict which customers are likely to stop shopping, using three years of its own purchase history held in spreadsheets.

The analytics team uses Excel and SQL, employs no data scientists, and does not want to write code.

What should they use?

- ☐ A SageMaker training jobs with a custom algorithm
- ☐ B Amazon Bedrock with prompt engineering
- ☐ C SageMaker Canvas
- ☐ D Amazon Q Business

◆ PAUSE AND DECIDE

What should they use?

- A

 Custom training jobs
The heavier option. No data scientists exist.
- B

 Bedrock
Generative models are poor at tabular prediction.
- C

 SageMaker Canvas
No-code prediction on their own tabular data.
- D

 Q Business
Answers questions from documents. Not prediction.

ON THE EXAM

Tabular prediction rules out Bedrock. 'No code' rules out custom training. Canvas sits exactly between.

A travel company's assistant confidently invents hotel amenities and cancellation terms that do not exist.

Accurate, current hotel data is available in the company's database and document store.

What is the most effective fix?

- ☐ A Raise the temperature
- ☐ B Use RAG to ground answers in the hotel data
- ☐ C Apply RLHF
- ☐ D Switch to a larger model

◆ PAUSE AND DECIDE

What is the most effective fix?

- A

 Raise temperature
Backwards. More randomness means more invention.
- B

 Use RAG
Grounds answers in accurate data the company already holds.
- C

 Apply RLHF
Shapes tone and preference. Supplies no facts.
- D

 Larger model
Still lacks the hotel data. Hallucination persists.

ON THE EXAM

Confident invention plus accurate data already available is the RAG signature. The data being mentioned is the clue.

A model asked to tag incoming emails as Complaint, Query or Feedback invents new categories such as Urgent, and formats its output differently each time.

The team has written a clear instruction describing each category.

What is the most effective next step?

- ☐ A Fine-tune on labelled emails
- ☐ B Add few-shot examples to the prompt
- ☐ C Increase max tokens
- ☐ D Use chain-of-thought prompting

◆ PAUSE AND DECIDE

What is the most effective next step?

- A

 Fine-tune
Would work, but costs money before trying the free option.
- B

 Few-shot examples
Examples pin down both the label set and the format.
- C

 Increase max tokens
Allows longer output. Unrelated to labels.
- D

 Chain-of-thought
Helps multi-step reasoning. This is a format problem.

ON THE EXAM

Instructions already exist and did not work. Examples succeed where descriptions fail.

An automated system extracts data from mortgage applications. About five percent of extractions return low confidence scores.

Regulation requires that no application is declined on the basis of an uncertain automated reading.

What should be added?

- ☐ A Amazon A2I
- ☐ B Retrain the model on more data
- ☐ C SageMaker Model Monitor
- ☐ D Bedrock Guardrails

◆ PAUSE AND DECIDE

What should be added?

- A

 Amazon A2I
Routes low-confidence predictions to a human reviewer.
- B

 Retrain on more data
May reduce uncertainty. Never eliminates it.
- C

 Model Monitor
Watches for drift. Does not review individual cases.
- D

 Guardrails
Filters generative content. Not confidence thresholds.

ON THE EXAM

Regulation demands a guarantee. Retraining offers an improvement, which is not the same thing.

Before indexing an S3 bucket of internal documents into a knowledge base, a bank must identify which files contain customer account numbers.

Four security services are proposed.

Which one does this?

- | | |
|---|--|
| ☐ A Amazon Inspector | ☐ B Amazon GuardDuty |
| ☐ C Amazon Macie | ☐ D AWS Trusted Advisor |

◆ PAUSE AND DECIDE

Which one does this?

- | | |
|---|---|
| <div data-bbox="173 482 234 539">A</div> Inspector
<i>Software vulnerabilities and network exposure.</i> | <div data-bbox="1304 482 1365 539">B</div> GuardDuty
<i>Threat detection from account activity.</i> |
| <div data-bbox="173 628 234 685">C</div> Macie
<i>Discovers and classifies sensitive data inside S3 objects.</i> | <div data-bbox="1304 628 1365 685">D</div> Trusted Advisor
<i>Account-level cost and best-practice advice.</i> |

ON THE EXAM

All four are real security services. Only Macie opens the objects and reads what is inside.

A team sends the same paragraph of English text to two different foundation models and finds each reports a different token count.

They ask whether one of the models is misconfigured.

What is the explanation?

- ☐ A One model is faulty
- ☐ B Each model has its own tokenizer
- ☐ C The temperature settings differ
- ☐ D One model truncated the input

◆ PAUSE AND DECIDE

What is the explanation?

- A

 One model is faulty
Both are behaving correctly.
- B

 Different tokenizers
Each model splits text differently. Same text, different counts.
- C

 Temperature differs
Affects output sampling, not input token counting.
- D

 Truncation
Would apply only if the context window were exceeded.

ON THE EXAM

Token counts are model-specific. This matters for cost estimates, which cannot be compared across models directly.

How did you do?

- ◆ **17 or more.** Comfortably on track. Focus revision on the domains you missed.
- ◆ **14 to 16.** Around the pass mark. Close, but tighten the weak domains before booking.
- ◆ **11 to 13.** More study needed. Rewatch the modules covering your misses.
- ◆ **10 or fewer.** Work through the core modules again before more practice.
- ◆ **Which domain did your misses cluster in?** That matters more than the total.

ON THE EXAM

The real exam scales scores, so this is a rough guide rather than a prediction. Pattern of misses beats raw count.