

MODULE 14 · EXAM STRATEGY - ALL DOMAINS

How AWS Writes These Questions

Exam mechanics, qualifier words, and the habits that turn a borderline pass into a comfortable one.

DOMAIN all five EXAM WEIGHT n/a MODULE 14 of 14

What you are walking into

DETAIL	VALUE
Questions	65 total — 50 scored, 15 unscored and unmarked
Time	90 minutes
Scoring	Scaled 100–1000. Pass mark 700
Guessing	No penalty. Unanswered counts as wrong
Formats	Multiple choice, multiple response, ordering, matching
Result	Pass or fail, with section-level feedback
Retake	14-day wait, unlimited attempts, full fee each time
Validity	Three years

ON THE EXAM

Confirm timing and cost on the current AWS exam guide before booking — these details do get revised.

Three consequences worth acting on

- ◆ **Fifteen questions do not count**, and you cannot tell which. Do not panic at an odd question.
- ◆ **Never leave anything blank**. No penalty for guessing means a guess is free.
- ◆ **Scaled scoring is not a percentage**. 700 does not mean exactly 70% correct.
- ◆ You do **not** need to pass each domain individually. Only the total matters.

ON THE EXAM

If a question looks unlike anything you studied, it may well be one of the fifteen. Answer it and move on.

The four question types, and how each behaves

- ◆ **Multiple choice** — one correct answer, three distractors. The majority.
- ◆ **Multiple response** — choose two or three. The stem always tells you how many.
- ◆ **Ordering** — arrange steps into the correct sequence.
- ◆ **Matching** — pair 3 to 7 items. **All pairs must be right** to score anything.
- ◆ Matching is all-or-nothing, so it deserves more time than a single multiple choice.

ON THE EXAM

Matching gives no partial credit. If you are certain of four pairs and unsure of one, that whole question hangs on the fifth.

What AWS actually means

THE PHRASE	WHAT IT IS REALLY ASKING
LEAST operational overhead	Which option is most managed? Prefer serverless and fully managed.
MOST cost-effective	Cheapest that still meets the requirement. Often batch or a smaller model.
LEAST development effort	Which needs least code? Prefer a finished service over building.
MOST secure	Which adds a real control? Often IAM, KMS or PrivateLink.
Fastest to implement	Managed service over custom build, every time.
Real-time / immediately	Rules out batch and asynchronous options.
Without writing code	Canvas, or a managed AI service.

ON THE EXAM

These phrases are not decoration. Each one eliminates options on its own, before you consider the services at all.

Read the **last sentence first**. It contains the qualifier, and the qualifier decides the answer.

How 'select two' questions behave

- ◆ The two correct answers usually do **different jobs**, not the same job twice.
- ◆ If two options are near-identical, they are often both wrong.
- ◆ Look for a pair that **combines**: one addresses detection, one addresses response.
- ◆ Selecting the wrong number of answers scores zero. Count before moving on.

ON THE EXAM

A good pair complements. If your two picks overlap heavily, reconsider — you have probably missed a second dimension.

The five shapes, condensed

- ◆ Right data, wrong verb.
- ◆ Right verb, wrong data.
- ◆ **The heavier option** — works, but more effort than needed.
- ◆ **The adjacent service** — close, but does a neighbouring job.
- ◆ **True but irrelevant** — a correct statement that does not answer the question.

ON THE EXAM

The heavier option is the most dangerous, because it is genuinely a working solution. Effort qualifiers exist to rule it out.

Ninety minutes, sixty-five questions

- ◆ That is roughly **80 seconds each**. Most take less; a few take much more.
- ◆ **Never spend more than two minutes** on one question. Flag it and move on.
- ◆ Aim to reach question 65 with fifteen minutes left for flagged items.
- ◆ Answer every flagged question before you review, so nothing is left blank.
- ◆ Do not change answers on review without a concrete reason. First instinct is usually better.

ON THE EXAM

Running out of time is a bigger risk than not knowing content. Pace is a skill you practise, not a thing you hope for.

Weight your revision to the domains

- ◆ Domain 3, Applications of Foundation Models — 28%. The largest by some margin.
- ◆ Domain 2, Generative AI Fundamentals — 24%.
- ◆ Domain 1, AI and ML Fundamentals — 20%.
- ◆ Domain 4, Responsible AI — 14%. And Domain 5, Security — 14%.
- ◆ Domains 2 and 3 together are **over half the exam**. Weight your time accordingly.

ON THE EXAM

RAG, prompting, model selection and Bedrock sit in Domains 2 and 3. That is where marks concentrate.

Carry these into the exam

- ◆ Read the last sentence first. The qualifier decides the answer.
- ◆ Never leave a question blank. Guessing is free.
- ◆ Matching is all-or-nothing. Give it extra time.
- ◆ 'Least effort' exists to eliminate the heavier option.
- ◆ Two minutes maximum, then flag and move on.
- ◆ Fifteen questions do not count. An odd question is not a crisis.

A question ends: 'Which solution meets these requirements with the LEAST operational overhead?'

Two options both solve the problem correctly. One uses a managed AWS service; the other builds the same capability on EC2 instances the team would maintain.

Which is correct, and why?

- ☐ A The EC2 build — it offers more control
- ☐ B The managed service — the qualifier asks for least overhead
- ☐ C Either; the qualifier is stylistic
- ☐ D Whichever costs less

◆ PAUSE AND DECIDE

Which is correct, and why?

- A

 The EC2 build
Control was not asked for. Overhead was.
- B

 The managed service
'Least operational overhead' is the deciding instruction.
- C

 Either
Qualifiers are never stylistic on this exam.
- D

 Whichever is cheaper
Cost was not the qualifier. Overhead was.

ON THE EXAM

When two options both work, the qualifier is the whole question. It is there precisely to separate them.

A question describes a nightly document classification job over 90,000 files, with results reviewed the next morning.

It ends: 'Which approach is MOST cost-effective?'

What does the qualifier point at?

- ☐ A A provisioned throughput endpoint
- ☐ B Batch inference
- ☐ C The largest available model for accuracy
- ☐ D A real-time endpoint with autoscaling

◆ PAUSE AND DECIDE

What does the qualifier point at?

- A

 Provisioned throughput
Hourly cost for capacity idle most of the day.
- B

 Batch inference
Roughly half price, and nobody is waiting.
- C

 Largest model
Accuracy was not the qualifier. Cost was.
- D

 Real-time endpoint
Always-on billing for an overnight job.

ON THE EXAM

Two signals combine: 'MOST cost-effective' plus 'reviewed the next morning'. Cost qualifier plus no urgency equals batch.

A multiple response question asks: 'Which TWO actions should the company take to secure its Bedrock workload?'

Two of the four options both describe restricting who can invoke the model, using slightly different wording.

What does that similarity suggest?

- ☐ A Both are correct; pick them
- ☐ B They likely overlap; look for a pair covering different dimensions
- ☐ C Pick the more detailed of the two
- ☐ D Choose three answers to be safe

◆ PAUSE AND DECIDE

What does that similarity suggest?

- A

 Pick both
Correct answers usually address different dimensions.
- B

 Look for a complementary pair
Access plus encryption, or access plus network isolation.
- C

 The more detailed one
Detail is not correctness.
- D

 Choose three
Wrong count scores zero. The stem said two.

ON THE EXAM

Near-identical options are a design signal. A well-built pair complements rather than duplicates.

A matching question asks you to pair five AWS services with five descriptions. You are confident about four pairs and genuinely unsure about the fifth.

Time remaining is comfortable.

What is the best approach?

- ☐ A Answer the four and leave the fifth blank
- ☐ B Work out the fifth by elimination — partial credit does not exist
- ☐ C Guess quickly and move on; it is only one question
- ☐ D Skip the whole question and return later

◆ PAUSE AND DECIDE

What is the best approach?

- A

 Leave the fifth blank
Scores zero for the entire question.
- B

 Eliminate to find the fifth
All pairs must be right; four correct pairs are worth nothing alone.
- C

 Guess and move on
Reasonable under time pressure. Time is comfortable here.
- D

 Skip entirely
Wastes the four you already know.

ON THE EXAM

With four pairs placed, the fifth is forced. Elimination usually solves matching questions outright.

A question asks which service detects sensitive data in an S3 bucket. One option reads: 'Amazon Macie uses machine learning to discover sensitive data.'

Another reads: 'Encrypting data at rest with AWS KMS protects it from unauthorised access.'

Why is the second option wrong?

- ☐ A The statement is factually incorrect
- ☐ B It is true, but does not answer the question asked
- ☐ C KMS does not work with S3
- ☐ D It is the heavier option

◆ PAUSE AND DECIDE

Why is the second option wrong?

- A

 Factually incorrect
It is entirely accurate. That is what makes it dangerous.
- B

 True but irrelevant
The question asked about detection, not protection.
- C

 KMS does not work with S3
It does. This claim is false.
- D

 The heavier option
A different shape. This one is simply off-topic.

ON THE EXAM

The hardest distractors are the ones you cannot fault. Check that a true statement answers the actual question.

Forty minutes into the exam you are on question 22, having spent five minutes on a question about vector store selection that you still cannot resolve.

You have 43 questions remaining.

What should you do?

- ☐ A Keep working; you are close
- ☐ B Guess, flag it, and move on immediately
- ☐ C Leave it blank and return at the end
- ☐ D Skip the next few easy questions to save time

◆ PAUSE AND DECIDE

What should you do?

- A

 Keep working
Five minutes already spent. Pace is now the bigger risk.
- B

 Guess, flag, move on
Secures a chance at the mark and protects the remaining 43.
- C

 Leave blank, return later
Risks forgetting. Blank scores zero guaranteed.
- D

 Skip easy questions
Sacrifices certain marks to chase an uncertain one.

ON THE EXAM

Do the arithmetic: 43 questions in 50 minutes is 70 seconds each. You are already behind, and one mark is one mark.

During review you return to a flagged question. Your original answer was Amazon Comprehend. On rereading, Amazon Textract now feels appealing because the scenario mentions documents.

You cannot identify anything you misread the first time.

Should you change it?

- ☐ A Yes — second readings are usually more careful
- ☐ B No — change only with a concrete reason
- ☐ C Yes — 'documents' clearly indicates Textract
- ☐ D Change it and flag it again

◆ PAUSE AND DECIDE

Should you change it?

- A

 Yes, second readings are better
Not when nothing new was noticed. That is doubt, not insight.
- B

 No, not without a reason
'Feels appealing' is not a reason. Keep the first answer.
- C

 Documents means Textract
Only for scanned images. Comprehend handles digital text.
- D

 Change and reflag
Adds churn without adding information.

ON THE EXAM

Change an answer when you spot something you missed. Do not change one because the alternative has started to feel better.

You have four weeks before the exam and roughly six hours a week to study.
Your practice scores are strongest in Domain 1 and weakest in Domains 2 and 3.

How should you allocate the time?

- ☐ A Evenly across all five domains
- ☐ B Mostly on Domains 2 and 3 — weakest and heaviest
- ☐ C Mostly on Domain 1, to protect your strength
- ☐ D On Domains 4 and 5, which are easiest to improve

◆ PAUSE AND DECIDE

How should you allocate the time?

- A

 Evenly
Ignores both weighting and your own weaknesses.
- B

 Domains 2 and 3
52% of the exam, and where you score worst. Both factors agree.
- C

 Protect Domain 1
Revising what you already know is comfortable, not useful.
- D

 Domains 4 and 5
Only 28% combined, and not your weak areas.

ON THE EXAM

Two factors point the same way here: weighting and weakness. When they agree, the decision is easy — study what is both heavy and weak.