

PRACTICE EXAM 3 OF 3 · EXAM STRATEGY - ALL DOMAINS

Practice Exam 3

The final twenty. Weighted to bring the full sixty into line with the real exam.

QUESTIONS 20 MIX Weighted PART 3 of 3

An engineer is tuning a summarisation prompt. They lower the temperature to 0.2 and simultaneously reduce top-p to 0.3, then find results hard to interpret.

A colleague suggests changing the approach.

What is the recommended practice?

- ☐ A Always set both to the same value
- ☐ B Adjust temperature or top-p, not both at once
- ☐ C Set both to their maximum
- ☐ D Neither affects output; adjust max tokens instead

◆ PAUSE AND DECIDE

What is the recommended practice?

- A

 Same value for both
They control different things. Matching them means nothing.
- B

 Adjust one, not both
They interact confusingly. Change one at a time.
- C

 Both at maximum
Maximum randomness. The opposite of tuning.
- D

 Neither affects output
Both directly shape token selection.

ON THE EXAM

Temperature reshapes the probabilities. Top-p truncates the candidate list. Moving both at once makes the effect of either impossible to read.

A supermarket chain has three years of transaction records with no labels attached. It wants to discover natural groupings of shopper behaviour to design new loyalty tiers. It has no predefined categories in mind.

Which technique applies?

- | | |
|---|---|
| ☐ A Classification | ☐ B Regression |
| ☐ C Clustering | ☐ D Reinforcement learning |

◆ PAUSE AND DECIDE

Which technique applies?

- A

 Classification
Supervised. Requires labels and predefined categories.
- B

 Regression
Predicts a number. Nothing numeric is being predicted.
- C

 Clustering
Unsupervised grouping with no labels and no predefined classes.
- D

 Reinforcement learning
Learns from rewards through action. No environment here.

ON THE EXAM

'No labels' plus 'no predefined categories' plus 'discover groupings' is the clustering signature, stated three ways.

A team is explaining to executives why foundation models can be adapted to many tasks without task-specific training data.

They need a one-sentence description of how these models were originally trained.

Which describes foundation model pre-training?

- ☐ A Supervised learning on labelled examples for each task
- ☐ B Self-supervised learning on vast unlabelled data
- ☐ C Reinforcement learning from human rankings
- ☐ D Clustering of similar documents

◆ PAUSE AND DECIDE

Which describes foundation model pre-training?

- A

 Supervised per task
That is traditional ML. Foundation models are not task-specific.
- B

 Self-supervised on unlabelled data
The data supplies its own labels by hiding parts of itself.
- C

 RLHF
A later alignment stage, not pre-training.
- D

 Clustering
Groups documents. Does not produce a generative model.

ON THE EXAM

Self-supervision is why no human labelling is needed at this scale, and why one model serves many tasks.

A national health service is legally required to keep all patient data, and all processing of it, inside its own country.

It plans to build a generative AI assistant on Amazon Bedrock.

What is the primary control?

- ☐ A Encrypt all data with a customer managed KMS key
- ☐ B Deploy in an in-country Region and use models available there
- ☐ C Apply Bedrock Guardrails
- ☐ D Restrict access with IAM policies

◆ PAUSE AND DECIDE

What is the primary control?

- A

 KMS encryption
Good practice. Encrypted data in the wrong country is still there.
- C

 Guardrails
Controls content, not geography.
- B

 In-country Region
Residency is a Region choice. Data stays where you put it.
- D

 IAM policies
Controls who accesses data, not where it sits.

ON THE EXAM

Every wrong option is a genuine best practice. None of them changes which country the data is in.

A firm must choose between three foundation models for contract review. Published benchmark scores are similar across all three.

The board wants evidence based on the firm's own contracts before committing.

What should the team do?

- ☐ A Pick the highest benchmark score
- ☐ B Run Bedrock model evaluation on their own prompts and documents
- ☐ C Pick the largest model
- ☐ D Compare perplexity across the three

◆ PAUSE AND DECIDE

What should the team do?

- A

 Highest benchmark
General capability. The board asked about their contracts.
- B

 Bedrock model evaluation
Compares candidates on the firm's own data and prompts.
- C

 Largest model
Size is not evidence, and costs more.
- D

 Perplexity
Measures fluency, not contract review quality.

ON THE EXAM

'Based on the firm's own contracts' rules out every general measure. Only task-specific evaluation answers it.

Engineers at a manufacturer waste hours searching twenty years of technical standards spread across SharePoint and S3. They want to ask a question and receive a specific answer. Access to some standards is restricted by department.

Which service fits?

- | | |
|--|---|
| ☐ A Amazon Comprehend | ☐ B Amazon Kendra |
| ☐ C Amazon Textract | ☐ D Amazon Personalize |

◆ PAUSE AND DECIDE

Which service fits?

- A

 Comprehend
Analyses text properties. Nobody is asking it a question.
- B

 Kendra
Connects to SharePoint and S3, and respects document permissions.
- C

 Textract
Extracts from scans. These are already digital.
- D

 Personalize
Recommends items. No search capability.

ON THE EXAM

Two signals: someone is asking a question, and permissions must be respected. Both point at Kendra.

A team adapts a vision model already trained on millions of general images so that it can recognise their own twelve product categories.

They have only a few hundred labelled product photographs.

Which concept describes this?

- | | |
|--|---|
| ☐ A Transfer learning | ☐ B Reinforcement learning |
| ☐ C Training from scratch | ☐ D Unsupervised learning |

◆ PAUSE AND DECIDE

Which concept describes this?

- A

Transfer learning
Reusing a pre-trained model's knowledge for a new related task.
- B

Reinforcement learning
Learning from rewards. No agent or environment here.
- C

From scratch
The scenario starts from an existing trained model.
- D

Unsupervised
They have labelled photographs.

ON THE EXAM

A few hundred examples achieving what would otherwise need millions is the tell. That efficiency is transfer learning.

A parts distributor's RAG assistant answers well on descriptive questions but fails when users search by exact part numbers such as HX-4471-B.

Semantic search alone is missing these.

What should be configured?

- ☐ A Increase top-k
- ☐ B Enable hybrid search
- ☐ C Reduce chunk size
- ☐ D Switch to a larger generation model

◆ PAUSE AND DECIDE

What should be configured?

- A

 Increase top-k
Retrieves more chunks. The right one is still not matching.
- B

 Hybrid search
Combines semantic with keyword matching for exact identifiers.
- C

 Reduce chunk size
Does not help an exact-string match problem.
- D

 Larger model
Generation is fine. Retrieval is failing.

ON THE EXAM

Semantic search matches meaning, and a part number has no meaning to match. Keyword search handles exact identifiers.

A hospital plans to let an AI system write preliminary radiology findings directly into patient records, visible to clinicians.

The governance board asks what must be in place before go-live.

What is required?

- ☐ A Model accuracy above 99%
- ☐ B A clinician review step before any finding is recorded
- ☐ C Bedrock Guardrails content filters
- ☐ D An AI Service Card

◆ PAUSE AND DECIDE

What is required?

- A

 99% accuracy
1 in 100 patients still receives a wrong finding on file.
- B

 Clinician review first
Controllability. The model assists; a clinician decides.
- C

 Content filters
Blocks harmful content. A clinical finding is not that.
- D

 AI Service Card
AWS documentation of an AWS service.

ON THE EXAM

Option A is the trap for treating safety as a threshold. In high-stakes decisions, oversight is structural, not statistical.

A wildlife charity has four hundred thousand unlabelled camera-trap images. It wants to train a model to identify species.

Humans must first categorise a sample so there is something to train on.

Which service does this?

- | | |
|--|---|
| ☐ A Amazon A2I | ☐ B SageMaker Ground Truth |
| ☐ C SageMaker Model Monitor | ☐ D Amazon Rekognition |

◆ PAUSE AND DECIDE

Which service does this?

- A

 Amazon A2I
Reviews predictions. No model exists yet.
- B

 Ground Truth
Humans label raw data so training can begin.
- C

 Model Monitor
Watches a deployed model for drift.
- D

 Rekognition
Pre-trained. Cannot label custom species categories.

ON THE EXAM

Both Ground Truth and A2I use human workforces. Ask when in the lifecycle the humans appear.

A high-volume application classifies short text snippets. No images or video are involved, quality needs are modest, and cost per request is the dominant concern.

The team is choosing from the Amazon Nova family.

Which model fits?

- | | |
|---|--|
| ☐ A Nova Micro | ☐ B Nova Pro |
| ☐ C Nova Premier | ☐ D Nova Canvas |

◆ PAUSE AND DECIDE

Which model fits?

- A

 Nova Micro
Text-only, lowest latency and cost. Exactly the requirement.
- B

 Nova Pro
Multimodal and more capable. Paying for unused capability.
- C

 Nova Premier
The most capable and most expensive option.
- D

 Nova Canvas
Image generation. No text classification.

ON THE EXAM

Text-only, modest quality, cost-dominant. Three constraints all pointing at the smallest model in the family.

A team wants to change the wording of prompts used in production without redeploying the application, and to compare versions before promoting one.

Prompts are currently hardcoded across several services.

What should they adopt?

- ☐ A Bedrock Prompt Management
- ☐ B Bedrock Guardrails
- ☐ C SageMaker Model Registry
- ☐ D Bedrock Knowledge Bases

◆ PAUSE AND DECIDE

What should they adopt?

- A

 Prompt Management
Stores prompts as versioned resources, editable without redeploying.
- B

 Guardrails
Safety controls on content, not prompt versioning.
- C

 Model Registry
Versions models, not prompts.
- D

 Knowledge Bases
Managed RAG. Nothing to do with prompt versions.

ON THE EXAM

Option C is the near miss: it does versioning and approval, but for models. Read what is being versioned.

A bank wants to analyse complaint emails for recurring themes using a foundation model. The emails contain names, account numbers and addresses.

Policy states personal data must be removed before any text reaches a model.

Which approach satisfies the policy?

- ☐ A Amazon Macie scanning the S3 bucket
- ☐ B Comprehend PII redaction before the model call
- ☐ C A Guardrails PII filter on the model output
- ☐ D KMS encryption of the email store

◆ PAUSE AND DECIDE

Which approach satisfies the policy?

- A

 Macie on the bucket
Discovers sensitive data. Does not redact before a call.
- B

 Comprehend redaction first
Removes PII before the text reaches the model.
- C

 Guardrails on output
Too late. The PII already reached the model.
- D

 KMS encryption
Protects at rest. Decrypted text still reaches the model.

ON THE EXAM

The word 'before' decides it. Option C is a real control in the wrong position — the subtlest kind of wrong answer.

A veterinary practice deploys an image model that runs a handful of times per hour during clinic sessions and not at all overnight.

A real-time endpoint runs continuously, and the bill is dominated by idle time.

Which deployment option fits?

- | | |
|--|---|
| ☐ A Serverless inference | ☐ B Batch transform |
| ☐ C A larger real-time endpoint | ☐ D Asynchronous inference |

◆ PAUSE AND DECIDE

Which deployment option fits?

- A

 Serverless inference
Scales to zero between requests. Idle time stops costing.
- B

 Batch transform
No endpoint, but the vet needs an answer during the appointment.
- C

 Larger endpoint
More idle capacity to pay for.
- D

 Asynchronous inference
For large payloads and long jobs. The problem is idleness.

ON THE EXAM

Intermittent traffic with real idle gaps is the serverless signature. Batch fails because someone is waiting.

An assistant confidently states that a regulation introduced three months ago does not exist. The model was trained before the regulation was published. No retrieval is configured.

What explains this?

- ☐ A Temperature is too low
- ☐ B The model's knowledge cutoff predates the regulation
- ☐ C The context window is too small
- ☐ D Prompt injection

◆ PAUSE AND DECIDE

What explains this?

- A

 Temperature too low
Affects variation, not what the model knows.
- C

 Context window
Limits request size, not stored knowledge.
- B

 Knowledge cutoff
The model knows nothing after training unless you supply it.
- D

 Prompt injection
An attack. Nothing was attacked.

ON THE EXAM

The fix is RAG or supplying the text in the prompt. Note the model was confident while wrong — that is characteristic.

A company has already trained a specialised model on SageMaker. It wants to serve that model through the Bedrock API so it can use Guardrails and the surrounding tooling.

The model uses a supported architecture.

What capability enables this?

- ☐ A Bedrock fine-tuning
- ☐ B Bedrock custom model import
- ☐ C Bedrock Knowledge Bases
- ☐ D SageMaker JumpStart

◆ PAUSE AND DECIDE

What capability enables this?

- A

 Fine-tuning
Adapts a Bedrock base model. The model already exists.
- B

 Custom model import
Brings your own trained weights into Bedrock.
- C

 Knowledge Bases
Managed RAG. Does not host your model.
- D

 JumpStart
A catalogue of pre-trained models to start from.

ON THE EXAM

Imported models are hosted on provisioned throughput, so the standing hourly cost applies here too.

A company documents how its AI system behaves when given deliberately misleading or malformed input, and how consistently it still produces correct output.

It is mapping this to AWS's responsible AI framework.

Which dimension is this?

- | | |
|--|--|
| ☐ A Controllability | ☐ B Veracity and robustness |
| ☐ C Transparency | ☐ D Governance |

◆ PAUSE AND DECIDE

Which dimension is this?

- A

 Controllability
Mechanisms to monitor and steer behaviour.
- B

 Veracity and robustness
Correct output even with unexpected or adversarial input.
- C

 Transparency
Helping stakeholders make informed choices.
- D

 Governance
Best practices across the AI supply chain.

ON THE EXAM

AWS's definition is 'correct system outputs, even with unexpected or adversarial inputs'.
The scenario is that sentence rephrased.

A regulator asks an insurer to show both that its pricing model does not disadvantage any age group, and that any individual premium can be explained.

The model runs on SageMaker.

Which service produces both?

- ☐ A SageMaker Clarify
- ☐ B SageMaker Model Monitor
- ☐ C Two services are needed
- ☐ D SageMaker Model Cards

◆ PAUSE AND DECIDE

Which service produces both?

- A

 Clarify
Bias detection and per-prediction explanation in one service.
- B

 Model Monitor
Production drift. Does not explain a decision.
- C

 Two services needed
Tempting, but Clarify covers both requirements.
- D

 Model Cards
Documents the model. Does not analyse it.

ON THE EXAM

Option C is the trap for half-remembering Clarify as either the bias tool or the explainability tool. It is both.

A search feature returns a document about 'reducing staff turnover' when a user searches for 'how to keep employees'. The two share almost no words.

A colleague asks how this is possible.

What makes this work?

- ☐ A Keyword matching with synonyms
- ☐ B Embeddings place similar meanings close together in vector space
- ☐ C The model was fine-tuned on HR documents
- ☐ D A higher temperature setting

◆ PAUSE AND DECIDE

What makes this work?

- A

 Keyword with synonyms
Would need every synonym defined in advance.
- C

 Fine-tuned on HR docs
Retrieval, not generation. No fine-tuning involved.
- B

 Embeddings and vector distance
Similar meanings sit close together, so distance measures similarity.
- D

 Higher temperature
An inference parameter for generation.

ON THE EXAM

This is why semantic search beats a keyword box, and why embeddings underpin RAG.

An agent can issue account credits through an internal API. Its execution role currently permits every action on the billing system, and no approval step exists.

A security reviewer flags two problems before launch.

Which pair of changes addresses them?

- | | |
|---|--|
| ☐ A Least privilege on the role, and return of control for approvals | ☐ B Lower the temperature, and add Guardrails |
| ☐ C Encrypt with KMS, and enable CloudWatch | ☐ D Multi-agent collaboration, and a larger model |

◆ PAUSE AND DECIDE

Which pair of changes addresses them?

- A

 Least privilege plus return of control
Scopes what the agent can do, and gates it on human approval.
- B

 Temperature and Guardrails
Guardrails filters content. Neither gates an action.
- C

 KMS and CloudWatch
Encryption and monitoring. Neither restricts the agent.
- D

 Multi-agent and larger model
More agents and more capability. Neither is a control.

ON THE EXAM

Two flagged problems need two matching controls. A good multi-part answer pairs complementary fixes, not two of the same kind.

Three sets of results — what now?

- ◆ **Consistently 17+ across all three.** Book the exam.
- ◆ **Improving each time.** A good sign. One more pass over your weakest domain.
- ◆ **Flat around 14 to 16.** Rewatch the two modules covering your most-missed domain.
- ◆ **The same domain missed in all three.** That is your gap. Nothing else matters as much.
- ◆ Remember Module 14: read the last sentence first, never leave a blank, two minutes maximum.

ON THE EXAM

Sixty questions across three papers gives you a trend, not a snapshot. Trust the pattern over any single score.