

INTRODUCTION · EXAM STRATEGY - ALL DOMAINS

AWS Certified Data Engineer

Associate. Exam code DEA-C01. A complete, original preparation course.

EXAM DEA-C01 QUESTIONS 65 in 130 min PASS 720 / 1000

Every explanation here comes from **AWS documentation**. Not one question is recycled.

What the exam actually validates

From the official exam guide

- ◆ **Implementing pipelines.** Ingest, transform, orchestrate.
- ◆ **Choosing stores.** Model data, catalog schemas, manage lifecycles.
- ◆ **Operating them.** Monitor, troubleshoot, optimise cost and performance.
- ◆ **Securing them.** Authentication, encryption, privacy, governance.
- ◆ Not tested: training models, language syntax, business conclusions.

ON THE EXAM

This is a build-and-operate exam, not a theory exam.

The format, in numbers

DETAIL

Questions	65 total. Only 50 are scored
Unscored	15, unmarked, mixed in. You cannot tell which
Time	130 minutes. Two minutes per question
Score	Scaled 100 to 1000. Pass at 720
Scoring	Compensatory. You need the overall mark, not each domain
Types	Multiple choice, and multiple response with five or more options

ON THE EXAM

No penalty for guessing. Never leave a question blank.

The four domains and their weights

- ◆ **Domain 1, Data Ingestion and Transformation.** 34 percent.
- ◆ **Domain 2, Data Store Management.** 26 percent.
- ◆ **Domain 3, Data Operations and Support.** 22 percent.
- ◆ **Domain 4, Data Security and Governance.** 18 percent.
- ◆ This course allocates its minutes in the same proportion.

ON THE EXAM

Domain 4 is roughly nine scored questions. Most candidates under-prepare it.

The exam guide was revised on the **twelfth of December, twenty twenty five.**

What version 1.1 added

None of this appears in older question banks

- ◆ **Open table formats.** Apache Iceberg, and Amazon S3 Tables.
- ◆ **Vector storage.** Index types like HNSW and IVF, Bedrock knowledge bases.
- ◆ **SageMaker Unified Studio.** Domains, domain units, projects.
- ◆ **SageMaker Catalog.** Business catalogs, data lineage, access through projects.
- ◆ **LLM integration** for data processing, and governance frameworks.

ON THE EXAM

Newly in scope: Aurora, Bedrock, Amazon Q, Kendra, Data Exchange, S3 Tables.

Services that changed while the banks did not

Answer from the documented catalogue, not from what you can buy today

- ◆ **Kinesis Data Analytics for SQL.** Applications deleted from January twenty twenty six.
- ◆ **Kinesis Data Analytics** became Managed Service for Apache Flink, 2023.
- ◆ **Kinesis Data Firehose** became Amazon Data Firehose, 2024.
- ◆ **AWS Glue for Ray.** Closed to new customers, April twenty twenty six.
- ◆ **Redshift Python UDFs.** Support ends June twenty twenty six.

ON THE EXAM

Where a service is retired but still examinable, this course teaches it and says so.

How the course is built

01

Foundations

This module, then the mental model that everything else hangs on.

02

Four domains

Nineteen modules, time allocated by the published weights.

03

Confusion pairs

Services that sound alike, put side by side.

04

Exam mechanics

How the questions are worded and what the qualifiers mean.

05

Practice exams

Three papers, sixty original questions, mixed domains.

ON THE EXAM

Twenty seven modules. Watch them in order; later ones assume earlier vocabulary.

Every scenario stops and asks you to
commit to an answer before the
walkthrough.

A retail company loads clickstream files into Amazon S3 every night. An analyst runs the same three aggregate queries each morning against the raw JSON files, and the queries have become slow and expensive.

The company wants to reduce query cost. Nobody is waiting on the results before nine in the morning, and the engineering team has no capacity for a new pipeline to operate.

Which change will reduce query cost with the least operational overhead?

- | | |
|---|---|
| A Load the data into an Amazon Redshift provisioned cluster and query it there | B Convert the files to Apache Parquet and partition by date, using an AWS Glue job |
| C Increase the Amazon Athena workgroup data scan limit | D Move the files to S3 Glacier Deep Archive between queries |

◆ PAUSE AND DECIDE

Which change will reduce query cost with the least operational overhead?

- A

 Load into a Redshift provisioned cluster
Would be fast, but adds a cluster and a load pipeline to operate. The question rules that out.
- B

 Convert to Parquet, partition by date
Columnar format plus partition pruning cuts bytes scanned, which is what Athena charges for.
- C

 Raise the workgroup scan limit
That is a guardrail against runaway spend. Raising it permits more cost, it does not reduce any.
- D

 Archive between queries
Cuts storage cost, not query cost, and Deep Archive cannot be queried without a restore.

THE SIGNAL

"No capacity to operate a new pipeline" eliminated A before you compared anything technical.

Two ways to use this course

WHAT FEELS PRODUCTIVE

Rewatch everything

- Start at module one, go again
- Scores stay comfortably high
- Revising material you already know
- Feels like progress. Moves nothing

WHAT ACTUALLY WORKS

Rewatch one domain

- Watch in order, once
- Sit practice exam one
- Find where the misses cluster
- Rewatch only that domain

THE TIEBREAKER

Which domain your misses cluster in matters far more than your score.

Before you start

- ◆ **Have the exam guide open.** It lists every in scope service.
- ◆ **Build something small.** One Glue job beats an hour of video.
- ◆ **Expect the qualifiers.** LEAST operational overhead appeared in one question in five.
- ◆ **Do not memorise the walkthroughs.** Learn the signal that decided each one.
- ◆ **Space it out.** Six to eight weeks beats two intense ones.

ON THE EXAM

Next: the mental model. How the pieces fit before we look at any one of them.