

MODULE 21 · EXAM STRATEGY - ALL DOMAINS

Confusion pairs

The services this exam deliberately puts side by side, and the one sentence that separates each pair.

DOMAIN All four EXAM WEIGHT — MODULE 21 of 27

A good distractor is **not wrong**. It is right for a question that was not asked.

How to use this module

- ◆ **Each pair has a deciding signal**, usually one phrase in the question.
- ◆ **Learn the signal, not the comparison.** You will not have time to weigh both sides.
- ◆ **Most pairs turn on one of four things:** existing technology, latency, frequency, or overhead.
- ◆ **The scenarios come in pairs**, where one changed constraint flips the answer.
- ◆ If you can recite a comparison but not the signal, you will still be slow.

ON THE EXAM

The exam gives you roughly two minutes a question. Recognising a signal takes seconds; comparing two services takes a minute.

Ingestion pairs

THE PAIR	WHAT DECIDES IT
Kinesis Data Streams or Amazon MSK	Existing Kafka, or Kafka APIs named. Otherwise Kinesis
Data Streams or Data Firehose	Do you need to process, replay or fan out? Streams. Just deliver? Firehose
Firehose or Managed Service for Apache Flink	Windowing, joins or aggregation in flight means Flink
DMS or AWS Glue	A database source with ongoing change means DMS. Files mean Glue
DataSync or Transfer Family	How your users connect. SFTP means Transfer Family
AppFlow or custom API code	A supported SaaS source. AppFlow removes the code entirely

ON THE EXAM

The first and last rows are the same rule. If something already exists, the answer is the option that does not replace it.

Kinesis Data Streams or Amazon Data Firehose?

DATA STREAMS

A durable log

- Retention, so records can be replayed
- Many independent consumers
- Sub-second, with enhanced fan-out
- You manage or scale the consumer

DATA FIREHOSE

A delivery pipe

- No consumer to write or run
- Buffers, converts format, delivers
- Latency measured in seconds
- One destination per stream

THE TIEBREAKER

Replay and multiple consumers are the two capabilities Firehose does not have. Either one in a question decides it.

Transformation pairs

THE PAIR	WHAT DECIDES IT
AWS Glue or Amazon EMR	A non-Spark engine named, or an existing Hadoop estate. Otherwise Glue
Glue Spark or Glue Python shell	Does it need to be distributed? Small and simple is Python shell
Glue or AWS Lambda	Under fifteen minutes and not distributed favours Lambda
EMR on EC2 or EMR Serverless	Spot pricing and cluster control, against no cluster to operate
Athena CTAS or a Glue job	One-off conversion, against recurring and incremental
ETL or ELT	Where the data already is. In the warehouse means transform there

ON THE EXAM

The first row is an elimination, not a comparison. Hive, HBase, Presto or Trino removes Glue before you read further.

AWS Glue or Amazon EMR?

AWS GLUE

Serverless Spark

- Spark and Python only
- Nothing to size, patch or scale
- Per-second billing, one minute minimum
- The default when overhead is mentioned

AMAZON EMR

The whole ecosystem

- Hive, HBase, Presto, Trino, Flink, Pig
- Spot Instances on task nodes
- Deep tuning and custom software
- EMR Serverless removes the cluster

THE TIEBREAKER

Three reasons choose EMR: an engine Glue lacks, an existing Hadoop estate, or Spot pricing on sustained work.

Migration and transfer pairs

THE PAIR	WHAT DECIDES IT
DataSync or the Snow Family	Bandwidth and time. Weeks over the wire means Snow
DataSync or Transfer Family	Whether your users need SFTP, FTPS or FTP
DMS or Snow for a database	Ongoing change means DMS. A one-off bulk move can use Snow
Full load or full load plus CDC	Can the source be read-only during the move?
Direct Connect or Site-to-Site VPN	Consistent bandwidth, against setting up quickly

ON THE EXAM

Availability warning: Snowball Edge closed to new customers in November 2025, and Snow support ends in commercial Regions on 31 December 2026.

A retailer streams point of sale events and needs them delivered into Amazon S3 as Parquet, partitioned by date, for later analysis. No processing is required between the producer and the destination, and a delay of a minute or two is acceptable.

The team has nobody available to build or operate a stream consumer.

What should the data engineer use?

- ☐ A Amazon Kinesis Data Streams with a Lambda consumer writing Parquet
- ☐ B Amazon Data Firehose with format conversion and dynamic partitioning
- ☐ C Amazon Managed Service for Apache Flink writing to Amazon S3
- ☐ D Amazon MSK with a Kafka Connect S3 sink

◆ PAUSE AND DECIDE

What should the data engineer use?

- A

 Data Streams with a Lambda consumer
A consumer to write, deploy and operate, which the question ruled out in its last sentence.
- B

 Firehose with conversion and partitioning
No processing needed, minutes of latency acceptable, no consumer to run. Format conversion and partitioning are configuration.
- C

 Managed Service for Apache Flink
A stream processing engine where nothing needs processing. Real capability, applied to a delivery problem.
- D

 MSK with a Kafka Connect sink
No existing Kafka is mentioned, and a connector is a component to configure and operate.

THE SIGNAL

No processing, tolerant latency, nobody to run a consumer. Three sentences, all pointing the same way.

The same retailer, the same point of sale events, the same Parquet destination in Amazon S3. The requirement has changed in one respect. A fraud team now needs to read the same events independently, in near real time, and a reporting team needs to reprocess the last three days of events after fixing a bug in their logic.

The delivery to Amazon S3 must continue as before.

What should the data engineer use now?

- ☐ A A second Firehose stream from the same producer
- ☐ B Kinesis Data Streams, with Firehose as one consumer and the fraud team as another
- ☐ C Firehose with an S3 destination, and have the fraud team read the S3 files
- ☐ D Increase the Firehose buffer interval so more data is available for reprocessing

◆ PAUSE AND DECIDE

What should the data engineer use now?

- A

 A second Firehose stream
Two independent deliveries, and neither retains anything. The three day reprocessing requirement is still unmet.
- B

 Data Streams, with Firehose as one consumer
Retention gives the reprocessing; multiple consumers give the fraud team a live feed; Firehose still handles S3.
- C

 Fraud team reads the S3 files
Near real time was stated. Files arrive after buffering, so the fraud team is always behind.
- D

 A longer buffer interval
Buffering is not retention. It delays delivery rather than making past records re-readable.

THE SIGNAL

Compare with scenario 1. One requirement changed and the answer moved, because replay and fan-out are what Firehose lacks.

Orchestration and workflow pairs

THE PAIR	WHAT DECIDES IT
Step Functions or MWAA	Existing Airflow DAGs, or reaching outside AWS, means MWAA
Step Functions or Glue workflows	Count the non-Glue steps. Zero means Glue workflows
EventBridge or Step Functions	Starting work, against coordinating a multi-step pipeline
Standard or Express workflows	Over five minutes, or needing full history, means Standard
Amazon SNS or Amazon SQS	Telling several parties, against handing work to one consumer
Retry or a dead-letter queue	Will it succeed next time? Transient retries; permanent routes

ON THE EXAM

The last row generalises well beyond queues. It is the same choice as failing a job or quarantining bad rows.

EventBridge or Step Functions?

AMAZON EVENTBRIDGE

Starts things

- Reacts to an event, or to a schedule
- Knows nothing about what happens next
- No retry of step three, no state
- Often triggers a state machine

AWS STEP FUNCTIONS

Coordinates things

- Knows the shape of the whole pipeline
- Retry and Catch per state
- Tracks where an execution is
- Waits, branches and runs in parallel

THE TIEBREAKER

They are usually both correct, in sequence. EventBridge starts a Step Functions execution.

A pipeline runs four AWS Glue jobs in sequence, followed by a Glue crawler. Each step must start only when the previous one has succeeded, and the team wants to be alerted if any step fails.

There are no other components. The team wants the least operational overhead.

What should the data engineer use?

- ☐ A An AWS Glue workflow with conditional triggers
- ☐ B An AWS Step Functions state machine invoking each job
- ☐ C Amazon MWAA with a DAG defining the dependencies
- ☐ D EventBridge rules chaining each job to the next

◆ PAUSE AND DECIDE

What should the data engineer use?

- A

A Glue workflow
Every component is a Glue job or crawler, so Glue workflows cover it with no additional service at all.
- B

Step Functions
Entirely capable, and it adds a service to orchestrate a pipeline that is already wholly inside Glue.
- C

MWAA
An Airflow environment for five Glue steps. No existing DAGs are mentioned to justify it.
- D

EventBridge chaining
EventBridge starts things. Chaining rules gives no view of the pipeline and no coordinated failure handling.

THE SIGNAL

Five steps, all Glue. Count the non-Glue components before anything else.

The same pipeline, with one addition. After the crawler completes, an AWS Lambda function must post a summary to a chat channel, and a Redshift stored procedure must run to refresh a set of aggregates.

The dependency and alerting requirements are unchanged, and the team still wants the least operational overhead.

What should the data engineer use now?

- ☐ A Keep the Glue workflow and trigger the Lambda function from a Glue job
- ☐ B An AWS Step Functions state machine coordinating all seven steps
- ☐ C Two Glue workflows with an EventBridge rule between them
- ☐ D Amazon MWAA, because the pipeline now spans several services

◆ PAUSE AND DECIDE

What should the data engineer use now?

- A

 Glue workflow plus a job that calls Lambda
Hides two steps inside a job, so the workflow cannot see or retry them. The failure handling requirement breaks.
- B

 Step Functions across all seven steps
Two non-Glue steps put the pipeline beyond Glue workflows. Step Functions coordinates all of it with retry and Catch throughout.
- C

 Two workflows joined by EventBridge
Splits one pipeline into two with a gap in the middle, and nothing coordinates failure across the join.
- D

 MWAA
Correct that coordination is needed, and it is the heaviest option with no existing Airflow to justify it.

THE SIGNAL

Two non-Glue steps. That is the entire difference from scenario 3, and it moves the answer.

Provisioned or serverless, across every service

PROVISIONED

Capacity you hold

- Billed whether used or not
- Reserved pricing available
- Cheaper at steady high utilisation
- Redshift nodes, Kinesis shards, MSK brokers

SERVERLESS

Work you perform

- Nothing between uses
- No commitment, no reservation
- Cheaper for spiky or unknown load
- Glue, Athena, Lambda, Redshift Serverless

THE TIEBREAKER

One decision, asked about eight services. The words unpredictable and steady do the work in every one of them.

Storage and table format pairs

THE PAIR	WHAT DECIDES IT
Intelligent-Tiering or lifecycle rules	Unknown or changing access patterns means Intelligent-Tiering
S3 Tables or Iceberg in a normal bucket	Who runs compaction? Nobody available means S3 Tables
Iceberg or plain Parquet	Does anything update or delete rows? If not, Iceberg adds only maintenance
Object Lock governance or compliance	An escape hatch, or a regulator. Compliance stops everyone
Versioning or replication	Recovering from a bad write, against surviving a Region
Glacier Instant or Flexible Retrieval	Milliseconds when you need it, against minutes to hours

ON THE EXAM

The third row is the one people over-apply. Append-only data with whole-partition expiry does not need a table format.

Catalog and format pairs

THE PAIR	WHAT DECIDES IT
A crawler or enableUpdateCatalog	Does the writing job already know the schema?
Partition projection or a crawler	Athena only, against every engine
Parquet or Avro	Reading few columns, against evolving schemas
Snappy or gzip	Splittability. Gzip cannot be divided
Glue Data Catalog or Hive metastore	An inherited estate can be federated

ON THE EXAM

A crawler discovers a schema. A job that wrote the data already knows it.

Data store pairs

THE PAIR	WHAT DECIDES IT
DynamoDB or Amazon Aurora	Key lookup at scale, against relational queries and joins
DocumentDB or DynamoDB	Existing MongoDB. Compatibility beats capability
Keyspaces or DynamoDB	Existing Cassandra, for the same reason
MemoryDB or ElastiCache	Durability. MemoryDB is a primary store, and is the one in scope
Neptune or a relational join	Relationships of varying depth means graph
Global secondary or local secondary index	An existing table needing a new pattern means GSI

ON THE EXAM

Three of these six rows are the same rule. An existing technology names its compatible service.

A company runs a self-managed Apache Cassandra cluster on Amazon EC2 supporting a high-write telemetry application. Around thirty services use Cassandra drivers. The operations team wants to stop managing the cluster.

Write throughput is high and predictable, and the access pattern is a lookup by partition key.

What should the data engineer recommend?

- | | |
|---|--|
| ☐ A Amazon DynamoDB, which handles high-write key-value workloads at scale | ☐ B Amazon Keyspaces, which is Cassandra compatible |
| ☐ C Amazon Aurora PostgreSQL with a wide table | ☐ D Amazon DocumentDB |

◆ PAUSE AND DECIDE

What should the data engineer recommend?

- A

DynamoDB
Genuinely suits the access pattern, and it means rewriting thirty services onto a different API and data model.
- B

Amazon Keyspaces
Cassandra compatible, so the thirty services keep their drivers and their queries, and AWS operates it.
- C

Aurora PostgreSQL
A relational engine for a high-write wide-column workload, plus rewriting every service.
- D

DocumentDB
MongoDB compatible, not Cassandra compatible. A different model entirely.

THE SIGNAL

Thirty services using Cassandra drivers is the constraint. The access pattern is a decoy that fits the wrong answer.

Query and warehouse pairs

THE PAIR	WHAT DECIDES IT
Athena or Amazon Redshift	Occasional and exploratory, against constant and concurrent
Redshift Spectrum or federated query	Files in S3, against a live operational database
Data sharing or zero-ETL	Warehouse to warehouse, against operational database into a warehouse
Spectrum or loading the data	How often is it queried? Rarely means leave it in S3
Materialized view or a plain view	Saving compute, against only saving typing
SPICE or direct query	Isolating the warehouse, against always-current figures

ON THE EXAM

The fifth row decides more cost questions than it should. A view runs the underlying query every single time.

Amazon Athena or Amazon Redshift?

AMAZON ATHENA

Per query

- Nothing running between queries
- Billed by bytes scanned
- Occasional, exploratory, ad hoc
- Concurrency shared with the account

AMAZON REDSHIFT

Per capacity

- A warehouse you keep
- Sort keys, distribution keys, materialized views
- Hundreds of concurrent users, sub-second
- Billed whether queried or not

THE TIEBREAKER

Frequency decides, not capability. Both can answer the query; only one is priced for how often it will be asked.

The four ways Redshift reaches outside

CAPABILITY	WHEN IT IS THE ANSWER
Redshift Spectrum	Cold history in S3, joined to warehouse tables in one query
Federated query	An occasional lookup against a live RDS or Aurora database
Data sharing	Another Redshift warehouse needs the same data, live, with no copy
Zero-ETL	An operational database's data must be in the warehouse continuously

ON THE EXAM

All four avoid building a pipeline. They differ in where the data lives and whether it moves.

An analytics team queries five years of transaction history. The most recent ninety days are queried by four hundred people through dashboards throughout the working day. Data older than a year is queried perhaps twice a quarter, for regulatory reporting.

Regulatory queries must be able to join old data to recent data in a single SQL statement.

What is the most appropriate design?

- A Keep everything in Amazon Redshift, sized for the full five years
- B Recent data in Redshift, older data in Amazon S3 as Parquet, queried together with Redshift Spectrum
- C Keep everything in Amazon S3 and query it all with Amazon Athena
- D Recent data in Redshift, older data in S3, with analysts using Athena for historical queries

◆ PAUSE AND DECIDE

What is the most appropriate design?

- A

 Everything in Redshift
Warehouse storage prices for four years of data queried twice a quarter. It works and it is the expensive option.
- B

 Redshift plus Spectrum over S3
Hot data gets warehouse performance for 400 concurrent users; cold data gets S3 prices; one statement spans both.
- C

 Everything in Athena
Four hundred concurrent dashboard users on a per-query engine. Module one's second scenario, repeated.
- D

 Two tools, Athena for history
Achieves the same saving and breaks the single-statement requirement. Athena and Redshift are separate engines.

THE SIGNAL

Two access patterns and a join requirement. That combination is what Spectrum exists for.

Operations pairs

THE PAIR	WHAT DECIDES IT
CloudWatch Logs or CloudTrail	What the application did, against who called the API
Concurrency scaling or table design	Queue wait time, against execution time
Anomaly detection or a fixed threshold	Does normal vary? Cyclical load needs anomaly detection
Alarm on failure or on absence	A job that never ran emits no failure metric
Logs Insights or OpenSearch	Investigating now, against dashboards and ongoing alerting
Athena over exported logs or CloudWatch	Years of retention and rare queries means export to S3

ON THE EXAM

The second row is the most reliably tested distinction in Domain 3. Waiting and running slowly have no answers in common.

Amazon CloudWatch Logs or AWS CloudTrail?

CLOUDWATCH LOGS

What happened inside

- Application and service output
- Exceptions, timings, row counts
- Searched with Logs Insights
- Nothing about who called an API

AWS CLOUDTRAIL

Who called what

- API calls, caller, time, parameters
- Management events by default
- Data events off by default, and charged
- Nothing about a job's internal behaviour

THE TIEBREAKER

CloudTrail records that a job was started. It cannot tell you what the job did to the data.

An Amazon Redshift cluster serves dashboards used by three hundred analysts each morning. During the busiest half hour, users report that dashboards are slow. Query history shows execution times unchanged at around one second, and queue wait times of two to four minutes.

An overnight ETL job sometimes overruns into that window.

What should the data engineer do?

- ☐ A Add sort and distribution keys to the dashboard tables
- ☐ B Enable concurrency scaling on the dashboard queue and set WLM query priority above the ETL workload
- ☐ C Run VACUUM and ANALYZE on the dashboard tables each morning
- ☐ D Move the dashboards to Amazon Athena

◆ PAUSE AND DECIDE

What should the data engineer do?

- A

 Sort and distribution keys
Improves execution time, which the question states is one second and unchanged.
- B

 Concurrency scaling plus WLM priority
Queue wait is the stated symptom. Priority puts dashboards ahead of ETL; concurrency scaling absorbs the burst.
- C

 VACUUM and ANALYZE
Maintenance that improves execution and query plans. Neither is the problem, and both add load during the peak.
- D

 Move to Athena
Three hundred concurrent users onto a per-query engine, to solve a queueing problem in the warehouse.

THE SIGNAL

The two numbers are given separately and deliberately. One second executing, four minutes waiting.

Security and governance pairs

THE PAIR	WHAT DECIDES IT
IAM or Lake Formation	Columns and rows. IAM cannot express either
Secrets Manager or Parameter Store	Does it rotate? That single word decides it
A role or a stored credential	Is the target AWS? Roles for AWS, secrets for everything else
SCP or permissions boundary	Capping an account, against capping a role
Gateway or interface endpoint	S3 and DynamoDB use gateway. Everything else is interface
Encryption or masking	Outsiders, against insiders who are allowed in

ON THE EXAM

The first and last rows are the two fastest eliminations in Domain 4. Both remove options without any technical reasoning.

IAM or AWS Lake Formation?

IAM

Resources

- Actions on ARNs
- Buckets, prefixes, tables, jobs
- Cannot name a column or a row
- The layer everything else sits on

LAKE FORMATION

Inside the table

- Databases, tables, columns, rows, cells
- Enforced by Athena, Spectrum, EMR and Glue
- LF-Tags for grants at scale
- Needs the catalog path to be the only path

THE TIEBREAKER

A column has no ARN. Any option proposing an IAM policy for column or row control can be discarded unread.

Secrets Manager or Parameter Store?

AWS SECRETS MANAGER

Secrets that rotate

- Automatic rotation on a schedule
- Rotates the credential and the stored value together
- Cross-account and cross-Region replication
- Charged per secret per month

PARAMETER STORE

Configuration

- Standard parameters are free
- SecureString encrypts with KMS
- No rotation capability at all
- Bucket names, thresholds, feature flags

THE TIEBREAKER

One word decides it. If the question says rotate, it is Secrets Manager and nothing else will do.

Encryption or masking?

ENCRYPTION

Against outsiders

- A stolen disk, an intercepted connection
- Transparent to anyone authorised
- Controlled by who holds the key
- Does nothing about over-broad access

MASKING

Against insiders

- An authorised person seeing too much
- Applied per role, per column
- Redshift masking, Lake Formation grants
- Does nothing if the storage is taken

THE TIEBREAKER

A support agent seeing a full card number is never an encryption problem. They decrypt correctly every time.

A healthcare analytics team must be prevented from seeing three columns in a patient table stored in Amazon S3 and queried through Amazon Athena. A separate requirement states that if a backup tape or disk were obtained by an outside party, the data must be unreadable.

Both requirements apply to the same table.

Which combination meets both?

- ☐ A An IAM policy denying access to the three columns, plus SSE-S3
- ☐ B Lake Formation column-level grants excluding the three columns, plus SSE-KMS with a customer managed key
- ☐ C SSE-KMS alone, since encryption prevents unauthorised reading
- ☐ D A redacted copy of the table for the analytics team, plus SSE-S3

◆ PAUSE AND DECIDE

Which combination meets both?

- A

 IAM column denial plus SSE-S3
IAM cannot express a column. The first half is impossible, whatever the second half does.
- B

 Lake Formation column grants plus SSE-KMS
Authorisation handles the insider requirement; encryption with a controlled key handles the outsider one.
- C

 Encryption alone
The analytics team decrypt legitimately every time they query. Encryption does nothing about an authorised reader.
- D

 A redacted copy
Meets the requirement by duplicating the dataset, with two pipelines and an inevitable drift between them.

THE SIGNAL

Two requirements, two layers. One is about insiders and one about outsiders, and neither control covers both.

Batch and streaming pairs

THE PAIR	WHAT DECIDES IT
Batch or streaming	Who is waiting, and what a faster answer is worth
Micro-batch or true streaming	Seconds are near real time. Milliseconds are not
Glue bookmarks or a full reload	Is the source append-only and identifiable?
Scheduled or event-driven	"Every night" against "as soon as it arrives"
At-least-once or exactly-once	Whether duplicates are tolerable downstream

ON THE EXAM

The first row is module one's rule. Streaming bills continuously, so it needs somebody who benefits continuously.

Pairs where the answer is both

- ◆ **EventBridge and Step Functions.** The first starts the second.
- ◆ **SNS and SQS.** A topic fanning out into several queues.
- ◆ **Data Streams and Firehose.** Firehose as one consumer of the stream.
- ◆ **Lake Formation and IAM.** IAM for the service, Lake Formation for the data.
- ◆ **The exam still makes you choose,** so pick the one the requirement names.

ON THE EXAM

Recognising that two services normally work together stops you rejecting the right answer for containing both.

Pairs that are not pairs

- ◆ **Athena and Redshift Spectrum** are two engines over one catalog definition.
- ◆ **DataZone and SageMaker Catalog** are the same capability, two names.
- ◆ **QuickSight, Quick Suite and Amazon Quick** are one service, renamed twice.
- ◆ **Kinesis Data Firehose and Amazon Data Firehose** are the same service.
- ◆ **Recognise the alias** rather than trying to distinguish two things that are one.

ON THE EXAM

Renaming is why several of these look like pairs. Treating an old and a new name as different services wastes time.

One failure, several services

FAILURE	WHERE IT APPEARS
Skew	Spark tasks, Kinesis shards, Redshift slices, DynamoDB partitions
Throttling with low average use	Kinesis, DynamoDB, and any rate-limited API
Silent wrongness	Type coercion, dropped joins, unregistered partitions
Waiting, not running	Redshift queue, Athena concurrency, a lock, a missing file
Cost from redundant work	Reprocessing, repeated queries, per-viewer computation
A default nobody changed	Glue DPU, log retention, CloudTrail data events, concurrency

ON THE EXAM

Adding capacity fixes none of these. It is offered as an option for all of them.

Pairs that turn on a number

- ◆ **Fifteen minutes.** Above it, Lambda is eliminated entirely.
- ◆ **One hundred partitions.** Athena CTAS cannot write more in one statement.
- ◆ **One hundred and twenty eight kilobytes.** The IA billing minimum per object.
- ◆ **Thirty, ninety and one hundred and eighty days.** Storage class minimums.
- ◆ **Twenty four hours.** Default Kinesis retention, and DynamoDB Streams retention.
- ◆ **A number in a question is rarely decoration.** It is usually there to cross a limit.

ON THE EXAM

When a scenario gives you a specific figure, check it against a limit before anything else.

The words that decide

PHRASE IN THE QUESTION	WHAT IT ELIMINATES
Least operational overhead	Anything with a cluster, an instance or a patch cycle
Existing Kafka, MongoDB, Cassandra, Redis	Anything requiring the applications to be rewritten
Unpredictable, spiky, variable	Provisioned capacity
A specific throughput number	On-demand, usually. Somebody measured it
Replay, reprocess, several consumers	Amazon Data Firehose
Columns, rows, cells	IAM policies
Rotate the credential	Systems Manager Parameter Store
Single SQL statement across both	Two separate query engines

ON THE EXAM

Read for these before you read the options. Each one removes at least one answer without any technical reasoning.

Using the pairs under time pressure

- ◆ **Read the last sentence of the scenario first.** The constraint often lives there.
- ◆ **Name the pair** before you read the options. It is usually one of the ones here.
- ◆ **Eliminate on the signal**, then compare only what survives.
- ◆ **Do not re-derive the comparison.** If you are weighing two services, you are already slow.
- ◆ **Two options that differ in one word** are the pair. That word is the answer.

ON THE EXAM

Next: how DEA-C01 questions are written, and how to read the ones this module has not prepared you for.