

MODULE 25 · EXAM STRATEGY - ALL DOMAINS

Practice exam 2

Twenty more questions, same weighting, different ground.

DOMAIN All four QUESTIONS 20 MODULE 25 of 27

Before you start

- ◆ If you took exam one, check whether your misses repeat a pattern.
- ◆ Time yourself at forty minutes for twenty questions.
- ◆ Read the question, then the scenario, then the constraint, then the options.
- ◆ Note the letter you chose as well as the answer, to spot a positional bias.
- ◆ Correct answers are spread evenly across A, B, C and D here too.

ON THE EXAM

A repeated pattern across two practice exams is a real finding. A single miss is often just a bad minute.

What changed since exam one

- ◆ **Same weighting**, seven, five, four and four across the domains.
- ◆ **Different ground within each domain**. No question repeats.
- ◆ **Services recur in new forms**. Glue appears three times, doing three different things.
- ◆ **Two multiple-response questions**, at different positions.
- ◆ **Answers rebalanced across the letters**, so a positional habit gains nothing.

ON THE EXAM

A service appearing in both exams is deliberate. Knowing what Glue is does not answer a question about bookmarks.

A payments company streams card transactions into Amazon Kinesis Data Streams. A new requirement is to detect, within seconds, when a single card is used more than five times in any two minute window, and to publish those cards to a downstream queue.

The team wants to avoid building and operating stream processing infrastructure.

What should the data engineer use?

- ☐ A Amazon Managed Service for Apache Flink with a two minute tumbling window
- ☐ B Amazon Data Firehose delivering to Amazon S3, with a scheduled Glue job checking for the pattern
- ☐ C An AWS Lambda consumer that counts transactions per card in memory
- ☐ D Amazon Athena querying the raw stream data every minute

◆ PAUSE AND DECIDE

What should the data engineer use?

- A

Managed Service for Apache Flink
Windowed aggregation over a stream is exactly what Flink does, and the managed service means no cluster to operate.
- B

Firehose plus a scheduled Glue job
Batch analysis of delivered files. The detection would arrive minutes or hours late, not within seconds.
- C

A Lambda consumer counting in memory
Lambda invocations are stateless and independent, so counts across a window would be lost between them.
- D

Athena on a schedule
An analytical engine over files, run every minute. It is neither fast enough nor a stream processor.

THE SIGNAL

"Within any two minute window" is the windowing signal. Counting across time is what separates Flink from Firehose.

An AWS Glue job reads a growing Amazon S3 prefix every night and writes aggregates to a curated table. Input volume has grown from two hundred gigabytes to four terabytes over eighteen months, and the job's runtime and cost have grown proportionally.

Files are only ever added to the prefix; existing files are never modified. Each run reprocesses the entire prefix.

What should the data engineer do?

- ☐ A Increase the number of workers so the job completes in the original time
- ☐ B Enable job bookmarks so each run processes only files added since the last run
- ☐ C Switch the job to Glue Flex execution
- ☐ D Partition the output table by date

◆ PAUSE AND DECIDE

What should the data engineer do?

A More workers

Restores the runtime by paying proportionally more. The job still reads four terabytes to process one night's data.

C Glue Flex

A discount on the same wasted work. It reduces the rate rather than the volume.

B Enable job bookmarks

Append-only input with unmodified files is exactly the bookmark case. Each run reads only what is new.

D Partition the output

Improves how the results are queried. It does nothing about how much input the job reads.

THE SIGNAL

Cost growing with total accumulated data rather than with daily volume is the reprocessing signature.

A company runs an Amazon EMR cluster on EC2 continuously. It executes a four hour Apache Hive job each night and is otherwise idle. Intermediate data is written to HDFS on the core nodes, and the team has a large body of existing Hive scripts.

Finance wants the cost reduced. The job must still finish overnight and the scripts must not be rewritten.

Which combination reduces cost the most?

- ☐ A Migrate the Hive scripts to AWS Glue and run them as Spark jobs
- ☐ B Move all core nodes to Spot Instances and keep the cluster running
- ☐ C Halve the cluster size and accept a longer runtime
- ☐ D Store data in Amazon S3 with EMRFS, run a transient cluster nightly, and use Spot Instances for task nodes

◆ PAUSE AND DECIDE

Which combination reduces cost the most?

- A

 Migrate Hive to Glue
Glue does not run Hive, so this means rewriting every script. The question forbids it explicitly.
- B

 All core nodes on Spot
Core nodes hold HDFS. An interruption loses data mid-job, and the cluster still runs all day.
- C

 Halve the cluster
Saves half of a bill you should not be paying. The waste is the idle time, not the node count.
- D

 EMRFS, transient cluster, Spot task nodes
Stops paying for twenty idle hours a day and discounts the four that remain. Hive runs unchanged.

THE SIGNAL

Twenty four hours of cluster for four hours of work. The waste is time, and Hive rules out Glue in one word.

A pipeline converts uploaded video metadata files into a normalised format. Each file takes between eight and forty minutes to process depending on size, and processing is single-threaded and cannot be distributed. About sixty files arrive each day at unpredictable times. The current implementation is an AWS Lambda function, and the longer files now fail.

What should the data engineer use?

- ☐ A Increase the Lambda function's timeout and memory allocation
- ☐ B An Amazon EMR cluster with a step per file
- ☐ C AWS Batch, running each file as a containerised job
- ☐ D An AWS Glue Spark job triggered per file

◆ PAUSE AND DECIDE

What should the data engineer use?

- A

 A longer Lambda timeout
Lambda's maximum is fifteen minutes and it cannot be raised. Files taking forty minutes will always fail.
- B

 An EMR cluster
A cluster to operate for sixty single-threaded jobs a day arriving unpredictably.
- C

 AWS Batch
No time limit, containerised, and it provisions compute per job then shuts down. It fills the gap above Lambda.
- D

 A Glue Spark job per file
Spark distributes work the question says cannot be distributed, and adds startup overhead to every file.

THE SIGNAL

Forty minutes crosses Lambda's hard fifteen minute ceiling. That number is in the question to eliminate Lambda.

A company needs to ingest opportunity and account records from Salesforce into Amazon S3 daily, with field mapping and light filtering applied. A developer has proposed writing a Lambda function that calls the Salesforce REST API, handles pagination, manages OAuth token refresh and retries, and writes JSON to Amazon S3.

The team wants to minimise the code they build and maintain.

What should the data engineer recommend?

- ☐ A Amazon AppFlow, with a scheduled flow performing the mapping and filtering
- ☐ B Build the Lambda function as proposed, with EventBridge Scheduler triggering it daily
- ☐ C AWS DMS with Salesforce as a source endpoint
- ☐ D AWS DataSync to transfer the Salesforce data

◆ PAUSE AND DECIDE

What should the data engineer recommend?

- A

Amazon AppFlow
A managed connector for Salesforce, with mapping and filtering as configuration. Authentication and pagination are handled.
- B

Build the Lambda function
Works, and means owning pagination, token refresh, retries and API changes forever. That is the code they said they want to avoid.
- C

AWS DMS
Migrates and replicates databases. Salesforce is a SaaS application, not a database endpoint.
- D

AWS DataSync
Moves files between file systems and object storage. It has no concept of a SaaS API.

THE SIGNAL

The proposal lists four things AppFlow does for you. A question describing that much plumbing is naming the managed service.

Four AWS Glue jobs run nightly, each triggered by its own cron schedule spaced ninety minutes apart. Job two occasionally overruns, and when it does, job three starts against incomplete data and produces incorrect aggregates. No error is raised.

The team wants the pipeline to be dependency-driven and to be alerted when any step fails.

Select TWO actions the data engineer should take.

- ☐ A Replace the cron schedules with an AWS Step Functions state machine invoking each job in sequence
- ☐ B Widen the gaps between the cron schedules to three hours
- ☐ C Add a Catch on each state routing failures to an SNS notification
- ☐ D Set each Glue job's maximum concurrent runs to one

◆ PAUSE AND DECIDE

Select TWO actions the data engineer should take.

- A

 A Step Functions state machine
Each job starts only when the previous one has succeeded. That is what dependency-driven means.
- B

 Widen the cron gaps
Guessing at timing more conservatively. The same failure returns on a slower night, and every night finishes later.
- C

 Catch routing to SNS
Delivers the alerting requirement, declared per state rather than written into each job.
- D

 Max concurrent runs of one
Stops a job overlapping itself. It says nothing about a different job starting on incomplete data.

THE SIGNAL

Two requirements stated, two answers. Dependency ordering and failure alerting are separate things.

A retailer writes order events to Amazon Kinesis Data Streams using store ID as the partition key. The chain has four hundred stores, but one flagship store generates around thirty percent of all events. Producers receive throttling errors during peak hours while total stream utilisation is around forty percent.

The team has already increased the shard count twice without improvement.

What should the data engineer do?

- ☐ A Increase the shard count again
- ☐ B Switch the stream to on-demand capacity mode
- ☐ C Use a composite partition key combining store ID with a random suffix
- ☐ D Enable enhanced fan-out for the consumers

◆ PAUSE AND DECIDE

What should the data engineer do?

- A

 More shards again
Twice already, without improvement. The flagship store's key still hashes to one shard however many exist.
- B

 On-demand capacity
Removes shard management and does not change how a key maps to a shard. A hot key still throttles.
- C

 A composite key with a random suffix
Spreads the hot store across several shards, so its traffic no longer concentrates on one. Consumers reassemble by store.
- D

 Enhanced fan-out
Gives each consumer dedicated read throughput. The throttling here is on the write side.

THE SIGNAL

Throttling at forty percent utilisation, unimproved by two capacity increases, is the hot partition signature.

A subscription business stores customer interaction records as Parquet in Amazon S3, partitioned by month, and queries them with Amazon Athena. Under privacy legislation it must erase all records for a customer within thirty days of a request, and receives around sixty such requests a month.

Currently an engineer identifies affected files, rewrites each without the customer's rows, and replaces them. This takes days and has twice corrupted a partition.

What should the company implement?

- ☐ A Convert the dataset to Apache Iceberg tables and use DELETE statements
- ☐ B Enable S3 Versioning so a failed rewrite can be rolled back
- ☐ C Repartition the data by customer identifier so whole files can be deleted
- ☐ D Apply S3 Object Lock so files cannot be corrupted during rewriting

◆ PAUSE AND DECIDE

What should the company implement?

- A

 Iceberg tables with DELETE
Row-level deletion becomes one atomic SQL statement. A failure leaves the table untouched rather than corrupted.
- B

 S3 Versioning
Makes the corruption recoverable without making the process faster or less manual. Days per request remains.
- C

 Partition by customer
The over-partitioning trap. Thousands of tiny partitions, and queries filter by month rather than by customer.
- D

 Object Lock
Makes objects undeletable, which directly conflicts with an erasure obligation.

THE SIGNAL

Row-level deletion inside shared files is precisely the problem open table formats exist to solve.

A company keeps six years of order history in an Amazon Redshift cluster, totalling ninety terabytes. The last six months are queried continuously by three hundred users. Older data is queried a handful of times each year for audits, and those audits must join old and recent data in a single query.

Managed storage is now the largest line on the Redshift bill.

What should the data engineer implement?

- ☐ A Reduce the cluster node count to lower the storage cost
- ☐ B Move the older data to a second Redshift cluster and use data sharing
- ☐ C Unload data older than six months to Amazon S3 as partitioned Parquet and query it with Redshift Spectrum
- ☐ D Unload older data to Amazon S3 and have auditors use Amazon Athena

◆ PAUSE AND DECIDE

What should the data engineer implement?

- A

 Fewer nodes
On managed storage, compute and storage are billed independently. Shrinking the cluster does not reduce the storage bill.
- B

 A second cluster with data sharing
Data sharing avoids copies between warehouses. It does not move data to cheaper storage.
- C

 Unload to Parquet, query with Spectrum
Cold data moves to S3 prices, and an external table joins to warehouse tables in one statement.
- D

 Auditors use Athena
Saves the same money and breaks the single-query requirement. Athena and Redshift are separate engines.

THE SIGNAL

"Join old and recent data in a single query" is the clause that makes it Spectrum rather than Athena.

A media company runs a self-managed MongoDB cluster on Amazon EC2 holding a content catalogue. Around eighty services query it using MongoDB drivers and aggregation pipelines. The operations team wants to stop managing replica sets and backups. Documents are deeply nested and the schema varies between content types.

What should the data engineer recommend?

- ☐ A Amazon Keyspaces, which is designed for flexible schemas
- ☐ B Amazon DynamoDB, which is fully managed and handles document workloads
- ☐ C Amazon Aurora PostgreSQL using JSONB columns
- ☐ D Amazon DocumentDB, which is MongoDB compatible

◆ PAUSE AND DECIDE

What should the data engineer recommend?

A Amazon Keyspaces

Cassandra compatible, not MongoDB compatible. Wide column is a different model from documents.

C Aurora PostgreSQL with JSONB

Can store documents, and every one of the eighty services would need rewriting to use SQL.

B DynamoDB

Fully managed, and a different API and data model. Eighty services and their aggregation pipelines would be rewritten.

D Amazon DocumentDB

MongoDB compatible, so the drivers and aggregation pipelines continue to work while AWS operates the cluster.

THE SIGNAL

Eighty services using MongoDB drivers is a compatibility requirement. Aggregation pipelines make it stricter still.

A company's Amazon S3 bill has grown steadily while the data the team knows about has stayed around twenty terabytes. S3 Storage Lens reports approximately seventy terabytes of billed storage.

Versioning was enabled three years ago. A nightly job overwrites a set of aggregate files. Several large uploads from an unreliable network link have failed over that period.

Which combination explains the gap and addresses it?

- ☐ A Enable S3 Intelligent-Tiering so infrequently accessed data moves to cheaper tiers
- ☐ B Enable Cross-Region Replication and delete the primary copies
- ☐ C Disable versioning to stop old versions accumulating
- ☐ D Add lifecycle rules to expire noncurrent versions and to abort incomplete multipart uploads

◆ PAUSE AND DECIDE

Which combination explains the gap and addresses it?

- A

 Intelligent-Tiering
Makes the invisible fifty terabytes cheaper without removing it. You would still be storing seventy.
- B

 Replication
Copies the problem to a second Region. Deleting primaries on a versioned bucket adds delete markers rather than freeing space.
- C

 Disable versioning
Stops new versions accumulating, leaves the existing ones, and discards the protection it was enabled for.
- D

 Expire noncurrent versions, abort incomplete uploads
A nightly overwrite on a versioned bucket retains every previous version, and failed uploads leave billed parts that never appear in a listing.

THE SIGNAL

Twenty terabytes known, seventy billed. Two planted details, versioning and failed uploads, explain the fifty.

A crawler has catalogued a set of pipe-delimited text files in Amazon S3. Analysts querying the resulting table in Amazon Athena find that every row appears in a single column named col0, containing the entire line as one string.

The files have a consistent structure with a header row.

What should the data engineer do?

- ☐ A Increase the crawler's sample size so it reads more files
- ☐ B Create a custom classifier defining the pipe delimiter, and rerun the crawler
- ☐ C Set the crawler's schema change policy to update the table
- ☐ D Convert the files to Parquet before crawling them

◆ PAUSE AND DECIDE

What should the data engineer do?

- A

A larger sample

Reading more files does not help when the crawler cannot parse the format of any of them.
- B

A custom classifier for the delimiter

The crawler could not recognise the delimiter, so it treated each line as one value. A classifier tells it how to read the file.
- C

Update the schema on change

Controls what happens when the schema changes. The schema was never inferred correctly in the first place.
- D

Convert to Parquet first

Would work, and requires a job to parse the files, which is the very thing that is failing.

THE SIGNAL

One column holding the whole line is the classic unrecognised-format symptom. A classifier is the fix.

An Athena table over five years of Parquet data is partitioned by a string column named event_date. Analysts filter with a function that extracts the date from an event_timestamp column instead. Results are correct.

Queries scan the entire table every time, and the Athena bill is far higher than expected.

What is the cause?

- ☐ A The table needs a Glue partition index
- ☐ B The workgroup's data scan limit has been set too high
- ☐ C The filter is applied to a derived value, so Athena cannot prune partitions
- ☐ D Parquet files are too small and need compaction

◆ PAUSE AND DECIDE

What is the cause?

- A

 A partition index
Speeds up looking partitions up in the catalog. Here every partition is being read, which is a pruning problem.
- B

 The scan limit
A guardrail that cancels expensive queries. It does not cause a query to read more than it needs.
- C

 The filter uses a derived value
The engine cannot connect a function's output to the partition column, so it prunes nothing and reads everything.
- D

 Small files
Adds request overhead. It does not explain scanning the entire table when a narrow date filter is applied.

THE SIGNAL

Correct results and a full table scan together is the silent pruning failure. Nothing errors.

An Amazon EMR transient cluster runs a nightly Spark job and terminates on completion. Last night the job failed after ninety minutes. The cluster terminated automatically and the operations team can see only that the step failed.

No application logs are available. The team needs to diagnose this and prevent the same blind spot recurring.

What should the data engineer do?

- ☐ A Rerun the job on a long-running cluster to reproduce the failure
- ☐ B Configure the cluster to archive logs to Amazon S3 and enable the persistent Spark history server
- ☐ C Enable AWS CloudTrail to record what the job did
- ☐ D Increase the cluster size so the job does not fail

◆ PAUSE AND DECIDE

What should the data engineer do?

- A

 Reproduce on a long-running cluster
May work for this one failure, at the cost of a rerun and a cluster, and it fixes nothing for next time.
- C

 Enable CloudTrail
Records the API calls that created the cluster and submitted the step. It contains nothing about why Spark failed.
- B

 Archive logs to S3, persistent history server
Logs survive termination and Spark's own view of the run is retained, so every future failure is diagnosable.
- D

 A bigger cluster
Capacity for a failure nobody has diagnosed. It might be skew, bad data or a code error.

THE SIGNAL

Log delivery is configured at cluster creation. After a transient cluster terminates, the decision has already been made.

A daily pipeline loads product data into a warehouse. AWS Glue Data Quality rules check that key columns are complete and that prices fall within an expected range. All rules pass every day.

Over three weeks, the proportion of products with a null description rose from under one percent to eleven percent, and nobody noticed until a catalogue page looked wrong.

What should the data engineer implement?

- ☐ A Add an IsComplete rule on the description column with a one hundred percent threshold
- ☐ B Run the ruleset twice daily instead of once
- ☐ C Increase the Glue job's worker count so more records are validated
- ☐ D Enable anomaly detection so statistics gathered across runs flag unexpected changes

◆ PAUSE AND DECIDE

What should the data engineer implement?

- A

IsComplete at one hundred percent
Description is evidently optional, since one percent were always null. This rule would fail every single day and be muted.
- B

Twice daily
Runs the same passing rules more often. A rule that does not fire is not improved by firing more frequently.
- C

More workers
Capacity for a data quality question. Every record is already being validated.
- D

Anomaly detection
A gradual shift breaks no fixed threshold. Detection over statistics gathered across runs is what catches a changing pattern.

THE SIGNAL

A gradual drift with all rules passing is exactly the failure static rules cannot catch. That is what anomaly detection is for.

An analytics team's Amazon Athena spend has grown to a level finance has queried. The data is uncompressed JSON in Amazon S3 in a single prefix, and queries filter on a date field and read eight columns out of seventy.

A proposal has been made to purchase Athena capacity reservations to make the cost predictable.

Select TWO changes the data engineer should make first.

- ☐ A Convert the data to Parquet with Snappy compression
- ☐ B Partition the data by date
- ☐ C Purchase Athena capacity reservations
- ☐ D Move the data to S3 Glacier Instant Retrieval

◆ PAUSE AND DECIDE

Select TWO changes the data engineer should make first.

- A

 Convert to Parquet with Snappy
Eight columns of seventy means reading roughly a ninth of the data, and compression reduces the bytes billed.
- B

 Partition by date
Queries filter on date against a single unpartitioned prefix, so every query reads everything. Partitioning prunes it.
- C

 Capacity reservations
Makes the cost predictable by fixing it, sized around queries that read seventy columns of raw JSON.
- D

 Glacier Instant Retrieval
Storage cost, plus a retrieval fee. Athena bills bytes scanned regardless of storage class.

THE SIGNAL

Reduce the work before you buy capacity for it. The word 'first' in the question is doing deliberate work.

A bank runs Glue jobs in private subnets. Compliance requires that no data traverses the public internet, and that jobs can reach only three named Amazon S3 buckets rather than any bucket in AWS.

The subnets currently route outbound traffic through a NAT gateway, and jobs reach Amazon S3 over its public endpoint.

What should the data engineer implement?

- ☐ A An IAM policy on the Glue role permitting only those three buckets
- ☐ B Bucket policies on each of the three buckets, keeping the NAT gateway
- ☐ C Move the Glue jobs into public subnets to simplify routing
- ☐ D An S3 gateway endpoint with an endpoint policy naming the three buckets, and remove the S3 route through the NAT gateway

◆ PAUSE AND DECIDE

What should the data engineer implement?

- A

An IAM policy on the role
Governs the identity rather than the network path, and anyone who can edit the role can widen it. Traffic still uses the internet.
- B

Bucket policies plus NAT
Bucket policies govern your buckets. They say nothing about traffic to somebody else's, and the internet path remains.
- C

Public subnets
Solves a routing inconvenience by placing the jobs on the internet, which is the opposite of the requirement.
- D

Gateway endpoint with an endpoint policy
Traffic leaves the internet path entirely, and the endpoint policy restricts which buckets are reachable through it.

THE SIGNAL

Two requirements: a private path, and a restricted set of destinations. Only an endpoint policy delivers the second.

A company has configured AWS Lake Formation column-level grants excluding three sensitive columns from an analytics team. The grants appear correctly in the Lake Formation console.

Testing shows the analytics team can still query all columns through Amazon Athena. The database was created before Lake Formation was adopted.

What is the most likely cause?

- ☐ A Column-level grants only apply to Amazon Redshift Spectrum, not to Athena
- ☐ B The IAMAllowedPrincipals setting is still present, so Lake Formation defers to IAM
- ☐ C The grants need up to twenty four hours to propagate
- ☐ D The table must be converted to Apache Iceberg for column grants to apply

◆ PAUSE AND DECIDE

What is the most likely cause?

- A

 Athena is not covered
False. Lake Formation permissions are enforced in Athena, Redshift Spectrum, EMR and Glue alike.
- B

 IAMAllowedPrincipals still present
That setting makes Lake Formation defer to IAM entirely, so grants are recorded and never consulted. Common on pre-existing databases.
- C

 Propagation delay
Invented. Grants take effect promptly; there is no twenty four hour window.
- D

 Convert to Iceberg
Lake Formation governs plain Parquet tables perfectly well. Iceberg is not a prerequisite.

THE SIGNAL

Grants that appear configured and do nothing is the classic Lake Formation surprise. The pre-existing database is the clue.

A company enabled Amazon Macie across its data lake nine months ago. The lake holds sixty terabytes across three hundred buckets, all encrypted with SSE-KMS using several customer managed keys. Macie reports almost no sensitive data findings.

A manual review of two datasets has found unredacted email addresses and payment card numbers.

What is the most likely explanation?

- ☐ A Macie lacks permission to use the KMS keys, so it can only report object metadata
- ☐ B Macie does not detect payment card numbers without a custom data identifier
- ☐ C Macie samples only the first megabyte of each object
- ☐ D The findings expired, because Macie retains results for ninety days

◆ PAUSE AND DECIDE

What is the most likely explanation?

- A

 No permission to use the KMS keys
Macie can only decrypt SSE-KMS objects if the key policy allows it. Without that it reports metadata and finds nothing.
- B

 No detector for card numbers
Payment card numbers and email addresses are both covered by built-in managed data identifiers.
- C

 First megabyte only
Invented. Macie samples intelligently across objects rather than truncating each one.
- D

 Findings expired
Detailed results expire after ninety days. That would not explain nine months of near-total silence.

THE SIGNAL

Near silence across three hundred encrypted buckets is a symptom, not a clean result. Ask what the tool could see.

A market research firm wants to sell three curated datasets to around one hundred and fifty external customers. Subscribers should receive quarterly revisions automatically, and access must end when a subscription lapses.

The firm does not want to build entitlement management or billing. A proposal is to create an S3 bucket per customer and copy the datasets into each after every revision.

What should the data engineer recommend instead?

- ☐ A A single bucket with a bucket policy listing all one hundred and fifty accounts
- ☐ B Automate the per-customer copying with a scheduled Glue job
- ☐ C AWS Lake Formation cross-account grants to each customer account
- ☐ D AWS Data Exchange, publishing the datasets with managed entitlements and revisions

◆ PAUSE AND DECIDE

What should the data engineer recommend instead?

- A

 One bucket policy naming 150 accounts
A policy that size is unmaintainable and hits length limits, and it manages no subscription lifecycle.
- B

 Automate the copying
One hundred and fifty copies, automated. The duplication, the storage and the revocation problem all remain.
- C

 150 Lake Formation grants
Technically workable for access, and it leaves billing, entitlement and lapse handling entirely to the firm.
- D

 AWS Data Exchange
Entitlement, revisions, subscription lifecycle and delivery are the product. None of it has to be built.

THE SIGNAL

Subscribers, automatic revisions, and access ending on lapse. Those three words are the Data Exchange signal.

Which module each question came from

QUESTIONS	REVISIT
1, 7	Module 3, streaming ingestion
2, 3, 4	Modules 2 and 5, Glue and transformation engines
5	Module 4, migration and connectivity
6	Module 6, orchestration
8, 9, 10, 11, 12	Modules 8 to 12, storage, Iceberg, Redshift, stores, catalog
13, 14, 15, 16	Modules 13 to 16, querying, monitoring, quality, cost
17, 18, 19, 20	Modules 17 to 20, network, authorisation, encryption, governance

ON THE EXAM

Compare with your exam one results. The same row failing twice is a genuine gap rather than a bad minute.

The distractor patterns used here

PATTERN	QUESTIONS WHERE IT APPEARED
Capacity for a non-capacity problem	2, 7, 15
Right answer, wrong question	9, 13, 14
Breaks a stated constraint	3, 5, 10
Buying capacity for wasted work	16
Asserting something false	4, 12, 18, 19
Automating a poor design	6, 20

ON THE EXAM

Four questions here contained an option asserting something a service cannot do. Those are the fastest eliminations available.

If you are running out of time

- ◆ **Twenty questions in forty minutes** is the real pace, and it is tighter than it sounds.
- ◆ **Bank the short ones**, which buys minutes for the long scenarios.
- ◆ **Three minutes on one question** means you have borrowed from two others.
- ◆ **Answer, flag, move on.** An answered flag is a chance; a blank is a zero.
- ◆ **If you finished with time to spare**, that is a real result too.

ON THE EXAM

If you ran over on this set, the reading order from module 22 is worth more practice than any additional service.

Scoring yourself

- ◆ **Sixteen or more** across both exams suggests you are ready.
- ◆ **Compare the two sets** rather than treating this one as a fresh measurement.
- ◆ **A repeated pattern** matters more than a repeated domain.
- ◆ **Review before you take exam three**, or it measures the same gaps again.
- ◆ **These are not calibrated** against the real exam. Use them to find gaps.

ON THE EXAM

Next: practice exam three, plus the scoring trend across all sixty questions and what to do with it.