

Exam mechanics: ordering, matching, qualifiers and 170 minutes

The new question types, the words that decide answers, and how to spend the time.

DOMAIN Cross-cutting MODULE 38 of 39 MINUTES 170

This is a reading exam with a time limit.
Most wrong answers are **correct answers**
to a slightly different question.

The exam, in numbers

Questions	65 total – 50 scored, 15 unscored and not identified
Time	170 minutes, plus 30 on request for non-native English speakers
Scoring	Scaled 100–1,000. Pass at 750
Model	Compensatory – you do not need to pass each domain
Unanswered	Scored as incorrect. There is no guessing penalty
Price	\$300 USD

ON THE EXAM

170 minutes over 65 questions is 2 minutes 37 seconds each. That is the number to plan around.

Five question types

- ◆ **Multiple choice** – one correct answer from four. **Multiple response** – two or more from five or more, and the question tells you how many.
- ◆ **Ordering** – arrange responses into the correct sequence. New since C02.
- ◆ **Matching** – pair items from two lists. Also new since C02.
- ◆ **Case study** – several questions about one scenario, so you read the scenario once. AWS introduced this alongside the other two.
- ◆ **Every type carries the same scoring value** as a multiple-choice question, and they are mixed in throughout.

ON THE EXAM

AWS's stated aim for the new types is to reduce reading time and test procedural knowledge more directly.

Ordering questions

- ◆ Arrange a small set of steps – typically three to five – into the correct sequence.
- ◆ **This course has been drilling these.** Isolate, capture, eradicate, recover. Snapshot, copy encrypted, restore.
- ◆ **Assume no partial credit.** AWS does not publish per-item rules, so a near-miss earning nothing is the safe assumption.
- ◆ **Anchor the ends first.** First and last steps are usually obvious; the middle then has few arrangements.
- ◆ Sequences here are about **evidence** or **dependency** – do not destroy what you need, do not act before a prerequisite.

ON THE EXAM

Every ordered process in this course is a candidate.

Matching questions

- ◆ Pair items from two lists – usually a requirement or symptom against a service or control.
- ◆ **Module 37 was a matching drill** – 40 pairs with the discriminator attached.
- ◆ **Assume no partial credit here either.** Three of four right scores as none.
- ◆ **Work the certain pairs first** and remove both items. Each removal shrinks the problem.
- ◆ If two items fit one slot, one fits a slot you have not read carefully. Go back.

ON THE EXAM

Matching rewards knowing the discriminator, not recognising the service.

Qualifiers, and what they eliminate

QUALIFIER	IT IS ASKING FOR	IT ELIMINATES
MOST secure	The strongest control	Cheaper or simpler options that also work
MOST cost-effective	The cheapest that meets the requirement	Over-engineering, however secure
LEAST operational overhead	Managed over self-built	Anything needing ongoing human effort
MOST operationally efficient	Fewest moving parts, most automation	Manual steps and bespoke tooling
FIRST	The immediate next action	Correct actions that come later
Which combination	Two or more that work together	Any single sufficient answer

ON THE EXAM

The same scenario with two different qualifiers has two different correct answers. Find the qualifier before the options.

Spending 170 minutes

- ◆ **2 minutes 37 seconds per question on average** – some take 40 seconds, some take five minutes.
- ◆ **First pass: answer what you know, flag the rest, leave nothing blank.**
- ◆ Finish the first pass by about 120 minutes, leaving 50 for flagged questions and a sweep.
- ◆ **On review, re-read the requirement and the qualifier**, not the whole scenario.
- ◆ Change an answer only if you can name what you missed.

ON THE EXAM

A five-minute question you get right costs two you never reached.

Reading a scenario

- ◆ **Read the question and qualifier first**, then the scenario – looking for something, not absorbing it.
- ◆ **Details are rarely decorative.** '1 Gbps hosted', 'no CVE existed', 'the grace period has not expired'.
- ◆ Note what the requirement **does not** ask for. That is what made Parameter Store correct in module 37.
- ◆ **Eliminate on a stated rule, not a feeling.**
- ◆ If two options survive, one is right for a qualifier you were not given.

ON THE EXAM

Every scenario in this course has at least one decisive clause. Real ones do too.

Practical logistics

- ◆ **ESL +30** – a 30-minute extension for non-native English speakers taking the exam in English. Request before scheduling.
- ◆ Or take a written exam in another language, with an English, Japanese or Spanish speaking proctor.
- ◆ Other accommodations go through Pearson VUE, in advance rather than on the day.
- ◆ **The score report classifies by domain**, so a fail tells you where to work.
- ◆ **Aim for 80%+ on practice exams before booking.**

ON THE EXAM

The ESL extension is a documented accommodation and it is free. Ask if it applies.

An ordering question presents four steps for recovering an S3 bucket whose policy denies every principal, in an organisation with centralised root access enabled.

The steps are: launch a root session scoped to S3UnlockBucketPolicy; delete the bucket policy; write a corrected bucket policy; call sts:AssumeRoot from the management account naming the member account.

Which sequence is correct?

- | | |
|--|--|
| <p>A Delete the policy, call AssumeRoot, launch the session, write the corrected policy</p> | <p>B Call AssumeRoot naming the account, launch the session with the task policy, delete the policy, write the corrected policy</p> |
| <p>C Launch the session, call AssumeRoot, delete the policy, write the corrected policy</p> | <p>D Write the corrected policy, call AssumeRoot, launch the session, delete the old policy</p> |

◆ PAUSE AND DECIDE

Which sequence is correct?

- A

 Delete first, then AssumeRoot
Puts the goal before the means. You cannot delete the policy until you hold credentials that the policy does not deny - which is the entire problem.
- C

 Session before AssumeRoot
Inverts the two privilege steps. AssumeRoot is what produces the session; the session cannot precede the call that creates it.
- B

 AssumeRoot, session, delete, write
Correct. Anchor the ends: the first step must obtain privilege and the last must restore service. The middle two follow by dependency.
- D

 Write the corrected policy first
You cannot write a bucket policy while the existing one denies every principal. The deny must go first.

THE TECHNIQUE

Anchor the ends. Obtaining privilege must be first; restoring service must be last. That leaves one arrangement.

A matching question pairs four requirements with four services. The requirements are: prove no external principal can reach a bucket; determine which internal principals can reach a table; find roles nobody has used in 90 days; and confirm a policy change does not widen access.

The services offered are: unused access analyzer; CheckNoNewAccess; external access analyzer; internal access analyzer.

Which pairing is correct?

- | | |
|--|---|
| <p>A External-prove no path; internal-unused roles; unused-widening check; CheckNoNewAccess-internal principals</p> | <p>B Internal-prove no external; external-internal principals; unused-widening; CheckNoNewAccess-unused roles</p> |
| <p>C External-widening check; internal-prove no external; unused-internal principals; CheckNoNewAccess-unused roles</p> | <p>D External-prove no external principal; internal-which internal principals; unused-roles unused 90 days; CheckNoNewAccess-policy widening</p> |

◆ PAUSE AND DECIDE

Which pairing is correct?

A External-no path; internal-unused; unused-widening; Check-internal
Gets one pair right and scrambles three. On an all-or-nothing assumption, one correct pair scores the same as none.

C External-widening; internal-external; unused-internal; Check-unused
Every pair misplaced. Worth reading only to notice that each service still sounds plausible against the wrong requirement.

B Internal-external; external-internal; unused-widening; Check-unused
Swaps external and internal, which is the single most likely mistake on this set - and then misplaces the other two as well.

D Each service with its own documented purpose
Correct. Start from the pairs you are certain of - CheckNoNewAccess is a policy check, unused access finds unused roles - and the remaining two resolve by name.

THE TECHNIQUE

Work the certain pairs first and remove both items. Two pairs removed leaves only external versus internal, which the names decide.

A company must give a partner read access to objects in one S3 prefix. The partner has an AWS account. Options under consideration are a cross-account role with an external ID, a bucket policy granting the partner account, presigned URLs generated hourly, and replicating the prefix to a bucket the partner owns.

The question is asked twice, with two different qualifiers.

Which answer changes, and why?

- A

 MOST secure favours the cross-account role with an external ID; MOST cost-effective favours the bucket policy
- B

 Both qualifiers favour presigned URLs, since they expire
- C

 MOST secure favours replication; MOST cost-effective favours the cross-account role
- D

 The qualifier does not change the answer; the requirement determines it alone

◆ PAUSE AND DECIDE

Which answer changes, and why?

- A** Secure favours the role with external ID; cost-effective favours the bucket policy
Correct. The role adds confused-deputy protection, scoping and session logging; the bucket policy costs nothing to run and meets the stated requirement.
- B** Both favour presigned URLs
Presigned URLs cannot be revoked and require generating and distributing them hourly forever. Neither the most secure nor the cheapest to operate.
- C** Secure favours replication
Replication copies your data into someone else's bucket, where you control neither its lifecycle nor its deletion. That is the least secure option here.
- D** The qualifier never changes the answer
The claim this example exists to disprove. Two qualifiers over one scenario routinely produce two different correct answers.

THE TECHNIQUE

Both A's answers are defensible – for their own qualifier. That is precisely why reading past the qualifier costs marks.

A company must ensure every EC2 instance across 200 accounts requires IMDSv2. Options are: account and Region metadata defaults plus an SCP; a Config rule with SSM Automation remediation; a scheduled Lambda that fixes non-compliant instances hourly; and updating every team's launch template.

The qualifier is LEAST operational overhead.

What does that qualifier select?

- ☐ A The Config rule with remediation, since AWS runs the evaluation
- ☐ B Updating every launch template, since it needs no ongoing service
- ☐ C Metadata defaults plus an SCP, because nothing runs and nothing needs maintaining
- ☐ D The scheduled Lambda, since it is fully automated

◆ PAUSE AND DECIDE

What does that qualifier select?

- A

 Config rule with remediation
Managed, and it still runs continuously, produces findings, and remediates – all of which is operational activity, plus per-evaluation cost.
- B

 Update every launch template
The largest overhead of the four: coordination across teams now, and a permanent obligation to catch every new template.
- C

 Metadata defaults plus an SCP
Correct. Both are settings rather than systems. Nothing executes, nothing evaluates, nothing needs a runbook – and new accounts inherit them.
- D

 Scheduled Lambda
Code you own: deployment, permissions, failure handling, monitoring, on-call. Automated is not the same as low-overhead.

THE TECHNIQUE

'Automated' and 'least operational overhead' are not synonyms. Code you own is overhead even when it runs itself.

Six things to carry forward

- ◆ **65 questions, 170 minutes, pass at 750, compensatory.** Never leave one unanswered.
- ◆ **Ordering and matching are new, and every type scores the same.** Assume no partial credit.
- ◆ **Anchor the ends on ordering. Remove certain pairs on matching.**
- ◆ **Find the qualifier before you read the options.** The same scenario with two qualifiers has two answers.
- ◆ **'Automated' is not 'least operational overhead'.** Code you own is overhead.
- ◆ **Eliminate on a stated rule.** If you cannot name why an option is wrong, you have not eliminated it.

ON THE EXAM

Next: what changed from SCS-C02, and where the console will not match what you were taught.