

MODULE 18 · EXAM STRATEGY - ALL DOMAINS

Confusion Pairs and Trigger Phrases

Every distinction the exam turns on, in one place, organised by the words that decide them.

COVERS All 5 domains FORMAT Rapid recall MODULE 18 of 22

Two options are usually both defensible.
One clause in the scenario separates them, and that clause is nearly always a requirement rather than a technology.

How to work with what follows

- ◆ **Read the left column first.** These are phrases as scenarios actually word them, not service names.
- ◆ **Cover the right column** and answer before revealing. Recognition is not recall.
- ◆ **Where a row surprises you**, that is a module to revisit before exam day.
- ◆ The five scenarios here each chain **several** distinctions, which is how the exam presents them.
- ◆ Nothing new is taught from here on. If a row is unfamiliar, the gap is in an earlier module.

ON THE EXAM

This is the module to return to the night before. Everything in it has appeared already, which is exactly what makes it useful as a final pass.

The phrase that names the throughput mode

WHEN THE SCENARIO SAYS	THE ANSWER IS
"No one is waiting", bulk, nightly, backfill	Batch inference. 50% off
Synchronous but tolerates extra latency	Flex tier. 50% off
Interactive and must stay fast under load	Priority tier
Steady, predictable, cannot be throttled	Reserved tier
The model is custom or fine-tuned	Provisioned Throughput
Bursty, unpredictable, "optimise cost"	On-demand. It is already correct

ON THE EXAM

The last row is the trap. Every reservation is a bet on sustained utilisation, so the word 'bursty' argues against all of them.

The phrase that names the customisation

WHEN THE SCENARIO SAYS	THE ANSWER IS
Content changes, needs citations, per-user access	RAG. Any one of the three is enough
Consistent tone or rigid output structure	Fine-tuning
"Similar accuracy at lower cost and latency"	Model Distillation
Does not understand the domain's jargon at all	Continued pre-training
Invocation logs already hold prompts and responses	Distillation. The dataset is free
Train a foundation model from scratch	Out of scope. That option is wrong

ON THE EXAM

The commonest wrong answer in Domain 1 is fine-tuning a model on documents so it 'knows' them. Distrust that option every time you see it.

The phrase that names the vector store

WHEN THE SCENARIO SAYS	THE ANSWER IS
Huge corpus, few queries, latency tolerated	S3 Vectors
Constant interactive traffic, hybrid search	OpenSearch Serverless
The filter joins to relational data	Aurora PostgreSQL with pgvector
"No single document contains that answer"	Neptune Analytics. GraphRAG
Text query must find a video or an image	Multimodal embeddings, not a store change
No vector database to provision at all	Managed Knowledge Base

ON THE EXAM

Query rate decides OpenSearch against S3 Vectors, not corpus size. A billion vectors queried twice an hour should not be renting a cluster.

The phrase that names the retrieval fix

WHEN THE SCENARIO SAYS

THE ANSWER IS

Part numbers or error codes fail, prose works

Hybrid search. A query-time change

Small chunks lose context, large lose precision

Hierarchical chunking

A compound or two-part question

Query decomposition

The right chunk returns but ranks low

Reranking

"I could not find information about that"

An over-restrictive metadata filter

Cites a document that was withdrawn

Deletion never propagated to the index

ON THE EXAM

Reranking fixes ordering and never coverage. If the answer was not in the candidate set, no amount of reranking will surface it.

Three places to version a thing

STORE	VERSIONS	THE TRIGGER
SageMaker Model Registry	Model artefacts and their metrics	Governed model promotion and rollback
Bedrock Prompt Management	Prompt text, variables and the model binding	Auditable prompt change control
AWS AppConfig	Runtime configuration, including model ID	Change behaviour without a deployment

ON THE EXAM

All three make rollback trivial for the same reason: the previous version still exists, so you repoint rather than rebuild. Which one depends on what changed.

Four places to keep configuration

STORE	BEST FOR	THE DISTINGUISHING FEATURE
AWS AppConfig	Feature flags, model IDs, behaviour	Gradual rollout and alarm rollback
SSM Parameter Store	Simple key-value settings	Cheap. No deployment strategy
Secrets Manager	Credentials, API keys, passwords	Automatic rotation
A prompt version	The model and prompt together	Model binding travels with it

ON THE EXAM

AppConfig against Parameter Store is the pair to learn. If a scenario mentions gradual rollout or automatic rollback, only AppConfig has those.

Four ways to orchestrate, in order of freedom

MECHANISM	WHO DECIDES THE PATH	REACH FOR IT WHEN
One Lambda function	You, in code	Two or three steps, under 15 minutes
Bedrock Flows	You, visually	Prompts, knowledge bases and branching
Step Functions	You, declaratively	Retries, human pauses, any AWS service
An agent	The model, at runtime	The path is genuinely unknown in advance

ON THE EXAM

Read top to bottom and freedom increases while determinism decreases. Pick the highest row that satisfies the requirement, not the most capable one.

The phrase that names the integration

WHEN THE SCENARIO SAYS	THE ANSWER IS
A legacy system that cannot be modified	Event-driven. EventBridge or change capture
"Up to three business days" for approval	Step Functions task token callback
Several agents need the same tool	An MCP server behind AgentCore Gateway
A tool runs long queries or holds state	Host the MCP server on ECS, not Lambda
Every participant sees the same stream	AppSync. Fan-out to subscribers
Two hours of agent execution	AgentCore Runtime. Lambda caps at 15 minutes

ON THE EXAM

The 15-minute Lambda ceiling eliminates options in several domains. Any scenario naming a duration longer than that has already discarded half its option set.

The phrase that names the API shape

WHEN THE SCENARIO SAYS	THE ANSWER IS
"Feels slow", users abandon the page	Stream it. ConverseStream
ThrottlingException, 429	Back off with jitter and retry
ServiceQuotaExceededException, 400	Raise the quota. Retrying never helps
Users may wait but must receive the result	Asynchronous. Job ID plus SQS
Routing by who is paying, not by content	API Gateway usage plans
Swap models without changing code	AppConfig, plus Converse

ON THE EXAM

The two exception rows look almost identical and behave oppositely. A 429 is momentary; a 400 quota error is a boundary that a retry loop will hit forever.

The phrase that names the control

WHEN THE SCENARIO SAYS	THE ANSWER IS
"However the user phrases it"	A denied topic, not a word filter
Must remain readable or usable	ANONYMIZE, not BLOCK
A claim the cited source does not support	Contextual grounding check
Verifiable compliance with a written policy	Automated Reasoning checks
A model outside Bedrock needs the same policy	The ApplyGuardrail API
A team must not be able to disable it	IAM enforcement, or an SCP

ON THE EXAM

Guardrails do not enforce business rules, know ownership, or act before invocation. Three earlier modules had a guardrail as the tempting wrong answer.

The phrase that names the security answer

WHEN THE SCENARIO SAYS	THE ANSWER IS
"Even a compromised container"	A VPC endpoint policy, not an IAM policy
"An account admin cannot change it"	A service control policy
Works in the owning account, fails elsewhere	The KMS key policy is missing kms:Decrypt
Prove which characters were redacted	Comprehend offsets, not an FM rewrite
Find sensitive data nobody knew existed	Macie. Discovery at rest
Which user's credentials did this?	AssumeRole with session tags

ON THE EXAM

Three of these six distinguish prevention from detection. Logging, Config rules and CloudTrail alerts all detect; only IAM, SCPs and endpoint policies prevent.

Three ways to put a human in the loop

MECHANISM	SHAPE	REACH FOR IT WHEN
Step Functions task token	Workflow pauses for a callback	Hours or days, inside a defined process
Amazon Augmented AI	Managed review workflow and UI	You want the review interface provided
API Gateway plus your own UI	You build the approval surface	Review lives inside an existing product

ON THE EXAM

A2I entered maintenance mode in June 2026 and remains a correct exam answer for managed human review. Task tokens are where new work goes.

The Amazon Q family, and what happened to it

NAME	WHAT IT IS	STATUS
Amazon Q Developer	Coding assistance and AWS troubleshooting	IDE plugins end of support Apr 2027
Kiro	Spec-driven agentic IDE	The named successor. GA Nov 2025
Amazon Q Business	Enterprise search and chat over documents	Closed to new customers Jul 2026
Amazon Quick	BI and dashboards, formerly QuickSight	Active. Q Business migrating into it

ON THE EXAM

QuickSight became Quick Suite, then Quick. Q Business is folding into it. Q Developer's console experience is unaffected by the IDE sunset. All four are examinable.

A media company's assistant answers questions about programme rights. It uses semantic-only search over a knowledge base built six weeks ago on 40 million chunks. Users report that queries containing contract reference codes such as "RGT-2019-4471" return the wrong contracts, while general questions work well.

Separately, the assistant sometimes quotes rights terms from contracts that were terminated last quarter. The nightly ingestion job reports success. Re-ingesting the corpus takes eleven hours and the team wants both problems fixed this week.

Which combination addresses both problems?

- | | |
|---|---|
| A Re-ingest with semantic chunking, and increase the sync frequency to hourly | B Switch retrieval to hybrid search, and fix deletion propagation so terminated contracts are removed from the index |
| C Re-embed with a larger embedding model, and add a metadata filter excluding terminated contracts | D Enable reranking, and move to event-driven synchronisation |

Which combination addresses both problems?

- A

 Semantic chunking plus hourly sync
Eleven hours of re-ingestion for a problem chunking cannot fix, and more frequent syncs will not remove chunks that deletion never removed.
- B

 Hybrid search plus deletion propagation
Reference codes failing while prose works is the hybrid signal, and it is a query-time change. A successful job that leaves terminated contracts retrievable is a deletion fault.
- C

 Larger embeddings plus a filter
Bigger embeddings still blur near-identical strings. A filter hides terminated contracts rather than removing them, and it is a re-index anyway.
- D

 Reranking plus event-driven sync
Reranking cannot promote a contract retrieval never returned, and the contracts were terminated a quarter ago, so frequency is not the fault.

THE SIGNAL

Two rows from two tables. "Reference codes fail, prose works" is hybrid search. "Job succeeds but withdrawn content is cited" is deletion propagation. And 'this week' rules out any re-index.

The phrase that names the cost or latency lever

WHEN THE SCENARIO SAYS

THE ANSWER IS

Which team or tenant spent this?

Application inference profiles, tagged

Users reword the same question

Semantic caching, not a key-value cache

Long shared context reused within minutes

Prompt caching. Stable prefix required

Independent sub-tasks running in sequence

Parallelise with a Step Functions Map state

Retrieved context dominates the token bill

Context pruning by relevance

Predictable daily spike at a known time

Scheduled scaling, not reactive autoscaling

ON THE EXAM

Application inference profiles aggregate per usage type per day. For per-request token detail you need per-request metadata in the invocation logs instead.

The phrase that names the measurement

WHEN THE SCENARIO SAYS	THE ANSWER IS
Real-time alerting on token thresholds	Native CloudWatch metrics
Attribute usage by application	Logs Insights over invocation logs
Which prompt text was in force?	Invocation logging. CloudTrail cannot
Agent step-level visibility, no instrumentation	CloudWatch generative AI observability
Is retrieval or generation at fault?	A retrieve-only evaluation job
The same input gives different answers	Output diffing, then lower temperature

ON THE EXAM

Metrics tell you that something changed. Logs tell you what changed. Evaluation tells you whether it is any good. Compliance questions usually need at least two of the three.

A regulated insurer must prove, for any customer interaction in the past three years, which prompt text was in force, which model produced the answer, who approved that prompt, and what the assistant actually said. It must also demonstrate that no business unit can disable the content safety policy.

Prompts are currently hardcoded in Lambda functions. Guardrails are configured per business unit and passed as an optional parameter.

Which combination meets all the requirements?

- ☐ A Enable CloudTrail across all accounts and retain logs for three years
- ☐ B Move prompts into Prompt Management with an approval gate, enable model invocation logging, define one versioned guardrail, and enforce it with IAM policy-based enforcement and cross-account safeguards
- ☐ C Store prompts as versioned S3 objects, enable CloudTrail, and publish a guardrail standard for teams to adopt
- ☐ D Enable model invocation logging and apply an SCP requiring a guardrail identifier on every call
- ☐ E Keep prompts in Git and rely on commit history

Which combination meets all the requirements?

- A** CloudTrail only
Records that a call happened, by whom and when. It holds no prompt text and no response, so three of the five requirements fail.
- B** Prompt Management, logging, one guardrail, IAM enforcement
Each clause maps to one mechanism: version for the text, approval gate for the sign-off, invocation logging for the content, enforcement for the safety policy.
- C** Versioned S3 plus a published standard
Versions the text with no approval workflow and no link from an invocation to the object used. A published standard is convention, not enforcement.
- D** Logging plus an SCP on guardrail identifier
Close. It enforces that some guardrail is present, not that it is the right one, and prompts remain unversioned and unapproved.
- E** Git history and code review
Records the change in a repository and cannot tie a specific customer interaction at a specific moment to the prompt version that served it.

THE SIGNAL

Five requirements, so score every option against all five. D is the strongest distractor because it is genuinely a preventive control, and it answers only one of the five.

Five distractor patterns that recur

- ◆ **Right service, wrong feature.** A guardrail for a business rule; IAM for a retrieval filter.
- ◆ **Detection offered where prevention was asked for.** CloudTrail, Config and alarms all detect.
- ◆ **Better engineering than the requirement.** An automated gate when "recorded human approval" was specified.
- ◆ **The expensive correct answer.** A re-index that works, when a query parameter would do.
- ◆ **Solving a different requirement well.** Fixing one of the three things the scenario listed.

ON THE EXAM

Every one of these describes an option that works. That is what makes them hard. Professional-level questions rarely include an option that simply does not function.

Answers that are wrong more often than right

- ◆ **Fine-tune the model on our documents** so it knows them. Almost always a retrieval problem.
- ◆ **Add an emphatic instruction to the prompt.** Prompts nudge; schemas and IAM enforce.
- ◆ **Increase retries or extend the timeout.** Rarely the fix, and it multiplies the cost of failing.
- ◆ **Build a Lambda that does it yourself.** On an exam that keeps saying "least operational overhead".
- ◆ **Add a disclaimer.** Allocates responsibility on paper and produces no evidence and no control.

ON THE EXAM

None of these is universally wrong, and each is wrong far more often than it is right. Treat them as options requiring positive justification rather than defaults.

A logistics platform runs an agent with nine tools that resolves shipment exceptions. Completion rate is stable at 84%. Costs have risen 70% in two months with flat ticket volume, and average tool calls per session have gone from 6 to 15.

Separately, the agent recently emailed a customer a summary containing another customer's shipment details. The email tool has no restrictions; the agent's instructions say it must only discuss the current customer's shipments.

Which combination addresses both issues?

- ☐ A Strengthen the agent's instructions about customer boundaries, and raise the iteration cap so sessions complete
- ☐ B Evaluate trajectories for tool selection and reasoning quality to find the degraded tool, and enforce the customer boundary in the tool and its IAM permissions
- ☐ C Apply a Bedrock guardrail with a denied topic covering other customers' data, and switch to a faster model
- ☐ D Add a reranker to improve tool result quality, and add a PII filter to the email tool

Which combination addresses both issues?

- A

 Stronger instructions, higher iteration cap
The instruction already stated the boundary and was not honoured. Raising the cap pays for more of the symptom and removes the signal.
- B

 Trajectory evaluation, plus tool and IAM enforcement
Stable completion on rising step count is degradation only trajectory evaluation explains, and a boundary that lived in the prompt has to move into the tool.
- C

 Denied topic plus faster model
Guardrails filter content categories and do not know whose data a record is. A faster model makes fifteen wasteful calls more quickly.
- D

 Reranker plus PII filter
Rerankers order retrieval results, not tool outputs. A PII filter masks personal data while leaving another customer's shipment details intact.

THE SIGNAL

Two independent faults from two modules: a trajectory problem invisible to outcome metrics, and a control that lived only in the prompt. Both have appeared before in this course.

Enterprise search: three answers, one question

OPTION	SHAPE	STATUS AND TRIGGER
Amazon Kendra	Managed enterprise search, many connectors	Closed to new customers Jul 2026. Still examinable
Bedrock Knowledge Bases	Managed RAG with a chosen vector store	The migration path AWS names. Where new work goes
OpenSearch directly	You build retrieval yourself	Maximum control, maximum operational overhead

ON THE EXAM

AWS recommends migrating Kendra applications to Bedrock Knowledge Bases. On the exam Kendra is still a correct answer for enterprise search across many source systems.

A European bank is deploying an assistant over customer records. A regulator requires that no customer data is processed outside the EU, and that the bank can demonstrate this rather than assert it. The bank has 40 AWS accounts and cannot rely on each team configuring correctly.

Traffic is spiky, peaking sharply each morning, and a single-Region deployment throttles at peak. The finance team separately wants monthly AI cost split by business unit.

Which combination meets all three requirements?

- A** Global inference profile for headroom, model invocation logging to prove where requests ran, and cost allocation tags on Lambda functions
- B** EU geographic inference profile, an SCP conditioned on `bedrock:InferenceProfileArn` at the organisation root, and tagged application inference profiles per business unit
- C** Single-Region deployment with a quota increase, CloudTrail alerts on non-EU invocations, and Cost Explorer filtered by account
- D** Provisioned Throughput in eu-west-1, AWS Config rules detecting non-EU profiles, and a monthly cost report exported to S3

Which combination meets all three requirements?

- A** Global profile, logging, Lambda tags

A global profile routes worldwide, breaching residency outright. Lambda tags attribute Lambda cost, not Bedrock cost.
- B** EU geographic profile, SCP, tagged application profiles

Geographic profile gives headroom inside the EU, the SCP makes it unbreakable by 40 teams, and tagged application profiles split cost in Cost Explorer.
- C** Quota increase, CloudTrail alerts, Cost Explorer by account

A quota keeps you inside one Region's capacity, and alerts detect a breach after data has already been processed.
- D** Provisioned Throughput, Config rules, exported report

Reserving capacity for a sharp daily peak on spiky traffic, plus Config rules that detect rather than prevent.

THE SIGNAL

Three requirements from three domains: capacity, enforcement and attribution. The two kinds of inference profile appear in one answer doing entirely different jobs.

Ten phrases that decide questions

- ◆ **LEAST operational overhead** means the managed service, not the thing you build.
- ◆ **Cannot be modified / cannot be changed by the team** means events, or an SCP.
- ◆ **Prevent rather than detect** eliminates CloudTrail, Config and every alarm.
- ◆ **Must remain readable / usable** means mask, never block.
- ◆ **However it is phrased** means semantic, never literal.
- ◆ **Fastest path / this week** eliminates every re-index and re-embedding.

ON THE EXAM

Underline the qualifier before reading the options. Half the option set usually falls away on that phrase alone, which is the cheapest time you will ever save.

And the ones about who is waiting

- ◆ **No one is waiting** means Batch. **May wait but must receive it** means asynchronous.
- ◆ **Feels slow / users abandon** means streaming, never a faster model.
- ◆ **Bursty or unpredictable** argues against every form of reserved capacity.
- ◆ **Even if compromised** assumes the caller's credentials are already stolen.
- ◆ **"First"** in a question means measure before you rebuild, almost every time.

ON THE EXAM

That last one is worth its own note. When a question asks what to do first, an option that commits weeks of work before establishing a cause is rarely the answer.

A healthcare provider's clinical assistant retrieves from approved protocols. Three problems have been raised. Clinicians cannot verify answers against the source text. Around 4% of answers include a dosage figure not present in any retrieved passage. And patient names sometimes appear in the assistant's logged prompts.

The provider needs all three addressed. Answers must remain complete and usable, and the assistant must keep discussing dosage.

Which combination addresses all three?

- A RetrieveAndGenerate for citations with document versions, contextual grounding checks blocking below a threshold, and Comprehend Medical PHI detection with masking before the prompt is built
- B A reasoning display explaining each answer, temperature set to zero, and a denied topic covering patient names
- C Model invocation logging for verification, a denied topic covering dosage, and guardrail sensitive information filters set to BLOCK
- D Citations from RetrieveAndGenerate, an LLM-as-a-judge evaluation for faithfulness, and Macie scanning the log bucket

Which combination addresses all three?

- A

 Citations with versions, grounding checks, Comprehend Medical masking
Verification, ungrounded claims and PHI each get the mechanism built for them, and masking keeps answers usable while dosage discussion continues.
- B

 Reasoning display, temperature zero, denied topic
An explanation is narrative, not evidence. Temperature zero makes it confidently wrong identically. A denied topic cannot enumerate patient names.
- C

 Logging, denied topic on dosage, BLOCK filters
Logging serves auditors, not clinicians at read time. A denied topic on dosage removes the feature, and BLOCK breaks the usability requirement.
- D

 Citations, judge evaluation, Macie on logs
Evaluation is a batch job and cannot stop an ungrounded answer reaching a clinician. Macie finds PHI in the logs after it has already been written.

THE SIGNAL

Three problems, three mechanisms, and two constraints that eliminate the tempting alternatives. "Must remain complete and usable" is mask over block; "keep discussing dosage" kills the denied topic.

What to carry into the exam room

- ◆ **Find the qualifier first.** Underline it before you read a single option.
- ◆ **Count the requirements.** Three stated requirements means three ticks per option.
- ◆ **Ask what is different about the failing cases.** The pattern beats the error message.
- ◆ **Distrust the five options** that are wrong more often than right.
- ◆ **Two options will both work.** The constraint the scenario stated is what separates them.

ON THE EXAM

Module 19 takes this further into pure technique: question anatomy, the time budget, multiple-response strategy and what to do when you genuinely do not know.