

MODULE 19 · EXAM STRATEGY - ALL DOMAINS

Exam Technique

Question anatomy, the time budget, elimination, multiple response, and what to do when you do not know.

AWS CONTENT None FORMAT Pure technique MODULE 19 of 22

At this level the binding constraint is not knowledge. It is **reading 244 words and deciding, 75 times, in 180 minutes**, without your judgement degrading.

What you are working with

FIGURE

Time	180 minutes
Questions	75, of which 65 are scored
Per question	2 minutes 24 seconds, on average
Passing score	750 of 1000, scaled
Scoring	Compensatory. Overall only, not per domain
Formats	Multiple choice, and multiple response

ON THE EXAM

Ten questions are unscored and unmarked, so treat all 75 as real. Compensatory scoring means a weak domain can be carried, which changes how you revise.

Where those 144 seconds actually go

- ◆ Median question, stem plus options, is **244 words** across the banks measured for this course.
- ◆ At a careful **200 words per minute**, that is **73 seconds of reading** before you think at all.
- ◆ Which leaves roughly **70 seconds to decide**, on a question with two defensible options.
- ◆ **You cannot afford to read twice.** Reading once, properly, is the whole technique.
- ◆ Some questions take 40 seconds. That is what funds the ones that take four minutes.

ON THE EXAM

Half your budget is consumed by reading before any thinking happens. That is why the reading method matters more than knowing one more service.

The five-step read

01

Last line

Read the actual question before the scenario.

02

Qualifier

Find LEAST, MOST, without, must not. Hold it.

03

Scenario

Now read it, collecting constraints, not detail.

04

Eliminate

Score each option against every constraint.

05

Commit

Answer, or flag. Do not linger between the two.

ON THE EXAM

Reading the question first changes what you notice in the scenario. Without it you absorb four paragraphs with no idea which facts are constraints and which are scenery.

How often the qualifier appears

- ◆ A **LEAST** qualifier appeared in 25 to 29% of questions across the two banks measured.
- ◆ "LEAST operational overhead" specifically: 18 to 21%. Roughly one question in five.
- ◆ **MOST** appeared in 7 to 17%, usually as most cost-effective or most secure.
- ◆ "without" appeared in 10 to 26%, and it is quietly one of the strongest constraints.
- ◆ Roughly **half of all questions** carry an explicit qualifier. Missing it is missing the question.

ON THE EXAM

'Without requiring code changes' or 'without affecting existing users' looks like phrasing and functions as a hard requirement. Treat 'without' as a qualifier in its own right.

"A retailer runs a customer service assistant on Amazon Bedrock. Traffic is highly variable. The team must reduce monthly spend and has asked for the option that requires the fewest changes to the existing application, without increasing response latency for customers."

Do not solve this. The drill is to identify what the question has already told you before you read a single option.

What are the constraints, and what has already been eliminated?

- ☐ A One constraint: reduce spend
- ☐ B Three constraints: reduce spend, fewest changes, no latency increase
- ☐ C Two constraints: reduce spend and fewest changes
- ☐ D No constraints; it is an open cost-optimisation question

◆ PAUSE AND DECIDE

What are the constraints, and what has already been eliminated?

A One constraint

Reads only the obvious goal and misses two hard requirements, which is how a defensible-looking wrong answer gets chosen.

C Two constraints

Catches the qualifier and misses "without increasing latency", which is the clause that kills Flex tier and smaller models.

B Three constraints

Reduce spend, fewest changes, and no latency increase. Highly variable traffic then eliminates every reserved-capacity option before you read one.

D No constraints

Every professional question states its constraints. An open optimisation question would have no single correct answer.

THE SIGNAL

"Highly variable" plus "fewest changes" plus "without increasing latency" removes reservations, Batch, Flex and any re-architecture. Caching and pruning survive, before you read an option.

Constraint or scenery?

- ◆ **A constraint restricts the answer.** "Cannot be modified", "must remain readable", "no console access".
- ◆ **Scenery is true and irrelevant.** "Case files reference one another heavily" when nothing asks a relationship question.
- ◆ **Test it:** would removing this sentence change which option is correct? If not, it is scenery.
- ◆ **Numbers are usually constraints.** 400 MB, twelve minutes, 2% of documents, one in fifteen.
- ◆ **Industry and company size are usually scenery,** unless a regulator is named.

ON THE EXAM

Scenery exists to make a wrong option attractive. Every distractor needs something in the scenario to point at, and an irrelevant true fact is the cheapest way to supply it.

Count the requirements, tick them off

- ◆ **Three stated requirements means three ticks per option**, not an overall impression.
- ◆ An option answering two of three is **wrong**, however elegantly it answers those two.
- ◆ This is the single highest-yield habit in this course. It converts judgement into arithmetic.
- ◆ It also defends against the option that is **better engineering** than what was asked for.
- ◆ Under time pressure people judge options holistically. That is exactly when the trap works.

ON THE EXAM

Multi-requirement questions are the professional tier's signature. Several scenarios in this course listed three requirements and offered options that each solved one beautifully.

"A pharmaceutical company with 4,000 employees across eleven countries runs a research assistant. The corpus contains 900,000 papers, many of which cite one another. Researchers ask questions such as 'what does the literature say about compound X'. Answers must cite the specific paper version. The team has two engineers."

The drill is to sort those sentences before deciding anything.

Which detail here is scenery rather than a constraint?

- | | |
|---|---|
| ☐ A "Answers must cite the specific paper version" | ☐ B "Many of which cite one another" |
| ☐ C "The team has two engineers" | ☐ D "900,000 papers" |

◆ PAUSE AND DECIDE

Which detail here is scenery rather than a constraint?

A Must cite the specific paper version

A hard constraint. It demands citations carrying a version, which eliminates any option returning prose without attribution.

C The team has two engineers

A constraint. It rules out options requiring significant operational work, and this exam states team size for exactly that reason.

B Many of which cite one another

True, interesting, and irrelevant. The stated question is single-topic semantic search. Nothing asks a multi-hop question, so GraphRAG is bait.

D 900,000 papers

A constraint. Corpus size interacts with store choice, index design and cost, and numbers in these scenarios almost always matter.

THE SIGNAL

B is module 4's scenario 1 trap, reused. An interesting true fact about the corpus, planted so that a graph database looks like the sophisticated answer.

Multiple response, and why it is different

- ◆ **All or nothing.** Four correct out of five with one miss scores exactly zero.
- ◆ In the bank measured for this course, **12% were multiple response**, all five options choosing two.
- ◆ **The count is given.** Select TWO means exactly two, so a third plausible option is wrong.
- ◆ **Find the two that cover different requirements.** Correct pairs usually complement rather than overlap.
- ◆ Two options doing the same job is a signal that one of them is not in the answer.

ON THE EXAM

The complementarity rule is the useful one. If a question wants real-time alerting and per-application attribution, the pair is one option for each, not two options that both alert.

A question asks for real-time alerting on token thresholds AND per-application usage attribution, with the LEAST operational overhead. Select TWO.

The five options are: build a Firehose pipeline into OpenSearch; use Logs Insights over invocation logs; use native CloudWatch metrics with alarms; configure a Grafana integration; write a Lambda that counts tokens into DynamoDB.

Which reasoning gets you to the pair fastest?

- ☐ A Pick the two options that sound most technically thorough
- ☐ B Match one option to each stated requirement, then apply the qualifier to break any tie
- ☐ C Eliminate anything unfamiliar, then choose between what remains
- ☐ D Choose the two cheapest options, since overhead is mentioned

◆ PAUSE AND DECIDE

Which reasoning gets you to the pair fastest?

- A

 Pick the most thorough-sounding two
Thoroughness is what distractors are built from. The Firehose pipeline sounds the most serious and is the most overhead.
- B

 One option per requirement, qualifier breaks ties
Two requirements, so one option each. Metrics and alarms answer alerting; Logs Insights answers attribution. The qualifier then eliminates the three pipelines you would build.
- C

 Eliminate the unfamiliar
A reasonable last resort and a poor first move. Unfamiliarity is about you, not about the question.
- D

 Choose the two cheapest
Cost and operational overhead are different things, and neither is stated for any option. You would be guessing at prices.

THE SIGNAL

Two requirements, two answers, one each. Three of the five options build a pipeline to reach data CloudWatch already holds, which the qualifier eliminates as a group.

Breaking the tie

- ◆ **Re-read the qualifier**, not the options. It is there precisely to separate the final two.
- ◆ **Which one requires less work?** On this exam, the lighter answer usually wins.
- ◆ **Which one is managed?** Build-it-yourself loses against a service that does the same thing.
- ◆ **Which one prevents rather than detects?** If the requirement used the word prevent, that decides it.
- ◆ **Which one solves a requirement nobody stated?** That is the wrong one, however good it is.

ON THE EXAM

Two defensible options is the designed state of a professional question, not a sign you have misread it. The tiebreaker is always in the scenario, never in your own preference.

You are 95 minutes in, on question 44. You have narrowed a question to two options and cannot separate them. You have re-read the scenario twice and spent nearly four minutes. You are currently five questions behind the pace you need.

What should you do?

- ☐ A Keep working on it. Getting it right matters more than the schedule
- ☐ B Pick the lighter of the two options, flag it, move on, and return only if time allows
- ☐ C Guess randomly, flag it, and move on
- ☐ D Skip it without answering and come back at the end

◆ PAUSE AND DECIDE

What should you do?

A Keep working on it

One mark, against the five questions this is costing you. Four minutes on a two-way choice has already passed the point of return.

C Random guess and flag

Correct on process and wasteful on information. You have already eliminated three options; use the tiebreaker rather than discarding it.

B Pick the lighter option, flag, move on

Converts a coin flip into a better-than-even guess using a real tiebreaker, banks a possible mark, and stops the bleeding.

D Skip without answering

The one genuinely bad option. There is no penalty for a wrong answer, so an unanswered question is a guaranteed zero.

THE SIGNAL

Never leave a question blank. And set a hard personal limit, around three minutes, after which you take your best guess and flag. Questions are worth the same regardless of difficulty.

What the unscored questions mean for you

- ◆ **Ten of the 75 do not count**, and AWS does not mark which. They are being trialled for future papers.
- ◆ **Practically: answer all 75 the same way.** You cannot identify them, so guessing which is wasted effort.
- ◆ **Psychologically: it is a relief.** A question that feels impossibly obscure may simply not count.
- ◆ That is one reason not to let a single brutal question shake your confidence in the next five.
- ◆ It is **not** a reason to relax. You still need 750 from the 65 that do count.

ON THE EXAM

Unscored questions are often the newest and roughest, which is exactly why one question can feel wildly harder than everything around it. Flag it, guess, and move on.

The checkpoints to hold in your head

- ◆ **Question 25 by minute 60.** Question 50 by minute 120. That leaves 60 minutes for the last 25.
- ◆ The back-loaded plan is deliberate. You will be slower late, and you want review time.
- ◆ **Check the clock at 25 and at 50 only.** Checking every question costs time and adds anxiety.
- ◆ **Behind at 25?** Tighten the three-minute limit to two and flag more aggressively.
- ◆ **Ahead at 50?** Do not slow down. Bank the time for the review pass.

ON THE EXAM

Two checkpoints, not fifteen. The purpose is to catch a systematic pacing problem early enough to correct it, not to monitor yourself continuously.

What to do with the time you saved

- ◆ **Review only flagged questions.** Re-reading confident answers wastes the time you banked.
- ◆ **Do not change an answer without a reason you can name.** "It feels wrong now" is not a reason.
- ◆ **A valid reason:** you spotted a constraint you had missed, or a later question reminded you of a fact.
- ◆ **Fresh eyes genuinely help** on questions you flagged early, before you had warmed up.
- ◆ If two flagged questions remain and five minutes are left, **make sure everything is answered first.**

ON THE EXAM

Changed answers are more often changed from right to wrong than the reverse, unless the change is driven by a specific realisation rather than a feeling.

When you genuinely do not know

- ◆ **Eliminate on process, not knowledge.** Options violating a stated constraint go, even if the service is unfamiliar.
- ◆ **Distrust the option that does the most work.** This exam rewards the lighter answer consistently.
- ◆ **Distrust absolutes.** "Always", "never", "guarantees" and "eliminates" are rarely in correct answers.
- ◆ **Prefer the managed service** over the assembled pipeline when nothing else separates them.
- ◆ **Answer anyway, and flag.** A reasoned guess on a professional exam is better than most people expect.

ON THE EXAM

These heuristics are a last resort, not a strategy. They exist for the four or five questions where the material genuinely is not there, and they will win you some of those.

Managing the second half

- ◆ **There is no scheduled break.** Plan for three continuous hours of concentration.
- ◆ **Accuracy falls before you notice it does.** Around question 50 is where technique has to carry you.
- ◆ **That is what the method is for.** Counting requirements works when you are tired; intuition does not.
- ◆ **Eat beforehand and hydrate sensibly.** Sensibly, given there is no scheduled break.
- ◆ **Ten seconds of looking away** after a hard question resets more than pushing straight on does.

ON THE EXAM

The reason this course insists on a mechanical method rather than judgement is fatigue. At question 60 you want a procedure to follow, not a feeling to consult.

The five errors people make to themselves

- ◆ **Reading the scenario before the question.** You absorb four paragraphs with no filter.
- ◆ **Answering the question you expected.** A familiar-looking scenario with an unfamiliar ask.
- ◆ **Choosing what you would build.** Your production experience is not the stated requirement.
- ◆ **Overrunning early.** Nine minutes on question eleven costs you four questions at the end.
- ◆ **Fighting a bad question.** Some questions are poorly worded. Flag, guess, move.

ON THE EXAM

The third one caught real scenarios in this course. Module 00 named it: answer the question that was asked, not the one your production experience wishes had been asked.

Test centre or online proctored?

PEARSON VUE TEST CENTRE

Someone else's problem

- No technical setup, no network risk
- A room designed for the purpose
- Travel, and a fixed appointment
- Fewer things that can go wrong

ONLINE PROCTORED

Your room, your risk

- No travel, wider slot availability
- System check, ID check, room scan
- Desk must be clear, nobody may enter
- A network problem is your problem

THE TIEBREAKER

If you take it online, run the system check on the actual machine and network days ahead, not on the morning. That check is the commonest avoidable failure.

What to do with the final seven days

- ◆ **Sit a full practice exam under timed conditions.** Not a quiz, the full 180 minutes.
- ◆ **Look at where your wrong answers cluster,** by domain and by reason. Not at the total score.
- ◆ **Rewatch only the weak domain.** Revising what you know feels productive and moves nothing.
- ◆ **Reread module 18** the day before. It is the module built for a final pass.
- ◆ **Do not learn new services in the last week.** Consolidate what is already there.

ON THE EXAM

Wrong answers cluster for two reasons: a domain gap, or a technique failure. Sort yours into those two piles, because they need completely different remedies.

What your wrong answers are telling you

YOU GOT IT WRONG BECAUSE	THAT IS A	FIX IT BY
You did not know the service	Domain gap	Rewatching that module
You missed a stated requirement	Reading failure	Counting requirements per option
You missed the qualifier	Reading failure	Question first, qualifier second
You picked what you would build	Judgement failure	Answering what was asked
You ran out of time	Pacing failure	The three-minute limit and flagging

ON THE EXAM

Only the first row is fixed by more study. The other four are fixed by this module, and in most practice results they outnumber the first.

The last 24 hours

- ◆ **Do not cram the night before.** Sleep contributes more to a 180-minute exam than one more module.
- ◆ **Reread your own notes**, not new material. Twenty minutes, then stop.
- ◆ **Arrive or log in early.** Rushing costs you the first five questions while you settle.
- ◆ **Know your two checkpoints** before you start: 25 by 60, 50 by 120.
- ◆ **Expect the first three questions to feel hard.** They always do. It is not a signal.

ON THE EXAM

That last point matters more than it sounds. Candidates who read the first question, panic, and re-read it three times have lost four minutes before they have settled.

The technique, in six lines

- ◆ **Question first, then qualifier, then scenario.** Reading once properly is the whole method.
- ◆ **Count the requirements and tick them off.** Judgement becomes arithmetic.
- ◆ **Multiple response is all or nothing.** Find two options covering different requirements.
- ◆ **Three-minute limit, then guess and flag.** Never leave anything blank.
- ◆ **Two checkpoints: 25 by 60, 50 by 120.** Review flagged questions only.

ON THE EXAM

Modules 20, 21 and 22 are three practice exams of 25 questions each. Sit the first one now, then use the results to decide what to rewatch.