

MODULE 20 · EXAM STRATEGY - ALL DOMAINS

Practice Exam 1

25 questions, domain-weighted. Nine-second holds. Sit it as timed practice, not as a lesson.

QUESTIONS 25 TIMED 60 minutes MODULE 20 of 22

How to sit this properly

- ◆ **Write every answer down** before the walkthrough. Recognition is not recall.
- ◆ **Give yourself 60 minutes** for the 25 questions. That is the real 2 minutes 24 per question.
- ◆ **Do not look anything up.** A score inflated by a search tells you nothing useful.
- ◆ **Note why you chose each answer**, not just what you chose. That is the diagnostic data.
- ◆ **Sort your errors afterwards** using module 19's five rows: domain gap, reading, judgement or pacing.

ON THE EXAM

Three multiple-response questions appear. They are marked "Select TWO" and are all-or-nothing, exactly as on the real paper.

A property surveyor receives inspection packets containing scanned condition reports, site photographs, and dictated voice notes. Every packet must yield the same nine fields, named identically regardless of which format the information arrived in.

The output feeds a Bedrock knowledge base. The engineering team has three people and has stated it will not build or maintain orchestration across multiple extraction services. Volume is 3,000 packets a day.

Which approach meets the requirements with the least operational overhead?

- ☐ A Amazon Textract, Amazon Rekognition and Amazon Transcribe orchestrated by AWS Step Functions, with Lambda functions normalising each output to the nine fields
- ☐ B Amazon Bedrock Data Automation with a project containing custom blueprints defining the nine fields for each modality
- ☐ C Send every item to a multimodal foundation model with a prompt requesting the nine fields
- ☐ D Convert all inputs to text, then use Amazon Comprehend custom entity recognition

◆ PAUSE AND DECIDE

Which approach meets the requirements with the least operational overhead?

- A

 Three services plus Step Functions
Works, and is the orchestration the team explicitly refused to maintain.
- B

 Bedrock Data Automation with blueprints
One service across all three modalities; blueprints enforce identical field names; it feeds a knowledge base directly.
- C

 Multimodal FM with a prompt
No guarantee of consistent field naming across 3,000 items a day, and the highest cost per item.
- D

 Comprehend custom entities after conversion
Still needs the conversion step this question is about, and photographs lose everything non-textual.

ON THE EXAM

Three modalities, one output schema, and a team refusing orchestration. The constraint is stated in words, so option A is eliminated by the scenario rather than by technology.

A retailer's assistant answers questions about products. The catalogue changes daily as lines are added, repriced and discontinued. Answers must state which catalogue entry they came from, and staff in different regions must only see products available in their own region.

A data scientist has proposed fine-tuning a foundation model on the full catalogue each week so the assistant responds without retrieval infrastructure.

How should the solution be built?

- ☐ A Fine-tune weekly on the full catalogue as proposed
- ☐ B Use RAG over the catalogue with metadata filtering on region, returning source citations
- ☐ C Use continued pre-training on the catalogue to teach the model retail vocabulary
- ☐ D Fine-tune one model per region and retrain each weekly

◆ PAUSE AND DECIDE

How should the solution be built?

- A

 Weekly fine-tuning
Daily changes against weekly training, no citations, and no per-region access. Three failures.
- C

 Continued pre-training
Teaches vocabulary rather than recall, and it is the most expensive rung on the ladder.
- B

 RAG with region filtering and citations
Freshness by re-indexing, citations from the retrieved chunk, and access control by filtering on a region attribute.
- D

 One fine-tuned model per region
Solves access control by brute force, multiplies the retraining bill, and still gives no citations.

ON THE EXAM

Three independent signals, any one of which rules out training: content changes daily, answers need attribution, and permissions differ by user.

An industrial supplier is building semantic search over 120 million product and specification chunks. The tool serves 400 concurrent users through the working day, with sustained query volume and a sub-second response requirement.

Many queries contain exact manufacturer part codes such as "BRG-88-4417K", and users report these currently return near-miss products. Others are descriptive, such as "bearing for a high-vibration environment".

Which vector store and search configuration should be used?

- ☐ A Amazon S3 Vectors with semantic search
- ☐ B Amazon OpenSearch Serverless with hybrid search
- ☐ C Amazon Neptune Analytics with graph traversal
- ☐ D Aurora PostgreSQL with pgvector and semantic search

◆ PAUSE AND DECIDE

Which vector store and search configuration should be used?

- A

 S3 Vectors, semantic
Built for sparse queries and latency tolerance. This is sustained traffic with a sub-second requirement, and S3 Vectors as a KB store is semantic-only.
- B

 OpenSearch Serverless, hybrid
Sustained interactive traffic justifies the OCU floor, and hybrid search is what makes exact part codes match alongside descriptive queries.
- C

 Neptune Analytics
Graph traversal answers relationship questions. Nothing here asks one.
- D

 Aurora pgvector, semantic
120 million vectors is demanding for pgvector, and semantic-only leaves the part-code problem unfixed.

ON THE EXAM

Two decisions in one question. Sustained high query volume picks the store; exact identifiers failing while prose works picks hybrid search.

A software vendor indexes its API reference documentation. Each page has a method signature, a parameters table, a description and a set of examples, under numbered section headings.

With 250-token chunks, retrieval returns parameter descriptions without the method they belong to, and the assistant attributes parameters to the wrong methods. With 2,000-token chunks the attribution errors stop but precision falls badly and answers become vague.

Which chunking strategy should the vendor adopt?

- ☐ A Semantic chunking, splitting at topic shifts detected by embedding similarity
- ☐ B Hierarchical chunking, searching child chunks and returning the parent chunk
- ☐ C Fixed-size chunking at 1,000 tokens as a compromise
- ☐ D No chunking, indexing each documentation page as a single chunk

◆ PAUSE AND DECIDE

Which chunking strategy should the vendor adopt?

- A

 Semantic chunking
Suits flowing prose. API reference documentation is explicitly structured under numbered headings, which semantic splitting ignores.
- B

 Hierarchical: search child, return parent
Precise matching on small children, with the parent supplying the method signature that scopes the parameters.
- C

 Fixed size at 1,000 tokens
Splits the difference between two known failures, which usually yields a little of both.
- D

 One chunk per page
One vector representing an entire page matches nothing precisely, and may exceed embedding input limits.

ON THE EXAM

Both failure modes are described for you: small chunks lose scope, large chunks lose precision. That pairing is the hierarchical signature.

A research tool loads a 90-page regulatory filing into the context window, then answers a series of analyst questions about it. Analysts typically ask between six and fifteen questions about the same filing within ten minutes.

Input tokens dominate the bill. The prompt is assembled as: request timestamp, analyst name, the full filing text, then the question. An engineer has proposed truncating the filing to its first 30 pages.

What is the better change?

- ☐ A Truncate the filing to its first 30 pages as proposed
- ☐ B Reorder the prompt so the filing text comes first as a stable prefix, then enable prompt caching
- ☐ C Enable prompt caching without changing the prompt order
- ☐ D Switch to a model with a larger context window

◆ PAUSE AND DECIDE

What is the better change?

- A

 Truncate to 30 pages
Cuts tokens by discarding two thirds of the document the tool exists to analyse.
- B

 Reorder for a stable prefix, then cache
Six to fifteen questions in ten minutes over identical content is the caching profile, and reordering is what makes the prefix cacheable at all.
- C

 Cache without reordering
The timestamp changes every call, so the prefix differs from the first character and no cache entry is ever reused.
- D

 Larger context window
Reprocesses the whole filing on every question and raises the ceiling on a bill that is already the problem.

ON THE EXAM

The prompt assembly order is stated deliberately. A timestamp at the front invalidates every cache entry, which makes C the trap for anyone who knows caching exists but not how it works.

An energy company loads a nightly file of 30 million meter readings into Amazon S3 and catalogues it in AWS Glue. A foundation model then generates consumption summaries.

Twice this year an upstream change silently altered a unit of measurement, and the model produced confident summaries wrong by three orders of magnitude. No job failed. The team wants to stop bad batches reaching the model, and to be alerted when the data's shape drifts even when every rule passes.

Which two actions should the team take? (Select TWO.)

- ☐ A A Glue Data Quality DQDL ruleset on value ranges and completeness
- ☐ B Glue Data Quality anomaly detection on historical profiles
- ☐ C Nightly record counts to CloudWatch with a drop alarm
- ☐ D Sample 1,000 records and ask a model whether the data looks right
- ☐ E S3 Versioning so bad batches can be rolled back

◆ PAUSE AND DECIDE

Which two actions should the team take? (Select TWO.)

- A

 DQDL ruleset on ranges and completeness
A unit change moves values by orders of magnitude, which a range rule catches deterministically. This is the gate for known failures.
- B

 Anomaly detection on historical profiles
The second requirement asks for alerting when shape drifts even though every rule passes. That is exactly what anomaly detection provides.
- C

 Record counts plus an alarm
A unit change does not alter the record count, so this alarm never fires.
- D

 Ask a model whether it looks right
A non-deterministic tool for a deterministic job, on a sample rather than the batch.
- E

 S3 Versioning
Recovery rather than prevention. The requirement is to stop bad batches reaching the model.

ON THE EXAM

Two requirements, so two answers, one each. Rules gate known failure modes; anomaly detection catches the ones nobody wrote a rule for. Neither alone satisfies the question.

A humanitarian organisation is building a knowledge base over field reports written in fourteen languages, including Arabic, Swahili and Tagalog. Staff query it in their own language and expect relevant reports regardless of the language those reports were written in.

The team is evaluating embedding models and has a fixed storage budget that limits it to roughly 40 million vectors.

Which embedding approach should be used?

- ☐ A Amazon Titan Text Embeddings V2 at 1,024 dimensions
- ☐ B Cohere Embed Multilingual v3
- ☐ C Translate all reports to English at ingestion, then use an English embedding model
- ☐ D Amazon Nova Multimodal Embeddings

◆ PAUSE AND DECIDE

Which embedding approach should be used?

- A

Titan Text Embeddings V2
A capable general model and the English-leaning default. Cross-language retrieval across fourteen languages is not what it is chosen for.
- B

Cohere Embed Multilingual v3
Purpose-built for multilingual corpora, so a Swahili query lands near an Arabic report about the same subject in one shared space.
- C

Translate everything to English first
Adds a translation pipeline, a per-document cost, and a lossy step before anything is embedded.
- D

Nova Multimodal Embeddings
The answer when text must retrieve images, video or audio. Everything here is text.

ON THE EXAM

Cross-language retrieval is a model capability, not a pipeline. Option D is the distractor for anyone who reads "multi" and stops there.

A Japanese financial services firm must ensure that all inference on customer data occurs within Japan. Its regulator requires evidence of the control rather than a policy statement. The firm operates 18 AWS accounts and cannot depend on each development team configuring correctly.

Traffic peaks sharply at market open, and a single-Region deployment currently throttles during that window.

Which combination should the firm implement?

- ☐ A A global cross-Region inference profile, with model invocation logging to evidence where requests ran
- ☐ B A jp geographic inference profile, plus a service control policy at the organisation root conditioned on bedrock:InferenceProfileArn
- ☐ C Single-Region deployment with a service quota increase, and CloudTrail alerts on invocations outside Japan
- ☐ D Provisioned Throughput in ap-northeast-1, and AWS Config rules detecting non-Japanese inference profiles

◆ PAUSE AND DECIDE

Which combination should the firm implement?

- A

 Global profile plus logging
A global profile routes to any commercial Region worldwide, breaching residency. Logging evidences the breach rather than preventing it.
- B

 jp geographic profile plus an SCP at the root
The profile gives headroom inside Japan, and the SCP makes the boundary unbreakable by any of the 18 account teams.
- C

 Quota increase plus CloudTrail alerts
A quota keeps you inside one Region's capacity, and alerts detect after the data has already been processed.
- D

 Provisioned Throughput plus Config rules
Reserving capacity for a sharp daily peak, and Config rules detect rather than prevent.

ON THE EXAM

"Cannot depend on each team configuring correctly" is what selects an SCP over an IAM policy. Only an SCP is a ceiling an account administrator cannot raise.

An airline processes ticket refund requests. Every request follows a fixed sequence: verify the booking exists, confirm the fare class permits a refund, check no chargeback is open, calculate the refund from a published fare table, and issue the payment.

An engineer has built this as a Bedrock agent with five tools. Finance has queried the cost, and audit has asked for evidence that every refund followed the required sequence of checks.

How should the airline rebuild this?

- ☐ A Keep the agent and add stronger instructions specifying the required order of checks
- ☐ B Rebuild as an AWS Step Functions state machine, invoking a foundation model only where judgement is required
- ☐ C Keep the agent and enable CloudWatch generative AI observability to produce the audit evidence
- ☐ D Split into five agents behind a supervisor agent, one per check

◆ PAUSE AND DECIDE

How should the airline rebuild this?

- A

 Stronger ordering instructions
Instructions influence a probabilistic loop. Audit asked for evidence the sequence was followed, and influence is not a guarantee.
- C

 Agent plus observability
Records what the agent did each run, which evidences that the order varied rather than that it did not.
- B

 Step Functions, model only where judgement is needed
A fixed published sequence is a workflow. The state machine enforces order structurally and its execution history is the audit evidence.
- D

 Five agents and a supervisor
More model calls, more cost, more non-determinism, plus a coordination layer that can also fail.

ON THE EXAM

"Every request follows a fixed sequence", then the sequence is listed. A known path is a workflow, and building it as an agent surrenders the guarantee audit asked for.

A university is adding an AI tutor to its existing student portal. When a student asks a question, the answer should begin appearing within a second rather than arriving complete after several seconds.

Each response is for one student only. The portal already authenticates students with Amazon Cognito, and the team wants request throttling per student and to reuse its existing REST API surface.

Which delivery mechanism fits best?

- ☐ A An API Gateway REST API with response streaming
- ☐ B An API Gateway WebSocket API with connection state in DynamoDB
- ☐ C AWS AppSync with a GraphQL subscription per student
- ☐ D A Lambda function URL with response streaming, called directly by the browser

◆ PAUSE AND DECIDE

Which delivery mechanism fits best?

- A

 API Gateway REST with response streaming
Streams to the one caller, keeps the existing REST surface, and brings Cognito authorisation and per-caller usage plans with it.
- B

 WebSocket API with DynamoDB state
Suits genuinely multi-turn sessions and adds connection state the scenario has no need for.
- C

 AppSync subscription per student
AppSync earns its place on fan-out to several subscribers. Each response here is for one student.
- D

 Lambda function URL
Streams well and does not directly support Cognito user pool authorisers or per-caller throttling.

ON THE EXAM

Three constraints, not one: one recipient, existing REST surface, and per-student throttling. Since November 2025 a REST API can stream, so it satisfies all three.

A bank's platform team is exposing a shared risk-calculation capability to six separate agent applications. The calculation opens a persistent connection to a mainframe gateway, holds session state between calls, and individual calculations take between two and eight minutes. Access must be governed per agent, and the platform team wants agents to discover the capability over a standard protocol rather than hand-writing an integration each.

How should the platform team expose this?

- ☐ A An MCP server on AWS Lambda, fronted by Amazon Bedrock AgentCore Gateway
- ☐ B An MCP server on Amazon ECS, fronted by Amazon Bedrock AgentCore Gateway
- ☐ C A shared Python library published to AWS CodeArtifact for each agent to import
- ☐ D An Amazon API Gateway REST API that each agent calls directly

◆ PAUSE AND DECIDE

How should the platform team expose this?

A MCP server on Lambda

Right protocol, wrong host. Persistent connections and held state do not suit a stateless function, though the 15-minute cap is not breached here.

C Shared library in CodeArtifact

Centralises code and not deployment. Every agent must upgrade its dependency, and there is no per-agent access control.

B MCP server on ECS behind AgentCore Gateway

Persistent connections and session state need a long-running host, and MCP behind Gateway gives discovery over a standard protocol with per-agent authorisation.

D REST API called directly

Solves the plumbing and not the contract. Each agent still hand-writes a tool definition, which the scenario rules out.

ON THE EXAM

Skill 2.1.7 names both hosts: Lambda for stateless lightweight tools, ECS for complex ones. Persistent connections and held state is the ECS half, stated almost verbatim.

A broadcaster self-hosts an open-weight model on Amazon SageMaker AI to generate subtitles from recorded programmes. Files average 600 MB and take eight to twenty minutes to process.

Producers upload files sporadically through the day, with long quiet periods overnight and at weekends. They accept waiting several minutes and expect the result to arrive without rescheduling anything. The current real-time endpoint runs continuously and errors on larger files.

Which SageMaker AI inference option should the broadcaster use?

- | | |
|---|---|
| ☐ A Real-time inference on a larger instance type with an increased timeout | ☐ B Serverless inference, so the endpoint scales to zero overnight |
| ☐ C Asynchronous inference, with payloads in Amazon S3 and autoscaling to zero when idle | ☐ D Batch transform on a schedule every twenty minutes |

◆ PAUSE AND DECIDE

Which SageMaker AI inference option should the broadcaster use?

- A

 Bigger real-time instance
The 6 MB payload cap and 60-second timeout belong to the inference option, not the instance. A larger instance changes neither.
- B

 Serverless inference
Scales to zero as required, and caps at 4 MB and 60 seconds. Both are breached long before the file finishes uploading.
- C

 Asynchronous inference with S3 payloads
1 GB payloads and up to an hour of processing cover 600 MB and twenty minutes, and it scales to zero through the quiet periods.
- D

 Scheduled batch transform
Would process the files and delivers on a timetable, which the scenario rules out.

ON THE EXAM

Four numbers decide this: 600 MB, twenty minutes, long idle periods, and results without rescheduling. Only asynchronous inference satisfies all four.

A customer-facing assistant on Amazon Bedrock fails intermittently during regional capacity events, and users see an error page.

The business states the assistant must keep providing some useful response during a disruption, and that a degraded answer is preferable to no answer. The team has already implemented exponential backoff with jitter, which reduced but did not eliminate the failures.

Which two additional measures should the team implement? (Select TWO.)

- ☐ A Increase the retry count and lengthen the maximum backoff interval
- ☐ B Invoke through a cross-Region inference profile to reach Regions with capacity
- ☐ C Fall back to a smaller model, then to a cached or templated response
- ☐ D Purchase Provisioned Throughput sized to peak demand
- ☐ E Move all traffic to Bedrock batch inference

◆ PAUSE AND DECIDE

Which two additional measures should the team implement? (Select TWO.)

- A** More retries, longer backoff

Backoff is already in place and the scenario says it was insufficient. More of a control that did not work is not an additional measure.
- B** Cross-Region inference profile

Regional capacity events are exactly what CRIS absorbs, routing to a Region that has capacity without any code change.
- C** Model fallback, then cached or templated response

"A degraded answer is preferable to no answer" is graceful degradation stated as a requirement.
- D** Provisioned Throughput at peak

Buys reserved capacity in one Region, which does not help during that Region's capacity event.
- E** Batch inference

Asynchronous with no completion guarantee, applied to a customer waiting for a reply.

ON THE EXAM

Two requirements again: absorb the regional shortfall, and degrade rather than fail. B answers the first and C answers the second.

A manufacturer runs a maintenance management system purchased in 2013. It writes every work order status change to a PostgreSQL database. The vendor no longer supports customisation and no internal developer remains who understands it.

The business wants AI-generated repair summaries attached to work orders when they reach "awaiting parts", and expects to add three further AI capabilities over the next eighteen months. Summaries need not be instant but must be reliable.

How should the AI capability be integrated?

- ☐ A Modify the maintenance system to call an Amazon API Gateway endpoint on each status change
- ☐ B Capture status changes from the database and publish them to Amazon EventBridge, with AI consumers subscribing to the events they need
- ☐ C Run a nightly batch job that queries the database for status changes
- ☐ D Replace the maintenance system with a platform that supports AI natively

◆ PAUSE AND DECIDE

How should the AI capability be integrated?

A Modify the system

The vendor does not support customisation and nobody remains who could do it. The scenario says this twice.

C Nightly batch job

Satisfies the first capability and commits every future one to a nightly cadence, while polling rather than reacting.

B Change capture into EventBridge

Reads a source the system already writes without touching it, and the three future capabilities subscribe rather than requiring new integrations.

D Replace the system

Replacing a maintenance platform to add a summary feature, on an eighteen-month roadmap that would not survive the migration.

ON THE EXAM

"Expects to add three further capabilities" is what selects an event bus specifically. One consumer needs no bus; four consumers and counting do.

A building society's assistant helps customers understand mortgage products. Compliance requires that it never provides anything that could be construed as personalised financial advice, however the customer phrases the request.

The team built a blocked-phrase list, and compliance has found that customers reach advice-like answers by rewording their questions in ways the list does not cover. Separately, transcripts shared with the quality team must have customer names and account numbers removed while remaining readable.

Which guardrail configuration meets both requirements?

- ☐ A A denied topic describing personalised financial advice, and sensitive information filters set to ANONYMIZE
- ☐ B An expanded word filter covering every observed phrasing, and sensitive information filters set to BLOCK
- ☐ C Content filters set to HIGH on all six categories, covering both requirements
- ☐ D A denied topic describing personalised financial advice, and sensitive information filters set to BLOCK

Which guardrail configuration meets both requirements?

- A** Denied topic plus ANONYMIZE
A denied topic is semantic, catching rewordings nobody listed. Anonymise removes the identifiers while keeping transcripts readable.
- B** Expanded word list plus BLOCK
Doubles down on the approach the scenario says already failed, and blocking on every name returns nothing usable.
- C** Content filters at HIGH
The six harm categories include neither financial advice nor PII. Maximum strength on an unrelated policy addresses neither requirement.
- D** Denied topic plus BLOCK
Correct on the first half. Blocking breaks the requirement that transcripts remain readable.

ON THE EXAM

Each half names its policy type. "However the customer phrases it" is semantic; "while remaining readable" is mask rather than block. D is the near miss.

A hospital group processes 2 million clinical notes a month. Before any note reaches a foundation model, all protected health information must be removed. The information governance team must also demonstrate precisely which parts of each note were redacted. An engineer has proposed prompting a foundation model to rewrite each note with patient details removed, arguing it handles edge cases better than a fixed detector.

How should the redaction step be built?

- ☐ A Prompt a foundation model to rewrite each note with patient details removed
- ☐ B Use Amazon Comprehend Medical to detect PHI entities and offsets, and redact those spans deterministically in a Lambda function
- ☐ C Apply a Bedrock guardrail with sensitive information filters to the notes
- ☐ D Use Amazon Macie to scan the notes in Amazon S3 and quarantine any object containing PHI

◆ PAUSE AND DECIDE

How should the redaction step be built?

- A

 FM rewrites each note
Returns new text with no record of what changed, so the demonstration requirement fails by construction.
- B

 Comprehend Medical offsets, redacted in Lambda
Offsets are an auditable record of exactly which spans were replaced, it is clinically aware, and it acts before the model sees anything.
- C

 Bedrock guardrail filters
Guardrails act at invocation, so the data has already reached Bedrock. The requirement is removal beforehand.
- D

 Macie scanning S3
Discovers and classifies sensitive data at rest. It reports on objects rather than redacting spans within a note.

THE SIGNAL

Two clauses decide it: "demonstrate precisely which parts were redacted" needs offsets, and "before any note reaches a model" rules out an invocation-time control.

An insurer runs its AI application on Amazon ECS tasks in private subnets with no internet gateway and no NAT gateway. The application invokes Amazon Bedrock and queries a Bedrock knowledge base whose source documents are in Amazon S3.

Security requires that no traffic traverses the public internet, and that a compromised container in this VPC cannot invoke any model other than the two the application is approved to use.

What should the insurer configure?

- ☐ A Interface endpoints for the Bedrock services, relying on the ECS task role's IAM policy to restrict which models can be invoked
- ☐ B Interface endpoints for bedrock-runtime and bedrock-agent-runtime, an S3 gateway endpoint, and endpoint policies restricting invocation to the two approved model ARNs
- ☐ C A NAT gateway with a security group permitting only the Bedrock service endpoints
- ☐ D An AWS Direct Connect circuit so Bedrock traffic avoids the public internet

◆ PAUSE AND DECIDE

What should the insurer configure?

- A

 Endpoints plus IAM only
The right idea in the wrong place. A compromised container runs as the task role and inherits everything that role may do.
- B

 Interface endpoints, S3 gateway, endpoint policies
Covers the runtime path, the knowledge base path and the S3 data path, and pins the approved models at the network boundary.
- C

 NAT gateway plus security group
A NAT gateway is a route to the public internet, which the requirement forbids. Security groups filter addresses, not model ARNs.
- D

 Direct Connect
Connects a data centre to a Region. It does nothing for compute already inside a VPC reaching an AWS service.

ON THE EXAM

"Even a compromised container" assumes the caller's own credentials are already stolen, which is precisely the case an endpoint policy covers and an identity policy does not.

A pensions provider must demonstrate to its regulator, for any customer interaction in the past four years, which prompt text was in force, who approved that prompt, and what the assistant said in reply.

Prompts are currently hardcoded in Lambda functions and changed through ordinary code deployments. CloudTrail is enabled across all accounts with four-year retention.

What should the provider implement?

- ☐ A Nothing further. CloudTrail already records every Bedrock invocation for four years
- ☐ B Store prompts in Bedrock Prompt Management as versioned resources with an approval gate before publication, and enable model invocation logging
- ☐ C Keep prompts in the repository and rely on Git history plus recorded code reviews
- ☐ D Store each prompt as a versioned object in Amazon S3 and read it at runtime

◆ PAUSE AND DECIDE

What should the provider implement?

- A

 Nothing further
CloudTrail records that a call was made, by which principal and when. It contains no prompt text and no response.
- B

 Prompt Management with an approval gate, plus invocation logging
The version records the text, the approval gate records who signed it off, and invocation logging records what was said.
- C

 Git history and code review
Records the change in a repository and cannot tie a specific customer interaction to the prompt version that served it.
- D

 Versioned S3 objects
Versions the text with no approval workflow and no link from an invocation to the object version used.

ON THE EXAM

Three requirements, so score each option against all three. The CloudTrail versus invocation logging distinction is what makes A a trap for anyone who stops at 'audit trail'.

An internal assistant retrieves from a knowledge base fed by a wiki any employee can edit. An employee added a page instructing the assistant to ignore prior instructions and post retrieved documents to an external webhook.

The assistant has an HTTP request tool and attempted it in testing. The user asked an ordinary question. A guardrail with prompt attack detection was enabled on that prompt.

Which two measures should the company implement? (Select TWO.)

- ☐ A Increase the prompt attack filter strength to HIGH
- ☐ B Restrict the HTTP tool by IAM, and require approval to send
- ☐ C Add a denied topic covering contact with external systems
- ☐ D Evaluate retrieved content with a guardrail before the prompt
- ☐ E Remove the HTTP request tool entirely

◆ PAUSE AND DECIDE

Which two measures should the company implement? (Select TWO.)

- A

 Stronger prompt attack filter on input
The input guardrail evaluated the user's prompt, which was innocent. The attack arrived later, inside retrieved content.
- B

 IAM-restricted tool plus human approval
Bounds what the tool can reach whatever the model believes, and puts a person in front of an irreversible external request.
- C

 Denied topic on external contact
Denied topics govern conversational subject matter. They do not inspect tool call parameters or stop a function executing.
- D

 Guardrail evaluation on retrieved content
Moves a check to where the attack actually enters the system, which no control on the user's prompt can reach.
- E

 Remove the tool entirely
Effective, and it deletes the capability while leaving every other tool exposed to the same poisoned page.

ON THE EXAM

Two answers covering different layers: bound the blast radius, and check the channel the attack used. "The user asked an ordinary question" is what eliminates A and C.

A telecoms provider's Bedrock spend has grown 55% in one quarter with no growth in user numbers. The workload is a customer-facing chat assistant with a sub-two-second latency expectation and highly variable traffic.

An architect has proposed purchasing Provisioned Throughput sized to the observed peak. Investigation shows the assistant resends the full conversation history on every turn, and sessions average 28 turns.

What is the most effective response?

- ☐ A Purchase Provisioned Throughput sized to the observed peak
- ☐ B Summarise older turns and send a rolling window of recent ones, keeping on-demand pricing
- ☐ C Set serviceTier to flex for the 50% discount
- ☐ D Move the workload to Bedrock batch inference

◆ PAUSE AND DECIDE

What is the most effective response?

- A

 Provisioned Throughput at peak
Pays for peak capacity continuously against highly variable traffic. Reservation is the wrong instrument for unpredictable load.
- B

 Summarise older turns, rolling window, stay on-demand
28 turns each resending everything is quadratic growth, which explains a 55% rise with flat user numbers better than anything about pricing.
- C

 serviceTier=flex
Trades latency under load for 50%, against a sub-two-second expectation, and Flex degrades precisely during peaks.
- D

 Batch inference
Asynchronous with no completion guarantee, applied to a customer waiting for a chat reply.

ON THE EXAM

Every pricing-mode option fails on a stated constraint, leaving the one option that reduces tokens rather than repricing them. The last sentence is the diagnosis.

A software company runs Bedrock applications for nine product teams from one account. Finance cannot determine which team drives the monthly bill and suspects two teams account for most of it.

Finance needs a monthly per-team cost breakdown in AWS Cost Explorer. A platform engineer separately wants to identify individual expensive requests to find inefficient prompt patterns.

Which combination meets both needs?

- ☐ A Create tagged application inference profiles per team for Cost Explorer attribution, and enable model invocation logging with per-request metadata for token-level detail
- ☐ B Create tagged application inference profiles per team, which provides both the monthly breakdown and per-request detail
- ☐ C Apply cost allocation tags to each team's Lambda functions
- ☐ D Build a Lambda wrapper that counts tokens per call and writes them to Amazon DynamoDB keyed by team

Which combination meets both needs?

- A

 Tagged profiles plus invocation logging metadata
Profiles give Cost Explorer the monthly per-team dollars; per-request metadata in the logs gives the token detail profiles cannot.
- B

 Profiles alone
The near miss. Application inference profiles aggregate to per usage type per day, so they answer finance and not the engineer.
- C

 Cost allocation tags on Lambda
Attributes Lambda cost, not Bedrock cost. The model spend stays a single number.
- D

 Lambda wrapper into DynamoDB
Builds a metering service AWS already provides, and its numbers will not reconcile with the actual bill.

ON THE EXAM

Two requirements at two granularities. Knowing that application inference profiles are daily aggregate rather than per-request is exactly what separates A from B.

A legal research tool analyses a matter across five independent dimensions: precedent, statute, procedural history, jurisdiction and risk. Each dimension retrieves its own documents and makes one model call. A final call composes a summary from the five.

The tool runs the five sequentially and takes 41 seconds end to end. Each dimension takes six to nine seconds. Lawyers describe it as unusable and do not read the intermediate results.

What is the most effective change?

- ☐ A Stream the final summary so text appears as it is generated
- ☐ B Run the five dimensions concurrently using a Step Functions Map state with bounded concurrency, then compose the summary
- ☐ C Switch to a smaller, faster model for all six calls
- ☐ D Combine all five dimensions into a single large prompt with one model call

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Stream the final summary
Improves time to first token on the last call only. The lawyer still waits roughly 35 seconds before that call begins.
- B

 Step Functions Map with bounded concurrency
The dimensions are independent, so total time becomes the slowest branch plus the summary. Roughly 41 seconds becomes about 15.
- C

 Smaller model everywhere
Reduces each of six sequential calls while keeping the sequence, trading analysis quality for a saving parallelism gives free.
- D

 One combined prompt
One enormous prompt across five document sets, with quality falling as attention divides, and the parallelism opportunity removed.

ON THE EXAM

"Five independent dimensions" run "sequentially" is the whole question. Independent work in sequence is latency you are choosing to pay.

A publisher is comparing four foundation models for generating article summaries. It has 1,200 articles, each with a human-written summary produced by its editorial team over several years.

The team wants an objective, repeatable comparison it can re-run whenever a new model appears, and wants to avoid paying editors to review output during model selection.

Which Bedrock Evaluations job type fits?

- ☐ A A human-based evaluation job with the editorial team scoring output from all four models
- ☐ B A programmatic evaluation job using the existing summaries as references, scored with ROUGE and BERTScore
- ☐ C An LLM-as-a-judge job scoring helpfulness and coherence without reference answers
- ☐ D A retrieve-and-generate RAG evaluation job

◆ PAUSE AND DECIDE

Which Bedrock Evaluations job type fits?

- A

 Human evaluation
The most trustworthy signal and the one thing the scenario excludes on cost. It is also not repeatable on demand.
- B

 Programmatic with existing summaries as references
1,200 human-written summaries are a reference dataset already paid for. ROUGE suits summarisation and BERTScore credits correct paraphrase.
- C

 LLM-as-a-judge without references
Would work, and discards the strongest asset available. Judging without references is what you do when references do not exist.
- D

 Retrieve-and-generate RAG evaluation
Evaluates a retrieval pipeline. There is no retrieval here; the model summarises a supplied article.

ON THE EXAM

"Each with a human-written summary" is the sentence. Reference answers already exist, which makes programmatic evaluation both cheaper and more objective.

A pharmaceutical research assistant has seen answer quality complaints rise sharply. The team holds a ground-truth dataset mapping 500 questions to the passages that should answer them.

Leadership wants to know whether to invest the next three weeks in the retrieval pipeline or in prompt and model work. The team has one engineer available and cannot do both.

What should the team do first?

- ☐ A Run a retrieve-only RAG evaluation against the ground-truth passages, measuring context relevance and coverage
- ☐ B Run a retrieve-and-generate evaluation to obtain an overall quality score
- ☐ C Re-ingest the corpus with hierarchical chunking and observe whether complaints fall
- ☐ D Run an LLM-as-a-judge model evaluation on the assistant's answers

◆ PAUSE AND DECIDE

What should the team do first?

- A

 Retrieve-only against ground truth
Answers the exact question leadership asked. Healthy coverage means the fault is downstream; poor coverage means it is not.
- B

 Retrieve-and-generate for an overall score
Confirms quality is poor, which everyone knows. One end-to-end number cannot separate the two halves.
- C

 Re-ingest with hierarchical chunking
Commits the engineer to a fix before establishing that retrieval is the problem at all.
- D

 LLM-as-a-judge on the answers
Scores generation quality against context that may already be wrong. A poor answer from poor context looks like a generation failure.

ON THE EXAM

"Ground-truth dataset mapping questions to the passages that should answer them" is the precondition for context coverage. The scenario hands you what retrieve-only needs.

A logistics assistant began returning "I don't have information about that" for roughly 12% of questions four days ago. It answered them before. No deployment has occurred and the knowledge base sync job succeeds nightly.

An engineer calls the Retrieve API directly with several failing questions and receives zero results. The same questions with the metadata filter removed return relevant chunks. A recent company restructure renamed several divisions.

What is the cause, and what should be checked?

- | | |
|--|--|
| <p>☐ A The embedding model has drifted; re-embed the corpus</p> | <p>☐ B The filter values supplied at query time no longer match the division values stored on the chunks; compare the two</p> |
| <p>☐ C Chunk boundaries are splitting the answers; re-ingest with hierarchical chunking</p> | <p>☐ D The nightly sync is too infrequent; move to event-driven synchronisation</p> |

◆ PAUSE AND DECIDE

What is the cause, and what should be checked?

- A

Embedding drift
Embedding models are fixed functions and do not drift. Removing the filter returns relevant chunks, which proves the index works.
- B

Filter values no longer match stored metadata
The bisect already isolated it. Two calls differing only by the filter, with opposite results, and a restructure that renamed divisions.
- C

Re-ingest hierarchically
Hours of work on chunk boundaries when the same chunks return perfectly once the filter is removed.
- D

Event-driven sync
The sync succeeds nightly and the content is retrievable without the filter, so frequency is not the fault.

ON THE EXAM

The engineer has already done the diagnosis. Two calls differing only by the filter isolates the fault completely, and the restructure is the change that broke it.

What your result actually means

- ◆ **20 or more of 25** is roughly on track. Under 17 and there is a domain to revisit.
- ◆ **Where your errors cluster matters far more than the total.** Count them by domain.
- ◆ **Then count them by reason**, using module 19's five rows. Domain gap, reading, judgement or pacing.
- ◆ **A near miss on a multiple-response question is a zero.** Score it as the exam would.
- ◆ **Rewatch only the weak domain**, then sit practice exam 2 on a different day.

ON THE EXAM

This is 25 questions, not 65, so treat the score as a signal rather than a prediction. The clustering of your errors is the reliable part.