

MODULE 21 · EXAM STRATEGY - ALL DOMAINS

Practice Exam 2

25 fresh questions, same weighting. Sit this on a different day from exam 1.

QUESTIONS 25 TIMED 60 minutes MODULE 21 of 22

Same rules as last time

- ◆ **Sixty minutes for 25 questions.** Time it properly, and note where you fall behind.
- ◆ **Write every answer down first,** and note why you chose it.
- ◆ **Nothing looked up.** The score is only useful if it reflects what you can recall.
- ◆ **Compare against exam 1 by domain,** not by total. Did the weak domain move?
- ◆ **Three multiple-response questions** appear again, marked "Select TWO" and all-or-nothing.

ON THE EXAM

The comparison that matters is per-domain against exam 1. A total that rose while the same domain stayed weak means you revised the wrong thing.

A software company invokes Amazon Bedrock from twelve microservices, each using InvokeModel with provider-specific request bodies. It wants to evaluate models from three different providers over the next year and adopt whichever performs best per use case. Each model swap currently requires rewriting the request and response handling in every affected service. The team wants to reduce that cost without giving up guardrails or tool use.

What should the company change?

- ☐ A Standardise on the Converse API, so switching model means changing only the model identifier
- ☐ B Build an internal library that wraps InvokeModel and normalises each provider's payload
- ☐ C Standardise on a single provider to avoid payload differences entirely
- ☐ D Use Bedrock Intelligent Prompt Routing to abstract the model choice

◆ PAUSE AND DECIDE

What should the company change?

- A

Standardise on Converse

One request shape across every messages-capable model, so a swap changes the modelId only. Guardrails and tool use are both native to it.
- B

Internal normalising library

Rebuilds what Converse already provides, and every new provider means a new adapter and a release across twelve services.
- C

Standardise on one provider

Solves the payload problem by abandoning the requirement to evaluate three providers.
- D

Intelligent Prompt Routing

Routes between two models in one family based on predicted difficulty. It is not an abstraction over providers.

ON THE EXAM

"Without giving up guardrails or tool use" is there to confirm Converse rather than a hand-rolled wrapper, since both are native to it.

A bank classifies 1.4 million transaction descriptions a day into 60 merchant categories using a frontier model. Accuracy is excellent and the business is satisfied with it.

The model spend is now the largest platform cost and P95 latency is affecting a downstream reconciliation job. Eighteen months of model invocation logs containing prompts and the accepted classifications are already retained. Accuracy must not measurably fall.

What should the bank do?

- ☐ A Run a Bedrock Model Distillation job using the invocation logs, with the frontier model as teacher and a small model as student
- ☐ B Switch to the smallest available model and adjust the prompt to compensate
- ☐ C Apply continued pre-training to a small model using eighteen months of transaction text
- ☐ D Purchase Provisioned Throughput for the frontier model to lower the per-token rate

◆ PAUSE AND DECIDE

What should the bank do?

- A

 Distillation seeded from invocation logs
Exactly its purpose. The logs already hold prompts and accepted responses, so the teacher never regenerates them, and the student inherits behaviour on this narrow task.
- B

 Smallest model plus prompt tuning
Abandons the stated accuracy constraint. A small model prompted cold will not match a frontier model on a 60-way classification.
- C

 Continued pre-training on transaction text
Teaches vocabulary. The model already reads transaction descriptions; what it must learn is the 60-category mapping, which needs labels.
- D

 Provisioned Throughput
Changes the billing shape rather than the unit economics, and does nothing for latency.

ON THE EXAM

"Similar accuracy at lower cost and latency" is distillation's stated purpose, and "logs containing prompts and the accepted classifications" is the dataset being handed to you.

A national archive is building semantic search over 60 years of digitised records, around 1.2 billion chunks once embedded. Thirty researchers use it, generating roughly 600 queries a day concentrated in the mornings.

A five-second response is acceptable. The archive costed a provisioned OpenSearch Serverless collection and found the monthly figure impossible to justify. The corpus grows slowly and is queried far more rarely than it is stored.

Which vector store should the archive use?

- ☐ A Amazon OpenSearch Serverless, accepting the OCU floor as the cost of low latency
- ☐ B Amazon S3 Vectors as the vector store for the knowledge base
- ☐ C Aurora PostgreSQL with pgvector at a low minimum ACU
- ☐ D Amazon Neptune Analytics, since records cross-reference one another

◆ PAUSE AND DECIDE

Which vector store should the archive use?

- A

OpenSearch Serverless
Buys low latency the archive says it does not need, and pays a per-hour floor for 600 queries. The cost objection is stated.
- B

S3 Vectors
Enormous corpus, sparse queries, latency tolerance. Storage-and-request billing matches a workload stored far more than queried.
- C

Aurora pgvector
Cheaper than OpenSearch, and 1.2 billion vectors is far beyond what pgvector serves comfortably.
- D

Neptune Analytics
Graph storage answers relationship questions. Cross-referencing is scenery here; the requirement is semantic search.

ON THE EXAM

Query rate rather than corpus size decides this. Option D is the scenery trap: an interesting true fact about the records with no question attached to it.

An investment firm's research assistant uses the RetrieveAndGenerate API over a knowledge base. Single-topic questions are answered well.

Compound questions fail. "Compare our energy sector exposure in 2024 with 2026 and explain what drove the change" returns material about one year only, and the comparison is absent or invented. Retrieval returns five chunks, all concerning the same year.

What should the firm configure?

- ☐ A Increase top-k from 5 to 30 so chunks about both years are returned
- ☐ B Enable query decomposition through orchestrationConfiguration so each sub-question is retrieved separately
- ☐ C Switch to the Retrieve API and assemble the context in application code
- ☐ D Add a reranking model to surface chunks from both years

◆ PAUSE AND DECIDE

What should the firm configure?

- A

 Raise top-k to 30
One embedded query is one point in the space. Thirty results from the same bad vector are thirty results from the same neighbourhood.
- B

 Enable query decomposition
The compound question becomes separate sub-queries, each retrieved on its own terms, so both years get guaranteed coverage.
- C

 Switch to Retrieve
Actively removes the orchestration step that solves this, since plain Retrieve has none.
- D

 Add a reranker
Reorders the five same-year chunks. Coverage, not ordering.

ON THE EXAM

A question containing two topics embeds into an average sitting near neither. Whichever year dominates the phrasing wins, and more results from that vector change nothing.

An insurer fine-tuned a foundation model on its underwriting guidelines. In testing, invoking the fine-tuned model by passing the base model ID to Converse returns a `ValidationException`.

An architect has proposed adding exponential backoff and retrying. The model must serve an internal tool during business hours at a steady, well-understood request rate.

What is the correct fix, and why did the call fail?

- ☐ A Retry with exponential backoff; the `ValidationException` indicates transient capacity pressure
- ☐ B Create a Provisioned Throughput deployment for the custom model and invoke it by that deployment ARN
- ☐ C Set `serviceTier` to `reserved`, which is required for fine-tuned models
- ☐ D Invoke through a global cross-Region inference profile to find available capacity

◆ PAUSE AND DECIDE

What is the correct fix, and why did the call fail?

A Retry with backoff

Misreads the error. ValidationException means the request itself is unsupported, and retrying an invalid request returns the same error forever.

C serviceTier=reserved

Reserved is a service tier for base models on the on-demand path. It does not make a custom model invocable.

B Provisioned Throughput deployment, invoked by ARN

A custom model must be deployed before it serves anything, and is then called by the deployment ARN. Steady business-hours load suits Provisioned Throughput.

D Global cross-Region profile

CRIS routes base models across Regions. A custom model exists only where you customised it, so there is nothing to route to.

ON THE EXAM

Two lessons: a custom model needs its own deployment and is called by that ARN, and ThrottlingException means retry while ValidationException means the request is wrong.

A publisher's assistant retrieves from a knowledge base of editorial policies. Two problems have emerged. It sometimes quotes a policy withdrawn two months ago, and searches containing internal reference codes such as "ED-2024-118" return the wrong policies.

The nightly ingestion job reports success every night. The knowledge base uses semantic-only search. Re-ingesting the 8 million chunks takes nine hours and the team needs both problems fixed this week.

Which two actions should the team take? (Select TWO.)

- ☐ A Switch retrieval to hybrid search
- ☐ B Re-ingest using semantic chunking
- ☐ C Fix deletion propagation so withdrawn policies leave the index
- ☐ D Re-embed the corpus with a larger embedding model
- ☐ E Move from nightly to event-driven synchronisation

◆ PAUSE AND DECIDE

Which two actions should the team take? (Select TWO.)

- A

 Switch to hybrid search
Reference codes failing while prose works is the hybrid signal, and it is a query-time parameter rather than a re-index.
- B

 Re-ingest with semantic chunking
Nine hours on chunk boundaries, which fix neither exact-code matching nor stale content.
- C

 Fix deletion propagation
A job succeeding nightly while withdrawn policies remain retrievable means additions work and removals do not.
- D

 Larger embedding model
Bigger embeddings still blur near-identical strings, and it is a full re-index for an unlikely gain.
- E

 Event-driven synchronisation
The policy was withdrawn two months ago, so roughly sixty successful syncs have already run. Frequency is not the fault.

ON THE EXAM

Two problems, two answers, and "this week" eliminates every re-index. The sixty successful nightly syncs since withdrawal is what rules out frequency.

A healthcare knowledge base retrieves the top 5 chunks for each question. Evaluation shows the passage that answers the question is retrieved within the top 20 about 94% of the time, but appears in the top 5 only 71% of the time.

Answer quality tracks that 71% closely. Re-ingesting takes twelve hours. The team wants the fastest improvement in answer quality.

What should the team do?

- ☐ A Retrieve the top 20 and enable reranking, then pass the reranked top 5 to the model
- ☐ B Re-ingest with hierarchical chunking to improve retrieval precision
- ☐ C Pass all 20 retrieved chunks to the model on every question
- ☐ D Re-embed the corpus with a higher-dimensional embedding model

◆ PAUSE AND DECIDE

What should the team do?

- A

 Retrieve 20, rerank, pass top 5
94% recall at 20 proves retrieval finds it; the loss is ordering, which is precisely what reranking fixes. A query-time change, live in minutes.
- B

 Re-ingest hierarchically
Twelve hours on chunk boundaries, when the evaluation already shows the right chunk is being retrieved.
- C

 Pass all 20 chunks
Would help and quadruples input tokens on every question, diluting the signal with fifteen mostly irrelevant chunks.
- D

 Higher-dimensional embeddings
A full re-index to improve recall that is already 94%. It targets the number that is not the problem.

ON THE EXAM

The two percentages are the diagnosis. High recall at 20 with low precision at 5 is an ordering problem, and ordering is the one thing reranking fixes.

A SaaS analytics platform serves 1,200 customers from a single Bedrock knowledge base. Every chunk carries a customerId attribute, and the web tier passes a metadata filter on customerId taken from a field in the client request.

A security review found that editing that field in the browser returns another customer's documents. The company needs isolation it can evidence to auditors and wants the smallest change that closes the gap.

What should the company do?

- ☐ A Validate the customerId field against the list of known customer identifiers before using it
- ☐ B Derive customerId server-side from the authenticated session and apply the filter in trusted backend code
- ☐ C Create a separate knowledge base and vector index for each of the 1,200 customers
- ☐ D Apply a Bedrock guardrail that blocks responses containing other customers' data

◆ PAUSE AND DECIDE

What should the company do?

- A

 Validate against known identifiers
The substituted value is a perfectly valid customer ID. Validation confirms the value is well-formed, never that it belongs to this caller.
- C

 1,200 separate knowledge bases
Genuinely isolates, and is a large architectural change with 1,200 indexes and pipelines. Far more than the smallest fix.
- B

 Derive from the authenticated session
The vulnerability is that the caller chooses their own filter. Deriving it from the verified session removes the choice entirely.
- D

 Guardrail on responses
Guardrails filter content categories, not ownership. They have no way to know which customer a document belongs to.

ON THE EXAM

A metadata filter is only an access control if the filter cannot be chosen by the caller. A is the sharpest distractor because it validates format while leaving the choice intact.

A consultancy is building an agent that investigates client compliance questions. Investigations run between forty minutes and two hours, calling internal systems and third-party regulatory databases on the individual consultant's behalf.

Each consultant's own credentials must be used against the third-party systems, and every action must be attributable to that consultant.

Which platform should the consultancy build on?

- ☐ A AWS Lambda invoking Bedrock in a loop, with credentials in AWS Secrets Manager
- ☐ B Amazon Bedrock AgentCore, using Runtime for the long execution and Identity for outbound credentials
- ☐ C Amazon Bedrock Agents Classic with action groups backed by Lambda
- ☐ D AWS Step Functions orchestrating Lambda functions with a shared service credential

◆ PAUSE AND DECIDE

Which platform should the consultancy build on?

- A

Lambda loop with Secrets Manager
Lambda caps at fifteen minutes, and a two-hour investigation cannot run inside one. Shared secrets also lose per-consultant attribution.
- B

AgentCore Runtime plus Identity
Runtime supports up to eight hours with session isolation, and Identity handles outbound credentials so the agent acts as the individual consultant.
- C

Agents Classic with Lambda action groups
Fulfilment still runs in Lambda under the fifteen-minute ceiling, and Classic is closed to new customers.
- D

Step Functions with a shared credential
Handles the duration well and a shared service credential destroys the attribution requirement.

ON THE EXAM

Two requirements name their own answers: "up to two hours" is AgentCore Runtime's eight-hour window, and per-user outbound credentials is AgentCore Identity.

A legal review API accepts a contract and returns a structured assessment. An Amazon API Gateway REST API fronts a Lambda function that calls Bedrock synchronously. The Lambda timeout is 180 seconds.

Short contracts work. Long contracts return a 504 to the client after about 29 seconds, while CloudWatch shows the Lambda running on and completing successfully at around 70 seconds. Bedrock is billed for every one of these. Clients may wait, but must receive the result.

What is the correct fix?

- ☐ A Reduce the Lambda timeout to 29 seconds so it fails when the gateway does
- ☐ B Make the endpoint asynchronous: return a job ID immediately, process via Amazon SQS, and deliver the result by polling or notification
- ☐ C Switch to a faster model so long contracts complete within 29 seconds
- ☐ D Increase the client's read timeout to 180 seconds

◆ PAUSE AND DECIDE

What is the correct fix?

A Lower the Lambda timeout

Makes the failure tidier and still a failure. The result is never delivered, which is the one thing that must happen.

C Faster model

Buys headroom against a fixed wall that longer contracts will breach again. It treats a design mismatch as a performance problem.

B Asynchronous with job ID and SQS

"Clients may wait, but must receive the result" is the definition of async. The work completes and the answer is delivered instead of computed and discarded.

D Raise the client read timeout

The 504 comes from the gateway, not the client. A longer client timeout waits longer for a response the gateway has already abandoned.

ON THE EXAM

Work that outlives its caller's timeout is asynchronous work, whatever the endpoint currently claims. The successful 70-second completion is the diagnosis.

Three applications in one AWS account share a Bedrock model. Combined average traffic sits well below the account's requests-per-minute quota, and daily average dashboards look healthy.

Every weekday between 08:00 and 08:25 all three peak together, and users of the customer-facing application see failures. CloudWatch shows ThrottlingException. Exponential backoff with jitter is already in place and reduced but did not eliminate the errors. The customer-facing application must never be the one that fails.

What should the team do?

- ☐ A Increase the retry count and lengthen the maximum backoff interval
- ☐ B Measure true peak RPM per application, request a quota increase sized to the summed peaks, and enforce per-application rate limits so the customer-facing app keeps its share
- ☐ C Move all three applications to Provisioned Throughput
- ☐ D Place an Amazon SQS queue in front of all three applications' model calls

What should the team do?

A More retries, longer backoff

Backoff smooths a transient spike. This is a sustained twenty-five-minute overload, so longer waits move the failure later.

C Provisioned Throughput for all three

Buys guaranteed capacity for a twenty-five-minute daily peak and pays for it around the clock, including for two internal applications.

B Quota sized to summed peaks, plus per-app rate limits

Fixes the capacity shortfall and decides who loses when capacity is short. Without the second half the noisiest app still starves the customer-facing one.

D SQS in front of all three

Correct for the batch applications and wrong for the customer-facing one, where a human is waiting and a queued response is a failed response.

ON THE EXAM

Averages hide coincident peaks. And "must never be the one that fails" is a prioritisation requirement, which extra capacity alone never satisfies.

A council wants a grant assessment workflow: classify an application, retrieve the relevant funding criteria, draft an assessment, then route applications above £50,000 to a human officer who must approve or reject before anything is sent.

Approval can take up to five working days. A rejection must return the workflow to the drafting step with the officer's comments. Every external call needs retry handling.

Which orchestration should the council use?

- ☐ A Amazon Bedrock Flows, using a condition node to branch on application value
- ☐ B AWS Step Functions, using a task token callback for the officer approval and retry policies on each state
- ☐ C A single Lambda function calling Bedrock repeatedly, polling for the approval
- ☐ D A Bedrock agent with tools for classification, retrieval, drafting and approval

◆ PAUSE AND DECIDE

Which orchestration should the council use?

- A

Bedrock Flows with a condition node
Handles classify, retrieve and draft well. A five-day human pause with a rejection loop is outside what a prompt-chain builder is for.
- B

Step Functions with a task token callback
The callback pattern exists for waiting on a human for days, and per-state retry policies give the error handling asked for.
- C

One Lambda polling for approval
Lambda has a fifteen-minute maximum. Five working days is not something a function can hold.
- D

A Bedrock agent with tools
The sequence is fixed and listed in order, so making it non-deterministic buys flexibility nobody asked for.

ON THE EXAM

"Up to five working days" eliminates any Lambda-waiting answer before you consider anything else. The rejection loop and retries then choose Step Functions over Flows.

An enterprise has fourteen teams calling Bedrock directly with their own IAM roles. Nobody can attribute spend to a team, guardrail configuration differs between teams, and three teams have hardcoded a model reaching end of life in five months.

A platform team of two has been asked to fix all three problems without requiring all fourteen teams to rewrite their applications at once.

Which two actions should the platform team take? (Select TWO.)

- ☐ A Build an internal GenAI gateway that every team points at, adopted team by team
- ☐ B Apply cost allocation tags to each team's IAM roles
- ☐ C Centralise guardrail and model selection behind that gateway
- ☐ D Apply an SCP denying bedrock:InvokeModel except for approved model IDs
- ☐ E Move each team into its own AWS account

◆ PAUSE AND DECIDE

Which two actions should the platform team take? (Select TWO.)

- A** Internal GenAI gateway, adopted incrementally
One place that gives cost attribution per caller, and pointing at a gateway is an endpoint change rather than a rewrite.
- B** Cost allocation tags on IAM roles
Attributes some spend and does nothing for guardrail drift or hardcoded models. One of three problems.
- C** Centralise guardrails and model selection behind it
Closes the other two gaps. Guardrails applied once for everyone, and a model change made in one place rather than fourteen.
- D** SCP restricting model IDs
Solves the lifecycle problem by breaking three teams' applications on the day it is applied.
- E** One account per team
Gives per-team bills, multiplies guardrail configuration by fourteen, and is a large migration the scenario's constraint forbids.

ON THE EXAM

The pair is complementary: the gateway is the mechanism, and centralising policy behind it is what makes the mechanism solve all three problems.

A government agency must build a document assistant over classified planning records. Its security policy states that no record, no extract, and no text derived from a record may leave its accredited facility under any circumstances.

The facility runs AWS Outposts. A team has proposed keeping the vector store on the Outpost and sending only retrieved passages to Amazon Bedrock in the Region over AWS Direct Connect.

Does the proposal meet the requirement, and if not, what does?

- | | |
|---|---|
| ☐ A Yes. Direct Connect is a private circuit, so passages never traverse the public internet | ☐ B No. Retrieved passages are derived text. Host an open-weight model on SageMaker AI on the Outpost so inference stays local |
| ☐ C No. Use a geographic inference profile so processing stays in the correct jurisdiction | ☐ D Yes, provided model invocation logging is disabled so no derived text is retained |

◆ PAUSE AND DECIDE

Does the proposal meet the requirement, and if not, what does?

- A

 Yes, Direct Connect is private
Confuses a private path with data not leaving. The passages still travel to a Region and are processed there.
- B

 No. Open-weight model on SageMaker on the Outpost
"No text derived from a record may leave" includes retrieved passages, so inference has to happen locally. That means giving up Bedrock.
- C

 Geographic inference profile
Bounds processing to a geography. The requirement names a facility, and a geographic profile still means a Region.
- D

 Disable invocation logging
Changes what is retained rather than what is transmitted. The data left the facility before anything was logged.

ON THE EXAM

"No text derived from a record" does all the work. Retrieved chunks are derived text, so the whole retrieval-in-Region pattern is excluded and Bedrock goes with it.

An engineering firm's assistant answers questions from a knowledge base of approved design standards. It cites sources correctly. In testing, roughly 5% of answers contain a load tolerance figure that appears in no cited document.

The wording is confident and the citation is real; only the number is invented. Safety requires that no answer containing an unsupported engineering claim reaches a designer, and requires evidence the check ran on every answer.

What should the firm add?

- ☐ A Raise content filter strength to HIGH across all six categories
- ☐ B Enable contextual grounding checks with a grounding threshold, passing the retrieved passages as the source, and block responses scoring below it
- ☐ C Set temperature to zero so the model stops inventing figures
- ☐ D Add a denied topic covering load tolerance information

◆ PAUSE AND DECIDE

What should the firm add?

A Content filters at HIGH

The six harm categories do not include "factually unsupported". An invented figure is not hateful, violent or sexual, so nothing fires.

C Temperature zero

Makes the model confidently wrong in the same way every time. Determinism is not accuracy.

B Contextual grounding against the retrieved passages

Exactly the failure it detects: a claim the cited source does not support. The assessment returns per call, which is the evidence safety wants.

D Denied topic on load tolerance

Stops the assistant discussing load tolerance at all, removing the feature rather than making it accurate.

ON THE EXAM

Hallucination is not a content safety problem. A figure appearing in no cited document is ungrounded, and only grounding compares an answer to its source.

A central security team owns a guardrail encrypted with a customer managed KMS key, shared with nine workload accounts. The guardrail's resource policy allows the organisation. Applications in the security account use it successfully.

Applications in every workload account fail at inference time with an access denied error. IAM identity policies in those accounts have been reviewed and correctly permit both the Bedrock action and the guardrail ARN.

What is the most likely cause?

- ☐ A The guardrail resource policy omits the workload accounts, despite allowing the organisation
- ☐ B The KMS key policy does not grant kms:Decrypt to the workload account principals
- ☐ C The workload accounts lack a VPC endpoint for bedrock-runtime
- ☐ D A service control policy is denying bedrock:ApplyGuardrail in the workload accounts

◆ PAUSE AND DECIDE

What is the most likely cause?

- A

 Resource policy omits the accounts
The scenario states the resource policy allows the organisation, which covers all nine. The first of the two required grants is in place.
- C

 Missing VPC endpoint
Produces a connectivity failure or timeout, not access denied, and would not spare the security account.
- B

 KMS key policy missing kms:Decrypt
An encrypted guardrail needs both grants. The security account works because it owns the key; nobody else can decrypt until the key policy says so.
- D

 SCP denying ApplyGuardrail
Possible in principle, and it would not explain why the pattern follows key ownership so exactly.

ON THE EXAM

"Works in the owning account, fails everywhere else" is the signature of a key policy gap. The owner has implicit key access; nobody else does until granted.

A multinational must guarantee that Australian customer data is processed only by models running in Australian Regions. Applications currently use an `au` geographic inference profile by convention, documented in an architecture standard.

An audit found one team using a global inference profile in a test account, having copied a code sample. The organisation wants a control making this impossible rather than detectable, applied across all 52 accounts, that a local account administrator cannot remove.

What should the organisation implement?

- ☐ A An IAM policy in each account denying invocation on non-Australian inference profile ARNs
- ☐ B A service control policy at the organisation root denying `bedrock:InvokeModel` unless `bedrock:InferenceProfileArn` matches Australian profiles
- ☐ C AWS Config rules detecting use of global inference profiles, with automatic remediation
- ☐ D CloudTrail alerts on invocations using non-Australian inference profiles

What should the organisation implement?

- A

 IAM policy in each account
The right mechanism at the wrong altitude. Fifty-two policies to maintain, and any account administrator can edit their own.
- B

 SCP on InferenceProfileArn at the root
One policy, every account, and no account administrator can grant past it. The condition key makes the rule specific rather than blanket.
- C

 Config rules with remediation
Detects after the fact and remediates after that. The data has already been processed outside Australia by then.
- D

 CloudTrail alerts
Pure detection, which is what the audit already achieved. The requirement asks for impossible rather than detectable.

ON THE EXAM

Two clauses point at an SCP: "impossible rather than detectable" and "a local administrator cannot remove it". Nothing else in AWS satisfies the second.

A local authority deploys an assistant helping caseworkers assess housing eligibility. The assistant retrieves policy and produces a recommendation; caseworkers make the final decision.

Counsel has asked what happens when a decision is challenged. The current design logs the caseworker's decision and the assistant's recommendation, but nothing about what informed the recommendation or whether the caseworker agreed with it.

What should the authority add?

- ☐ A Remove the recommendation so the assistant only retrieves policy
- ☐ B Record the retrieved chunks with document versions, the prompt and model version, the recommendation, whether the caseworker accepted or overrode it, and their stated reason
- ☐ C Enable model invocation logging and retain it for the statutory period
- ☐ D Add a disclaimer stating the recommendation is advisory and the caseworker is responsible

◆ PAUSE AND DECIDE

What should the authority add?

A Remove the recommendation

Avoids the governance question by deleting the feature. A legitimate option in some contexts, and not what a challenge process needs.

C Invocation logging only

Captures prompt and completion, which is one link. It records nothing about which policy version informed it or what the caseworker did.

B Record the full chain including accept or override

Completes every link, and the accept-or-override record is what evidences meaningful human oversight rather than rubber-stamping.

D A disclaimer

Allocates responsibility on paper and produces no evidence. When a decision is challenged it answers none of the questions asked.

ON THE EXAM

The override record is the point. Human oversight is only meaningful if you can show humans sometimes disagreed, and a 100% acceptance rate is not oversight.

A media group runs generative AI in seven AWS accounts. Some teams use Bedrock, one self-hosts an open-weight model on SageMaker AI, and one calls a third-party model API outside AWS. Content safety policy differs everywhere and two teams have none.

Central security must apply one policy everywhere, prove it is applied, and prevent a team disabling it. They cannot require simultaneous application rewrites.

Which two actions should security take? (Select TWO.)

- ☐ A Define one versioned guardrail and enforce it on Bedrock calls with IAM policy-based enforcement
- ☐ B Have the non-Bedrock teams call the ApplyGuardrail API against that same guardrail
- ☐ C Publish a standard guardrail configuration for teams to adopt in their next release
- ☐ D Route all seven accounts' traffic through one central Bedrock account
- ☐ E Apply an SCP denying bedrock:InvokeModel unless a guardrail identifier is present

◆ PAUSE AND DECIDE

Which two actions should security take? (Select TWO.)

- A** One guardrail, IAM policy-based enforcement
Covers the Bedrock teams and makes the policy non-optional rather than a convention an application can omit.
- B** ApplyGuardrail for the non-Bedrock teams
Two teams do not use Bedrock inference. ApplyGuardrail evaluates content independently of the model, which no Bedrock-only control can reach.
- C** Publish a standard for adoption
Convention, not enforcement. No proof of application and nothing stops a team disabling it.
- D** Central Bedrock account
A large re-architecture that still cannot reach the third-party API, and the scenario forbids simultaneous rewrites.
- E** SCP requiring a guardrail identifier
Nearly right. It enforces that some guardrail is present rather than the correct one, and does nothing for the two non-Bedrock teams.

ON THE EXAM

The pair is complementary: enforcement for the Bedrock path, ApplyGuardrail for everything else. Two teams not using Bedrock is what makes B necessary.

A council's resident helpline assistant answers benefit questions. Analysis shows around 65% of questions are variations of about 50 underlying topics, phrased differently by each caller. Responses do not vary by caller and change only when policy is updated.

Cost and latency must both fall. An engineer implemented an ElastiCache lookup keyed on a hash of the question text and reports a 3% hit rate.

What should the team implement?

- ☐ A Increase the ElastiCache TTL so entries survive longer
- ☐ B A semantic cache: embed each question, match against stored questions above a similarity threshold, and return the cached answer on a hit
- ☐ C Enable Bedrock prompt caching on the system prompt
- ☐ D Pre-compute answers to the 50 topics and serve them from Amazon CloudFront

◆ PAUSE AND DECIDE

What should the team implement?

- A

 Longer TTL
The 3% hit rate is not an expiry problem. Entries are not matching because no two callers type the same string.
- C

 Prompt caching on the system prompt
Saves the shared prefix and still invokes the model on every question. It avoids no invocations.
- B

 Semantic cache on question embeddings
65% of traffic collapses onto 50 topics once you match on meaning rather than characters, and caller-independent answers make reuse safe.
- D

 Pre-compute 50 answers at the edge
Appealing, and serving them requires mapping an arbitrary question onto one of 50. That mapping is a semantic match, so you have built B with extra steps.

ON THE EXAM

"Phrased differently by each caller" plus a 3% hit rate on a hash is the semantic-cache signal stated twice. "Responses do not vary by caller" is what makes caching safe.

A news organisation's assistant generates long article briefings. Users complain it is slow and many navigate away before an answer appears. CloudWatch shows median InvocationLatency of 5.2 seconds, no throttling, and TimeToFirstToken is not being measured. The team already reduced maxTokens, which did not reduce complaints and produced briefings that editors called incomplete.

What should the team change?

- ☐ A Switch from Converse to ConverseStream and stream tokens to the browser as they are generated
- ☐ B Move the workload to Provisioned Throughput to guarantee capacity
- ☐ C Enable a global cross-Region inference profile to reduce latency
- ☐ D Increase Lambda memory and enable provisioned concurrency

◆ PAUSE AND DECIDE

What should the team change?

- A

 Switch to ConverseStream
Time to first token drops to a fraction of a second. Total latency is unchanged, but the wait stops being a blank screen.
- B

 Provisioned Throughput
Fixes throttling. The scenario states there is none, so this buys reserved capacity nobody needs.
- C

 Global cross-Region inference profile
Cross-Region inference is a throughput mechanism rather than a latency one, and can add a network hop rather than remove one.
- D

 Lambda memory and provisioned concurrency
Attacks cold starts, worth milliseconds against a five-second model call.

ON THE EXAM

"Navigate away before an answer appears" is a perceived-latency complaint. That TTFT is not being measured is the clue that nobody has looked at the right metric.

An agent built on Amazon Bedrock AgentCore with seven tools responds in anywhere between three and fifty seconds, apparently at random. The team cannot tell whether the variation comes from the number of tool calls, one slow tool, retrieval, or the model itself.

They need step-level visibility into what the agent did on each run, correlated with the user's request, and want to avoid building custom instrumentation.

What should the team enable?

- ☐ A Amazon CloudWatch Container Insights on the compute running the agent
- ☐ B CloudWatch generative AI observability, giving connected traces across AgentCore runtime, tools, gateway and memory
- ☐ C Model invocation logging, analysing prompt lengths with CloudWatch Logs Insights
- ☐ D Custom CloudWatch metrics emitted from each tool using the embedded metric format

◆ PAUSE AND DECIDE

What should the team enable?

A Container Insights

Infrastructure health for compute that is largely idle while waiting on tool calls. It cannot see a single agent step.

C Invocation logging plus Logs Insights

Shows the model calls and not the tool executions between them. Variable agent latency is usually the loop.

B GenAI observability across AgentCore components

Purpose-built for this: step-level traces across runtime, tools, gateway and memory, with no instrumentation to write.

D Custom metrics from each tool

Would work, and is the build-it-yourself answer the scenario rules out. It also misses orchestration time between tools.

ON THE EXAM

"Three to fifty seconds at random" across an agent loop is a variable iteration count, so you need step-level tracing. "Avoid custom instrumentation" then eliminates D.

A team runs a nightly LLM-as-a-judge evaluation on its assistant. The evaluator model and the model being evaluated come from the same provider and family. Scores have sat around 0.93 for five months.

A separate manual review by subject matter experts rated a sample of the same responses far lower, finding frequent unsupported claims. The team is unsure which number to believe.

What should the team change?

- ☐ A Trust the manual review and stop running automated evaluation
- ☐ B Switch the evaluator to a model from a different family, add faithfulness as an explicit criterion, and calibrate the judge against the expert-rated sample
- ☐ C Raise the judge's scoring threshold so fewer responses pass
- ☐ D Run the evaluation more frequently to increase the sample size

◆ PAUSE AND DECIDE

What should the team change?

- A

 Abandon automated evaluation
Throws away the only mechanism that scales. Expert review of every response is not affordable, which is why the judge existed.
- B

 Cross-family evaluator, faithfulness criterion, calibrate
Addresses all three faults: self-preference from same-family judging, a criterion that never checked grounding, and no human calibration.
- C

 Raise the threshold
Moves the line on a score that is not measuring the right thing. A biased 0.93 does not become trustworthy at 0.96.
- D

 Evaluate more often
More samples of a systematically biased measurement. Frequency improves precision, never accuracy.

ON THE EXAM

Same-family judging, a five-month flat score, and experts disagreeing are three symptoms of one cause. D is the classic error: treating a bias problem as a sample-size problem.

A regulatory classification assistant assigns one of twelve compliance categories. Spot checks find it usually assigns the correct category, but two reviewers testing the same document on the same day received different categories.

The team cannot reproduce the discrepancy on demand. Temperature is 0.8. No exception is ever raised, and the golden set pass rate sits at 91%, unchanged for months.

What should the team do first?

- ☐ A Run output diffing: submit the same documents repeatedly, measure category variance, then set temperature near zero and constrain output with a schema
- ☐ B Fine-tune the model on labelled classification examples to improve accuracy
- ☐ C Add a reasoning display so reviewers can see why each category was chosen
- ☐ D Increase the golden set size so the discrepancy is caught

◆ PAUSE AND DECIDE

What should the team do first?

- A

 Output diffing, then temperature and schema
The complaint is inconsistency, which diffing measures directly. Temperature 0.8 on a classification task is the rule broken outright.
- B

 Fine-tune for accuracy
Accuracy is not the complaint; it usually assigns the correct category. Training it to be inconsistent more accurately does not help.
- C

 Reasoning display
Shows reviewers a plausible explanation for each of two conflicting answers. Explanations are narrative, not evidence.
- D

 Larger golden set
A golden set scored once per run will not surface variance. You need the same input many times, not more inputs once.

ON THE EXAM

"The same document, different categories" is a variance problem, and no accuracy measure detects it. That the 91% has been flat for months is the confirmation.

A production assistant is producing poor answers. Four theories are proposed in one meeting: the model was superseded last month, someone edited the system prompt, the knowledge base sync has been failing silently, and a new document set was ingested with different chunking.

All four changed within the same four-week window. The team has one engineer and needs to identify the cause quickly.

What is the most effective approach?

- ☐ A Revert all four changes simultaneously, confirm quality recovers, then reapply them one at a time
- ☐ B Call the Retrieve API directly with failing questions to establish whether retrieval or generation is at fault, then bisect within that half
- ☐ C Run a full retrieve-and-generate evaluation to measure the overall quality drop
- ☐ D Re-ingest the corpus with the original chunking configuration and observe

◆ PAUSE AND DECIDE

What is the most effective approach?

- A

 Revert everything, reapply one by one
Rigorous and slow. Reverting a model version and a re-ingestion is days of work before the first data point arrives.
- B

 Retrieve directly, then bisect within the half
One API call eliminates two of the four theories. Sync and chunking are retrieval-side; model and prompt are generation-side.
- C

 Full end-to-end evaluation
Quantifies a drop everyone already agrees on. One end-to-end score cannot separate four candidate causes.
- D

 Re-ingest with original chunking
Acts on one of four theories with no evidence, spending the single engineer's time on a long operation.

ON THE EXAM

Two theories are retrieval-side and two are generation-side. A single Retrieve call splits the candidates in half before any change is made.

Reading this result against exam 1

- ◆ **Compare per domain, not by total.** Did the domain you revised actually improve?
- ◆ **A total that rose while the weak domain stayed weak** means you revised the wrong thing.
- ◆ **Sort errors by reason again.** Reading and pacing failures should fall between papers; domain gaps need study.
- ◆ **Score multiple-response strictly.** One miss is a zero, as it is on the day.
- ◆ **Leave exam 3 until last**, ideally in the final week, as a dress rehearsal.

ON THE EXAM

Two papers give you a trend rather than a snapshot. The direction of travel per domain is far more informative than either score on its own.