

MODULE 22 · EXAM STRATEGY - ALL DOMAINS

Practice Exam 3

The dress rehearsal. Sit this in your final week, under proper conditions.

QUESTIONS 25 TIMED 60 minutes MODULE 22 of 22

Make this one count

- ◆ **One sitting, sixty minutes, no interruptions.** Practise the stamina as well as the knowledge.
- ◆ **Use the two checkpoints:** question 8 by minute 20, question 16 by minute 40.
- ◆ **Apply the three-minute limit** and flag. Do not let one question eat the paper.
- ◆ **Answer everything.** There is no penalty for a wrong answer, and a blank is a guaranteed zero.
- ◆ **Do not stop to look things up.** That habit is what makes a practice score meaningless.

ON THE EXAM

Scaled to 25 questions, the pacing checkpoints are question 8 by minute 20 and question 16 by minute 40, which mirrors the back-loaded plan from module 19.

A conglomerate has eleven business units each building generative AI independently. Architectures differ widely, security reviews take months because every submission is unfamiliar, and two units have shipped without guardrails.

The CTO wants a consistent way to assess and build generative AI across all eleven units, and a mechanism that makes the approved pattern the easiest path for a team to take.

Which combination should the conglomerate adopt?

- ☐ A The AWS Well-Architected Generative AI Lens for assessment, with approved patterns published through AWS Service Catalog and CloudFormation
- ☐ B A written architecture standard circulated to all eleven units
- ☐ C An SCP denying `bedrock:InvokeModel` in every account until a review is passed
- ☐ D A central platform team that builds every generative AI application on behalf of the units

◆ PAUSE AND DECIDE

Which combination should the conglomerate adopt?

- A** GenAI Lens plus Service Catalog and IaC
The Lens gives one assessment framework across eleven units, and Service Catalog makes the approved pattern the path of least resistance rather than an obligation.
- B** A written standard
Consistency by convention. It gives no assessment framework and nothing makes the approved pattern easier than ignoring it.
- C** SCP blocking Bedrock until review
Governance that blocks experimentation gets routed around, and it addresses neither assessment consistency nor build consistency.
- D** Central team builds everything
Achieves consistency by creating a bottleneck across eleven business units. Reviews taking months would become builds taking months.

ON THE EXAM

Two requirements: a consistent way to assess, and a mechanism making the approved pattern easiest. The Lens answers the first and Service Catalog answers the second.

A maritime insurer's assistant handles queries containing dense industry terminology: charter party clauses, laytime, demurrage, general average. The base model consistently misinterprets these terms, treating "general average" as an arithmetic mean.

Retrieval works correctly and returns the right documents, but the model then misreads them. The insurer holds 40 years of unlabelled maritime correspondence and reports.

Which customisation approach fits?

- ☐ A Fine-tuning on labelled question-and-answer pairs written by underwriters
- ☐ B Continued pre-training on the unlabelled maritime corpus
- ☐ C Model Distillation using a frontier model as teacher
- ☐ D Adding a glossary of maritime terms to the system prompt

◆ PAUSE AND DECIDE

Which customisation approach fits?

A Fine-tuning on labelled pairs

Teaches a task format or tone. The problem is that the model does not understand the vocabulary before any task begins.

C Model Distillation

Transfers a teacher's behaviour to a smaller model. There is no teacher here that already reads maritime terminology correctly.

B Continued pre-training on the unlabelled corpus

The one case that justifies the most expensive rung: the model does not understand the domain's language, and 40 years of unlabelled text is exactly its input.

D A glossary in the system prompt

Worth trying first and unlikely to be sufficient. Dozens of terms with contextual meanings will not fit as reliable definitions.

ON THE EXAM

This is the rare scenario where the top of the ladder is right. Retrieval works, so it is not a knowledge problem, and the failure is comprehension of the vocabulary itself.

A professional services firm wants an assistant over documents held in SharePoint, Confluence and OneDrive. Consultants must only see documents they already have permission to open in the source system.

The firm has no vector database expertise, does not want to operate one, and wants the fastest route to a working system.

Which approach fits best?

- ☐ A A custom Bedrock knowledge base on Amazon OpenSearch Serverless, with a custom connector per source system
- ☐ B An Amazon Bedrock Managed Knowledge Base using its native SharePoint, Confluence and OneDrive connectors with document-level access control
- ☐ C Export all three sources to Amazon S3 nightly and index from there
- ☐ D Amazon Kendra with connectors to all three systems

◆ PAUSE AND DECIDE

Which approach fits best?

A Custom KB with custom connectors

Requires operating a vector store the firm has no expertise in, and hand-building three connectors that already exist.

C Nightly export to S3

Flattens three permission models into a bucket, so the source access control the requirement depends on is lost.

B Managed Knowledge Base with native connectors

No vector database to provision, native connectors for all three sources, and document-level access control that honours source permissions.

D Amazon Kendra

Would have fitted well and closed to new customers in July 2026. AWS names Bedrock Knowledge Bases as the migration path.

ON THE EXAM

Three constraints all point one way: no vector store to operate, three enterprise sources with native connectors, and permissions honoured per document.

A startup is about to ingest 200 million chunks into a new knowledge base for a multi-tenant product. It plans to add per-customer isolation, recency filtering and document-level source attribution "in a later phase once the core search works".

An architect has raised a concern about this sequencing.

What is the concern, and what should the startup do?

- ☐ A No concern. Metadata can be added to existing chunks at any time
- ☐ B Metadata is attached at ingestion, so adding tenant, timestamp and source attributes later means re-ingesting all 200 million chunks. Design the metadata schema before loading anything
- ☐ C The concern is index size; reduce embedding dimensions before ingesting
- ☐ D The concern is cost; use S3 Vectors rather than OpenSearch to keep the re-ingestion cheap

◆ PAUSE AND DECIDE

What is the concern, and what should the startup do?

- A

 No concern
Metadata rides with a chunk from ingestion. Retro-fitting attributes across an existing index is a re-ingestion, not an update.
- C

 Reduce dimensions first
A real consideration and not the concern here. Dimension choice affects cost and memory, not whether filtering is possible later.
- B

 Design the metadata schema before loading
Tenant isolation, recency filtering and source attribution are all metadata-dependent, and all three become a 200-million-chunk re-ingestion if deferred.
- D

 Use S3 Vectors to make re-ingestion cheap
Optimises the cost of a mistake rather than avoiding it, and the ingestion cost is the embedding, not the store.

ON THE EXAM

All three deferred features are metadata features. Tenant isolation especially, since it is the security control, and discovering it late means re-ingesting before you can ship safely.

A council digitised 30 years of planning applications as scanned PDFs. Many contain site plans, tables of measurements and handwritten annotations. Ingested into a knowledge base with default parsing, retrieval quality is poor and tables come back as unreadable runs of numbers.

The council needs the tabular data preserved so the assistant can answer questions about specific measurements.

What should the council change?

- ☐ A Use Amazon Bedrock Data Automation as the parser for the knowledge base
- ☐ B Increase the chunk size so tables are not split
- ☐ C Switch to a multimodal embedding model
- ☐ D Enable hybrid search so numeric values match exactly

◆ PAUSE AND DECIDE

What should the council change?

- A

 BDA as the knowledge base parser
Default parsing flattens a scanned table into a run of numbers. BDA can act as the parser, extracting document structure before anything is chunked.
- B

 Larger chunks
Keeps more of the mangled table together. The damage happened at parsing, before chunking was reached.
- C

 Multimodal embeddings
The answer when a text query must retrieve an image or video. The requirement is to read the table's contents as data.
- D

 Hybrid search
Helps match exact values that are correctly extracted. These values are not correctly extracted.

ON THE EXAM

The failure is at parsing, upstream of chunking, embedding and search. Three of these options tune stages that only ever saw already-damaged text.

A retailer deploys its assistant through AWS CodePipeline. Prompts live in Bedrock Prompt Management. Last month a prompt version was published directly in the console, and quality dropped for three days before anyone noticed.

The pipeline was green throughout, because no application code had changed. The team wants prompt changes governed as strictly as code, without slowing routine releases.

Which two actions should the team take? (Select TWO.)

- ☐ A Restrict console publishing so only the pipeline can publish prompt versions
- ☐ B Add a pipeline stage that scores a golden set and gates promotion on the result
- ☐ C Add unit tests asserting the prompt text matches an expected string
- ☐ D Enable CloudTrail on Prompt Management so console publications are recorded
- ☐ E Move prompts back into the Lambda source code

◆ PAUSE AND DECIDE

Which two actions should the team take? (Select TWO.)

- A

 Restrict console publishing
Closes the path that bypassed the pipeline entirely. Without this, any gate can be walked around.
- B

 Golden set scored as a pipeline gate
The pipeline had no quality signal for prompts, so even a properly promoted bad prompt would have gone green.
- C

 Assert the prompt text in unit tests
Tests that the prompt is what you wrote, not that it works. Any deliberate change updates the assertion and it passes again.
- D

 CloudTrail on Prompt Management
Detective, and detection was never the failure. Three days passed with a full audit trail available and nobody looking.
- E

 Prompts back into source code
Fixes governance by discarding versioning, the approval gate and the model binding.

ON THE EXAM

Two problems, two answers: a path that bypassed the pipeline, and a pipeline with no quality signal. Fixing only one leaves the other wide open.

A knowledge base holds policy documents tagged with an effective date and a status of draft, current or superseded. Users ask questions in natural language such as "what does the current expenses policy say about rail travel" and "which policies changed after April".

The application currently applies no filters. Answers frequently cite superseded policies, and date-scoped questions are answered from the wrong period.

What should the team implement?

- ☐ A Explicit metadata filtering on status, applied server-side, plus implicit filtering so date phrases in the question become date filters
- ☐ B Implicit metadata filtering only, so the model infers both status and date from the question
- ☐ C Delete superseded policies from the knowledge base entirely
- ☐ D Add "only use current policies" to the system prompt

◆ PAUSE AND DECIDE

What should the team implement?

- A

 Explicit on status, implicit on dates
Status is a known constraint the application should assert; the date is only expressed in the user's prose, which implicit filtering infers.
- B

 Implicit only
Leaves the status constraint to inference when the application already knows it should exclude superseded policies on every query.
- C

 Delete superseded policies
Destroys the historical record the second question type depends on. "Which policies changed after April" needs them.
- D

 Prompt instruction
The model can only exclude superseded policies it can see. Retrieval has already put them in the context.

ON THE EXAM

Two constraints of two kinds. Status is known to the application and belongs in an explicit filter; the date exists only in the user's wording, which is what implicit filtering is for.

A logistics firm fine-tuned a model two years ago to produce structured shipment summaries in a fixed format. It runs on Provisioned Throughput at a fixed monthly cost. The base model has just moved to Legacy status with an end-of-life date in seven months.

An engineer proposes starting a new fine-tuning run immediately on the newest base model, reusing the original training dataset.

What should the firm do first?

- ☐ A Start the fine-tuning run immediately, since seven months is a short window
- ☐ B Evaluate the newest base model with prompt engineering against the original dataset before deciding whether to fine-tune at all
- ☐ C Request an end-of-life extension from AWS Support
- ☐ D Migrate the existing custom model to an on-demand deployment to cut cost while planning

◆ PAUSE AND DECIDE

What should the firm do first?

- A

 Start fine-tuning immediately
Assumes the answer. It recreates a maintenance obligation and a Provisioned Throughput bill that may no longer be needed.
- B

 Evaluate the new base model with prompting first
The training dataset is already a golden dataset. Testing whether a plain prompt now clears the bar costs days and may remove the custom model entirely.
- C

 Request an EOL extension
EOL dates are published service lifecycle decisions, and this defers a migration that must happen anyway.
- D

 Migrate to on-demand deployment
A sensible cost move on a model going away in seven months. It optimises something you are about to delete.

ON THE EXAM

The word "first" is doing the work. Structured output is exactly what newer model generations improved most, so measure before you rebuild.

A telecoms agent handles account queries with thirteen tools. It answers correctly about 76% of the time. The commonest failure is choosing between "get_plan" and "get_subscription", whose descriptions are three words each.

The team proposes splitting into four specialist agents behind a supervisor, expecting each to reason better over fewer tools. They have not measured which tool confusions cause which failures.

What should the team do first?

- ☐ A Build the four-agent supervisor architecture as proposed
- ☐ B Instrument tool selection to identify the failing pairs, then rewrite the tool descriptions and parameter schemas
- ☐ C Switch to a larger foundation model with stronger reasoning
- ☐ D Reduce the tool set to the seven most-used tools

◆ PAUSE AND DECIDE

What should the team do first?

- A

 Build the supervisor architecture
Adds a routing model and coordination cost, while carrying the three-word descriptions into whichever agent inherits both tools.
- C

 Larger model
A stronger model still cannot distinguish two tools described in three words each. It pays more per call for the same ambiguity.
- B

 Instrument, then rewrite the tool contracts
The stated failure is tool selection, and selection happens from the description alone. It costs days and the diagnosis survives any later architecture change.
- D

 Reduce to seven tools
Removes capability to work around a description problem, and the two confused tools are both core.

ON THE EXAM

"Descriptions are three words each" states the fault, and "they have not measured" tells you the proposed architecture change is a guess.

A trading platform generates a market commentary each time a threshold is breached. Every subscribed desk should see the commentary appear in real time, in the same panel, as it is generated.

Up to 90 users may be subscribed simultaneously. The platform authenticates with Amazon Cognito and the team does not want to manage connection state itself.

Which delivery mechanism fits?

- ☐ A A Lambda function URL with response streaming, called by each client
- ☐ B AWS AppSync, with clients subscribing to a GraphQL subscription keyed by instrument
- ☐ C An API Gateway WebSocket API storing each connection ID in Amazon DynamoDB and broadcasting
- ☐ D An API Gateway REST API with response streaming, polled by each client

◆ PAUSE AND DECIDE

Which delivery mechanism fits?

- A

 Function URL per client
Ninety clients calling means ninety generations of the same commentary, at ninety times the cost.
- B

 AppSync GraphQL subscription
Generate once, fan out to every subscriber. AppSync manages the connections and integrates with Cognito natively.
- C

 WebSocket with connection IDs in DynamoDB
Achieves the same outcome by building the fan-out yourself, which the scenario rules out in words.
- D

 REST streaming with polling
Polling a stream is a contradiction. Each poller re-triggers generation or receives a partial response.

ON THE EXAM

"Every subscribed desk sees the same commentary" is fan-out, and "does not want to manage connection state" then chooses AppSync over a WebSocket API.

A research institute hosts nine fine-tuned open-weight models, each on its own SageMaker AI real-time endpoint on a single GPU instance. Eight serve fewer than 300 requests a day; the ninth serves steady production traffic.

The GPU bill is dominated by idle capacity on the eight quiet endpoints, which cannot be shut down because researchers query them unpredictably. A few seconds of response time is acceptable for those eight. The production endpoint's performance must not be affected.

What should the institute do?

- ☐ A Consolidate all nine models onto one endpoint using inference components
- ☐ B Consolidate the eight low-traffic models onto one endpoint using inference components, leaving the production model on its own endpoint
- ☐ C Move all nine models to serverless inference endpoints
- ☐ D Move the eight low-traffic models to asynchronous inference endpoints

◆ PAUSE AND DECIDE

What should the institute do?

- A

 Consolidate all nine
Cheapest, and it places production on shared instances with eight unpredictable neighbours. The scenario forbids affecting it.
- B

 Consolidate the eight, leave production alone
Eight idle GPUs become one shared endpoint, and the isolation requirement is honoured by not touching the production model.
- C

 All nine to serverless
Serverless does not support GPU-backed hosting of models this size, and caps at 60 seconds.
- D

 The eight to asynchronous inference
Would cut idle cost, at the price of turning every researcher query into a submit-and-collect flow for a few seconds of work.

ON THE EXAM

Two requirements pull opposite ways: consolidate for cost, isolate production for performance. The answer that splits the estate honours both; the cheapest honours one.

A retailer routes every customer question to a frontier model. Analysis shows 82% are simple order-status lookups the smallest model answers identically, and the returned JSON already carries a confidence value.

The remaining 18% are genuinely complex. Cost must fall substantially with no measurable drop in answer quality.

Which two approaches are appropriate here? (Select TWO.)

- ☐ A Cascade: answer with the small model, re-run on the frontier model below a confidence threshold
- ☐ B Switch entirely to the smallest model and accept the quality loss
- ☐ C Route order-status questions to the small model by intent classification
- ☐ D Enable prompt caching on the system prompt and change nothing else
- ☐ E Purchase Provisioned Throughput for the frontier model

◆ PAUSE AND DECIDE

Which two approaches are appropriate here? (Select TWO.)

- A

 Cascade on returned confidence
The output already carries a confidence value, so the check costs nothing to build. 82% stops at the small model and the hard 18% still escalates.
- C

 Deterministic intent routing to the small model
Order-status lookups are a fixed, recognisable intent, so classifying them cheaply avoids the frontier model entirely for most traffic.
- B

 Smallest model only
Explicitly violates the no-measurable-quality-drop constraint on the complex 18%.
- D

 Prompt caching only
Caches a short instruction block while every question still runs through the frontier model. The dominant cost is untouched.
- E

 Provisioned Throughput
Changes the billing shape rather than the unit economics. The expensive model is still doing all the work.

ON THE EXAM

Two complementary tiering strategies: route what you can recognise up front, and cascade what you cannot. Both keep the hard 18% on the frontier model.

A bank runs an existing Amazon Bedrock Agents Classic agent for internal servicing. Its action groups must trigger payments through the bank's core banking platform, which enforces its own authorisation, dual control and audit trail.

Security will not permit a Lambda function to hold credentials for the core platform, and requires that every payment be executed by the existing platform under its own controls.

How should the agent trigger payments?

- ☐ A Use return of control, so the agent returns the predicted action to the bank's application, which executes it through the core platform
- ☐ B Use Lambda fulfilment with credentials stored in AWS Secrets Manager
- ☐ C Give the agent an MCP tool that calls the core platform directly
- ☐ D Remove payments from the agent and handle them in a separate manual process

◆ PAUSE AND DECIDE

How should the agent trigger payments?

- A

 Return of control
The agent decides and hands the predicted action back to the bank's own application, which executes it under the core platform's existing authorisation and audit.
- B

 Lambda fulfilment with Secrets Manager
Security explicitly forbids a Lambda function holding those credentials, and it bypasses the platform's dual control.
- C

 MCP tool calling the platform directly
Same objection in a different wrapper. Something outside the bank's application still holds credentials and calls the platform.
- D

 Remove payments entirely
Deletes the capability rather than integrating it, when a mechanism designed for exactly this case exists.

ON THE EXAM

Return of control is the Agents Classic feature people forget. It lets the agent decide while your existing system, with its own auth and audit, does the acting.

A charity runs a document assistant on Amazon Bedrock from Lambda functions. New models appear every few weeks and the team wants to adopt them quickly. Each should roll out to a small share of traffic first, and revert automatically if the error rate rises.

No change may require a code deployment, and the team of two wants the least operational overhead.

What should the charity implement?

- ☐ A Store the model ID in AWS AppConfig, retrieve it at runtime, and use AppConfig deployment strategies with a CloudWatch alarm for rollback
- ☐ B Store the model ID in AWS Systems Manager Parameter Store and read it on each invocation
- ☐ C Publish a Lambda alias per model and shift traffic with weighted routing in AWS CodeDeploy
- ☐ D Store the model ID in Amazon DynamoDB with an admin Lambda function to update it

◆ PAUSE AND DECIDE

What should the charity implement?

- A

 AppConfig with deployment strategies and alarm rollback
Runtime configuration, gradual rollout and automatic rollback are all native features rather than things you build.
- B

 Parameter Store
Holds the value cheaply and has no deployment strategy, so no gradual rollout and no automatic rollback.
- C

 Lambda aliases and weighted CodeDeploy routing
Gives gradual rollout and rollback, and shifting between aliases means publishing a version, which is a deployment.
- D

 DynamoDB plus an admin Lambda
Solves it, and you now own a configuration service. Rollout and rollback logic are yours to write and maintain.

ON THE EXAM

AppConfig against Parameter Store is the pair to know. Both hold a value; only AppConfig has gradual rollout and alarm-triggered rollback.

An insurer's assistant explains policy terms to customers. Compliance has written the eligibility rules as a formal set of conditions and requires that any statement the assistant makes about eligibility be verifiably consistent with those rules, with an auditable record of the check.

Contextual grounding checks are already enabled against the policy documents, and unsupported claims have stopped. Eligibility statements remain a concern because they combine several rules.

What should the insurer add?

- ☐ A Raise the contextual grounding threshold so more responses are blocked
- ☐ B Add Automated Reasoning checks, encoding the eligibility rules as a formal policy
- ☐ C Add a denied topic covering eligibility statements
- ☐ D Route all eligibility questions to a human agent

◆ PAUSE AND DECIDE

What should the insurer add?

- A

 Raise the grounding threshold
Grounding asks whether a source supports a claim. It cannot check whether combining several rules produces a logically consistent conclusion.
- B

 Automated Reasoning checks on a formal policy
Formal verification against encoded rules, deterministic and auditable, which is exactly what "verifiably consistent" and "auditable record" describe.
- C

 Denied topic on eligibility
Removes the capability rather than verifying it. The assistant exists to explain policy terms.
- D

 Route everything to a human
Safe and abandons the automation for the question type customers most want answered.

ON THE EXAM

The two hallucination policies answer different questions. Grounding asks "does the source support this?"; Automated Reasoning asks "is this logically consistent with the rules?"

A European retailer receives an erasure request from a former customer. Their data appears in support documents indexed in a Bedrock knowledge base, in nine months of assistant conversation history in Amazon DynamoDB, and in model invocation logs written to Amazon S3.

The team has deleted the support documents from the source S3 bucket and considers the request fulfilled. Legal wants evidence the person's data can no longer appear in any assistant answer.

What still needs to be done?

- ☐ A Nothing further. Deleting the source documents removes them from the knowledge base automatically
- ☐ B Run an ingestion sync so deletions propagate to the vector index, delete the relevant conversation history and invocation log records, verify with a retrieval test, and record the evidence
- ☐ C Delete and rebuild the entire knowledge base from the current source documents
- ☐ D Add a metadata filter excluding the former customer from retrieval results

What still needs to be done?

A Nothing further

Deleting a source object does not remove its chunks from the vector index until a sync runs, and never touches DynamoDB or the S3 logs.

C Delete and rebuild the knowledge base

Clears the index by brute force at significant cost, and still leaves conversation history and invocation logs untouched.

B Sync, delete the other copies, verify, record evidence

Reaches all three stores the scenario names, and the retrieval test plus the record are what satisfy legal's actual request.

D Metadata filter excluding them

Hides the data from retrieval while retaining it. Erasure means deletion, and a filter is a configuration somebody can remove.

ON THE EXAM

The scenario lists three stores in its first sentence and the team addressed one. Erasure questions are completeness questions: count the copies before you count the options.

A retailer discovered that a shared Amazon S3 bucket feeding a knowledge base contained spreadsheets of customer records nobody knew were there. They had been ingested and were appearing in answers.

The team must find any other sensitive data across 280 buckets, prevent unknown sensitive files being ingested in future, and ensure any personal data reaching a prompt is masked rather than blocking the request.

Which combination addresses all three requirements?

- ☐ A Amazon Macie to scan and classify the buckets, an Amazon Comprehend validation step in the ingestion pipeline, and Bedrock guardrail sensitive information filters set to ANONYMIZE
- ☐ B Bedrock guardrail sensitive information filters set to BLOCK on every request
- ☐ C Amazon Macie to scan the buckets, with findings used to remove sensitive files manually
- ☐ D A Lambda function scanning documents at ingestion with regular expressions for common PII formats

Which combination addresses all three requirements?

- A

 Macie, Comprehend in the pipeline, guardrail ANONYMIZE
One control per requirement, each at the right moment: discovery at rest, prevention at ingestion, masking at invocation.
- B

 Guardrail filters set to BLOCK
Acts only at invocation, so it finds nothing across 280 buckets and prevents no ingestion. It also blocks when the requirement says mask.
- C

 Macie plus manual removal
Answers discovery well and leaves the other two unaddressed. Manual removal does not prevent the next unknown file.
- D

 Lambda with regular expressions
Covers ingestion only, and regular expressions miss names and addresses that have no fixed format.

ON THE EXAM

Three requirements at three different moments. No single service acts at all three, which is why every one-service option fails.

An assistant built on Amazon Bedrock AgentCore must query Amazon DynamoDB tables owned by 30 business units, each in its own AWS account. Business unit data boundaries must be strictly maintained, and audit requires knowing which user's credentials performed each action.

The assistant serves 40,000 interactions a day, so the design must scale with account count.

Which two elements should the design include? (Select TWO.)

- ☐ A Assume a role in each target account using sts:AssumeRole with session tags carrying the end user's identity
- ☐ B Condition the target accounts' resource policies on those session tags using aws:PrincipalTag
- ☐ C Create one IAM role per end user in every business unit account
- ☐ D Use a single shared cross-account role for all queries
- ☐ E Replicate all business unit tables into the assistant's own account

◆ PAUSE AND DECIDE

Which two elements should the design include? (Select TWO.)

- A

 AssumeRole with session tags
Carries the end user's identity into the temporary credentials, so CloudTrail records which user's context performed each action.
- B

 Resource policies conditioned on aws:PrincipalTag
Turns those tags into an enforced boundary rather than only an audit annotation, which is what maintains the business unit separation.
- C

 One role per user per account
Preserves identity by brute force and does not scale. Roles per user across 30 accounts becomes unmanageable.
- D

 A single shared cross-account role
Crosses the boundary correctly and loses the person, so every action is logged as the shared role.
- E

 Replicate the tables centrally
Destroys the data boundary the requirement exists to protect.

ON THE EXAM

The pair is complementary: session tags carry the identity, and a principal-tag condition makes that identity enforceable rather than merely recorded.

A bank's data platform holds customer tables in AWS Glue Data Catalog. A new generative AI team wants to build a knowledge base over a subset of that data, and the risk function is concerned it will read columns containing salary and health information that the team is not entitled to see.

The bank wants the restriction enforced before any data reaches an embedding model.

What should the bank implement?

- ☐ A AWS Lake Formation column-level permissions restricting which columns the AI team's pipeline may read
- ☐ B Bedrock guardrail sensitive information filters on the resulting assistant
- ☐ C A metadata filter on the knowledge base excluding salary and health chunks
- ☐ D Amazon Macie scanning the Glue tables and alerting on sensitive columns

◆ PAUSE AND DECIDE

What should the bank implement?

- A

 Lake Formation column-level permissions
Governs what the pipeline may read at all, so a forbidden column never becomes a chunk and never reaches an embedding model.
- B

 Guardrail sensitive information filters
Acts at invocation, long after the data was embedded and indexed. The requirement is enforcement before embedding.
- C

 Metadata filter on the knowledge base
Hides chunks at query time that should never have been created. The data is already embedded and stored.
- D

 Macie alerting on sensitive columns
Detective. It reports that the team read something it should not have, after the fact.

ON THE EXAM

"Before any data reaches an embedding model" places the control upstream of the whole pipeline, which is the one place Lake Formation sits and the others do not.

A publisher must generate SEO descriptions for a back catalogue of 3.1 million articles already stored in Amazon S3. Descriptions are needed within two weeks and nothing consumes any individual result as it is produced.

A pilot on 8,000 articles using synchronous on-demand invocation worked correctly, but finance has rejected the projected cost of the full run. Models must not change and quality must not degrade.

Which approach meets the requirements at the lowest cost?

- | | |
|--|--|
| <p>A Keep synchronous invocation with serviceTier set to flex, spread over several days</p> | <p>B Submit the articles as a Bedrock batch inference job reading from and writing to Amazon S3</p> |
| <p>C Purchase Provisioned Throughput for one month and run the job at full rate</p> | <p>D Switch to a smaller model and keep synchronous on-demand invocation</p> |

◆ **PAUSE AND DECIDE**

Which approach meets the requirements at the lowest cost?

- A

 Synchronous with serviceTier=flex
Also 50% off and defensible. It keeps 3.1 million synchronous calls the team must orchestrate and retry itself.
- B

 Bedrock batch inference via S3
50% off, the data is already in S3, and Bedrock owns the orchestration and retries. Nothing waits on any result, which is Batch's precondition.
- C

 Provisioned Throughput for a month
Buys guaranteed capacity at a premium hourly rate for a job with no latency requirement to protect.
- D

 Smaller model
Directly violates the stated constraint that quality must not degrade and models must not change.

ON THE EXAM

A and B both give 50%. Batch wins on operational overhead: the data is already in S3 and no result is time-sensitive. When price ties, ask who runs the loop.

An internal assistant has run for eleven months. Every CloudWatch dashboard is green: latency stable, no throttling, error rate near zero, token usage flat. Nothing has been deployed for eight weeks.

Helpdesk ticket volume has risen 35% over two months and staff say the assistant "used to be better". The team cannot reproduce a failure and has no data showing anything wrong.

What should the team put in place?

- ☐ A Increase CloudWatch metric resolution and add more infrastructure alarms
- ☐ B Run the golden set on a schedule and publish pass rate as a metric, sample live responses for LLM-as-a-judge scoring, and track empty-retrieval and escalation rates
- ☐ C Enable AWS X-Ray tracing to find where latency is being introduced
- ☐ D Redeploy the most recent known-good version of the application

◆ PAUSE AND DECIDE

What should the team put in place?

- A

 More resolution, more infrastructure alarms
More precision on metrics that are already correct. The infrastructure is healthy; the answers are not.
- B

 Golden set, sampled judging, retrieval and escalation rates
Builds the quality layer the team never had. Rising tickets is escalation rate showing up in somebody else's dashboard.
- C

 X-Ray tracing
Excellent for finding where time goes. Nobody reported a latency problem and the metrics confirm there is not one.
- D

 Redeploy a known-good version
Nothing has been deployed for eight weeks, so there is no regression to roll back to.

ON THE EXAM

Green dashboards prove nothing about quality. "Nothing deployed for eight weeks" plus rising complaints is drift in the corpus, the users or the questions.

A media company's Bedrock costs jumped 300% over one weekend with no release and no traffic increase. CloudWatch shows InputTokenCount rising sharply from Saturday morning. The team has static alarms on invocation count, which did not fire because invocation count is unchanged.

They want to detect this class of event automatically in future, and to know quickly when spend departs from normal.

Which combination should the team implement?

- | | |
|--|--|
| ☐ A CloudWatch anomaly detection on token metrics, plus AWS Cost Anomaly Detection at the billing layer | ☐ B Static CloudWatch alarms on InputTokenCount at a fixed threshold |
| ☐ C Weekly review of Cost Explorer by the platform team | ☐ D AWS Config rules checking that token limits are configured on every application |

◆ PAUSE AND DECIDE

Which combination should the team implement?

- A

 Anomaly detection on tokens, plus Cost Anomaly Detection
Anomaly detection learns the normal band rather than needing a threshold nobody could have set, and the billing layer catches what application metrics miss.
- B

 Static threshold on InputTokenCount
Better than nothing, and someone must guess a number. Too low and it fires constantly; too high and the next event slips under it.
- C

 Weekly Cost Explorer review
A weekend event would run for days before anyone looked, which is the situation they are trying to fix.
- D

 Config rules on token limits
Checks that a configuration exists rather than detecting that consumption has changed.

ON THE EXAM

The alarm on invocation count is the lesson. Token spend rose while invocation count held, so an alarm on the wrong metric never fires however well tuned it is.

A medical device manufacturer is validating an assistant that summarises adverse event reports for regulatory submission. Accuracy is safety-critical, there are no existing reference summaries, and the regulator expects evidence of expert review.

The team has 40 clinical specialists available for a limited period and 6,000 reports to validate against.

How should the manufacturer structure the evaluation?

- ☐ A Human-based evaluation on every one of the 6,000 reports
- ☐ B Human-based evaluation on a stratified sample to calibrate an LLM-as-a-judge evaluation, which then scores the full 6,000 with specialists reviewing anything below threshold
- ☐ C LLM-as-a-judge evaluation on all 6,000 with no human involvement
- ☐ D Programmatic evaluation using ROUGE against the original report text

◆ PAUSE AND DECIDE

How should the manufacturer structure the evaluation?

- A

 Human evaluation on all 6,000
The strongest evidence and it does not fit the specialists' limited availability. It also gives nothing repeatable for the next release.
- B

 Human calibration, judge at scale, human review below threshold
Expert judgement where it is most valuable, scale from the judge, and a documented calibration that evidences expert review to the regulator.
- C

 Judge only, no humans
The regulator expects evidence of expert review, and an uncalibrated judge on safety-critical content has no established agreement with people.
- D

 ROUGE against the original report text
Rewards word overlap with the source. A good summary paraphrases, so this scores the wrong property.

ON THE EXAM

The three job types are complements, not alternatives. Humans calibrate the judge and handle the hard cases; the judge provides scale and repeatability.

A procurement agent resolves supplier queries using ten tools. Task completion rate has held at 89% for four months and the team considers it healthy.

Costs have risen 80% over the same period with no increase in query volume. Average tool calls per query have climbed from 4 to 16, and some queries now hit the iteration cap before completing.

What does this indicate, and what should the team evaluate?

- ☐ A Nothing is wrong; completion rate is the metric that matters and it is stable
- ☐ B The agent is degrading. Evaluate trajectories for tool selection accuracy and reasoning quality, and identify which tool's outputs changed
- ☐ C The model has become slower; switch to a faster model
- ☐ D Raise the iteration cap so more queries complete

◆ PAUSE AND DECIDE

What does this indicate, and what should the team evaluate?

- A

 Nothing is wrong
Completion rate held while tool calls quadrupled. The agent reaches the same outcome by a far worse route, and some queries no longer finish.
- B

 Degradation. Evaluate trajectories, find the changed tool
Step count leads and completion rate lags. Something is returning worse information and the agent is compensating by trying harder.
- C

 Faster model
The problem is more calls, not slower calls. A faster model makes sixteen wasteful calls more quickly.
- D

 Raise the iteration cap
Hides the symptom by paying for more of it, and removes the one signal that revealed the problem.

ON THE EXAM

Outcome metrics lag; trajectory metrics lead. A stable success rate on a rising step count is exactly the shape of an agent quietly degrading.

A tender analysis tool retrieves clauses and returns a structured JSON assessment. Most tenders work. Long tenders fail to parse and the application shows an error.

Logs show a successful Bedrock invocation with no exception. The response ends mid-object. Input token counts for the failing requests are well within the model's context window.

What is happening, and what is the fix?

- ☐ A Context window overflow. Reduce the number of retrieved clauses
- ☐ B Output truncation at maxTokens. Check stopReason, raise maxTokens, and request a more compact output schema
- ☐ C The model is producing invalid JSON. Add an instruction requiring valid JSON
- ☐ D A network timeout is cutting the response. Increase the client read timeout

◆ PAUSE AND DECIDE

What is happening, and what is the fix?

- A

 Input overflow; retrieve fewer clauses
The scenario states input tokens are well within the window. This applies the input-side fix to an output-side failure.
- B

 Output truncation at maxTokens
Ends mid-object, no exception, and longer tenders produce longer assessments. stopReason confirms it, and a compact schema addresses the cause.
- C

 Instruct it to produce valid JSON
The JSON was valid until it stopped. The model made no syntax error; it was cut off mid-generation.
- D

 Client read timeout
A timeout aborts the connection and usually raises an exception. Bedrock logged a successful invocation.

ON THE EXAM

Input and output overflow both correlate with document length, which is why they are confused. "Ends mid-object" with no exception is the output side.

Where to go from here

- ◆ **Compare all three papers by domain.** A flat weak domain across three attempts is where the last week goes.
- ◆ **Re-read module 18** the night before. It exists for exactly that pass.
- ◆ **Re-read module 19** the morning of. Two checkpoints, three-minute limit, answer everything.
- ◆ **Check the exam guide and service updates** before your date. This catalogue moves fast.
- ◆ **Build something small in Bedrock** if you have not. One retrieval pipeline teaches what ten hours of video cannot.

ON THE EXAM

That is the whole course: nine hours across 23 modules, roughly 175 original questions, every explanation sourced from AWS documentation. Good luck on exam day.