MODULE 38 · EXAM STRATEGY - ALL DOMAINS

The maintenance-mode map

What is closed, what replaced it, and what to answer in the exam anyway.

MODULE 38 of 43 DOMAIN Exam skills WEIGHT revision

admin@sthithapragna.com

Ten SageMaker AI features closed to new customers on 30 July 2026, in one announcement.

What changed on 30 July 2026

FEATURE	STATUS
Ground Truth	Maintenance. Closed to new customers
Mechanical Turk	Maintenance
Augmented AI (A2I)	Maintenance
Clarify	Maintenance
Model Monitor	Maintenance
Debugger	Maintenance
Role Manager	Maintenance
Geospatial	Maintenance
Studio Lab	Maintenance
Profiler	Sunset, with a termination date
Ground Truth Plus	End of support, 30 June 2026

Three different statuses

- **Maintenance mode:** closed to new customers. Existing customers continue as normal.
- AWS continues investing in **security and availability**, and plans **no new features**.
- **Sunset:** a termination date has been announced. Migration is not optional.
- **End of support:** already gone.
- Only one of those three forces action immediately. Maintenance is a signal, not a deadline.

> ON THE EXAM
>
> The distinction matters at work: a maintenance-mode feature you already use is safe to keep running, and unwise to build new work on.

What did not change

- **Training, inference, Studio and Canvas** are all unaffected.
- Pipelines, the Model Registry, Feature Store and JumpStart are unaffected.
- So **every domain of this exam still has a current, supported answer** for its core tasks.
- What was pruned is a layer of managed conveniences above stable primitives.
- The primitives stayed. The wrappers went.

> ON THE EXAM
>
> This matters for perspective: nine features is a lot of names and a small proportion of the platform.

Removed from MLA-C02 entirely

REMOVED FROM SCOPE	WHY IT MATTERS HERE
SageMaker Neo	Compilation moved to the inference optimization toolkit. Module 22
Bring your own container	Out of scope in C02. Prebuilt containers cover it
Model pruning	Not in the C02 guide. Quantization and distillation are
Ordering and matching question types	C02 uses multiple choice and multiple response only

ON THE EXAM

Different from maintenance mode: these services exist and are simply not examined. Module 40 covers them for C01 candidates.

The wider AWS wave

- The same June announcement moved roughly **twenty services and features** to maintenance.
- Five entered **sunset**, with termination dates.
- An earlier wave in March had already moved several services, so this is a pattern rather than an event.
- **Consolidation** is the stated logic: several older services folding into newer lines.
- Expect more. Any course or certification will lag a wave like this by months.

> ON THE EXAM
>
> The useful conclusion is not which ten. It is that a service name is a fact with a shelf life, and a capability usually is not.

The exam consequence

- **Every one of these remains examinable** if the exam guide names it. Learn how it works.
- The exam asks what a service does, and a service in maintenance still does it.
- **Do not answer 'it is deprecated'.** That is not one of the options, and it is not what is being tested.
- The status matters at work, where building new on a closed feature creates a migration.
- This course has flagged each one as it appeared, which is why this module is a map rather than news.

> ON THE EXAM
>
> Two different questions: what does this service do, and should I build on it? The exam asks the first. Your employer asks the second.

What replaced what

CLOSED	WHAT AWS POINTS YOU TO
Ground Truth	Bedrock Data Automation, or a third-party labelling service
Mechanical Turk	A vendor workforce, or your own
Augmented AI (A2I)	Your own review workflow, with Step Functions and a queue
Clarify	Published bias formulas, the SHAP library, Bedrock Evaluations
Model Monitor	AWS open-source monitoring solutions, MLflow Apps, Evidently AI
Debugger	SageMaker AI MLflow, TensorBoard on SageMaker, CloudWatch
Role Manager	IAM directly, with policies in infrastructure as code
Bedrock Agents	AgentCore, with a published migration mapping

ON THE EXAM

Two patterns in this table: some were replaced by a newer service, and others by 'do it yourself with the open tooling underneath'.

Ground Truth, Mechanical Turk and A2I

- **Ground Truth** labelled datasets before training. **A2I** reviewed predictions after deployment.
- That distinction is still examinable, and module 37 lists it as a confusion pair.
- **Mechanical Turk** was one workforce option among private and vendor workforces.
- **Bedrock human evaluation uses Ground Truth work teams**, so that path may be unavailable too. Module 18.
- The **concepts** outlive the services: active learning, confidence thresholds, consensus labelling.

> ON THE EXAM
>
> The human-in-the-loop idea is not going anywhere. What changed is which managed service orchestrates it.

Clarify

- **Clarify's mathematics is public.** Pre-training metrics are published formulas over label counts.
- Post-training metrics are standardised functions of **per-group confusion-matrix counts**.
- **Feature attribution uses the SHAP library, which is the engine Clarify itself is built on.**
- AWS points to those, plus **Bedrock Evaluations**, logged to a SageMaker AI MLflow App.
- So you lose the packaging, not the capability. That is unusually clean.

> ON THE EXAM
>
> Of all ten, this is the one where nothing is actually lost: the formulas were always published and SHAP was always open source.

Debugger and Profiler

- **Debugger is maintenance. Profiler is sunset**, with a termination date. Different urgency.
- Debugger watched tensors during training and raised rules for vanishing gradients and poor GPU use.
- AWS points to **SageMaker AI MLflow, TensorBoard on SageMaker, and CloudWatch**.
- Migration is concrete: **remove DebuggerHookConfig, Debugger TensorBoardOutputConfig and rules**.
- Both remain examinable, and you should know what they did.

> ON THE EXAM
>
> This is the clearest migration instruction AWS published: three specific configuration objects to delete from your estimator.

Model Monitor

- Four dimensions: **data quality, model quality, bias drift, feature attribution drift**. All examinable.
- AWS points to **open-source monitoring solutions** it publishes, running in your own account.
- Built on SageMaker AI, Athena, Lambda, EventBridge, SQS, SNS and QuickSight, with **MLflow Apps**.
- They use **Evidently AI**, and include a pattern for reconciling **delayed ground truth by inference ID**.
- For lighter cases: publish custom metrics to CloudWatch and use **anomaly detection alarms**.

> ON THE EXAM
>
> The inference-ID reconciliation is the notable part: it solves module 29's label-latency problem, which Model Monitor handled with a merge job.

Role Manager

- **Role Manager is on the list**, and module 34 presented it without a maintenance flag.
- That was an omission in this course, and this slide is the correction.
- It built execution roles from personas rather than from a blank policy document.
- It remains examinable, and it remains available to existing customers.
- The replacement is **IAM directly**, with policies in infrastructure as code. Modules 28 and 34.

> ON THE EXAM
>
> Flagged here rather than quietly patched, because knowing a course got something wrong is more useful than a silently corrected slide.

Geospatial and Studio Lab

- **Geospatial capabilities** handled satellite and map data in SageMaker AI.
- **Studio Lab** was a free notebook environment requiring no AWS account.
- Neither appears in the MLA-C02 exam guide, and neither was taught in this course.
- They are listed for completeness, because you will see them in the same announcement.
- If a question mentions either, it is testing something else.

> ON THE EXAM
>
> Included so the list is complete. Recognising a name and knowing it is out of scope is itself useful under time pressure.

Bedrock Agents and the wider shift

- **Bedrock Agents is now Agents Classic**, in maintenance, with a published migration mapping.
- **AgentCore** is the successor: Runtime, Gateway, Memory, Identity, Observability, Policy.
- The same announcement moved **Kendra toward Bedrock Knowledge Bases**.
- So the pattern is **consolidation**: several older services folding into the Bedrock and AgentCore line.
- Module 27 covers the agent mapping in detail, including the multi-agent gap.

> ON THE EXAM
>
> This half of the wave is a generational shift rather than a pruning. The capability moved rather than being handed back to you.

What to say, and where

QUESTION	THE ANSWER
Exam: which service labels a training dataset?	Ground Truth
Work: what should we use to label a new dataset?	Not Ground Truth. It is closed to new customers
Exam: what detects feature attribution drift?	Model Monitor with Clarify
Work: how do we monitor drift on a new project?	The open-source solutions, or CloudWatch anomaly detection
Exam: what explains an individual prediction?	Clarify, using SHAP
Work: how do we explain predictions now?	SHAP directly. Same engine, no service wrapper

ON THE EXAM

Three pairs, and the exam answer and the work answer differ in every one. Holding both is the whole point of this module.

What to do on Monday

- **Inventory:** which of the ten does your organisation actually use?
- **Separate** workloads already using them from new work. Only the second is urgent.
- **Profiler first**, because it has a termination date and the others do not.
- Read the availability change page for each one you use. The migration guidance is specific.
- Add a check to your design review: is every named service still open to new customers?

> ON THE EXAM
>
> The last bullet is the durable fix. One question in a review catches the next wave without anybody reading an announcement.

Spotting a stale source

- **Check the date on anything you read**, and treat a rate or a limit as dated information.
- This course found several: a 6 MB payload limit, a 15-minute async ceiling, 'Bedrock has no API keys'.
- **Each of those was true when written.** That is how a stale fact survives.
- Prefer the service documentation over a blog post, and prefer a page over an announcement.
- An announcement is a dated document by design. It is never revised.

> ON THE EXAM
>
> The 2021 async announcement still says 6 MB and 15 minutes, correctly, because it describes what was true in 2021.

How to track this yourself

- The **AWS Service Availability Updates** page lists maintenance, sunset and end of support together.
- Each affected service gets an **availability change page** in its own documentation.
- Those pages carry the **replacement guidance**, which is usually the most useful part.
- The service's own features page carries the notice too, dated.
- Check before designing anything new on a service you have not used this year.

> ON THE EXAM
>
> The availability change pages are genuinely good: they say what replaced it, and often exactly what to remove from your code.

What this means for the material you are studying

- This course was built after the announcement and **flags each affected feature where it appears**.
- **Module 34 missed one.** Role Manager was taught without a flag, and this module corrected it.
- Several figures in this course are dated: payload limits, prices, coverage lists.
- Where a fact could not be confirmed from AWS documentation, the course said so.
- Treat any specific number here as accurate **on the day it was recorded**.

> ON THE EXAM
>
> A course that claims to be permanently current is making a promise about a moving target. The honest version dates its claims.

Learn the capability, and check the status separately. They change at different speeds.

> An exam question describes a team needing to detect that the features their deployed model relies on have shifted over time, and asks which AWS capability provides this.
>
> The options include Model Monitor with Clarify, CloudWatch anomaly detection, Bedrock Evaluations, and building it with Evidently AI.

What should you answer?

A CloudWatch anomaly detection, since Model Monitor is closed to new customers

B Model Monitor with Clarify, which is the named AWS capability for attribution drift

C Bedrock Evaluations, which covers model quality measurement

D Build it with Evidently AI, which AWS now recommends

◆ PAUSE AND DECIDE

What should you answer?

A CloudWatch anomaly detection, since Model Monitor is closed to new customers

A correct fact applied to the wrong question. The exam asks which capability does this, and anomaly detection on custom metrics is not feature attribution drift.

B Model Monitor with Clarify, which is the named AWS capability for attribution drift

Feature attribution drift is one of Model Monitor's four documented dimensions, with Clarify supplying the attribution. Maintenance status does not change what it does.

C Bedrock Evaluations, which covers model quality measurement

Evaluates foundation models and RAG systems. It has nothing to do with feature attribution on a tabular model.

D Build it with Evidently AI, which AWS now recommends

A real replacement path, and it answers 'what should we build' rather than 'which AWS capability provides this'.

THE SIGNAL

Knowing the maintenance status is useful and it is not the answer. The exam asks what a service does, and a service in maintenance still does it.

A new team in a newly created AWS account needs to label 40,000 images for a classification project. They read a tutorial recommending SageMaker Ground Truth with a private workforce.

They cannot create a labelling job and ask whether their permissions are wrong.

What is happening, and what should they do?

A. A missing IAM permission. Add the Ground Truth managed policy

B. A Region issue. Ground Truth is not available in every Region

C. A service quota. Request a labelling job quota increase

D. Ground Truth is closed to new customers. Use another labelling route

◆ PAUSE AND DECIDE

What is happening, and what should they do?

A. A missing IAM permission. Add the Ground Truth managed policy

A reasonable first instinct, and no policy grants access to a feature closed to the account. Module 34's phantom-user lesson in reverse: permissions are not the only gate.

B. A Region issue. Ground Truth is not available in every Region

Regional availability is a real constraint for many services. It is not what a newly created account hits here.

C. A service quota. Request a labelling job quota increase

Quotas cap how much you can use. They do not open a feature that is closed to new customers.

D. Ground Truth is closed to new customers. Use another labelling route

Closed since 30 July 2026. A new account cannot sign up, and the tutorial predates the change.

THE SIGNAL

The tell is 'a newly created AWS account' plus 'a tutorial'. Tutorials outlive the availability of what they teach.

> A team has an existing production workload using Model Monitor for data quality, set up eighteen months ago. It works correctly and raises alarms as expected.
>
> A new engineer reads that Model Monitor is in maintenance mode and proposes migrating immediately to the open-source solutions.

How should you respond?

A Existing customers continue as normal. Plan a migration, and do not treat it as urgent

B Migrate immediately. Maintenance mode means the service will stop working

C Do nothing and remove the monitoring, since the capability is being withdrawn

D Migrate to Bedrock Evaluations, which replaces Model Monitor

◆ PAUSE AND DECIDE

How should you respond?

A Existing customers continue as normal. Plan a migration, and do not treat it as urgent

Maintenance mode is not sunset. AWS continues investing in security and availability, and no termination date has been announced.

B Migrate immediately. Maintenance mode means the service will stop working

Confuses maintenance with sunset. Profiler has a termination date; Model Monitor does not.

C Do nothing and remove the monitoring, since the capability is being withdrawn

Removes a working control on a production model, which is strictly worse than either migrating or waiting.

D Migrate to Bedrock Evaluations, which replaces Model Monitor

Bedrock Evaluations assesses foundation models and RAG. It does not monitor a tabular model's input distributions.

THE SIGNAL

Maintenance and sunset are different words with different consequences. Only one of them comes with a date.

> A regulated organisation needs per-prediction explanations for a credit model, in a new AWS account. Their architecture document specifies SageMaker Clarify with online explainability.
>
> Compliance requires the explanations to be versioned and auditable alongside the training runs.

What should the revised design be?

A Request an exception so Clarify can be enabled for this account

B Drop per-prediction explanations and report global feature importance instead

C Compute SHAP directly in a pipeline, logging results to a SageMaker AI MLflow App

D Use Bedrock Evaluations, which provides explainability for any model

◆ PAUSE AND DECIDE

What should the revised design be?

A Request an exception so Clarify can be enabled for this account

Closed to new customers is a lifecycle decision, not a quota or a permission. There is no exception to request.

B Drop per-prediction explanations and report global feature importance instead

Abandons a regulatory requirement because a service is unavailable. Global importance cannot explain one applicant's decision.

C Compute SHAP directly in a pipeline, logging results to a SageMaker AI MLflow App

SHAP is the engine Clarify is built on, so the capability is unchanged. MLflow supplies the versioning and lineage compliance asked for.

D Use Bedrock Evaluations, which provides explainability for any model

Evaluates foundation models and RAG systems. It does not produce SHAP attributions for a tabular credit model.

THE SIGNAL

Clarify is a packaging of public mathematics and an open-source library. Losing the packaging does not lose the capability.

What to take from this module

- **Ten features closed to new customers on 30 July 2026.** Core SageMaker AI is untouched.
- **Maintenance is not sunset.** Only Profiler has a termination date.
- **Answer the exam question with the named service.** Status is a work answer, not an exam answer.
- **Clarify's replacement is the cleanest:** public formulas and the SHAP library it was built on.
- **Check the date on everything.** A stale fact is usually a fact that was once true.

> ON THE EXAM
>
> Module 39 covers how MLA-C02 questions are constructed, and how to work through them under time pressure.