MODULE 43 · EXAM STRATEGY - ALL DOMAINS

Practice exam 3

The hardest set. Twenty questions, more judgement, and several where the right answer is to push back.

MODULE 43 of 43 DOMAIN Practice WEIGHT 20 questions

admin@sthithapragna.com

Before you start

- **Two minutes a question, timed.** This set is closest to the real thing in difficulty.
- **Several questions have no comfortable answer.** That is deliberate and it is realistic.
- Watch for options that are **true statements answering a different question**.
- Watch for facts that **used to be true**. At least one appears here.
- Score all three sets together afterwards. Sixty questions is a real sample.

> ON THE EXAM
>
> If this set feels harder than the other two, that is the intent. It is better to meet these questions here than on the day.

A data team runs ad hoc exploratory queries over 200 TB of event data in S3, perhaps forty queries a week, with no predictable pattern. They also need sub-second dashboard queries over a curated 400 GB subset, refreshed nightly, for 300 business users.

They want to minimise total cost and operational effort.

What should they build?

A Athena for both workloads

B Redshift for both workloads

C Athena over S3 for exploration, and Redshift for the curated dashboard subset

D EMR with Presto for both workloads

◆ PAUSE AND DECIDE

What should they build?

A Athena for both workloads

Right for exploration, and Athena is not built for sub-second concurrent dashboard queries across 300 users.

B Redshift for both workloads

Loading 200 TB into a warehouse to run forty unpredictable queries a week is enormous cost for very low utilisation.

C Athena over S3 for exploration, and Redshift for the curated dashboard subset

Two workloads with opposite shapes. Athena's per-query billing suits sporadic exploration; a warehouse suits small, hot, concurrent queries.

D EMR with Presto for both workloads

The most capable and the most to operate, and the scenario names operational effort as a thing to minimise.

THE SIGNAL

Two workloads are described with opposite profiles. The trap is assuming one service must serve both.

> A data scientist encodes a high-cardinality merchant identifier, with 90,000 distinct values, by replacing each merchant with the mean fraud rate observed for that merchant across the full dataset. Cross-validated performance is excellent; production performance is poor.
>
> Some merchants appear only two or three times.

What is the problem?

A One-hot encoding should have been used for 90,000 categories

B Target encoding computed on the full dataset leaks the label, especially for rare categories

C The merchant identifier should be dropped as it has no predictive value

D 90,000 categories is too many for any encoding scheme

◆ PAUSE AND DECIDE

What is the problem?

A One-hot encoding should have been used for 90,000 categories

One-hot on 90,000 values produces an unusable feature space. Target encoding is the right family of technique; the leak is the fault.

B Target encoding computed on the full dataset leaks the label, especially for rare categories

The encoded value contains the label. For a merchant appearing three times, the encoding nearly reveals that row's own outcome.

C The merchant identifier should be dropped as it has no predictive value

Merchant identity is usually highly predictive of fraud. The signal is real; the encoding is contaminated.

D 90,000 categories is too many for any encoding scheme

Hashing, embeddings and correctly-fitted target encoding all handle high cardinality. Cardinality is not the failure.

THE SIGNAL

Target encoding uses the label. Computed before the split, or without smoothing on rare categories, it hands the answer to the model.

A new team in a new AWS account must label 25,000 images for a classification project. Their architecture document, written eighteen months ago, specifies SageMaker Ground Truth with a private workforce.

They cannot create a labelling job. Their deadline is six weeks away.

What should the revised plan be?

A Request a service quota increase for labelling jobs

B Add the Ground Truth managed policy to their execution role

C Switch the project to unsupervised learning to avoid labelling

D Use a third-party labelling service or an in-house workflow, since Ground Truth is closed to new customers

◆ PAUSE AND DECIDE

What should the revised plan be?

A Request a service quota increase for labelling jobs

Quotas cap how much of a feature you can use. They do not open a feature closed to new customers.

B Add the Ground Truth managed policy to their execution role

No policy grants access to a feature the account cannot use. Permissions are not the gate here.

C Switch the project to unsupervised learning to avoid labelling

Changes the problem to fit the tooling. Image classification needs labels, and the business requirement has not changed.

D Use a third-party labelling service or an in-house workflow, since Ground Truth is closed to new customers

Ground Truth entered maintenance on 30 July 2026 and is closed to new customers. An eighteen-month-old document predates that.

THE SIGNAL

'A new AWS account' plus 'a document written eighteen months ago' is the maintenance-mode tell, and both are stated.

> A knowledge base must cover 30,000 engineering drawings and specification documents. Roughly 40 percent of the useful content is inside tables and annotated diagrams rather than prose. A text-only extraction pipeline currently loses most of it.
>
> Answers must cite the source document and page.

What should they use for ingestion?

A. Bedrock Data Automation for multimodal parsing during ingestion

B. Amazon Textract for OCR only, then standard chunking

C. Larger chunks, so tables are not split across boundaries

D. A higher numberOfResults at query time

◆ PAUSE AND DECIDE

What should they use for ingestion?

A Bedrock Data Automation for multimodal parsing during ingestion

Parses tables and diagrams into text a retriever can index, which is the only option that addresses content the text pipeline cannot see.

B Amazon Textract for OCR only, then standard chunking

OCR recovers text from images and does not interpret an annotated diagram's meaning. Better than nothing, and not the full answer.

C Larger chunks, so tables are not split across boundaries

Helps a table that is being split. It cannot help content that was never extracted in the first place.

D A higher numberOfResults at query time

Retrieves more of what was indexed. If 40 percent was never indexed, retrieving more of the remainder does not reach it.

THE SIGNAL

The failure is at extraction, not retrieval. Options C and D both tune retrieval over content that was never captured.

> A retrieval system uses 300-token chunks and returns the top five. Users ask questions whose answers span several pages of a procedure, and evaluation shows context coverage is the weakest metric while relevance is high. The generating model has a large context window and cost is not a constraint.
>
> Latency tolerance is generous.

Which TWO changes should they try?

A Reduce chunk size to 150 tokens

B Increase chunk size, so a chunk holds more of a procedure

C Increase numberOfResults above five

D Switch from semantic to keyword-only search

◆ PAUSE AND DECIDE

Which TWO changes should they try?

A. Reduce chunk size to 150 tokens

Smaller chunks hold less of each procedure, which worsens coverage. It would help precision, which is already fine.

B. Increase chunk size, so a chunk holds more of a procedure

Answers span several pages, so a larger chunk carries more of the procedure per retrieved passage. Cost and latency are explicitly not constraints.

C. Increase numberOfResults above five

Low coverage with high relevance means what comes back is right and there is not enough of it. More passages is the direct fix.

D. Switch from semantic to keyword-only search

Relevance is already high, so retrieval is finding the right material. Keyword-only would likely reduce it.

THE SIGNAL

Low coverage with high relevance means 'right material, not enough of it'. Both correct answers increase how much material reaches the model.

A European bank must build a document assistant. Legal requires that customer documents never leave the EU, that no data crosses Regional boundaries during inference, and that the strongest available prompt-attack filtering is applied.

The Standard guardrail tier requires opting in to cross-Region inference.

How should they proceed?

A Enable Standard tier, since stronger safety outranks residency

B Enable Standard tier and document the cross-Region flow as an accepted risk

C Abandon guardrails, since the tier they need is unavailable

D Use Classic tier within the EU Region, and compensate with tighter content filters and application controls

◆ PAUSE AND DECIDE

How should they proceed?

A Enable Standard tier, since stronger safety outranks residency

A legal residency requirement is not traded against a technical preference. It is a constraint, not a priority.

B Enable Standard tier and document the cross-Region flow as an accepted risk

Documenting a breach of a legal requirement does not make it acceptable. Risk acceptance does not apply to regulatory constraints.

C Abandon guardrails, since the tier they need is unavailable

Classic tier guardrails are available and useful. Having less than the maximum is not a reason to have none.

D Use Classic tier within the EU Region, and compensate with tighter content filters and application controls

Residency is a hard constraint, so the design works within it and strengthens what it can elsewhere.

THE SIGNAL

A stated legal requirement is a constraint, not a preference. Two options here trade it away, and one abandons safety entirely.

A team has a well-structured tabular dataset of 400,000 rows and 60 features, and a churn prediction problem. Their data scientist proposes a six-week programme to build and tune a custom neural network.

There is no existing baseline, and the business wants a first result within two weeks.

How should you respond?

A Run Autopilot first to establish a baseline, then decide whether custom work is justified

B Approve the neural network, since deep learning outperforms tree models on large datasets

C Reduce the feature count to speed up the custom build

D Use a foundation model with the tabular data in the prompt

◆ PAUSE AND DECIDE

How should you respond?

A Run Autopilot first to establish a baseline, then decide whether custom work is justified

Delivers a result inside two weeks, and it gives the six-week proposal something to beat. Without a baseline, nobody can say whether it was worth it.

B Approve the neural network, since deep learning outperforms tree models on large datasets

Not true for structured tabular data, where gradient-boosted trees are usually the stronger and far cheaper choice.

C Reduce the feature count to speed up the custom build

Sixty features is modest. Cutting them to save time weakens the model without addressing the timeline properly.

D Use a foundation model with the tabular data in the prompt

The wrong tool for a 400,000-row structured classification problem, and far more expensive per prediction.

THE SIGNAL

The absence of a baseline is the planted detail. A six-week build with nothing to compare against cannot be evaluated even when it finishes.

A team runs long training jobs on Spot instances with checkpointing every epoch. They also enable early stopping, configured to halt when validation loss has not improved for three epochs.

After enabling Spot, some jobs now stop much earlier than expected and produce weaker models.

What is the likely interaction?

A Spot interruptions corrupt checkpoints, so training resumes from bad weights

B Early stopping is incompatible with Managed Spot Training

C On resumption the early-stopping counter state must be restored with the checkpoint, or it may halt prematurely

D Spot instances have less memory, so fewer epochs complete

◆ PAUSE AND DECIDE

What is the likely interaction?

A Spot interruptions corrupt checkpoints, so training resumes from bad weights

SageMaker AI writes checkpoints to S3 and resumes from the last complete one. Corruption is not the described symptom.

B Early stopping is incompatible with Managed Spot Training

They are not incompatible. They interact, which is a different and more useful statement.

C On resumption the early-stopping counter state must be restored with the checkpoint, or it may halt prematurely

Early stopping keeps state: the best loss so far and epochs since improvement. A checkpoint that saves weights but not that state resumes with a reset or misleading counter.

D Spot instances have less memory, so fewer epochs complete

Spot is the same hardware at a different price. Capacity is identical.

THE SIGNAL

Checkpointing is not only about weights. Any training state the loop depends on has to survive a restart, or the loop behaves differently after one.

A model flags transactions for manual review. Reviewing costs about four pounds per case. A missed fraud costs about nine hundred pounds on average. The team has a precision-recall curve and must pick an operating threshold.

Review capacity is 1,200 cases a day; the model currently flags 400 at the proposed threshold.

How should they choose the threshold?

A Pick the threshold that maximises F1, as a balanced default

B Lower the threshold to catch more fraud, since capacity allows and a miss costs 225 times a review

C Raise the threshold to improve precision and reduce review cost

D Use the default 0.5 threshold, which is calibrated by construction

◆ PAUSE AND DECIDE

How should they choose the threshold?

A. Pick the threshold that maximises F1, as a balanced default

F1 weights precision and recall equally. The scenario gives costs that are emphatically not equal, so balance is the wrong objective.

B. Lower the threshold to catch more fraud, since capacity allows and a miss costs 225 times a review

The costs and the spare capacity are both stated. Nine hundred against four is a 225-to-1 ratio, and there are 800 unused review slots a day.

C. Raise the threshold to improve precision and reduce review cost

Saves four pounds per avoided review at the risk of a nine-hundred-pound miss. The arithmetic runs the wrong way.

D. Use the default 0.5 threshold, which is calibrated by construction

0.5 is a default, not a calibrated value, and it ignores every cost the scenario supplies.

THE SIGNAL

Every number needed is stated: both error costs and the spare capacity. The question is arithmetic dressed as judgement.

A document assistant sends a 6,000-token policy manual and a 200-token system prompt with every request, followed by a short user question. It handles 90,000 requests a day. Input token cost now dominates the bill.

The manual changes roughly monthly. Answers are short.

What is the MOST effective change?

A Reduce maxTokens to shorten the answers

B Switch to a cheaper model for all traffic

C Move the manual into a knowledge base and retrieve only relevant sections

D Enable prompt caching on the stable prefix, and consider retrieval as a second step

◆ PAUSE AND DECIDE

What is the MOST effective change?

A. Reduce maxTokens to shorten the answers

Output is already short and input dominates the bill. Trimming the smaller half cannot fix the larger one.

B. Switch to a cheaper model for all traffic

Reduces the rate and not the volume, and it changes answer quality for every request including the ones that were fine.

C. Move the manual into a knowledge base and retrieve only relevant sections

A genuinely good answer and a larger change: an index to build, sync and evaluate. Caching gets most of the saving with almost no work.

D. Enable prompt caching on the stable prefix, and consider retrieval as a second step

6,200 tokens repeat identically on every one of 90,000 daily requests, and the stable part is already at the front. That is the ideal caching shape.

THE SIGNAL

The manual is identical on every request and sits at the front of the prompt. That is precisely the shape prompt caching is built for.

A regulator asks a bank to produce, for a deployed credit model, its intended use, the data it was trained on, its performance across demographic subgroups, its known limitations, and who approved it for production.

The team has the model in a registry and the training code in a repository.

What should they produce?

A The training job's CloudWatch logs and hyperparameters

B A SHAP summary plot for the most recent 1,000 predictions

C A model card, with the registry's approval record and subgroup evaluation attached

D The pipeline definition from the repository

◆ PAUSE AND DECIDE

What should they produce?

A The training job's CloudWatch logs and hyperparameters

Evidence of how it was trained. It says nothing about intended use, limitations or who approved it.

B A SHAP summary plot for the most recent 1,000 predictions

Explains what the model relies on. It is one input to the answer, not the answer to five questions.

C A model card, with the registry's approval record and subgroup evaluation attached

A model card is the artefact designed to hold exactly these fields, and the registry supplies the approval record the regulator asked for.

D The pipeline definition from the repository

Shows how the model is built. It does not document intended use, subgroup performance or approval.

THE SIGNAL

Five separate questions are asked, and only one artefact is designed to answer all five in one document.

A retailer re-scores its entire customer base of 31 million records every Sunday night to produce segment assignments for the week ahead. The complete input file is available by 22:00 and results are needed by 06:00.

Nothing consumes an individual customer's score before the load completes.

Which option is MOST cost-effective?

A Batch transform

B Asynchronous inference, submitting 31 million requests

C A real-time endpoint with auto-scaling, driven by a Lambda loop

D Serverless inference, invoked per record

◆ PAUSE AND DECIDE

Which option is MOST cost-effective?

A Batch transform

The dataset exists in full before processing starts, nobody waits for an individual score, and instances terminate when the job ends. No idle charge at all.

B Asynchronous inference, submitting 31 million requests

Async suits requests arriving individually over time. Here the whole dataset exists up front, and 31 million queued requests adds enormous overhead.

C A real-time endpoint with auto-scaling, driven by a Lambda loop

Builds orchestration to feed an endpoint that then sits idle all week. It is the most expensive way to do a weekly batch job.

D Serverless inference, invoked per record

Per-invocation overhead 31 million times, with no benefit over a job designed for exactly this shape.

THE SIGNAL

'The complete input file is available' is the batch tell. Set 1 asked the opposite way round: requests arriving individually meant async.

A CPU-based recommendation model serves 2,000 requests a second from a fleet of x86 instances. Inference is a straightforward matrix operation with no framework-specific custom kernels. The team wants better price-performance and has a containerised build pipeline.

Latency requirements are unchanged and accuracy must not move.

What should they try first?

A Quantize the model to INT8

B Rebuild the container for Arm64 and move to Graviton instances

C Move to GPU instances

D Compile the model with the inference optimization toolkit

◆ PAUSE AND DECIDE

What should they try first?

A. Quantize the model to INT8

Quantization can move accuracy, and the scenario states accuracy must not move. It is also a bigger change than an instance swap.

B. Rebuild the container for Arm64 and move to Graviton instances

Graviton improves price-performance for CPU inference, needs no model change, and cannot affect accuracy. The containerised pipeline makes the rebuild cheap.

C. Move to GPU instances

A GPU at 2,000 small CPU-bound requests per second is likely to be under-utilised and dearer, not cheaper.

D. Compile the model with the inference optimization toolkit

Accuracy-neutral and worth doing, and it is more work than an instance family change and does not itself reduce the hourly rate.

THE SIGNAL

'Accuracy must not move' rules out quantization immediately. Graviton is the only option that cannot change model output at all.

> An endpoint serves a service whose traffic is flat overnight, then rises sharply from 08:00 to a plateau by 09:00. Target tracking auto-scaling is configured on invocations per instance. Every morning users see elevated latency and some throttling for about twelve minutes.
>
> Scaling eventually settles correctly and the endpoint is not undersized at plateau.

What should they change?

A Increase the target value so scaling triggers later

B Switch to a larger instance type and remove auto-scaling

C Lengthen the scale-in cooldown

D Add a scheduled scaling action that raises minimum capacity before 08:00

◆ PAUSE AND DECIDE

What should they change?

A Increase the target value so scaling triggers later

Triggering later makes the morning gap worse. The problem is that scaling reacts too slowly, not too eagerly.

B Switch to a larger instance type and remove auto-scaling

Sizes for the peak all day, which pays for the overnight hours at plateau capacity. It solves latency by abandoning the saving.

C Lengthen the scale-in cooldown

Slows scaling down, which helps after a dip. The failure is on the way up, in the morning.

D Add a scheduled scaling action that raises minimum capacity before 08:00

Reactive scaling cannot pre-empt a known daily ramp. A schedule provisions capacity before the traffic arrives, and target tracking handles the rest.

THE SIGNAL

The ramp is predictable and happens at the same time daily. Reactive scaling always lags a ramp; only a schedule can lead it.

A compliance knowledge base indexes policy documents. A policy was withdrawn on Monday and deleted from the source S3 bucket. On Wednesday a user received an answer citing that withdrawn policy as current.

The ingestion job runs on the first of each month.

What is the problem, and what should change?

A The generating model hallucinated the policy. Enable contextual grounding

B Chunk overlap caused duplication. Reduce the overlap percentage

C The index still holds the deleted document until the next sync. Increase sync frequency and alarm on sync recency

D The user's query was ambiguous. Add query rewriting

◆ PAUSE AND DECIDE

What is the problem, and what should change?

A. The generating model hallucinated the policy. Enable contextual grounding

The model was faithful to what it retrieved. Grounding checks the answer against the retrieved passages, which genuinely contained the withdrawn policy.

B. Chunk overlap caused duplication. Reduce the overlap percentage

Overlap affects how passages are split, not whether deleted source documents remain indexed.

C. The index still holds the deleted document until the next sync. Increase sync frequency and alarm on sync recency

Deleting from the source does not remove it from the vector index. The document stayed retrievable, and a monthly sync leaves a window of up to a month.

D. The user's query was ambiguous. Add query rewriting

The query found exactly the document it asked about. The document should no longer have been there.

THE SIGNAL

This is a disclosure risk, not a quality problem. Everything worked correctly over an index that was out of date.

> A team wants five specialist agents, each owning a domain, with a coordinator that routes each request to the right specialist and combines their outputs. They plan to use AgentCore and expect routing between agents to be a configuration option.
>
> They have six weeks and no framework engineering capacity.

How should you advise them?

A Proceed as planned. AgentCore provides routing-mode multi-agent collaboration out of the box

B Warn them: AWS describes multi-agent collaboration as limited, and full collaboration needs custom framework code

C Use Agents Classic instead, which supports multi-agent collaboration natively

D Build all five specialisms into a single agent with five tools

◆ PAUSE AND DECIDE

How should you advise them?

A Proceed as planned. AgentCore provides routing-mode multi-agent collaboration out of the box

AWS's own documentation calls multi-agent collaboration limited. A supervisor pattern via agents-as-MCP-tools works; routing mode is not straightforward.

B Warn them: AWS describes multi-agent collaboration as limited, and full collaboration needs custom framework code

The honest answer. The supervisor pattern is achievable, and the configuration-driven routing they are expecting is not available, which their timeline depends on.

C Use Agents Classic instead, which supports multi-agent collaboration natively

Agents Classic is in maintenance mode and is not the place to start new work, whatever its feature set.

D Build all five specialisms into a single agent with five tools

A reasonable simplification worth considering, and it is not what they asked for and does not address the false expectation.

THE SIGNAL

The question is about an expectation that will not hold. Correcting it inside six weeks is worth more than any architecture suggestion.

> Users of a streaming assistant report it 'sits there doing nothing' before replying, though the dashboard shows invocation latency well within target. The client is a vendor application the team cannot modify or instrument.
>
> They need an alarm and an SLA on what users actually experience.

What should they use?

A The TimeToFirstToken CloudWatch metric

B Invocation latency, at the 99th percentile

C A synthetic canary calling the model every minute

D Model invocation logs, timing request and response entries

◆ PAUSE AND DECIDE

What should they use?

A The TimeToFirstToken CloudWatch metric

Measures request-sent to first-token-received on the streaming APIs, and AWS states it needs no client-side instrumentation. Alarms and SLA baselines are its stated purpose.

B Invocation latency, at the 99th percentile

Measures total call duration, which the scenario says is within target. For a streaming response that is not what the user perceives.

C A synthetic canary calling the model every minute

Measures synthetic traffic and builds infrastructure to reproduce a metric that already exists.

D Model invocation logs, timing request and response entries

Logs record the invocation, not the moment the first token arrived, and it means building metrics from log parsing.

THE SIGNAL

'Cannot modify or instrument the client' is the constraint TimeToFirstToken was released for, in March 2026.

A pilot knowledge base serves around 350 questions a month across a 1,500-document corpus, on OpenSearch Serverless. Finance queries a line of several hundred pounds a month that does not move when usage does. Token costs are a few pounds.

The pilot may or may not be extended.

What is happening?

A Provisioned Throughput was purchased in error

B Embedding every query is the dominant cost

C The OpenSearch Serverless minimum applies at zero traffic. Consider S3 Vectors or a dev-test collection

D Retrieved chunks are inflating input tokens. Reduce numberOfResults

◆ PAUSE AND DECIDE

What is happening?

A Provisioned Throughput was purchased in error

Would also produce a fixed charge, and nothing suggests it was bought. The knowledge base's own store explains it without inventing a purchase.

B Embedding every query is the dominant cost

Query embedding is a couple of hundred input tokens. At 350 questions a month it is negligible, and it would vary with usage.

C The OpenSearch Serverless minimum applies at zero traffic. Consider S3 Vectors or a dev-test collection

A two-OCU minimum produces a continuous floor regardless of queries, which matches a large charge that does not move with usage.

D Retrieved chunks are inflating input tokens. Reduce numberOfResults

A real lever, and the scenario states token costs are a few pounds. Halving a few pounds does not address several hundred.

THE SIGNAL

'Does not move when usage does' is the diagnostic phrase. Only a fixed floor behaves that way, and at 350 questions it dominates everything.

A benefits assistant must never state that a claimant qualifies for a payment when the documented rules say otherwise. Eligibility depends on age bands, contribution history and residency duration, all written down as formal criteria.

When a response is rejected the team needs an explanation of which rule it violated, for the audit trail.

Which safeguard fits?

A Content filters at HIGH strength on the output

B Denied topics covering eligibility statements

C Contextual grounding checks, with the rules document as the grounding source

D Automated Reasoning checks, with the eligibility rules authored as a policy

◆ PAUSE AND DECIDE

Which safeguard fits?

A. Content filters at HIGH strength on the output

Content filters address harmful content. A wrong eligibility statement is incorrect, not harmful by category, and passes every filter.

B. Denied topics covering eligibility statements

Blocks the assistant from discussing eligibility, which is the one thing it exists to do.

C. Contextual grounding checks, with the rules document as the grounding source

Probabilistic scoring of support. A wrong eligibility claim phrased in the source's own language can score as well supported, and it cannot explain which rule was broken.

D. Automated Reasoning checks, with the eligibility rules authored as a policy

Formal logic against authored rules, and it can explain why a claim is wrong. Age bands and residency durations are exactly the logically verifiable case.

THE SIGNAL

Two signals: rules that are genuinely formal, and a requirement for an explanation rather than a score.

> A retraining pipeline registers a new model whenever training accuracy exceeds 0.90. Every weekly run passes. Production accuracy has fallen for four consecutive months while the pipeline reports success each week.
>
> The team proposes raising the threshold to 0.95.

Which TWO changes should you recommend?

A. Gate on a metric computed on a held-out evaluation set, not on training accuracy

B. Raise the training accuracy threshold to 0.95 as proposed

C. Trigger retraining on drift or new data rather than on a weekly schedule

D. Remove the ConditionStep, since it is not preventing anything

◆ PAUSE AND DECIDE

Which TWO changes should you recommend?

A Gate on a metric computed on a held-out evaluation set, not on training accuracy

Training accuracy measures memorisation, so any model passes. A held-out metric is the only version of this gate that can ever fail.

B Raise the training accuracy threshold to 0.95 as proposed

Raising a broken gate makes it slightly harder to pass and no more meaningful. An overfitted model clears 0.95 as easily as 0.90.

C Trigger retraining on drift or new data rather than on a weekly schedule

A weekly schedule retrains when nothing changed and misses degradation between runs. Four months of decline with weekly retraining is that failure.

D Remove the ConditionStep, since it is not preventing anything

Correct that it prevents nothing, and the fix is to make it meaningful rather than to delete the only quality gate in the pipeline.

THE SIGNAL

Two faults are described: a gate measuring the wrong thing, and a trigger unrelated to whether retraining was warranted. Each needs its own fix.

Scoring sixty questions

- **Combine all three scoresheets.** Sixty questions is a real sample; twenty is not.
- **Look for the domain with the most misses**, then the task area within it.
- Count knowledge errors against reading errors across all sixty. The ratio tells you what to do next.
- Mostly reading errors means practise the working order, not more content.
- Mostly knowledge errors in one domain means rewatch three or four modules, not the whole course.

> ON THE EXAM
>
> This is a diagnostic, not a predictor. It tells you where the next few hours should go.

Where to go from here

- **Module 37** is the highest-value revision here. Twenty-four pair tables, pure discrimination.
- **Module 38** the day before, so the maintenance-mode answers are fresh.
- **Module 39** for the format, and check whether you booked the beta or the general release.
- **Do not learn anything new in the last two days.** Consolidate what is there.
- Every specific figure in this course is accurate as recorded. Check anything you will act on at work.

> ON THE EXAM
>
> Forty-four modules, four domains, and one habit worth more than all of them: read the ask, then the constraints, then the options.