

Escalation, Confidence and Human Review

Knowing when to stop, handing over usefully, and the errors a confidence threshold will never catch.

The goal is not to escalate **less**. It is to escalate at the **right moment**, with the right information.

What should always escalate

- ◆ **It falls outside policy.** The agent should not be deciding whether to make an exception.
- ◆ **The action is irreversible or high-value.** Refunds above a limit, deletions, anything that cannot be undone.
- ◆ **The customer is distressed or explicitly asks for a person.** That request is itself the criterion.
- ◆ **The agent has failed twice on the same thing.** A third attempt is rarely the one that works.
- ◆ **It lacks the capability.** No tool exists for what is needed.
- ◆ **And a catch-all: anything not resolved.** Module 4's orchestration check, doing its job.

ON THE EXAM

The first five are cases you can name. The sixth catches the ones you did not think of.

Criteria that generalise

- ◆ **Describe the property, not the case.** "Requests an exception to stated policy", not "asks about the 30-day rule".
- ◆ **A rule written from one incident catches that incident.** The next one arrives in a different shape.
- ◆ **Group by why it needs a human,** not by what the customer said.
- ◆ **Test a criterion by asking what it misses,** not by checking it catches the example you started from.

ON THE EXAM

This is module 12's lesson again. "Most effective" almost always means the answer that generalises.

Retry, or escalate?

RETRY

Something transient went wrong

- A timeout, a rate limit, a 503
- The next attempt has a real chance
- Nothing about the request was wrong
- Cap the attempts, then escalate anyway

ESCALATE

The situation needs judgement

- Outside policy, or beyond the agent's authority
- The same failure has happened twice
- The customer asked for a person
- A retry would repeat, not resolve

THE TIEBREAKER

Deterministic failures do not become resolvable on the third attempt. Module 6 said this about errors; it is true about cases too.

An escalation that arrives as a raw transcript is a **ticket**. An escalation with structure is a **handover**.

What the human actually needs

- ◆ **Why it escalated**, in one line. Which criterion fired.
- ◆ **What the customer wants**, stated plainly, separate from how they said it.
- ◆ **What has already been established**. Order found, account verified, policy checked.
- ◆ **What has already been said to them**, so the human does not contradict it.
- ◆ **What remains to be decided**, and the options if there are any.
- ◆ **The transcript as an attachment, not as the handover**. Available, not required.

ON THE EXAM

Everything above is derivable from the session. Nobody derives it unless the escalation is designed to.

A confidence score is not a probability. It is a **signal you have to calibrate** before it means anything.

Using confidence properly

- ◆ Ask for it alongside every output, per record or per field.
- ◆ Filter in your code, not in the prompt. Module 12's rule; the model filtering hides what it dropped.
- ◆ Compute a review flag from the score, so routing is a rule rather than a judgement.
- ◆ Calibrate the threshold against a labelled set before you trust any particular number.

ON THE EXAM

The score is an input to your routing logic. It is not the routing logic.

Calibration, and what the threshold trades

- ◆ **Take a few hundred outputs and label them correct or not.** Then see where the errors actually sit.
- ◆ **The threshold trades reviewer load against misses.** Lower it and you review more; raise it and you miss more.
- ◆ **Choosing 0.8 because it sounds right is guessing** with a number attached.
- ◆ **Recalibrate when the inputs change.** A threshold tuned on last year's documents is a threshold tuned on last year.

ON THE EXAM

Calibration is unglamorous and it is the difference between a confidence score and a decoration.

Reviewing low-confidence output catches uncertain errors. It **never** catches confident ones.

Two review queues, not one

BELOW THE THRESHOLD

Everything uncertain

- Catches errors the model hesitated on
- Volume is manageable and predictable
- The obvious queue, and necessary
- Blind to anything scored highly

SAMPLED FROM ABOVE IT

A slice of the confident output

- Catches errors the model was sure about
- A small percentage is enough to detect a pattern
- Also tells you whether the threshold is right
- The queue almost nobody builds

THE TIEBREAKER

The first queue protects the cases you already suspected. The second tells you what you are wrong about.

Sampling that finds things

- ◆ **Stratify.** Sample across document types, sources and confidence bands, not uniformly at random.
- ◆ **A small percentage is enough** to detect a systematic problem, though not to measure it precisely.
- ◆ **Feed what you find back into the criteria,** the schema or the threshold.
- ◆ **Sampling only the easy category tells you the easy category is fine.**

ON THE EXAM

You are sampling to discover unknown failure modes, not to produce a headline accuracy number.

When a dependency is down

- ◆ **Say what is unavailable and what it means**, in terms the person cares about.
- ◆ **Offer what you can still do**. Partial help beats a blank refusal.
- ◆ **Capture the request and commit to a follow-up**, so the interaction still ends somewhere.
- ◆ **Never answer from memory as a substitute for the tool**. A plausible invented answer is worse than an outage.

ON THE EXAM

Degrading gracefully still means ending resolved or escalated. An outage is not permission to end nowhere.

Reading the stem: escalation and review questions

WHAT THE STEM SAYS	WHAT IT WANTS
"the human has to read the whole transcript"	A structured handover, not a ticket
"errors are getting through despite review"	Sample the high-confidence output too
"the agent was confident and wrong"	A confidence threshold cannot catch this
"we picked 0.8 as the threshold"	Calibrate against a labelled set
"it keeps trying the same thing"	Cap the retries and escalate
"the backend was down and it made something up"	Explain, offer, capture. Never substitute memory for a tool.

ON THE EXAM

The second and third rows are the same insight. Confidence-based review has a blind spot, by construction.

Module 17 in five lines

- ◆ Escalate at the right moment with the right information. The rate is not the metric.
- ◆ Write criteria as properties, not as the incident that prompted them.
- ◆ A handover names why, what is established, what was said, and what remains.
- ◆ Low-confidence review is blind to confident errors. Sample above the threshold too.
- ◆ Calibrate the threshold on labelled data, and recalibrate when the inputs change.

ON THE EXAM

That closes the domain modules. Six scenarios, then the three cross-cutting modules.

An extraction pipeline routes any record scoring below 0.75 confidence to human review. Reviewers correct roughly 30% of what they see, so the queue is clearly earning its place.

An audit of records that were never reviewed finds an error rate of about 4%, concentrated in one document format the model handles confidently and incorrectly.

The team is considering lowering the threshold.

What is the most effective change?

- ☐ A Lower the threshold to 0.85 so more records are reviewed
- ☐ B Add a sampled review of records scoring above the threshold, stratified by document format
- ☐ C Retrain reviewers to be more thorough on the records they already see
- ☐ D Ask the model to be more conservative when assigning confidence

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Lower the threshold
The errors are confident ones. Raising the bar reviews more uncertain records and still misses these.
- B

 Sample above the threshold
Confident errors are only findable by looking at confident output. Stratifying by format targets the known cluster.
- C

 Better reviewers
Those records never reach a reviewer. Thoroughness on the wrong set changes nothing.
- D

 More conservative scoring
It is not hesitating. It has misunderstood a format, which does not feel uncertain from the inside.

THE SIGNAL

"Confidently and incorrectly" rules out every answer that adjusts the confidence mechanism. That is the blind spot, not a tuning problem.

A support agent escalates about 12% of conversations. Escalations arrive in the human queue as the full conversation transcript and nothing else.

Agents report that handling an escalated case takes almost as long as handling it from scratch. In several cases a human has contradicted something the agent already told the customer.

The escalation rate itself is considered appropriate.

What should change?

- | | |
|---|---|
| <p>A Reduce the escalation rate so agents handle fewer cases</p> | <p>B Replace the transcript with a structured handover: reason, established facts, what was said, what remains</p> |
| <p>C Have the agent write a one-paragraph summary at the top of the transcript</p> | <p>D Train human agents to read transcripts more efficiently</p> |

◆ PAUSE AND DECIDE

What should change?

A Escalate less

The stem says the rate is appropriate. This fixes the symptom by handling fewer cases correctly.

C A summary paragraph

A real improvement, and prose is not a field. "What was said to the customer" can still be omitted.

B A structured handover

Named fields, including what was already said, which is the specific cause of the contradictions.

D Read faster

Puts the cost on humans permanently rather than removing it once.

THE SIGNAL

The contradictions are the tell. They happen because "what was already said" is not a named field anywhere.

A returns agent has clear policy in its system prompt. After a customer explained a family bereavement, the agent granted an out-of-policy refund outside the returns window.

The team writes a new escalation rule: "escalate any request mentioning bereavement".

Two months later, the agent grants an out-of-policy exception for a customer who was in hospital.

What is the underlying problem with their approach?

- ☐ A The rule should have listed more circumstances, including hospitalisation
- ☐ B The criterion was written from an instance. It should be "any request for an exception to stated policy"
- ☐ C The system prompt needs stronger wording about the returns window
- ☐ D The agent should not handle emotionally sensitive conversations at all

◆ PAUSE AND DECIDE

What is the underlying problem with their approach?

- A

 List more circumstances
The same mistake a third time. There is always another circumstance nobody listed.
- B

 Write the property, not the instance
The property is "an exception is being requested". The reason for the request is irrelevant to whether a human decides.
- C

 Stronger policy wording
The policy was already clear. Persuasion failing against a sympathetic case is the situation, not the fix.
- D

 Exclude sensitive conversations
Over-broad, and it still needs a definition of sensitive, which is the same problem again.

THE SIGNAL

Both incidents share one property: a customer asked for an exception. Neither shares the circumstance. Write the rule at the shared level.

An agent looks up order status through a warehouse API. During a two-hour outage the API returned errors on every call.

Reviewing the outage, the team finds the agent told several customers their orders had shipped. Those statements were plausible and, in four cases, wrong.

The system prompt tells the agent to be helpful and to keep customers informed.

What is the most important change?

- ☐ A Add an instruction never to guess order status
- ☐ B On tool failure, state what is unavailable, offer what can still be done, and capture the request for follow-up
- ☐ C Retry the API more aggressively during outages
- ☐ D Cache recent order statuses so answers are available during an outage

◆ PAUSE AND DECIDE

What is the most important change?

A Never guess

A prohibition with no alternative. The agent still has a customer to answer and a prompt telling it to be helpful.

C Retry harder

A two-hour outage does not respond to retries. It also delays the customer without helping.

B Explain, offer, capture

Gives the agent a correct thing to do when the tool fails, so the interaction still ends somewhere.

D Cache statuses

Useful, and a cached status during an outage is stale by definition. Confidently wrong, from a different source.

THE SIGNAL

The agent had no acceptable behaviour available. Telling it not to do the wrong thing, without supplying a right thing, rarely works.

A team is launching a document classification system. They plan to route anything below 0.8 confidence to human review, having picked 0.8 as a sensible-looking figure.

They have 400 labelled documents from a pilot, with known correct classifications, which they have not used.

Reviewer capacity is limited and they are unsure whether 0.8 will overwhelm it.

What should they do first?

- ☐ A Launch with 0.8 and adjust based on reviewer feedback
- ☐ B Run the pilot set through the system and plot errors against confidence to choose the threshold
- ☐ C Set the threshold to whatever produces a review volume matching reviewer capacity
- ☐ D Route everything to review for the first month

◆ PAUSE AND DECIDE

What should they do first?

A Launch and adjust

Reviewer feedback tells you about volume, not about what is slipping past. The misses stay invisible.

C Fit to capacity

Sets a quality control by staffing level. It may sit far above or below where the errors actually are.

B Use the labelled set

They already have the data that turns a guess into a decision. Errors against confidence shows what each threshold costs.

D Review everything

Safe and unsustainable, and it never produces the threshold. It defers the decision rather than informing it.

THE SIGNAL

"400 labelled documents, which they have not used" is the planted detail. The answer is sitting on their disk.

A billing agent has clear escalation criteria and a structured handover format. In production, 3% of conversations end with a final message and no escalation, leaving the customer unresolved.

Investigation shows these are cases where the agent believed it had resolved the issue.

The team has already improved the escalation criteria twice.

What closes the gap?

- ☐ A Improve the escalation criteria a third time
- ☐ B An orchestration check after each conversation that escalates anything not marked resolved
- ☐ C Lower the confidence threshold at which the agent escalates
- ☐ D Add examples of unresolved conversations to the system prompt

◆ PAUSE AND DECIDE

What closes the gap?

- A

 Better criteria again
Criteria are evaluated by the agent. It believed it had resolved these, so no criterion fires.
- B

 A post-run check
Runs whatever the agent concluded. It does not ask the agent to recognise its own failure.
- C

 Lower the threshold
Confidence-based escalation needs the agent to be uncertain. It was confident.
- D

 Examples in the prompt
Context. It improves the rate on a problem defined by the agent not noticing.

THE SIGNAL

Every option except B asks the agent to detect a failure it does not know it is having. That is the whole shape of this question.