

MODULE 18 · EXAM STRATEGY - ALL DOMAINS

Telling Near-Identical Answers Apart

Twenty-eight pairs, one discriminator each. Almost no new content, which is the entire point.

MODULE 18 of 20 PAIRS 28 NEW MATERIAL almost none

Understanding a distinction and **recognising** it under time pressure are different skills.

How recognition fails

- ◆ **You recognise the topic but not the pair.** "This is about sessions" is not yet an answer.
- ◆ **You recall both halves and not the discriminator.** The one sentence that separates them is the thing to memorise.
- ◆ **You pick the more powerful mechanism.** Hooks over permissions, agents over workflows, because they sound more capable.
- ◆ **You accept a reversed answer.** Grep for filenames reads perfectly plausibly at speed.

ON THE EXAM

The third and fourth failures cost the most marks, because both feel like competence rather than error.

How to use the next six slides

- ◆ **Read the pair, then cover the right-hand column** and say the discriminator out loud.
- ◆ **If you cannot, note the module number** and go back to it before the practice exams.
- ◆ **The discriminator is the memorisable unit**, not the definitions.
- ◆ **Twenty-eight pairs.** Expect to be shaky on three or four, which is exactly what this is for.

ON THE EXAM

Being shaky on a pair here is useful information. Being shaky on it in the exam is a mark.

Architecture and orchestration

PAIR	THE DISCRIMINATOR
Workflow / agent	Could you draw the flowchart in advance?
Chaining / dynamic decomposition	Same steps every time, or do they depend on what step two returns?
Parallelisation / orchestrator-workers	Could you write the subtask list before seeing the input?
Subagent / more context	Does the work need its own window, or does the parent already hold it?
Fork / subagent	Fork branches a decision. A subagent isolates context.
Terminator / guardrail	end_turn is completion. Hitting a cap is failure.

ON THE EXAM

Modules 1 to 4. The last row decides at least one question on most papers.

Sessions, state and context

PAIR	THE DISCRIMINATOR
Session / checkpointing	Sessions persist the conversation. Checkpoints persist the files.
continue / resume / fork	Most recent / a named one / a branch. Resume never creates.
Resume / fresh with summary	What changed while you were gone? Nothing means resume.
Refresh in place / start again	Can you list what changed? If not, that uncertainty is the answer.
Clearing / compaction	Clearing loses content, keeps it re-fetchable. Compaction keeps substance, loses wording.
In context / on disk	Keep what the next decision needs. Write down what a later one might.

ON THE EXAM

Modules 4, 15 and 16. Row one is the pair people get wrong most confidently.

Tools and MCP

PAIR	THE DISCRIMINATOR
Split / merge a tool	One tool, three behaviours: split. Success is always three calls: merge.
Resource / tool	Needs to know it before deciding: resource. Needs to do it: tool.
isError / protocol error	The tool ran and failed, or the request never reached it.
Empty result / error	A search matching nothing succeeded.
Glob / Grep	Glob is for names. Grep is for contents.
Adopt / build a server	Adopt for a common system. Build only for your own.

ON THE EXAM

Modules 5 to 7, plus Glob and Grep from 16. Row five is the one most often tested reversed.

Configuration and enforcement

PAIR	THE DISCRIMINATOR
CLAUDE.md / settings.json	Context versus enforcement. Only settings has a defined winner.
Import / path-scoped rule	Imports always load. Path scoping is what reduces context.
Permission / hook	Permissions are flat. Hooks can evaluate a condition.
PreToolUse / PostToolUse	Prevention versus reaction. PostToolUse cannot undo.
disable-model-invocation / user-invocable: false	Removes Claude as invoker / removes you as invoker.
--append-system-prompt / --system-prompt	Adds to the default / replaces it, including what you relied on.

ON THE EXAM

Modules 8, 9 and 11. Row five is two similarly named fields with opposite effects.

Prompting and structured output

PAIR	THE DISCRIMINATOR
Criteria / examples	Easier to define, or easier to demonstrate?
One prompt / two prompts	"We improved X and Y got worse" means split it.
Filter in prompt / downstream	A model asked to filter hides what it discarded.
Required / optional field	Required plus unknowable equals fabricated.
Blind / informed retry	Identical failures mean send the error back, not the request.
Batch / synchronous	Is anybody waiting for this answer?

ON THE EXAM

Modules 12 to 14. Row two is the clearest single signal in Domain 4.

Reliability and review

PAIR	THE DISCRIMINATOR
Retry / escalate	Transient, or needs judgement? Deterministic failures do not resolve.
Low-confidence review / sampling above	Confident errors are invisible to a threshold.
Prompt guarantee / code guarantee	"Guarantee" eliminates every prompt-shaped option.
Same-session / fresh-session review	The author holds the misunderstanding that caused the bug.
Plan mode / asking for a plan	A boundary, or a request.

ON THE EXAM

Modules 6, 10, 11 and 17. Row three is the highest-value line in the whole course.

Three questions resolve most pairs: **who decides, can it refuse, and what changed.**

The three questions, applied

QUESTION	PAIRS IT RESOLVES
Who decides?	Workflow/agent, resource/tool, prompt filter/downstream, the two invocation fields
Can it refuse?	CLAUDE.md/settings, permission/hook, plan mode/asking, prompt/code guarantee
What changed?	Resume/fresh, refresh/restart, retry/escalate, blind/informed retry

ON THE EXAM

When two options both look right, work out which of these three the stem is really asking about.

Reversed answers

- ◆ An option with the two halves swapped reads **plausibly** when you are moving fast.
- ◆ **Glob** for contents, **Grep** for filenames. Batch the urgent ones. Local scope for the team.
- ◆ Say the discriminator before you read the options, so you are matching rather than evaluating.
- ◆ If an option feels obviously right, check the direction once. One second, and it catches this whole class.

ON THE EXAM

You cannot spot a reversal by reading harder. You spot it by knowing the answer before you look.

Module 18 in five lines

- ◆ The discriminator is the memorisable unit. The definitions you already have.
- ◆ Who decides, can it refuse, what changed. Three questions, most of the pairs.
- ◆ The more powerful mechanism is not the default answer.
- ◆ Know the answer before reading the options, so reversals cannot land.
- ◆ A pair you cannot recall now is a module to rewatch before the practice exams.

ON THE EXAM

Ten rapid scenarios next. Short stems, four seconds each. Say the pair, then the discriminator, then the letter.

An agent edited three files during a session, then the session was resumed after the engineer reverted those files on disk.

The engineer expects the agent to know the files are back to their original state.

Is that expectation correct?

- ☐ A Yes. Resuming restores the session's view of the filesystem
- ☐ B No. Sessions persist the conversation, not the files. It still believes it edited them
- ☐ C Yes, if checkpointing is enabled for the session
- ☐ D No, unless the engineer runs /context first

◆ PAUSE AND DECIDE

Is that expectation correct?

- A

 Resuming restores the filesystem
Conflates the two mechanisms. A session has no view of the filesystem to restore.
- B

 Conversation, not files
The conversation records the edits. Reverting the files does not revert the record of editing them.
- C

 If checkpointing is on
Checkpointing could revert the files. It does not update the conversation either.
- D

 Run /context first
Reports what memory files loaded. Nothing to do with filesystem state.

THE PAIR

Pair: session versus checkpointing. Sessions persist the conversation; checkpoints persist the files. Module 4.

A team needs to find every file whose name matches `report_*.py`, and separately every file containing the string `legacy_auth`.

Which tools?

- ☐ A Grep for the filenames, Glob for the string
- ☐ B Glob for the filenames, Grep for the string
- ☐ C Glob for both
- ☐ D Read the directory, then scan each file

◆ PAUSE AND DECIDE

Which tools?

- A

 Grep names, Glob string
The reversal. Reads plausibly at speed, which is exactly why it is offered.
- B

 Glob names, Grep string
Names are paths. Strings are contents.
- C

 Glob for both
Glob cannot look inside a file.
- D

 Read everything
Loads every file to answer what two searches answer for nothing.

THE PAIR

Pair: Glob versus Grep. Glob is for names, Grep is for contents. Module 16.

A rule must hold whenever any tool writes to a path under /production, and it must hold whether or not the agent was told about it.

Where does the rule go?

- ☐ A The system prompt, stated firmly
- ☐ B CLAUDE.md, so every session loads it
- ☐ C A PreToolUse hook that inspects the path and blocks
- ☐ D A PostToolUse hook that alerts on production writes

◆ PAUSE AND DECIDE

Where does the rule go?

- A

 System prompt
Context. "Whether or not it was told" rules it out by definition.
- B

 CLAUDE.md
The same thing by a different route. Also context.
- C

 PreToolUse block
Fires on the event, inspects the argument, and refuses. Conditional, so a permission cannot express it.
- D

 PostToolUse alert
The write already happened. An alert about an outage is not prevention.

THE PAIR

Two pairs at once: context versus enforcement, and PreToolUse versus PostToolUse. Modules 6, 8 and 9.

A review prompt covers accessibility and performance. Emphasising accessibility reduced performance findings; emphasising performance reduced accessibility findings. Total findings stayed flat.

What does this indicate?

- ☐ A max_tokens is truncating the response
- ☐ B One prompt is doing two jobs. Split it
- ☐ C The model needs a larger context window
- ☐ D The criteria need to be more explicit

◆ PAUSE AND DECIDE

What does this indicate?

- A

 Truncation
Total findings stayed flat. Truncation would reduce the total, not shift the mix.
- C

 Bigger window
Nothing here is about capacity.
- B

 Two jobs, one prompt
Improving one at the cost of the other is the signature. Attention is being traded, not lost.
- D

 More explicit criteria
Would help coverage within a category. It does not stop two categories competing.

THE PAIR

Pair: one prompt versus two. "We improved X and Y got worse" is the signal. Module 12.

An agent needs to know the valid status codes and the database schema before it can plan any query. Both change about twice a year.

How should this be exposed?

- ☐ A A tool the agent calls to fetch the schema
- ☐ B An MCP resource holding the schema and status codes
- ☐ C An MCP prompt the user invokes
- ☐ D A subagent that looks the schema up

◆ PAUSE AND DECIDE

How should this be exposed?

- A

 A tool
Costs a call and a turn to learn something it needed before it could decide anything.
- C

 A prompt
Prompts are user-invoked templates. Nobody is asking for a schema.
- B

 A resource
Stable, referenced constantly, required before deciding. The resource signature exactly.
- D

 A subagent
A round trip and a context load to fetch something small and stable.

THE PAIR

Pair: resource versus tool. Needs to know it before deciding means resource. Module 7.

A validation failure returns the same error on all three retries. The failing field is a date in the wrong format.

What should change?

- ☐ A Add exponential backoff between retries
- ☐ B Increase the retry count
- ☐ C Send the validation error back with the retry
- ☐ D Loosen the schema to accept the format

◆ PAUSE AND DECIDE

What should change?

- A

 Backoff
For transient failures. Identical repeats mean deterministic.
- B

 More retries
More identical failures at more cost.
- C

 Send the error back
*The model learns what was wrong and corrects that field.
The loop converges.*
- D

 Loosen the schema
Accepts bad data to make the error disappear. A loud failure becomes a quiet one.

THE PAIR

Pair: blind versus informed retry. "Identical failures" means deterministic, so no retry policy helps. Modules 6 and 13.

A nightly reconciliation job of 30,000 records is reviewed by a team each morning. Finance wants the bill reduced.

What is the first lever?

- ☐ A Run the records in parallel
- ☐ B Send them through the Batch API
- ☐ C Reduce max_tokens on each response
- ☐ D Add more subagents to share the work

◆ PAUSE AND DECIDE

What is the first lever?

- A

 Parallel
Buys wall-clock time. Costs exactly the same.
- B

 Batch
Reviewed each morning means nobody is waiting. Half price for a scheduling change.
- C

 Shorter responses
A small share of the bill on a job like this.
- D

 More subagents
Parallelism again, with extra context loads. More expensive, not less.

THE PAIR

Pair: batch versus synchronous, and the parallelism trap. Parallelism buys time; batch buys money. Modules 3 and 14.

A skill deploys to production. Claude has begun invoking it on its own when users ask hypothetical questions about deployment.

The team still wants to run it manually.

Which field?

- ☐ A user-invocable: false
- ☐ B disable-model-invocation: true
- ☐ C allowed-tools, restricted to read tools
- ☐ D context: fork

◆ PAUSE AND DECIDE

Which field?

- A

 user-invocable: false
The reversal. Removes the slash command and leaves Claude as the only invoker.
- B

 disable-model-invocation
Removes Claude as an invoker; the slash command stays. Both requirements, one field.
- C

 allowed-tools
Restricts what the skill may do, not who may start it.
- D

 context: fork
Isolates context. Nothing to do with invocation.

THE PAIR

Pair: the two invocation fields. Similar names, opposite effects. Read which invoker the stem wants removed. Module 9.

An extraction system routes anything below 0.7 confidence to review. An audit finds a cluster of errors in records scoring above 0.9.

What is the fix?

- ☐ A Raise the threshold to 0.9
- ☐ B Sample records above the threshold for review
- ☐ C Ask the model to score more conservatively
- ☐ D Increase reviewer capacity

◆ PAUSE AND DECIDE

What is the fix?

- A

 Raise to 0.9
The errors score above 0.9. Raising the bar reviews more and still misses them.
- B

 Sample above
Confident errors are only findable by looking at confident output.
- C

 Score conservatively
It is not hesitating. It misunderstood something, which does not feel uncertain.
- D

 More reviewers
More capacity aimed at the same blind spot.

THE PAIR

Pair: threshold review versus sampling above it. A threshold is blind to confident errors by construction. Module 17.

Every pull request review runs the same three independent checks in the same order. Reviewers cannot tell whether a missing finding means clean code or a skipped step.

What is the right structure?

- ☐ A One agent with all three checklists and a stronger prompt
- ☐ B Prompt chaining: three sequential calls, each with one check
- ☐ C One agent with three tools, one per check
- ☐ D Three subagents run by a coordinator

◆ PAUSE AND DECIDE

What is the right structure?

- A

 One agent, stronger prompt
Reduces skipping without eliminating it. Reviewers need certainty, not a better rate.
- B

 Prompt chaining
Same steps, same order, independent. Your code runs all three, so a missing finding provably means clean code.
- C

 Three tools
Better, and the agent still owns the path, so it can still skip one.
- D

 Three subagents
Parallelism where the requirement is certainty. The coordinator still decides what to spawn.

THE PAIR

Pair: workflow versus agent. Same steps every time means you could draw the flowchart, so no decision-maker is needed. Module 2.