

MODULE 19 · EXAM STRATEGY - ALL DOMAINS

How This Exam Writes Its Questions

Item anatomy, the qualifier that decides it, four distractor families, and where two minutes actually goes.

MODULE 19 of 20 QUESTIONS 60 MINUTES 120

The exam, in numbers

- ◆ **60 questions, 120 minutes.** Two minutes each, with no reading time on top.
- ◆ **Four options, one correct.** The practice bank was four options in 133 of 134 items.
- ◆ **Some items may ask for more than one answer.** None appeared in the bank; the format allows it, so read the instruction.
- ◆ **Scaled to 1000, pass at 720.** Not a raw percentage, so do not compute your score in the room.
- ◆ **Four of the six scenarios, drawn at random.** You cannot predict which.

ON THE EXAM

Two minutes per question, sixty times, with no break in the middle. Pace is a skill you can practise.

Anatomy of an item

PART	WHAT IT DOES	TYPICAL SIZE
The setup	Establishes a production system and its constraints	2 to 3 sentences
The symptom	What is going wrong, in observable terms	1 to 2 sentences
The planted detail	The one fact that decides the answer	Often the last sentence
The qualifier	"Most effective", "root cause", "guarantee"	In the question line
Four options	Two plausible, one weak, one reversed or invented	About 19 words each

ON THE EXAM

Measured on the practice bank: stems ran 34 to 150 words, median 101. Options averaged around 19.

About 180 words to read, per question, sixty times. That is roughly **45 seconds** before you have thought at all.

Where the two minutes goes

- ◆ **About 45 seconds reading**, if you read everything once at a normal pace.
- ◆ **About 15 seconds forming your own answer** before looking at the options.
- ◆ **About 30 seconds matching and eliminating.**
- ◆ **About 30 seconds spare**, which is your buffer for the hard ones.
- ◆ **Re-reading the stem a third time is where the buffer goes.** Read it twice, deliberately.

ON THE EXAM

The average is not the target. Easy items should take 60 seconds so hard ones can take three minutes.

A reading order that saves time

- ◆ **Read the question line first**, before the stem. Now you know what you are reading for.
- ◆ **Then the stem, watching for the planted detail.** It is usually the last sentence.
- ◆ **Then say your own answer**, in a phrase, before looking at the options.
- ◆ **Then read the options, matching rather than evaluating.**
- ◆ **Reading the options first invites them to frame the stem for you.**

ON THE EXAM

Committing to an answer before reading the options is what makes a reversed distractor visible.

The qualifiers, by how often they appear

QUALIFIER	SHARE OF BANK STEMS	WHAT IT MEANS
"most effective"	40%	The answer that generalises
"BEST" or "best"	11%	The same thing, in fewer letters
"most reliable" / "appropriate"	5%	Fewer ways to fail
"root cause"	rare	A workaround is still wrong
"guarantee" / "regardless"	rare	Every prompt-shaped option dies

ON THE EXAM

From the 134-question practice bank, not Anthropic. The bottom two are rare and decisive.

What "most effective" is really asking

- ◆ Which option fixes the class of problem, not the instance in front of you.
- ◆ Options that patch the stated case are the classic near-miss. They work, and they work once.
- ◆ Prefer the change that makes the next occurrence impossible over the one that handles this occurrence.
- ◆ Two options both work? Ask which one you would still be happy with in six months.

ON THE EXAM

Almost half the paper turns on this word. It is worth more attention than any single technical fact.

Distractors are not random. There are four shapes, and each one is **wrong in a predictable way.**

The four distractor families

FAMILY	HOW TO SPOT IT
Right idea, wrong layer	The correct principle applied in the prompt instead of in code
Symptom, not cause	Genuinely reduces the failure without addressing why it happens
Over-engineered	A heavier mechanism than the problem needs, often the impressive one
Reversed or invented	The two halves swapped, or a capability that does not exist

ON THE EXAM

Every scenario in this course has had at least one of these. Now you get to name them.

Family one: right idea, wrong layer

- ◆ **The rule belongs in code and the option puts it in a prompt.** Or in CLAUDE.md, which is the same thing.
- ◆ **The check belongs in the tool and the option asks the agent to remember.**
- ◆ **A parameter the model sets, offered as a control.** confirmed: true is not a control.
- ◆ **The tell:** you agree with the principle and the mechanism cannot enforce it.

ON THE EXAM

This is the most common family on this exam, because the principle in the option is always correct.

Family two: symptom, not cause

- ◆ It **reduces the failure rate** without changing why the failure happens.
- ◆ **Common shapes:** more retries, a bigger context window, a stricter prompt, a shorter output.
- ◆ **Test it:** if this option ships, does the same failure return at a different scale?
- ◆ **Watch for it under "root cause" and "most effective".** Both explicitly exclude this family.

ON THE EXAM

"It would help" is not the standard. Ask whether it would still be helping in six months.

Family three: over-engineered

- ◆ **An agent where a workflow would do.** Subagents where one context would do.
- ◆ **Building a server where adopting one would do.** A model where a rule already works.
- ◆ **It sounds like good architecture**, which is precisely why it is offered.
- ◆ **The tell:** it adds a decision-maker to something that had no decision to make.

ON THE EXAM

The more sophisticated option is not the default answer. Anthropic's own guidance is to find the simplest thing that works.

Family four: reversed or invented

- ◆ **Reversed: the two halves of a pair swapped.** Glob for contents, batch the urgent ones.
- ◆ **Invented: a mechanism that does not exist.** Subagents messaging each other, an SDK that reads your policy.
- ◆ **Both are easy marks if you know the answer before you read.**
- ◆ **Invented options often sound the most technical.** Specificity is not evidence.

ON THE EXAM

An invented mechanism is not a trick question. It is testing whether you know the real mechanism.

What the data says about guessing

- ◆ **Answer letters were near-uniform in the bank.** A 35, B 33, C 33, D 33. There is no lucky letter.
- ◆ **The longest option was correct 38% of the time**, against 25% by chance. A weak signal, not a strategy.
- ◆ **Eliminating one option raises a guess from 25% to 33%.** Two raises it to 50%.
- ◆ **Always answer.** There is no penalty for a wrong answer, so a blank is a guaranteed zero.

ON THE EXAM

Measured on the practice bank. Use elimination, which is real, rather than surface heuristics, which are noise.

Flagging, and the second pass

- ◆ **Answer everything on the first pass**, even when you are unsure. Then flag it.
- ◆ **Flag on a rule**, not a feeling: two options you could not separate, or a term you did not recognise.
- ◆ **Do not flag more than about ten**. A second pass over twenty items is not a second pass.
- ◆ **On review, change an answer only with a reason you can state**. Vague doubt is not a reason.

ON THE EXAM

A flagged item with an answer already in it costs you nothing. An unanswered one costs you the whole mark if you run out of time.

The six scenarios, and what each really tests

SCENARIO	WHAT IT IS ACTUALLY ABOUT
Customer support agent	Escalation, guaranteed outcomes, tool limits
Code generation with Claude Code	Memory hierarchy, plan mode, enforcement layers
Multi-agent research	Delegation, context isolation, coordinator responsibilities
Developer productivity tools	Built-in tools, exploration, MCP integration
Claude Code in CI/CD	Headless flags, classifying endings, fresh-session review
Structured data extraction	Schema design, fabrication, batch and retry

ON THE EXAM

The scenario is the setting. The domain being tested is often a different one, so read for the decision.

Exam-day practicalities

- ◆ **Check the current fee, booking route and access rules yourself.** Those move, and this course cannot be current on them.
- ◆ **If you sit it online, test your setup before the day,** not thirty minutes before.
- ◆ **Know where the clock is** and glance at it every ten questions, not every question.
- ◆ **At question 30 you should have about 60 minutes left.** That is the only checkpoint you need.

ON THE EXAM

One mid-point check is enough. Watching the clock constantly costs more time than it saves.

Module 19 in five lines

- ◆ Read the question line first, then the stem, then commit before reading the options.
- ◆ "Most effective" is on nearly half the paper and means the answer that generalises.
- ◆ Four distractor families: wrong layer, symptom not cause, over-engineered, reversed or invented.
- ◆ Answer everything, flag sparingly, and change an answer only with a stated reason.
- ◆ Easy items in a minute so hard ones can have three.

ON THE EXAM

Five scenarios next. Each one names the distractor family behind every wrong option.

A support agent must never issue a refund above £1,000 without manager approval. The system prompt states this and the tool description repeats it.

Over four months, six refunds above the limit were issued. In each case the customer had given a compelling reason.

Audit requires the team to demonstrate the limit cannot be exceeded.

What is the most effective change?

- ☐ A Add examples to the system prompt showing correct refusals above £1,000
- ☐ B Reject amounts above £1,000 inside the refund tool, returning an error naming the approval tool
- ☐ C Add an approved_by_manager parameter the agent must set for amounts above £1,000
- ☐ D Route all refunds through a manager approval subagent

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Examples in the prompt
Symptom, not cause. Improves the rate against a requirement that says demonstrate it cannot happen.
- B

 Reject inside the tool
Runs on every call, cannot be persuaded, and can be shown to an auditor.
- C

 approved_by_manager parameter
Right idea, wrong layer. The flag is set by the agent you are constraining.
- D

 Approval subagent for everything
Over-engineered. It puts a subagent in front of every small refund to police six large ones.

THE FAMILIES

Three families in one item: symptom, wrong layer, over-engineered. Naming them is faster than reasoning through each.

A research coordinator delegates to four subagents. The synthesis subagent consistently omits findings that the data subagent produced.

Logs confirm the data subagent completed and returned a detailed result to the coordinator. The synthesis subagent's prompt reads: "Synthesise the research into a report."

What is the root cause?

- ☐ A The synthesis subagent's maxTurns is too low to gather everything
- ☐ B The data subagent should pass its findings to the synthesis subagent directly
- ☐ C The coordinator is not including the findings in the synthesis prompt
- ☐ D The synthesis subagent needs a more capable model

◆ PAUSE AND DECIDE

What is the root cause?

- A

 Turn cap too low
Symptom, not cause. More turns to search a context that will never contain the findings.
- B

 Direct subagent messaging
Invented. Subagents have no channel to each other; the coordinator is the only route.
- C

 Findings not in the prompt
The prompt is quoted and contains no findings. The prompt is the only channel in.
- D

 Bigger model
Symptom, not cause. A larger model omits data it never received just as reliably.

THE FAMILIES

The quoted prompt is the planted detail. When a stem quotes something in full, read it as evidence, not as colour.

A team processes 200,000 documents monthly for archival classification. Results feed a quarterly report.

The pipeline sends each document as a synchronous request on arrival. The bill has doubled year on year.

No result is consumed before the quarterly report is compiled.

What is the most cost-effective change?

- ☐ A Process documents in parallel to reduce total runtime
- ☐ B Send the documents through the Message Batches API
- ☐ C Build a multi-agent pipeline with specialised classifier subagents
- ☐ D Reduce max_tokens so each classification response is shorter

◆ PAUSE AND DECIDE

What is the most cost-effective change?

- A

 Parallel processing
Reversed understanding. Parallelism buys wall-clock time and costs exactly the same.
- B

 Batch API
Nothing is consumed before the quarterly report. Half price for a scheduling change.
- C

 Multi-agent pipeline
Over-engineered, and more expensive. Subagents multiply tokens; the stem asked to reduce cost.
- D

 Shorter responses
Symptom, not cause. Output is a small share of a classification bill.

THE FAMILIES

Option C is the family worth watching: an impressive architecture offered as an answer to a cost question.

An engineer needs the agent to locate every file whose name ends in .config.ts, then find which of those reference a deprecated helper.

The repository has about 4,000 files.

What is the most efficient approach?

- ☐ A Read every file in the repository, then filter
- ☐ B Glob for the filename pattern, then Grep those files for the helper
- ☐ C Grep for the filename pattern, then Glob for the helper
- ☐ D Spawn a subagent to explore the repository and report back

◆ PAUSE AND DECIDE

What is the most efficient approach?

- A

 Read everything
Symptom of not knowing the search tools. Four thousand files into context to answer about a handful.
- B

 Glob then Grep
Names first with Glob, then contents within that set using Grep. Neither loads a file.
- C

 Grep then Glob
Reversed. Glob cannot look inside files, so the second step cannot work.
- D

 Subagent
Over-engineered. A round trip and a fresh context load for two cheap searches.

THE FAMILIES

Reversed and over-engineered in the same item. Both are eliminable in seconds if you formed your answer first.

An extraction schema requires a `policy_number` field. About 15% of submitted documents have no policy number.

For those, extraction returns a plausible number that does not exist in any system. The prompt already instructs the model not to invent values.

Downstream systems have created records against these numbers.

What is the most effective fix?

- ☐ A Strengthen the instruction and add an example with no policy number
- ☐ B Make `policy_number` nullable, documenting that null means not present
- ☐ C Validate extracted policy numbers against the source text before accepting
- ☐ D Lower the temperature so the model invents less

◆ PAUSE AND DECIDE

What is the most effective fix?

- A

 Stronger instruction
Wrong layer. The prompt says do not invent; the schema says produce a value. The schema wins.
- B

 Make it nullable
Removes the conflict. Absence becomes expressible, so there is nothing to invent.
- C

 Validate against source
Symptom, not cause. Catches it after the fact, on every document, forever.
- D

 Lower temperature
Symptom, not cause. More consistent inventions, not fewer.

THE FAMILIES

"The prompt already instructs" is the eliminator, and it is placed deliberately. When a stem tells you an instruction exists and is failing, stop writing instructions.