

MODULE 20 · EXAM STRATEGY - ALL DOMAINS

What Changed Since the Exam Guide

Every divergence this course flagged, in one place, ranked by whether it changes your answer.

MODULE 20 of 20 DIVERGENCES 13 CHANGE YOUR ANSWER 5

You need **two** answers for a handful of questions: the one the exam wants, and the one you would build with.

How big is the gap, really

- ◆ **Thirteen divergences across twenty modules.** Five change an answer; eight only affect reading the docs.
- ◆ **Nothing in the exam's framing became wrong.** Several things became incomplete.
- ◆ **One technique was withdrawn**, which is the only genuine correction in the list.
- ◆ **The architecture, the loop and the enforcement principles are unchanged.** That is most of the course.

ON THE EXAM

Proportionally: about a fifth of the syllabus has drifted, and the drift is almost all additive.

Tier one: these change your answer

TOPIC	EXAM ANSWER	CURRENT REALITY
Schema-conformant JSON	Tool input_schema + tool_choice	Native structured outputs, GA
Prefill	A recommended technique	Withdrawn on the latest models
Commands and skills	Two mechanisms	One system, same slash command
The delegation tool	Task	Renamed Agent; Task is an alias
"Which memory files loaded?"	/memory	/context reports what loaded

ON THE EXAM

Give the middle column in the exam room. Row two is the only outright withdrawal.

Structured outputs, since it is the largest

- ◆ **Two forms.** `output_config.format` constrains what Claude says; strict on a tool constrains how it calls.
- ◆ **Both use grammar-constrained sampling.** Invalid output is unreachable, not merely discouraged.
- ◆ **They solve different problems and combine.** Valid tool arguments and a structured final answer.
- ◆ **Supported keywords are a subset.** No numeric bounds, no string length, no recursion.
- ◆ **You still need a validation layer.** Shape is guaranteed; meaning is not.

ON THE EXAM

The exam's tool technique is still good engineering. Structured outputs is simply better, where it is available.

One correction to make

- ◆ **Module 12 listed prefill in the technique ladder.** That is the exam-era guidance, and it is on the paper.
- ◆ **Prefill is no longer supported on the latest models.** Do not reach for it in new work.
- ◆ **For format control, use structured outputs,** or examples, which are unaffected.
- ◆ **Everything else on that ladder is current.** Be clear and direct is still the first and largest lever.

ON THE EXAM

This is the only place where following the course's exam-era guidance at work would leave you using something withdrawn.

Tier two: these will confuse you in the docs

TOPIC	WHAT THE GUIDE SAYS	WHAT YOU WILL FIND
Large toolsets	Consolidate	Plus a Tool Search Tool, loading on demand
Filling context window	Summarise manually	Context editing and a memory tool primitive
Claude Code in CI	Config loads normally	--bare skips hooks, skills, MCP, CLAUDE.md
Entering plan mode	Shift+Tab	Also /plan, a CLI flag, a settings default
Memory files	CLAUDE.md hierarchy	Plus auto memory and AGENTS.md
Hooks	A handful of events	Around thirty; skills carry their own

ON THE EXAM

None of these makes a guide-era answer wrong. All of them make it incomplete, which is why they read as surprises.

Tier three: moving, and not worth studying

TOPIC	STATUS
MCP going stateless; roots, sampling and logging deprecated	A 2026 release candidate with a long removal window
Subagent nesting to five levels	Documented, and reports differ on which subagent types get it
Auto mode as the default permission mode	Shipped for new sessions on paid plans; plan mode is unaffected

ON THE EXAM

Answer the exam version for all three. These are here so the documentation does not unsettle you the week before.

Everything that makes this exam **hard** is unchanged. The drift is in the surface, not the reasoning.

What has not changed at all

AREA	STATUS
The agentic loop and stop_reason	Unchanged. Still the loop condition.
Hub and spoke; subagents inherit nothing	Unchanged. Siblings still cannot talk.
isError versus protocol error	Unchanged, and clarified in favour of isError for validation.
settings.json precedence; managed policy wins	Unchanged.
Prompts persuade; code enforces	Unchanged, and the highest-value idea in the course.
Batch at 50%; required fields cause fabrication	Unchanged.

ON THE EXAM

If you know this table, the drift in the other three is a footnote rather than a problem.

Holding both answers

- ◆ **In the exam room, answer from the guide.** If an option describes something the guide predates, it is not the answer.
- ◆ **At your desk, build from today's documentation.** The exam is not a style guide.
- ◆ **Do not try to reconcile them.** They are answers to different questions, asked by different people.
- ◆ **The five tier-one rows are the only ones worth rehearsing.** Everything else you will simply recognise.

ON THE EXAM

Knowing which mode you are in is the whole skill. It takes a second and it costs nothing.

Check these yourself before booking

- ◆ **The fee, and whether it has changed.**
- ◆ **Access rules.** Availability has been partner-gated, and gating changes.
- ◆ **The exam guide version.** If Anthropic has published a new one, it supersedes this module.
- ◆ **Anthropic's certification pages are the authority.** Not this course, and not a blog post.

ON THE EXAM

This course cannot stay current on commercial details. It can tell you which ones to check.

Module 20 in five lines

- ◆ Thirteen divergences. Five change your answer; eight only affect reading the docs.
- ◆ Structured outputs is the largest gap. Answer the tool technique; build with the real feature.
- ◆ Prefill is the only withdrawal. Do not use it in new work.
- ◆ The loop, the hierarchy, and persuade-versus-enforce have not moved at all.
- ◆ Exam room: the guide. Desk: the docs. Know which room you are in.

ON THE EXAM

Three scenarios next, then the practice exams. All three test whether you can pick the exam answer over the newer one.

An exam item asks how to guarantee that a document extraction returns JSON conforming to a defined schema, every time.

One option describes defining a tool whose `input_schema` is the target shape and forcing it with `tool_choice`. Another describes passing the schema in an output configuration so generation is grammar-constrained.

Which do you select?

- ☐ A The grammar-constrained option, since it is the stronger guarantee
- ☐ B The tool `input_schema` option, since that is the technique the exam is scored against
- ☐ C Neither. Add a validation-and-retry loop instead
- ☐ D Whichever the stem's scenario has already established as in use

◆ PAUSE AND DECIDE

Which do you select?

- A

Grammar-constrained

Technically the better answer, and it postdates the guide. An option the guide predates is not what the paper is scored against.
- B

Tool input_schema

The exam's answer, and still sound engineering. This is the mode switch working.
- C

Validation and retry

A safety net, not a guarantee. The question asked how to guarantee conformance.
- D

Match the scenario

Sensible instinct in general. It does not resolve a question about which technique is correct.

THE SIGNAL

This is uncomfortable and it is the right answer. You are being scored against a published blueprint, not against current best practice.

A colleague preparing for the same exam reads the current Claude Code documentation and finds that slash commands and skills are one system, and that the delegation tool is called Agent.

They are now unsure whether the course material on subagents and slash commands is wrong.

They sit the exam next week.

What should you tell them?

- ☐ A Restudy from the current documentation, since it is authoritative
- ☐ B The mechanisms are the same; only the names and packaging changed. Answer Task and treat commands and skills as the guide does
- ☐ C Ignore the documentation entirely until after the exam
- ☐ D Contact Anthropic to confirm which version the exam uses

◆ PAUSE AND DECIDE

What should you tell them?

- A

 Restudy from current docs
Overcorrection. The behaviour did not change, so there is nothing to restudy, and it costs a week.
- B

 Same mechanism, new label
Task still works as an alias and skills still produce a slash command. Nothing they learned is wrong.
- C

 Ignore the docs
Overcorrects the other way. Reading the docs is how they found this, which is useful.
- D

 Ask Anthropic
Reasonable, and a week is short. The answer is knowable without asking.

THE SIGNAL

Both renames are aliases or mergers, not behaviour changes. Panic costs more than the divergence does.

An engineer passes the exam and is asked to build a production extraction pipeline the following month.

They implement the tool input_schema technique, manual context summarisation, and a prefilled response for format control, because those are the techniques they revised.

Which choice is actually a problem?

- ☐ A The tool input_schema technique
- ☐ B Manual context summarisation
- ☐ C The prefilled response
- ☐ D All three should be replaced with current features

◆ PAUSE AND DECIDE

Which choice is actually a problem?

- A

 Tool input_schema
Still works and is still sound. Structured outputs is better where available; this is not broken.
- B

 Manual summarisation
Still works. Context editing is more convenient, not a correction.
- C

 The prefilled response
The only withdrawal in the list. Unsupported on the latest models, so this one genuinely fails.
- D

 Replace all three
Two of them are fine. Rewriting working code because a newer feature exists is not an improvement.

THE SIGNAL

The distinction that matters: superseded is not the same as withdrawn. Only one row in this module is a withdrawal.