

PRACTICE EXAM 2 OF 4 · EXAM STRATEGY - ALL DOMAINS

Practice Exam 2

Twenty new questions. No topic repeats from exam 1. One question pairs with an earlier one deliberately.

QUESTIONS 20 REPEATS none PAUSE 9 seconds

Before you start

- ◆ **Leave a day or two since exam 1 if you can.** Two papers in one sitting measures stamina, not knowledge.
- ◆ **Have your exam 1 domain tally to hand.** The comparison is the point.
- ◆ **If you rewatched a domain, watch whether it moved.** If it did not, the problem is not coverage.
- ◆ **Two papers gives you a direction.** Not yet a trend, but better than a snapshot.

ON THE EXAM

You are looking for whether the same domain misses twice. That is the finding, not the score.

A document-processing agent occasionally never returns. The loop continues calling the API long after a complete answer has been produced, until the turn cap aborts it.

Logs show the final responses contain only text and no tool calls. The loop implementation exits when the response text contains the phrase "processing complete".

What is the root cause?

- ☐ A The turn cap is set too low for large documents
- ☐ B The loop determines completion from response text rather than from `stop_reason`
- ☐ C The agent needs a system prompt instructing it to say "processing complete"
- ☐ D Tool results are being appended without the preceding assistant turn

◆ PAUSE AND DECIDE

What is the root cause?

- A

 Turn cap too low
The cap is what stops the runaway. It is the limiter, not the cause.
- C

 Prompt it to say the phrase
Makes correctness depend on the model choosing a phrase. That is the bug.
- B

 Completion read from text
No tool calls means stop_reason is end_turn. The loop had its exit and ignored it.
- D

 Missing assistant turn
That raises an API error immediately, not a loop that runs too long.

THE SIGNAL

"Only text and no tool calls" is the planted detail. No tool calls means end_turn. Module 1.

A project CLAUDE.md states that all new modules must use dependency injection. An engineer's user-level CLAUDE.md states a personal preference for constructor arguments instead.

That engineer reports Claude usually uses dependency injection, sometimes uses constructor arguments, and once used both in the same file.

They want the project standard to apply reliably.

What should they do?

- ☐ A Nothing. Project memory overrides user memory, so the standard already wins
- ☐ B Remove the conflicting rule from the user file, and enforce the pattern with a review step
- ☐ C Move the project rule into CLAUDE.local.md so it is read last
- ☐ D Add "this overrides personal preferences" to the project CLAUDE.md

◆ PAUSE AND DECIDE

What should they do?

- A

 Project overrides user
The claim the docs never make. Memory files are concatenated, and the behaviour disproves it.
- C

 Read last wins
Load order is not precedence. It also puts a team standard in a personal gitignored file.
- B

 Remove the conflict, then enforce
Ends the contradiction, then puts the standard somewhere with a defined outcome.
- D

 Stronger wording
Two contradictory instructions in one context window. Emphasis does not resolve that.

THE SIGNAL

"Usually, sometimes, once both" is concatenation with no override rule. Module 8.

A prompt asks Claude to "review this API specification carefully for problems".

Across two hundred specifications it reliably flags missing error responses, sometimes flags inconsistent naming, and has never flagged a missing rate-limit definition.

The team has strengthened the wording twice.

What is the most effective change?

- ☐ A Strengthen the wording further and note that rate limits are especially important
- ☐ B Replace the adjectives with a named list of checks, including rate-limit definitions
- ☐ C Add three example specifications with their correct reviews
- ☐ D Use a more capable model, since API review is demanding

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Strengthen further
Already tried twice. It also patches one category rather than fixing coverage.
- C

 Example reviews
Helps format and judgement. Cannot guarantee coverage of a category the examples omit.
- B

 Named checks
Turns a variable sweep into a fixed checklist you can extend when a gap appears.
- D

 Better model
It is not failing to understand APIs. Rate limits were never in scope.

THE SIGNAL

"Never flagged" is the tell. A capability problem produces sometimes; a scope problem produces never. Module 12.

An engineer resumes a session after a deployment. Two files the agent had read were changed by that deployment.

They ask the agent to re-read both files before continuing. It does. It then generates code calling a method that was removed in the deployment.

Nothing was truncated and no error occurred.

What is happening?

- ☐ A The re-read returned cached contents
- ☐ B Both versions are now in the conversation, and nothing marks the older one as superseded
- ☐ C The agent cannot read files modified since the session began
- ☐ D The context window filled and the newer read was dropped

◆ PAUSE AND DECIDE

What is happening?

- A

 Cached contents
Invents a caching layer. The stem says it re-read them.
- B

 Both versions present
A conversation is append-only. The new read sits after the old one; neither is marked current.
- C

 Cannot read changed files
Invented limitation. Reading has nothing to do with when a file changed.
- D

 Window filled
That raises an error, and the stem rules out truncation and errors.

THE SIGNAL

Re-reading adds a version. It never removes one. Say the earlier read is superseded, or start fresh. Module 16.

A due-diligence system analyses acquisition targets. For every target it runs exactly four analyses: financial, legal, technical and market. The four never vary and are independent.

An architect proposes an orchestrator that decides at run time which analyses each target needs.

Turnaround time is the main complaint.

What should you advise?

- ☐ A Build the orchestrator; run-time decisions will reduce unnecessary work
- ☐ B Parallelise the four fixed analyses and aggregate the results
- ☐ C Run the four analyses sequentially with a shared context
- ☐ D Merge the four analyses into one prompt to reduce overhead

◆ PAUSE AND DECIDE

What should you advise?

- A

 Orchestrator
Over-engineered. You could write the subtask list before seeing the input, so nothing needs deciding.
- B

 Parallelise the fixed four
The list never varies, so write it at author time and run the four concurrently.
- C

 Sequential
Four times the wall-clock time for no benefit. The complaint was turnaround.
- D

 One merged prompt
Four jobs, one attention budget. All four get worse.

THE SIGNAL

Could you write the subtask list before seeing the input? Yes, so parallelise. Module 2.

A tool called `manage_subscription` takes a required action parameter accepting "view", "upgrade", "downgrade" and "cancel", plus eleven other parameters used by different actions. Two problems recur. The agent occasionally calls it with action "cancel" when a customer asked only what cancelling would cost. And the support team cannot grant view access without also granting cancel.

What is the most appropriate change?

- ☐ A Split into separate tools, one per action, each with only its own parameters
- ☐ B Add a confirm parameter the agent must set for cancel and downgrade
- ☐ C Improve the description to explain when each action applies
- ☐ D Add a hook that blocks cancellations the customer did not request

◆ PAUSE AND DECIDE

What is the most appropriate change?

- A

 Split by action
Fixes both symptoms. Distinct tools are selected more reliably, and permissions become per-tool.
- B

 Confirm parameter
The agent sets that flag too. A parameter the model controls is not a control.
- C

 Better description
May reduce wrong-mode selection. Does nothing about the permissions boundary.
- D

 Hook on cancel
Real enforcement, and it cannot reliably infer what the customer requested.

THE SIGNAL

Two symptoms, one cause. A permissions boundary you cannot draw means the tool boundary is wrong. Module 5.

An internal assistant answers questions against a 60,000-token compliance manual. Every request sends the full manual followed by a short question.

The manual is revised annually. Around 5,000 questions are asked per day and users expect an answer within seconds. Input tokens dominate the bill.

This pairs with exam 1 question 7, where batch was correct.

What is the most effective change?

- ☐ A Move the workload to the Batch API for the 50% discount
- ☐ B Cache the manual prefix so repeated questions read from the cache
- ☐ C Summarise the manual and send the summary
- ☐ D Reduce max_tokens on the response

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Batch it
Users expect an answer in seconds. Batch requires the one thing this workload cannot give.
- B

 Cache the prefix
Large, stable, shared prefix in front of a short question, reused 5,000 times a day.
- C

 Send a summary
Cheaper and lossy. Compliance answers depend on exact wording.
- D

 Shorter responses
The stem says input dominates. This targets the smaller half.

THE SIGNAL

The pair: in exam 1 nobody was waiting, so batch. Here users wait, so cache. One constraint changed the answer.

A team wants a subagent that audits infrastructure code for misconfigurations and reports findings. It must never be able to modify any file, regardless of how it is prompted or what an engineer asks it to do.

An engineer proposes a system prompt beginning "You are a read-only auditor. Never edit or write files."

What is the most reliable enforcement?

- ☐ A Use the proposed system prompt and repeat the constraint in CLAUDE.md
- ☐ B Define the subagent with read tools only, so write tools are not in its toolset
- ☐ C Have the coordinator inspect the subagent's output and reject modifications
- ☐ D Set the subagent's turn cap to 1 so it has no opportunity to write

◆ PAUSE AND DECIDE

What is the most reliable enforcement?

- A

 Prompt plus CLAUDE.md
Both are context. "Regardless of how it is prompted" rules this out by definition.
- B

 Read tools only
A tool it does not have is a tool it cannot call. The restriction holds under any prompt.
- C

 Coordinator inspects output
Detection after the fact, and the coordinator sees only the final message.
- D

 maxTurns of 1
One turn is enough to call a write tool, and it breaks an auditor that must read many files.

THE SIGNAL

"Regardless of how it is prompted" eliminates every prompt-shaped option before you finish reading them. Module 3.

A team requires that every SQL file Claude Code writes is checked by their linter immediately afterwards, so violations are visible before commit.

The rule currently sits in CLAUDE.md. Compliance is around 85%; during large multi-file changes the lint step is sometimes skipped.

They do not want Claude to have to remember.

What should they implement?

- ☐ A A PreToolUse hook on Write and Edit that blocks unlinted content
- ☐ B A PostToolUse hook on Write and Edit that runs the linter
- ☐ C A skill called /lint that Claude invokes after editing
- ☐ D A permissions rule denying Write on SQL files unless linted

◆ PAUSE AND DECIDE

What should they implement?

- A

 PreToolUse blocking
Blocks the write mid-task and breaks multi-step reasoning. A lint step is not a prohibition.
- B

 PostToolUse linter
The file is written, then linted, every time. Claude never has to remember.
- C

 A /lint skill
Still relies on Claude choosing to invoke it. That is the 15% they already have.
- D

 Conditional permission
Permissions match on the tool, not on content. They cannot express "unless linted".

THE SIGNAL

"Also do this" is additive, so PostToolUse. Prohibitions go in PreToolUse. Module 9.

To answer "when does the Henderson account renew?", an agent calls `search_accounts` to find the account, reads the account ID from the result, calls `get_contract` to fetch the contract, reads the contract ID, then calls `get_renewal_date`.

Three round trips per question. Traces show the intermediate IDs are never used for anything else.

What is the most effective change?

- ☐ A Run the three calls in parallel to reduce latency
- ☐ B Provide a tool taking an account name and returning the renewal date, resolving IDs internally
- ☐ C Cache the intermediate IDs for repeated questions
- ☐ D Add a system prompt example showing the correct three-call sequence

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Run in parallel
Impossible. Each call needs an ID the previous one produced. This is a chain.
- B

 One tool, internal resolution
Success is always the same three calls, so the sequence belongs inside the tool.
- C

 Cache the IDs
Helps repeat questions about Henderson. The first question about any account still takes three hops.
- D

 Show the sequence
The agent already performs it correctly. Teaching it what it does is not a fix.

THE SIGNAL

"Intermediate IDs never used for anything else" is the merge signal. Make the tool take what the agent has. Module 5.

An agent escalates about 9% of conversations. Escalations arrive in the human queue as the full transcript and nothing else.

Agents report that an escalated case takes nearly as long as one handled from scratch. In several cases a human contradicted something the agent had already told the customer.

The escalation rate is considered appropriate.

What should change?

- ☐ A Reduce the escalation rate so humans handle fewer cases
- ☐ B Replace the transcript with a structured handover naming the reason, established facts, what was said, and what remains
- ☐ C Have the agent add a summary paragraph at the top of the transcript
- ☐ D Train human agents to read transcripts faster

◆ PAUSE AND DECIDE

What should change?

- A

 Escalate less
The stem says the rate is appropriate. This handles fewer cases correctly.
- B

 Structured handover
Named fields, including what was already said, which is what causes the contradictions.
- C

 Summary paragraph
A real improvement, and prose is not a field. The missing item can go missing again.
- D

 Read faster
Puts the cost on humans permanently rather than removing it once.

THE SIGNAL

The contradictions are the tell: "what was already said" is not a named field anywhere. Module 17.

An agent has been working through a schema migration with an engineer for three hours. The engineer leaves for a meeting.

While they are away, a colleague merges a branch that renames six tables the agent has read and reasoned about.

The engineer returns. The remaining work depends on those six tables.

What is the most reliable way to continue?

- ☐ A Resume the session so the agent keeps its full understanding of the migration
- ☐ B Resume and ask the agent to re-read the six affected files
- ☐ C Start a fresh session with a summary of decisions made and the current schema
- ☐ D Fork the session so the pre-merge state is preserved for comparison

◆ PAUSE AND DECIDE

What is the most reliable way to continue?

- A

 Resume
Restores six stale reads and every conclusion drawn from them.
- B

 Resume and re-read
The strongest distractor, and it leaves both versions in an append-only history.
- C

 Fresh session with a summary
Carries the decisions forward and leaves the stale reads behind.
- D

 Fork
Copies the stale history into a second session. Two copies of the wrong picture.

THE SIGNAL

"While they were away" is the planted phrase. Something changed, so start fresh with what you know now. Module 4.

A lease extraction schema has a `monthly_rent` field defined as a single number.

Several leases have been varied. The original states £4,200; a variation deed changes it to £4,650. The extracted value is sometimes one and sometimes the other, with no indication both exist.

The property team need both figures and which is current.

What should change?

- ☐ A Instruct the model to always take the value from the most recent variation
- ☐ B Model the field as a list of entries, each with a value, an effective date and a source reference
- ☐ C Add a second field for the varied rent
- ☐ D Add a boolean `has_been_varied` flag alongside the number

◆ PAUSE AND DECIDE

What should change?

- A

 Take the most recent
Discards the original, which the team explicitly need. It also makes the model arbitrate a legal question.
- B

 A list with effective dates
Both values survive, each dated, and downstream decides which applies.
- C

 A second field
Works for one variation. A twice-varied lease breaks it.
- D

 A varied flag
Tells you something was lost without telling you what.

THE SIGNAL

If a document can legitimately contain two answers, a scalar field is the wrong shape.
Module 13.

A pipeline generates a feature, then in the same session asks Claude to review its own work against the team's standards before opening a pull request.

Reviews are consistently positive. Human reviewers later find defects the automated review missed, several stemming from a misreading of an existing utility's contract.

What is the most effective change?

- ☐ A Give the review step a longer, more detailed checklist
- ☐ B Run the review in a separate session that sees only the diff and the standards
- ☐ C Raise max_turns on the review step so it can look more carefully
- ☐ D Instruct the review step to assume the code is wrong

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Longer checklist
A checklist cannot catch a misreading. The reviewer shares it and ticks the box.
- B

 Separate session
The misunderstanding lives in the authoring context. A reviewer without it reads the code as written.
- C

 More turns
More time spent agreeing with itself. Duration was never the constraint.
- D

 Assume it is wrong
Adjusts tone, not information. Different remarks about code it still misunderstands.

THE SIGNAL

"Stemming from a misreading" rules out every answer that asks the reviewer to work harder. Module 11.

A weekly job runs an agent to summarise changes across a set of repositories. Each run happens on a fresh CI container.

The team wants each run to remember the previous week's summary so recurring themes are recognised. They have added the continue flag.

Every run behaves as though it is the first.

Why, and what should they do?

- ☐ A `--continue` only works interactively. Use `--resume` with the session ID
- ☐ B No session exists on this host. Persist the transcript or a session store, or carry conclusions forward in a file
- ☐ C Sessions expire after seven days, so the previous week's session is gone
- ☐ D `--continue` requires the previous run to have ended cleanly

◆ PAUSE AND DECIDE

Why, and what should they do?

- A

 Use `--resume` with the ID
Closer, and incomplete. Resume also reads a local transcript this container does not have.
- C

 Seven-day expiry
Invented. There is no such expiry.
- B

 No session on this host
Transcripts are local files. A fresh container starts with nothing, so something must be carried across.
- D

 Unclean termination
Invented, and it does not explain why even the second run behaves as a first.

THE SIGNAL

"A fresh CI container" is the whole question. Session state is a file on a machine, and this machine is new. Module 4.

A compliance agent uses a policy search tool. When a query matches no policies, the tool returns an error reading "Search failed: no matching policies."

Engineers notice the agent runs several progressively broader searches after an empty result, then stops using the tool for the rest of the session and answers from memory.

What is the most effective change?

- ☐ A Return a successful empty result stating what was searched and how to broaden it
- ☐ B Instruct the agent in the system prompt that empty searches are normal
- ☐ C Have the tool automatically broaden the query when nothing matches
- ☐ D Return an empty array with no message

◆ PAUSE AND DECIDE

What is the most effective change?

- A

 Successful empty result with context
Fixes the miscategorisation and gives the agent something to do next.
- B

 Explain it in the prompt
Asks the agent to disbelieve its own tool. The tool says failed; the prompt says ignore that.
- C

 Auto-broaden
The agent cannot tell an exact match from a widened one. It hides the boundary.
- D

 Empty array, no message
Better than an error, and it gives no context. Half the fix.

THE SIGNAL

"Stops using the tool" is the giveaway. Repeated failures taught it the tool is unreliable.
Module 6.

A booking agent checks availability through a partner API. During a ninety-minute outage the API returned errors on every call.

Reviewing the outage, the team finds the agent told several customers that rooms were available on their requested dates. Those statements were plausible and, in three cases, wrong.

The system prompt tells the agent to be helpful and keep customers informed.

What is the most important change?

- ☐ A Add an instruction never to state availability without a successful lookup
- ☐ B On tool failure, state what is unavailable, offer what can still be done, and capture the request for follow-up
- ☐ C Retry the partner API more aggressively during outages
- ☐ D Cache recent availability so answers exist during an outage

◆ PAUSE AND DECIDE

What is the most important change?

- A

 Never state availability
A prohibition with no alternative. It still has a customer and a prompt telling it to be helpful.
- B

 Explain, offer, capture
Gives the agent a correct thing to do when the tool fails, so the interaction ends somewhere.
- C

 Retry harder
A ninety-minute outage does not respond to retries.
- D

 Cache availability
Cached availability during an outage is stale by definition. Confidently wrong, different source.

THE SIGNAL

Telling it not to do the wrong thing, without supplying a right thing, rarely works. Module 17.

An automated review runs on every pull request with `claude -p` and posts the result as a comment. On large pull requests the comment sometimes stops mid-sentence or covers only the first files changed.

The workflow reports success every time. The JSON output shows a stop reason indicating the turn cap was reached. The cap is set to 4.

What should they change?

- ☐ A Remove `--max-turns` so large reviews run to completion
- ☐ B Raise the cap, and have the workflow check the stop reason before posting
- ☐ C Switch to text output so the comment reads more naturally
- ☐ D Split large pull requests into smaller ones

◆ PAUSE AND DECIDE

What should they change?

- A

 Remove the cap
In print mode there is no default, so removing it means no ceiling at all.
- B

 Raise it, and check the stop reason
Fixes the truncation and the silent success separately, because they are two faults.
- C

 Text output
Discards the stop reason, which is the only reason they can diagnose this.
- D

 Smaller pull requests
Sensible practice, and a capped run is still posted as a complete review.

THE SIGNAL

Two faults: a cap too low, and a workflow publishing a capped run as finished. Module 11.

A contract analysis prompt ends with: "Report only risks you are highly confident about, to avoid overwhelming the legal team."

An audit finds material risks present in contracts where the model reported nothing. There is no record of what it considered and discarded.

Lawyers say they would rather see uncertain risks than miss real ones.

What should change?

- ☐ A Lower the bar in the prompt to "reasonably confident"
- ☐ B Report all risks with a confidence score and filter in the pipeline against a tuned threshold
- ☐ C Remove the confidence instruction so everything is reported
- ☐ D Ask the model to explain why it discarded low-confidence risks

◆ PAUSE AND DECIDE

What should change?

- A

 Lower the bar in words
*Moves an invisible threshold to a different invisible place.
Still no record.*
- B

 Score, then filter downstream
*Everything recorded, threshold is a number you can tune,
uncertain risks route to review.*
- C

 Report everything
*Solves the misses by creating the overwhelm the
instruction existed to prevent.*
- D

 Explain the discards
*Asks the model to report reliably on a suppression it
already made silently.*

THE SIGNAL

"No record of what it considered and discarded" is the sentence. Any answer leaving the filter in the model keeps the blind spot. Module 12.

An MCP server for meeting transcripts needs to produce short abstracts as part of one of its tools.

The team would prefer not to hold their own API key, not to manage their own rate limits, and not to make the model a configuration option customers must set.

Which MCP capability fits?

- ☐ A Elicitation, asking the user to supply an abstract
- ☐ B Sampling, requesting a completion from the host's model through the client
- ☐ C Roots, granting the server access to the model's context
- ☐ D A resource containing the transcript so the agent summarises it itself

◆ PAUSE AND DECIDE

Which MCP capability fits?

A Elicitation

Elicitation asks the user for input. Nobody wants to be asked to write the abstract.

C Roots

Roots are filesystem boundaries. This description is invented.

B Sampling

The server borrows the host's model: no key, no rate limits, no model configuration.

D A resource

Moves the work into the main conversation, which is the context cost the server was avoiding.

THE SIGNAL

Three constraints, one capability. No key, no rate limits, no model choice is the definition of sampling. Module 7.

Scoring exam 2

- ◆ **Compare domains with exam 1 before you compare totals.** The same domain twice is the finding.
- ◆ **A domain that improved after rewatching** means the method works. Repeat it for the next weak one.
- ◆ **A domain that missed twice despite rewatching** is not a coverage problem. Work the scenarios, not the slides.
- ◆ **A total that moved but with misses in new places** is normal variance at twenty questions.

ON THE EXAM

Forty questions is enough to see a pattern. It is still not enough to predict a score.

Which domain did you miss?

QUESTION NUMBERS	DOMAIN	MODULES TO REWATCH
1, 5, 8, 12, 15	D1 Agentic architecture	1 to 4
6, 10, 16, 20	D2 Tool design and MCP	5 to 7
2, 9, 14, 18	D3 Claude Code	8 to 11
3, 7, 13, 19	D4 Prompt and structured output	12 to 14
4, 11, 17	D5 Context and reliability	15 to 17

ON THE EXAM

Add these to your exam 1 tally. Two papers, same weak domain, means work the scenarios rather than rewatching.