PRACTICE PAPER 3 OF 3 · ALL DOMAINS

Practice Paper 3

The last seventeen questions, completing a full 53-question sitting. One echoes paper 1 with a single detail flipped. The paper ends by reading your trend across all three papers.

QUESTIONS 17 HOLD 9 s each MODULE 36 of 36

admin@sthithapragna.com

Before you start

- **Keep your paper 1 and paper 2 tallies** by domain to hand.
- **Try it as a timed block:** about 38 minutes for 17 questions, exam pace.
- **One echo question:** find the detail that flips the answer.
- **Read the qualifier first** on every question: first, best, cheapest.
- **The final slides** combine all 53 questions into one picture.

> THE HABIT
>
> Exam pace is about 2 min 15 s per question. Most should take less.

> Every night, a service summarises the day's 30,000 recorded sales calls for a dashboard managers open at 9am.
>
> Today it streams each summary in real time, overnight, at full price.

What should the team change?

A Keep streaming, since streaming is the modern default

B Turn on fast mode to finish the night sooner

C Submit the night's calls as a Message Batch

D Add more real-time workers to finish earlier

◆ PAUSE AND DECIDE

What should the team change?

A Keep streaming, since streaming is the modern default
Streams to nobody.

B Turn on fast mode to finish the night sooner
Pays more for unneeded speed.

C Submit the night's calls as a Message Batch
Nobody waits: half price.

D Add more real-time workers to finish earlier
Full price, more rate-limit pressure.

THE SIGNAL

The echo: nobody is waiting, so the answer flips from streaming to batch.

> A developer's chatbot forgets everything said earlier in a conversation.
>
> The code sends only the user's latest message with each API call.

What is the fix?

A. Send the conversation history with every request

B. Move the chatbot to a model with a larger window

C. Raise max_tokens so the model remembers more

D. Turn on prompt caching so earlier turns are stored

◆ PAUSE AND DECIDE

What is the fix?

A Send the conversation history with every request

The API is stateless.

B Move the chatbot to a model with a larger window

An empty window, made bigger.

C Raise max_tokens so the model remembers more

Output length is not memory.

D Turn on prompt caching so earlier turns are stored

A discount, not storage.

THE SIGNAL

The Messages API is stateless. The conversation is whatever you send.

> An internal agent built on the Claude Agent SDK runs in CI. It should only read files and run the test suite.
>
> Nobody is present to approve anything, and it must never edit code.

Which two settings fit? Choose two.

A allowedTools listing only the read and test tools

B The dontAsk permission mode, so anything else is denied

C bypassPermissions, so the run never stalls

D A system prompt telling it never to edit files

E Running the agent as an administrator account

◆ PAUSE AND DECIDE

Which two settings fit? Choose two.

A allowedTools listing only the read and test tools
Pre-approves exactly what it needs.

B The dontAsk permission mode, so anything else is denied
Denies the rest, without prompting.

C bypassPermissions, so the run never stalls
Allows everything.

D A system prompt telling it never to edit files
Guidance, not a lock.

E Running the agent as an administrator account
More privilege, not less.

THE SIGNAL

Locked down and unattended: allow exactly what is needed, and deny everything else without asking.

> A developer forced JSON by prefilling the assistant's reply with an opening brace.
>
> After moving to Opus 5.5, every such request returns HTTP 400.

What should the developer do?

A Prefill with an opening square bracket instead

B Move back to an older model permanently

C Retry each failed request until one succeeds

D Structured outputs with a JSON schema instead

◆ PAUSE AND DECIDE

What should the developer do?

A Prefill with an opening square bracket instead
Still a prefill.

B Move back to an older model permanently
Permanent avoidance.

C Retry each failed request until one succeeds
Every retry fails the same way.

D Structured outputs with a JSON schema instead
The supported replacement.

THE SIGNAL

Prefill is rejected on current models. Structured outputs guarantee the shape instead.

> A team must replace a home-grown HTTP client for Claude, used in 60 files across three services, with the official SDK.
>
> The services are in production, and the client has no tests.

What is the best approach?

A Replace it in all sixty files in one release

B Characterisation tests, then small batches

C Rewrite all three services from scratch on the SDK

D Leave the old client in place indefinitely

◆ PAUSE AND DECIDE

What is the best approach?

A Replace it in all sixty files in one release
Big-bang change.

B Characterisation tests, then small batches
Pin behaviour, then small steps.

C Rewrite all three services from scratch on the SDK
Over-engineering.

D Leave the old client in place indefinitely
Permanent avoidance.

THE SIGNAL

Large refactors: pin current behaviour with tests first, then move in small, verified steps.

> Users have coaxed an assistant into revealing its system prompt.
>
> The prompt contains an internal API key that the assistant was told to use in links.

What is the right fix?

A Add a line telling it never to reveal the prompt

B Move the assistant to a more capable model

C Remove the key from the prompt, and revoke it

D Rotate the key weekly, but keep it in the prompt

◆ PAUSE AND DECIDE

What is the right fix?

A Add a line telling it never to reveal the prompt
Prompts can leak anyway.

B Move the assistant to a more capable model
Fewer leaks, not none.

C Remove the key from the prompt, and revoke it
Nothing secret left to leak.

D Rotate the key weekly, but keep it in the prompt
Still leaks, just more often.

THE SIGNAL

Assume the system prompt can leak. Secrets never belong in it.

A review assistant checks each pull request for security issues, style problems and missing tests.

The three checks are independent, and their findings are combined into one comment.

Which design fits best?

A One long sequential agent that does all three in turn

B A chain where each check's output feeds the next

C An autonomous agent that decides which checks to run

D Three focused calls in parallel, then aggregate

◆ PAUSE AND DECIDE

Which design fits best?

A One long sequential agent that does all three in turn

Slower, crowded context.

B A chain where each check's output feeds the next

No dependency to chain.

C An autonomous agent that decides which checks to run

Autonomy the task does not need.

D Three focused calls in parallel, then aggregate

Independent work, run in parallel.

THE SIGNAL

Independent subtasks with a combined result: parallelise, then aggregate.

> An app uses Anthropic's server-side web search tool. Some responses end with a stop reason of pause_turn.
>
> The app shows those responses to users as if they were finished.

Which two changes fit? Choose two.

A Treat pause_turn as an error and show an apology

B Turn off web search for the whole app

C Send the assistant content back to continue the turn

D Raise max_tokens so searches finish in one response

E Check the stop reason before showing any response

◆ PAUSE AND DECIDE

Which two changes fit? Choose two.

A Treat pause_turn as an error and show an apology

Gives up a resumable turn.

B Turn off web search for the whole app

Removes the feature.

C Send the assistant content back to continue the turn

The documented continuation.

D Raise max_tokens so searches finish in one response

A different limit entirely.

E Check the stop reason before showing any response

200 is not the same as finished.

THE SIGNAL

pause_turn means the server-tool loop paused. Send the content back, and check stop reasons first.

> Requests sharing a long cached prefix arrive about every 20 minutes during the working day.
>
> Usage shows the cache is written again on almost every request.

What should the team change?

A Use the one-hour cache duration for that prefix

B Send a dummy request every few minutes to keep it warm

C Turn prompt caching off, since it is not helping

D Move these requests to a larger model

◆ PAUSE AND DECIDE

What should the team change?

A Use the one-hour cache duration for that prefix
Outlives the gap between requests.

B Send a dummy request every few minutes to keep it warm
Pays for traffic nobody needs.

C Turn prompt caching off, since it is not helping
Gives up the saving.

D Move these requests to a larger model
Unrelated to cache expiry.

THE SIGNAL

Requests further apart than five minutes: choose the one-hour cache, written at a higher price, read cheaply.

> An assistant can create and delete calendar events for users. Creating is low risk.
>
> A few deletions last month removed the wrong meetings.

What should the developer change?

A Require the user to confirm every tool call

B Require confirmation for deletions only

C Remove the delete tool from the assistant

D Ask the assistant in its prompt to be careful

◆ PAUSE AND DECIDE

What should the developer change?

A Require the user to confirm every tool call
Approval fatigue.

B Require confirmation for deletions only
Approval sized to risk.

C Remove the delete tool from the assistant
Loses a needed capability.

D Ask the assistant in its prompt to be careful
Guidance, not a gate.

THE SIGNAL

Gate the risky, irreversible action, and leave the safe ones fast.

> A web front end calls the Claude API directly from the browser, with the organisation's API key in the JavaScript bundle.
>
> Anyone can read the bundle.

What should the team do?

A Obfuscate the key inside the JavaScript bundle

B Ask each user to paste in their own API key

C Rotate the exposed key every day

D Revoke it; call Claude from the backend only

◆ PAUSE AND DECIDE

What should the team do?

A Obfuscate the key inside the JavaScript bundle
Still readable.

B Ask each user to paste in their own API key
Pushes it to users.

C Rotate the exposed key every day
Still exposed, daily.

D Revoke it; call Claude from the backend only
Keys never reach the browser.

THE SIGNAL

A key in the browser is a published key. Calls go through your backend, which holds the key.

> A team changed five things in a prompt at once: wording, order, two examples and an output format.
>
> Eval scores rose on some cases and fell on others, and nobody can tell why.

What should the team do?

A Keep every change, since the average went up

B Revert, then one change at a time on evals

C Add more changes until every case improves

D Move to a larger model to smooth out the results

◆ PAUSE AND DECIDE

What should the team do?

A Keep every change, since the average went up
Ships regressions blind.

B Revert, then one change at a time on evals
Attributable changes.

C Add more changes until every case improves
More confusion.

D Move to a larger model to smooth out the results
Swap the model.

THE SIGNAL

Iterative refinement: one change at a time, measured on the eval set.

> A team wants Claude Code to review every pull request in CI and post its findings in a fixed format the pipeline can parse.
>
> It must not change any code.

How should the team run it?

A Headless with -p, a JSON schema, and read-only tools

B An interactive session left open on a shared server

C A session with write access so it can fix issues too

D Developers paste each diff into a chat window

◆ PAUSE AND DECIDE

How should the team run it?

A Headless with -p, a JSON schema, and read-only tools

Scriptable, parseable, read-only.

B An interactive session left open on a shared server

Not automatable.

C A session with write access so it can fix issues too

Breaks the must-not.

D Developers paste each diff into a chat window

Manual, inconsistent.

THE SIGNAL

CI means headless: -p, structured output, and only the tools the job needs.

> A policy assistant answers confidently about rules that do not appear in the documents it is given.
>
> Staff have acted on two invented answers.

What should the developer do first?

A Lower the temperature to stop invented answers

B Move the assistant to the largest model

C Answer only from the documents, with quotes

D Fine-tune a model on the policy documents

◆ PAUSE AND DECIDE

What should the developer do first?

A Lower the temperature to stop invented answers

A dial, now rejected.

B Move the assistant to the largest model

Fewer, not none.

C Answer only from the documents, with quotes

Anthropic's grounding guidance.

D Fine-tune a model on the policy documents

Train a custom model.

THE SIGNAL

Hallucination is fluent prediction. Ground it in sources, require quotes, and permit uncertainty.

> A team runs twelve Claude features in production, on several models.
>
> It has no record of which feature uses which model, and no alerts beyond total spend.

Which two practices should it add? Choose two.

A. A model inventory with owners, checked against deprecations

B. Never pin a model, so features always get the newest

C. Upgrade every feature on the day a model launches

D. Delete request logs daily to keep storage small

E. Alerts per feature on errors, latency and cost

◆ PAUSE AND DECIDE

Which two practices should it add? Choose two.

A A model inventory with owners, checked against deprecations

Know what retires, and who acts.

B Never pin a model, so features always get the newest

Untested changes by default.

C Upgrade every feature on the day a model launches

Big-bang upgrades.

D Delete request logs daily to keep storage small

Removes the evidence.

E Alerts per feature on errors, latency and cost

Spot problems per feature.

THE SIGNAL

Operate what you run: know your models and owners, and watch each feature.

> Every file an agent edits must be formatted with the team's formatter.
>
> The instruction is in the agent's system prompt, and it sometimes forgets.

What is the best control?

A Repeat the instruction at the end of every user turn

B Ask reviewers to reformat files during code review

C Run the formatter in CI only, after the pull request

D A post-tool-use hook that formats each edited file

◆ PAUSE AND DECIDE

What is the best control?

A Repeat the instruction at the end of every user turn

Louder guidance.

B Ask reviewers to reformat files during code review

Pushes it to people.

C Run the formatter in CI only, after the pull request

Too late in the loop.

D A post-tool-use hook that formats each edited file

Deterministic, every edit.

THE SIGNAL

Every time, without exception, is a job for a hook, not an instruction.

> Four downstream services parse a triage feature's JSON output. The team wants to add a field and rename another.
>
> The services are owned by other teams and release on their own schedules.

What is the safest approach?

A Rename and add both now, and tell the other teams

B Keep a separate schema for each consumer forever

C Add the field; keep the old name for a while

D Stop validating the output so changes cannot break it

◆ PAUSE AND DECIDE

What is the safest approach?

A Rename and add both now, and tell the other teams

Breaks consumers on release.

B Keep a separate schema for each consumer forever

Four schemas to maintain.

C Add the field; keep the old name for a while

Backward-compatible change.

D Stop validating the output so changes cannot break it

Removes the control.

THE SIGNAL

Contracts change additively first. Removals wait until consumers have moved.

Answer key, paper 3

QUESTIONS 1-6	QUESTIONS 7-12	QUESTIONS 13-17
1: C (D2)	7: D (D1)	13: A (D3)
2: A (D5)	8: C+E (D2)	14: C (D5)
3: A+B (D1)	9: A (D5)	15: A+E (D2)
4: D (D6)	10: B (D8)	16: D (D1)
5: B (D2)	11: D (D2)	17: C (D2)
6: C (D7)	12: B (D6)	

SCORING

Choose-two items count only when both letters match.

All 53 questions, by domain

DOMAIN	EXAM WEIGHT	QUESTIONS ACROSS PAPERS
D2 Applications and Integration	33.1%	18
D5 Model Selection	16.8%	9
D1 Agents and Workflows	14.7%	8
D6 Prompt and Context	11.0%	6
D8 Tools and MCPs	10.6%	5
D7 Security and Safety	8.1%	4
D3 Claude Code / D4 Debugging	3.1% / 2.6%	2 / 1

ON THE EXAM

Weight revision by exam weight times misses. A D2 gap costs most.

Reading your trend

- **Misses falling from paper 1 to paper 3:** your revision is working. Keep going.
- **The same domain missing in all three:** rewatch it, then retake only its questions.
- **Misses spread thinly:** look for one rule or trap behind them, often the partial answer.
- **Both echo questions right:** you read constraints, not keywords.
- **About 45 of 53 or better** is a strong sign. It is a guide, not the scaled score.

> THE HABIT
>
> Revising what you already know feels productive and moves nothing. Spend time where you miss.

The last week before the exam

- **Rewatch your weakest domain,** then module 31 (confusion pairs) and module 33 (question craft).
- **Re-read module 32** and check the release notes for anything newer.
- **Retake the papers a week apart,** not back to back, so they test recall.
- **On the day:** read the ask first, find the qualifier, name each trap.
- **Trust the ten rules.** They break most ties the exam sets.

> FINALLY
>
> Good luck. Build carefully, test honestly, and let the evidence decide.