

MODULE 32 · EXAM STRATEGY - ALL DOMAINS

# The naming migration

What got renamed, what got replaced, and which old material is merely old-fashioned rather than wrong.

COVERS All 7 domains    TYPE Revision    MODULE 32 of 36

Three categories, and only one is safe.  
**Renamed but unchanged. Renamed and  
changed. Replaced entirely.**

# The four platform renames

| OLD NAME                    | CURRENT NAME                         |
|-----------------------------|--------------------------------------|
| Databricks Asset Bundles    | Declarative Automation Bundles       |
| Databricks Repos            | Databricks Git Folders               |
| Delta Live Tables (DLT)     | Lakeflow Spark Declarative Pipelines |
| Databricks Workflows / Jobs | Lakeflow Jobs                        |

## ON THE EXAM

All four are in the safe category. Existing code, YAML and CLI commands run unchanged, and the DAB acronym was deliberately kept.

# Why it happened

- ◆ **Lakeflow is an umbrella** covering ingestion, transformation and orchestration as one product family.
- ◆ **Lakeflow Connect** for ingestion, **Lakeflow Spark Declarative Pipelines** for transformation, **Lakeflow Jobs** for orchestration.
- ◆ **The old names grew separately**, so Workflows and Delta Live Tables sounded unrelated despite doing adjacent jobs.
- ◆ **Knowing the family helps you guess** when you meet a Lakeflow name you have not seen.

## ON THE EXAM

Traces of the old names survive in places: classic SKUs still begin with DLT, and pipeline event log schemas still use dlt.

# Code-level renames

| OLD                     | CURRENT                             |
|-------------------------|-------------------------------------|
| import dlt              | from pyspark import pipelines as dp |
| @dlt.expect             | @dp.expect                          |
| @view                   | @temporary_view                     |
| APPLY CHANGES INTO      | AUTO CDC INTO                       |
| hasTag / hasTagValue    | has_tag / has_tag_value             |
| Single user access mode | Dedicated access mode               |
| Shared access mode      | Standard access mode                |

## ON THE EXAM

Every old form here still works. Recognise both, because the exam may show you either.

# One code change that is not just a rename

- ◆ `@table` now creates streaming tables.
- ◆ A new `@materialized_view` decorator creates materialised views.
- ◆ So the decorator's meaning shifted, not merely its prefix.
- ◆ Old material showing `@dlt.table` for a materialised view is describing different behaviour, not just an old name.

## ON THE EXAM

This is the clearest example of the second category. The name changed and so did what it does.

The exam guide itself uses old and new names, sometimes in the same sentence. **You must recognise both, because the guide does not choose.**

# Where the guide uses which

- ◆ **Databricks Repos** — the old name, in the Domain 5 development workflow objective.
- ◆ **Automation Bundle and Declarative Automation Bundles** — both, each with 'formerly Databricks Asset Bundles'.
- ◆ **Lakeflow Jobs** — the new name, used consistently across Domain 4.
- ◆ **Lakeflow Spark Declarative Pipelines** — the new name.

## ON THE EXAM

Note the pattern: where the guide expects confusion it supplies 'formerly X' itself. That is a hint about what it thinks is unfamiliar.

# Renamed, and nothing else changed

| TOPIC                     | OLD MATERIAL IS                                             |
|---------------------------|-------------------------------------------------------------|
| Bundles / Asset Bundles   | Still correct. YAML, CLI and the DAB acronym are unchanged. |
| Git Folders / Repos       | Still correct on workflow, branching and pull requests.     |
| Lakeflow Jobs / Workflows | Still correct on tasks, dependencies and the DAG.           |
| Dedicated / single user   | Still correct on when you need it: Runtime ML, GPU, RDD, R. |
| Standard / shared         | Still correct on Unity Catalog support and isolation.       |

## ON THE EXAM

For everything in this table, translate the vocabulary and trust the content.

# Actually changed — old material is wrong

| TOPIC                    | WHAT CHANGED                                                                   |
|--------------------------|--------------------------------------------------------------------------------|
| Failed casts             | ANSI mode is on by default from DBR 17. cast throws instead of returning null. |
| @table in pipelines      | Now creates streaming tables. Materialised views need @materialized_view.      |
| DBFS root and mounts     | Deprecated. New accounts are provisioned without access.                       |
| Personal access tokens   | Labelled legacy. OAuth M2M is the recommendation for CI/CD.                    |
| Table maintenance        | Predictive optimisation runs OPTIMIZE, VACUUM and ANALYZE automatically.       |
| Partitioning and Z-ORDER | Superseded by liquid clustering, which PO is built around.                     |

## ON THE EXAM

This is the short, dangerous list. Old material on these topics is not dated vocabulary — it is incorrect behaviour.

# Replaced rather than renamed

- ◆ **Hive metastore table access control** still exists, and Unity Catalog is a different model with no DENY.
- ◆ **Z-ORDER** still works, and liquid clustering is what predictive optimisation is designed around.
- ◆ **Directory listing** still works in Auto Loader, and file events is the recommendation.
- ◆ **cloudFiles.useIncrementalListing** is deprecated from Databricks Runtime 17.3 LTS.

## ON THE EXAM

The pattern here is that the old thing keeps working. A question offering it is not automatically wrong — it may just be the worse answer.

# What did not change

- ◆ **Delta Lake, the transaction log, and time travel.**
- ◆ **The medallion architecture**, and everything about bronze, silver and gold.
- ◆ **Spark itself:** lazy evaluation, wide and narrow transformations, shuffles, skew and spill.
- ◆ **The three-level namespace** and Unity Catalog's core object model.

## ON THE EXAM

Guard against over-correcting. Most of what you learnt from older material is still exactly right.

# Quick translation

| IF YOU READ                          | READ IT AS                                       |
|--------------------------------------|--------------------------------------------------|
| DLT, Delta Live Tables               | Lakeflow Spark Declarative Pipelines             |
| Workflows                            | Lakeflow Jobs                                    |
| Repos                                | Git Folders                                      |
| Asset Bundles, DABs                  | Declarative Automation Bundles                   |
| Single user / shared access mode     | Dedicated / standard                             |
| APPLY CHANGES INTO                   | AUTO CDC INTO                                    |
| <code>import dlt, @dlt.expect</code> | <code>pyspark.pipelines as dp, @dp.expect</code> |

## ON THE EXAM

Seven translations. Read them in both directions, because the exam may use either side.

A candidate is working through a 2024 question bank. One question describes configuring a Delta Live Tables pipeline with a target schema, a pipeline mode, and expectations, and asks which settings are required.

None of the terminology matches the current exam guide.

### How should they treat this question?

- ☐ A Discard it; the terminology is out of date so the content cannot be trusted
- ☐ B Translate DLT to Lakeflow Spark Declarative Pipelines and treat the content as still valid
- ☐ C Discard it; pipelines were removed from the current exam
- ☐ D Treat it as valid without translation, since the exam still uses DLT terminology

---

#### ◆ PAUSE AND DECIDE

## How should they treat this question?

- A

 Discard it entirely  
*Over-correction. The rename was non-breaking, so the pipeline settings it describes are still accurate.*
- B

 Translate the name, trust the content  
*Correct. This is the safe category: renamed, nothing else changed. Existing pipeline code runs unchanged.*
- C

 Pipelines were removed from the exam  
*They were de-emphasised — the old exam had a whole pipelines section, and the current guide folds expectations into one Domain 3 objective. Not removed.*
- D

 No translation needed  
*The current guide uses the new name. Recognising both is required, but the exam will not present it as DLT.*

### THE SIGNAL

The rename was non-breaking and existing code runs unchanged. That is the definition of the safe category.

A practice question shows a code sample using `@dlt.table` to define what it calls a materialized view, and asks what the decorator produces.

The candidate knows about the dlt to dp rename and translates it mentally to `@dp.table`.

### Is the translation sufficient?

- ☐ A Yes; `@dp.table` is simply the new spelling of `@dlt.table`
- ☐ B No; `@table` now creates streaming tables, and materialized views need `@materialized_view`
- ☐ C No; the decorator was removed entirely
- ☐ D Yes, but only on Databricks Runtime 17 and above

### ◆ PAUSE AND DECIDE

## Is the translation sufficient?

- A

Just a new spelling

*The prefix change is a rename; the decorator's meaning also changed, which a spelling translation does not capture.*
- C

Removed entirely

*It exists and is widely used. What changed is what it produces.*
- B

@table now creates streaming tables

*Correct. This is the second category – renamed and changed. A new @materialized\_view decorator exists for materialised views.*
- D

Sufficient on DBR 17+

*Invents a version condition. The decorator's behaviour is not gated on runtime version in that way.*

### THE SIGNAL

He did the translation correctly and it was not enough. A correct rename can still leave you with wrong behaviour.

A study guide from 2024 states: 'If a value cannot be converted, CAST returns NULL rather than failing, so casting is safe to use on untrusted input.' The candidate has no reason to doubt it.

Nothing in the sentence uses an outdated product name.

### What is wrong here?

- ☐ A Nothing; this is still correct behaviour
- ☐ B ANSI mode is on by default from DBR 17, so cast throws an error; try\_cast returns null in both modes
- ☐ C CAST was renamed to TRY\_CAST
- ☐ D The statement is only wrong on serverless compute

### ◆ PAUSE AND DECIDE

## What is wrong here?

A Still correct

*It was correct for years, which is why it is dangerous. From DBR 17 the default changed and cast raises CAST\_INVALID\_INPUT.*

C CAST was renamed

*Both exist and always did. try\_cast is a separate function, not a new name for cast.*

B ANSI mode changed the default

*Correct. try\_cast is the function that returns null in both modes, and is the safe choice on untrusted input.*

D Only wrong on serverless

*The default follows the runtime, not the compute type. DBR 17 and above behaves this way generally.*

### THE SIGNAL

No outdated name appears anywhere in the sentence. That is precisely why this category is the dangerous one.

Translate the vocabulary, then ask the harder question: **has the behaviour changed too?** An old name is a warning. A familiar name is not a guarantee.