

MODULE 33 · EXAM STRATEGY - ALL DOMAINS

How the exam is written

The format, the time budget, how scenarios encode their answers, and the guide's own sample question worked through.

COVERS Exam technique TYPE Revision MODULE 33 of 36

The format

	VALUE
Questions	45 scored
Time	90 minutes
Options	Four, A to D
Cost	USD 200
Pass mark	Not published
Validity	Two years
Domains	Seven, weighted

ON THE EXAM

Four options is the current format. A five-option question is from the previous version, which is a useful way to date any practice material.

Where the marks are

DOMAIN	WEIGHT	ROUGHLY
Transform (Domain 3)	22%	10 questions
Ingest (Domain 2)	21%	9 questions
Jobs (Domain 4)	16%	7 questions
Governance (Domain 7)	15%	7 questions
CI/CD (Domain 5)	10%	4–5 questions
Troubleshoot (Domain 6)	10%	4–5 questions
Platform (Domain 1)	6%	3 questions

ON THE EXAM

Domains 2 and 3 together are 43% — nearly nineteen questions. That is where preparation time earns the most.

The time budget

- ◆ **Ninety minutes for forty-five questions is two minutes each.**
- ◆ **Most questions take well under that.** A definition or a single-fact question takes twenty seconds.
- ◆ **Scenario questions are where the time goes,** and they are the ones worth spending it on.
- ◆ **Two minutes is an average, not an allowance.** Bank time on the short ones.

ON THE EXAM

If you are still on question 20 at the halfway mark, that is a signal to speed up rather than a crisis.

A well-written scenario contains **one detail that decides the answer**. Your job is to find it, not to evaluate all four options equally.

Where the deciding detail hides

- ◆ **In a number.** 18–24 minute runtime on a 15-minute schedule. 37 minutes against a 28-second percentile.
- ◆ **In a negation.** 'has never failed', 'nothing upstream changed', 'no spill is reported'.
- ◆ **In a timeframe.** 'three weeks later', 'intermittently', 'has worked for two years'.
- ◆ **In a constraint you might skim.** 'must not stop', 'the row count must be correct', 'they must not reach the base tables'.

ON THE EXAM

Numbers and negations are the two richest sources. A stated negative is almost always there to eliminate an option.

Words that carry weight

PHRASE	WHAT IT USUALLY MEANS
'intermittently'	A race condition or a timing dependency, not a logic error
'has worked for years'	Something outside the code changed, or it was always quietly wrong
'to be safe' / 'for margin'	A guess is being made where a mechanism exists
'must not' / 'must never'	A hard requirement. Any option that risks it is out.
'the count must be correct'	Rules out anything that removes rows
'nobody was alerted'	Notification design, not necessarily a failure
'no error was raised'	A silent failure — look for a same-type or positional mistake

ON THE EXAM

These are not exam tricks. They are how a competent question-writer signals which mechanism is being tested.

Count the requirements

- ◆ **Many scenarios state two or three requirements**, spread across the paragraphs.
- ◆ **Test each option against every requirement**, not against the one you noticed first.
- ◆ **A distractor usually satisfies all but one.** That is what makes it tempting.
- ◆ **If an option fails any stated requirement, it is out**, however good it looks otherwise.

ON THE EXAM

Sensor readings that must not reach gold, a dashboard that must not stop, and a fault that must be counted. One action does all three.

The four shapes of wrong answer

SHAPE	HOW TO SPOT IT
The other member of a pair	Plausible, related, wrong on one specific detail
The pair stated backwards	Names both halves and swaps them
An invented limit or rule	'Not supported above', 'requires a minimum', 'must be re-applied'
True but irrelevant	A correct statement that does not answer what was asked

ON THE EXAM

Module 31 covered the first three. The fourth is the hardest, because checking whether a statement is true is not the same as checking whether it answers the question.

Sometimes more than one option is correct. The question is then asking which is **best**, and best usually means **simplest, cheapest, or most automatic.**

The guide's sample questions

- ◆ **The exam guide contains five sample questions**, each labelled with the objective it aligns to.
- ◆ **They are retired from a previous version of the exam.** The guide says so explicitly.
- ◆ **But the objectives they are labelled with are the current ones**, so Databricks has matched them to the new outline.
- ◆ **They are not one per domain.** Two of the five are Domain 6.

ON THE EXAM

Retired questions matched to current objectives is still the best available evidence of style, difficulty and distractor design.

A data engineer notices a batch job's duration has doubled after a new data source was onboarded. In the Spark UI, the longest stage shows that most tasks finish in under 30 seconds, but one task takes over 10 minutes.

The stage's task summary shows Min/Median shuffle read near 400 MB while Max shuffle read exceeds 5 GB.

Which solution reduces the job runtime?

- ☐ A Increase cluster size to add more executors so the slow task finishes faster
- ☐ B Confirm adaptive query execution with skew join handling is active to automatically split the oversized partition at runtime
- ☐ C Reduce `spark.sql.shuffle.partitions` to coalesce more work into fewer tasks
- ☐ D Manually repartition the dataset using a salt key before the join to distribute skewed keys evenly

◆ PAUSE AND DECIDE

Which solution reduces the job runtime?

- A

 Add more executors
More idle capacity. One task holds the stage open, so the other executors already have nothing to do.
- B

 Confirm AQE skew join handling is active
Correct. AQE splits the oversized partition at runtime, it is on by default, and confirming it costs nothing.
- C

 Reduce shuffle.partitions
Fewer, larger tasks. It makes an over-large partition worse, and AQE coalesces downwards anyway.
- D

 Salt the key manually
Genuinely fixes skew, and it is code to write, test and maintain. Ask this only after AQE has failed to help.

THE SIGNAL

Two signals, both from module 24: max duration far above the median, and max shuffle read far above the median. Skew, twice over.

What that question tells you

- ◆ **It gives the same signal twice**, through duration and through shuffle read per task.
- ◆ **Two of the four options are wrong for reasons you can state**, not merely wrong.
- ◆ **The fourth option actually works**. The question is asking which is best.
- ◆ **It names a specific mechanism**, AQE skew join handling, rather than describing a symptom.

ON THE EXAM

Four scenarios in this course were built on the same discriminator. That was not a coincidence; it is the shape the exam favours.

When you genuinely do not know

- ◆ **Eliminate first.** Invented limits and true-but-irrelevant options can be removed without knowing the answer.
- ◆ **Prefer the option naming a mechanism** over the one restating the problem.
- ◆ **Be wary of absolutes** in options: 'always', 'never', 'the only way'.
- ◆ **Answer everything.** There is no penalty for a wrong answer, and a blank is certainly wrong.

ON THE EXAM

Eliminating two options and guessing gives you a coin flip on a question you would otherwise have lost entirely.

On the day

- ◆ **Read the whole question before the options.** Options are designed to steer you.
- ◆ **Flag and move on** rather than stalling. A question you return to often reads differently.
- ◆ **Do not change an answer without a reason you can state.** A vague second-guess is usually worse.
- ◆ **Check the halfway mark.** Around question 22 at 45 minutes means you are on pace.

ON THE EXAM

The most common time mistake is not going too slowly. It is stalling on one hard question early and losing composure.

A daily job must load a supplier file that arrives between 02:00 and 09:00. The data must be available as early as possible. The job must not run when no file has arrived, because empty runs are being billed. And the operations team must be alerted if no file arrives by 10:00.

Four designs are proposed.

Which design meets every stated requirement?

- | | |
|--|--|
| <div data-bbox="160 711 950 768">A A file arrival trigger on the volume</div> | <div data-bbox="1302 711 2356 882">B A file arrival trigger, plus a separate scheduled job at 10:00 that alerts if the table was not updated</div> |
| <div data-bbox="160 911 1223 1025">C A schedule every 30 minutes from 02:00, exiting early when no file is present</div> | <div data-bbox="1302 911 2356 1025">D A schedule at 09:15, with a failure notification configured</div> |

◆ PAUSE AND DECIDE

Which design meets every stated requirement?

- | | |
|---|--|
| <p>A A file arrival trigger alone</p> <p><i>Meets the first two requirements well and provides no alert. A trigger that never fires produces no run and therefore no notification.</i></p> | <p>B File arrival trigger plus a 10:00 check job</p> <p><i>Correct. Early availability, no empty runs, and an alert on absence — which needs a separate scheduled check, because absence is not an event.</i></p> |
| <p>C Poll every 30 minutes</p> <p><i>Fourteen runs a day, most of them empty, which fails the billing requirement explicitly.</i></p> | <p>D Schedule at 09:15</p> <p><i>Fails 'as early as possible' on every day the file arrives at 02:00.</i></p> |

THE TECHNIQUE

Option A is the trap, and it is the right answer to the first two requirements. The third one is what it cannot do.

A nightly job's gold aggregation has slowed from 12 minutes to 40. Investigation shows the silver table it reads has grown to several thousand small files, written by an hourly append job.

The table is a Unity Catalog managed table.

What is the best first action?

- ☐ A Schedule a nightly OPTIMIZE job on the silver table
- ☐ B Verify predictive optimisation is enabled on the account, since it runs OPTIMIZE on managed tables automatically
- ☐ C Rewrite the hourly job to append less frequently
- ☐ D Increase the gold job's cluster size

◆ PAUSE AND DECIDE

What is the best first action?

- A

 Schedule a nightly OPTIMIZE
Genuinely fixes it, and Databricks recommends disabling scheduled OPTIMIZE jobs where predictive optimisation runs. You would be building what may already exist.
- B

 Verify predictive optimisation is enabled
Correct. The table is UC managed, PO runs OPTIMIZE automatically, and checking costs nothing before you build anything.
- C

 Append less frequently
Reduces file creation and changes the data's freshness to solve a maintenance problem. A large change for a side effect.
- D

 A bigger cluster
Pays permanently to read a badly laid-out table faster, rather than fixing the layout.

THE TECHNIQUE

'Unity Catalog managed table' is placed deliberately. It is the fact that makes B available and A redundant.

Find the deciding detail. Count the requirements. **When two options both work, choose the one that needs no new machinery.**