MODULE 37 · FOUNDATIONS AND EXAM SKILLS - ALL DOMAINS

Practice exam 1

Twenty questions, weighted like the real exam, in mixed order. Nine seconds to answer each one.

DOMAIN Practice exam WEIGHT - MODULE 37 of 39

How to take this paper

- Answer **every** question, even when unsure. A guess costs nothing.
- Commit before the explanation appears. The nine-second hold is deliberate.
- Write down A, B, C or D. Score yourself at the end, by domain.
- Note **why** you chose, not just what. Wrong reasoning with a right answer is still a gap.
- Below 15 out of 20 means restudy the weak domain before the next paper.

> ON THE EXAM
>
> Treat it as the real thing: no pausing to look things up, and no going back to change answers.

> A developer must build and test agents in an existing Foundry project, including running evaluations and uploading files. They should not be able to create projects or assign roles to others.
>
> Your security team requires least privilege.

Which role should you assign, at project scope?

A Foundry Project Manager, so they can manage the project they work in.

B Foundry User, which grants project-level data actions.

C Azure AI Developer, the role named for developers.

D Cognitive Services User, which covers inference calls.

◆ PAUSE AND DECIDE

Which role should you assign, at project scope?

A Foundry Project Manager.

Can manage projects and assign Foundry User: more than least privilege here.

B Foundry User.

Microsoft's guidance for day-to-day development: project data actions, no management.

C Azure AI Developer.

Microsoft says not to use it for Foundry projects, despite the name.

D Cognitive Services User.

A legacy service role, not the Foundry project data role.

ON THE EXAM

Domain 1. Foundry User for building; Project Manager to manage. The name 'Developer' is a trap.

> A European bank must keep all inference data processing inside the EU. Traffic is variable, with unpredictable peaks, and the team wants to pay only for what they use.
>
> Latency requirements are ordinary.

Which deployment type fits?

A Global Standard, for the widest capacity and lowest price.

B Provisioned, to guarantee throughput during the peaks.

C Batch, which processes work at half price overnight.

D Data Zone Standard, keeping processing in the zone.

◆ PAUSE AND DECIDE

Which deployment type fits?

A Global Standard.

Cheapest and widest, but processing can happen anywhere.

B Provisioned.

Reserved capacity billed whether used or not: the opposite of pay-per-use.

C Batch.

Asynchronous and cheap, but this is interactive traffic.

D Data Zone Standard.

Pay-per-token, with processing kept inside the EU data zone.

ON THE EXAM

Domain 1. Residency plus variable traffic plus pay-per-use points at a data zone standard deployment.

A retailer runs a provisioned deployment sized for normal traffic. During sales events, requests exceed the provisioned capacity and users receive 429 errors.

The team wants overflow handled automatically, without rewriting the application.

What should they configure?

A Spillover to a standard deployment in the same resource.

B A second provisioned deployment, sized for peak traffic.

C Client-side retries with exponential backoff and a longer timeout.

D A quota increase request submitted before each sales event.

◆ PAUSE AND DECIDE

What should they configure?

A Spillover.

Overflow goes to a standard deployment in the same resource, automatically.

B Bigger provisioned deployment.

Pays peak prices all year for a few days of traffic.

C Retries.

Helps briefly, but capacity is genuinely exhausted during the peak.

D Quota requests.

Manual, slow, and not automatic overflow.

ON THE EXAM

Domain 1. Spillover is the overflow hotel: provisioned fills, standard takes the rest.

> A support assistant must block responses with medium or high severity violent content, while still allowing low severity mentions, such as describing an accident report.
>
> You are configuring the content filter for the deployment.

How should you set the violence threshold?

A Low, which blocks the lowest severity and everything above it.

B High, so that only the most severe content is permitted through.

C Medium, the lowest severity that should be blocked.

D Off, with a custom blocklist listing violent terms instead.

◆ PAUSE AND DECIDE

How should you set the violence threshold?

A Low.

Blocks low, medium and high, so accident reports are blocked too.

B High.

Blocks only high, letting medium severity content through.

C Medium.

The threshold is the lowest severity blocked, so medium and high are blocked.

D Off plus blocklist.

Turning filters off needs approval, and exact terms can't judge severity.

ON THE EXAM

Domain 1. The threshold names the lowest severity that gets blocked, not the highest allowed.

> A deployed RAG assistant shows healthy latency, low error rates and stable token usage on its dashboards. Users report that answers are increasingly wrong, citing documents that don't support them.
>
> Leadership wants early warning of this in future.

What should the team add?

A Alerts on token usage, to detect unusual generation volumes.

B More detailed platform metrics, sampled at a higher frequency.

C A larger model deployment, to improve answer quality overall.

D Continuous evaluation of groundedness on live traffic.

◆ PAUSE AND DECIDE

What should the team add?

A Token alerts.

Measures volume, which is already stable, not answer quality.

B More platform metrics.

The platform layer is healthy. Bad answers live a layer above it.

C Bigger model.

An expensive guess with no measurement to confirm improvement.

D Continuous evaluation of groundedness.

Scores live answers against their retrieved context, catching ungrounded claims.

ON THE EXAM

Domain 1. Healthy platform metrics can hide bad answers. The answers layer needs its own measurement.

> A government customer requires that no traffic between their Foundry agents and their storage and search resources traverses the public internet. Agents must also reach an internal API on a private network.
>
> They are designing the deployment now.

What should they use?

A The standard agent setup, with VNet injection.

B The basic agent setup, with a firewall rule allowing only their IP range.

C Public endpoints secured with Entra ID and conditional access policies.

D A separate Foundry resource per environment, each with its own keys.

◆ PAUSE AND DECIDE

What should they use?

A Standard setup, VNet injection.

Agent traffic runs in your network, with private endpoints to dependencies.

B Basic setup with firewall rules.

Basic doesn't support private networking for dependencies.

C Public endpoints plus Entra.

Strong identity, but traffic still crosses the public internet.

D Resource per environment, keys.

Separation isn't isolation, and keys move the design backwards.

ON THE EXAM

Domain 1. Identity and network are separate axes. This requirement is network, so it needs the standard setup.

> A compliance assistant answers questions about regulations that are amended most months. Answers must always reflect the current text, and cite the clause they came from.
>
> The team has a large archive of past regulation versions.

What should they build?

A Fine-tune a model monthly on the amended regulation text.

B Fine-tune once on the full archive, including every past version.

C Retrieve the current regulations at request time, with citations.

D Put the current regulations in the system message of each request.

◆ PAUSE AND DECIDE

What should they build?

A Monthly fine-tuning.
Fine-tuning teaches behaviour, not reliable current facts, and can't cite.

B Fine-tune on the archive.
Mixes superseded versions into the model's behaviour.

C Retrieval with citations.
Current text at request time, with clause-level citations.

D Everything in the system message.
Doesn't scale, and costs tokens on every request.

ON THE EXAM

Domain 2. Knowledge that changes belongs in retrieval; fine-tuning is for behaviour.

> An extraction service returns JSON for insurance claims. Some claims genuinely have no policy number, and the team finds the model inventing plausible ones.
>
> Downstream systems require a consistent JSON shape for every claim.

What should the schema do?

A Mark every field required, so the shape never varies between claims.

B Allow the policy number to be null, keeping strict validation.

C Drop structured outputs and parse the model's free text with regex.

D Lower the temperature to zero so the model stops inventing values.

◆ PAUSE AND DECIDE

What should the schema do?

A Everything required.

Forces a value where none exists: an invitation to fabricate.

B Nullable field, strict schema.

Shape stays consistent, and 'missing' becomes a valid answer.

C Regex over free text.

Abandons the guaranteed shape downstream systems need.

D Temperature zero.

More deterministic invention is still invention.

ON THE EXAM

Domain 2. A required field with no null option is an invitation to fabricate.

> A team switches a summarisation endpoint to a reasoning model. Requests now sometimes return status incomplete with no visible text, and the logs show high output token counts.
>
> The prompt and inputs are unchanged.

What is the most likely cause?

A The token cap is too low, and reasoning tokens consumed the budget.

B The model is rate limited, and partial responses are being returned.

C The content filter is removing the visible portion of the answer.

D Structured outputs are truncating the response before completion.

◆ PAUSE AND DECIDE

What is the most likely cause?

A Cap too low for reasoning tokens.

Reasoning tokens bill as output and can exhaust the cap before any visible text.

B Rate limiting.

Returns 429 errors, not incomplete responses with high token counts.

C Content filter.

Returns a filtered response with a reason, not silent incompleteness.

D Structured outputs.

Constrains shape; it doesn't consume the token budget invisibly.

ON THE EXAM

Domain 2. Reasoning tokens are invisible, billed as output, and share the cap with the answer.

> A procurement agent researches suppliers with search tools and can also submit purchase orders through an MCP tool. Buyers must approve every purchase order.
>
> Research queries must not interrupt anyone.

How should approvals be configured?

A Require approval for every tool call the agent makes, for consistency.

B Require no approvals, and review submitted purchase orders afterwards.

C Instruct the agent, in its system message, to ask before submitting.

D Require approval only on the purchase order tool.

◆ PAUSE AND DECIDE

How should approvals be configured?

A Approve everything.

Interrupts research constantly; reviewers start approving without reading.

B Review afterwards.

The money has already moved by the time anyone looks.

C Instruction to ask.

A job description, not a gate. Nothing enforces it.

D Approval on the consequential tool.

Microsoft's guidance: scope gates to specific tool invocations.

ON THE EXAM

Domain 2. Scope approval to the consequential tool call; low-risk steps flow.

> A team enables tracing for an agent and sees spans for every call in Application Insights, but the prompts and responses are missing from the traces.
>
> They are debugging in a development environment and want to see message contents.

What should they do?

A Grant the Log Analytics Reader role to the developers' accounts.

B Enable message content recording, which is off by default.

C Switch from client-side tracing to server-side agent tracing.

D Increase the Application Insights sampling rate to capture more data.

◆ PAUSE AND DECIDE

What should they do?

A Log Analytics Reader.

Needed to read traces, but they can already see spans.

B Enable content recording.

Message content capture defaults to false; enable it in development only.

C Server-side tracing.

Changes where spans come from, not whether content is recorded.

D More sampling.

Captures more spans, still without their message content.

ON THE EXAM

Domain 2. Content capture is off by default, deliberately: prompts and responses often contain personal data.

> An agent with six tools frequently answers correctly but calls the wrong tool on the way, sometimes passing arguments that don't match the user's request.
>
> The team wants a metric that isolates this specific failure before they start changing instructions.

Which evaluator should they use?

A Groundedness, to check answers against the retrieved context.

B Task Adherence, to check the agent follows its instructions.

C Tool Call Accuracy, to check tool choice and arguments.

D Coherence, to check the answers read consistently.

◆ PAUSE AND DECIDE

Which evaluator should they use?

A Groundedness.
Measures support from context, not which tool was chosen.

B Task Adherence.
Broader. it may fall for the right reason or the wrong one.

C Tool Call Accuracy.
Scores whether the right tool was called with the right arguments.

D Coherence.
Measures readability, which isn't the problem here.

ON THE EXAM

Domain 2. Agent evaluators isolate agent failures: intent, tool calls, task adherence, completion.

A team has an existing multi-agent application written with Agent Framework, including custom handoff logic and its own state store. They want Foundry identity, scaling, tracing and versioned releases.

Rewriting the orchestration is out of scope.

How should it be deployed?

A As a prompt agent, with the handoff rules written into instructions.

B As connected agents, with one agent routing to the others.

C As a Foundry workflow, rebuilt in the visual designer.

D As a hosted agent, with the code packaged into a container.

◆ PAUSE AND DECIDE

How should it be deployed?

A Prompt agent.
Configuration only: it can't run custom orchestration code.

B Connected agents.
Not available in the new Foundry Agent Service.

C Rebuild as a workflow.
A rewrite, into a designer retiring on 1 December 2026.

D Hosted agent.
Runs their existing code with Foundry identity, scaling and tracing.

ON THE EXAM

Domain 2. Custom orchestration code means a hosted agent. Connected agents no longer exist.

A design team generates icons that must sit on coloured website backgrounds without a visible rectangle behind them. They currently request JPEG output and describe a transparent background in the prompt.

Every icon still arrives with a solid background.

What should they change?

A Set background to transparent and output format to PNG.

B Set output compression to zero so no background artefacts appear.

C Request WebP output, which supports transparency and small files.

D Describe the transparent background more precisely in the prompt.

◆ PAUSE AND DECIDE

What should they change?

A background transparent plus PNG.

The documented combination: transparency needs both the parameter and PNG.

B No compression.

A higher-quality JPEG still can't store transparency.

C WebP.

Supports transparency generally, but isn't supported in Azure OpenAI here.

D Better prompt wording.

Transparency is a parameter, not something a prompt switches on.

ON THE EXAM

Domain 3. Transparency is a parameter plus a format. Neither alone works, and WebP isn't available.

> Marketing has an approved eight-second product video whose framing, motion and transitions are exactly right. They now want the same video with a winter colour palette.
>
> Regenerating from an edited prompt keeps changing the framing.

What should they do?

A Generate four new videos and pick whichever framing is closest.

B Generate a new video using the first frame as an input reference.

C Remix the approved video by its ID, changing only the palette.

D Extend the duration to twelve seconds so more framing options appear.

◆ PAUSE AND DECIDE

What should they do?

A Four new generations.
Four fresh videos, none of them the approved one.

B First frame as reference.
Anchors the opening frame; motion and transitions are generated again.

C Remix by video ID.
Keeps framework, transitions and layout, changing one stated thing.

D Longer duration.
Changes length, not fidelity to the approved video.

ON THE EXAM

Domain 3. Remix edits a finished video; a reference image only anchors a new one's first frame.

A bank must share call transcripts with an internal analytics team. Names, account numbers and addresses must be removed, but analysts need to see what type of information was removed at each position.

Transcripts are otherwise read as normal text.

Which redaction policy fits?

A characterMask, replacing each value with a repeated character.

B entityMask, replacing each value with its entity type label.

C syntheticReplacement, substituting realistic fake values.

D noMask, returning the detected entities without redacted text.

◆ PAUSE AND DECIDE

Which redaction policy fits?

A characterMask.
Hides values and preserves length, but hides the type as well.

B entityMask.
Each value becomes its type, so analysts see what was removed where.

C syntheticReplacement.
Realistic substitutes look like real data, hiding that anything was removed.

D noMask.
Returns no redacted transcript at all.

ON THE EXAM

Domain 4. Masks hide, labels explain, synthetic values keep text natural. The requirement names the middle one.

> A marketplace translates roughly forty million short product reviews a month into six languages, for search and browsing. Cost is the dominant constraint.
>
> Reviews are informal, and no specific tone or gender handling is required.

Which approach fits?

A Text translation with the standard NMT engine.

B Text translation with an LLM deployment, for the best quality.

C Adaptive custom translation, using reference sentence pairs.

D Document translation in batch mode, grouping reviews into files.

◆ PAUSE AND DECIDE

Which approach fits?

A NMT text translation.
Fast, per-character billing, built for exactly this volume.

B LLM translation.
Token billing and higher latency for controls this job doesn't need.

C Adaptive custom translation.
Adds terminology control, and needs an LLM deployment.

D Batch document translation.
For files that must keep their layout, not short strings.

ON THE EXAM

Domain 4. NMT for volume and cost; LLM translation when tone, gender or terminology matter.

> A legal team indexes 60,000 scanned contracts held as PDFs in Blob Storage. The indexer runs successfully, but searches return almost nothing, and indexed documents contain only file names and dates.
>
> The PDFs display correctly when opened.

What is missing?

A. A vectorizer on the index, so queries can be embedded at query time.

B. A deletion detection policy, so stale documents stop crowding results.

C. A larger search tier, because 60,000 documents exceed the current one.

D. OCR in the skillset, because scanned pages carry no text layer.

◆ PAUSE AND DECIDE

What is missing?

A A vectorizer.

Embeds queries, but there's no document text to match against yet.

B Deletion policy.

Removes deleted documents; it doesn't add missing content.

C Bigger tier.

Capacity isn't the issue: the indexer ran successfully.

D OCR in the skillset.

Scanned pages have no text layer, so without OCR only metadata is indexed.

ON THE EXAM

Domain 5. Retrieval can only find what ingestion put there. Scans need OCR or Content Understanding.

> A newsroom publishes breaking stories that must be searchable within seconds. Its Azure AI Search indexer currently runs on the most frequent schedule available.
>
> Editors still report delays of several minutes before new stories appear.

What should the team change?

A Add a second indexer running on an offset schedule to halve the gap.

B Enable the enrichment cache so indexing runs complete more quickly.

C Push documents to the index directly from the publishing system.

D Increase the batch size so each indexer run processes more stories.

◆ PAUSE AND DECIDE

What should the team change?

A Second offset indexer.

Halves a five-minute gap at best, and doubles the moving parts.

B Enrichment cache.

Saves reprocessing cost; it doesn't make new content appear sooner.

C Push model.

The publishing system sends documents as they're published: near-instant.

D Bigger batches.

Changes throughput per run, not how often runs happen.

ON THE EXAM

Domain 5. Indexers pull on a schedule, as often as every five minutes. Seconds means the push model.

> A skillset extracts organisations from each chunk of a report. The skill runs correctly, and a debug session shows the organisations node populated under each chunk.
>
> The index field for organisations is empty for every document.

What is most likely wrong?

A The skill's context is set to /document instead of the chunk collection.

B The output field mapping from the enriched node to the index is missing.

C The field mapping from the blob metadata to the index is misconfigured.

D The index field type is a string rather than a collection of strings.

◆ PAUSE AND DECIDE

What is most likely wrong?

A Wrong context.

Would change how often the skill runs; the debug session shows it running per chunk.

B Missing output field mapping.

Enrichments are temporary: only mapped nodes reach the index.

C Field mapping.

Field mappings carry raw source fields, not enriched nodes.

D Field type.

A type mismatch usually errors; it doesn't silently empty every document.

ON THE EXAM

Domain 5. 'The skill ran but the field is empty' is almost always a missing output field mapping.

How did you do?

- **Answer key:** B D A C D A, C B A D B C D, A C, B A, D C B.
- **18 to 20:** exam-ready on this sample. Move to practice exam 2.
- **15 to 17:** solid. Restudy the domains where you lost marks, then continue.
- **Below 15:** revisit those domain modules before paper 2, not after.
- Check **why** you chose each answer. Right answer, wrong reason still counts as a gap.

> ON THE EXAM
>
> Questions 1-6 Domain 1, 7-13 Domain 2, 14-15 Domain 3, 16-17 Domain 4, 18-20 Domain 5.

Marks lost to **mechanics** are the cheapest to recover. Marks lost to **confusion pairs** are the next cheapest.